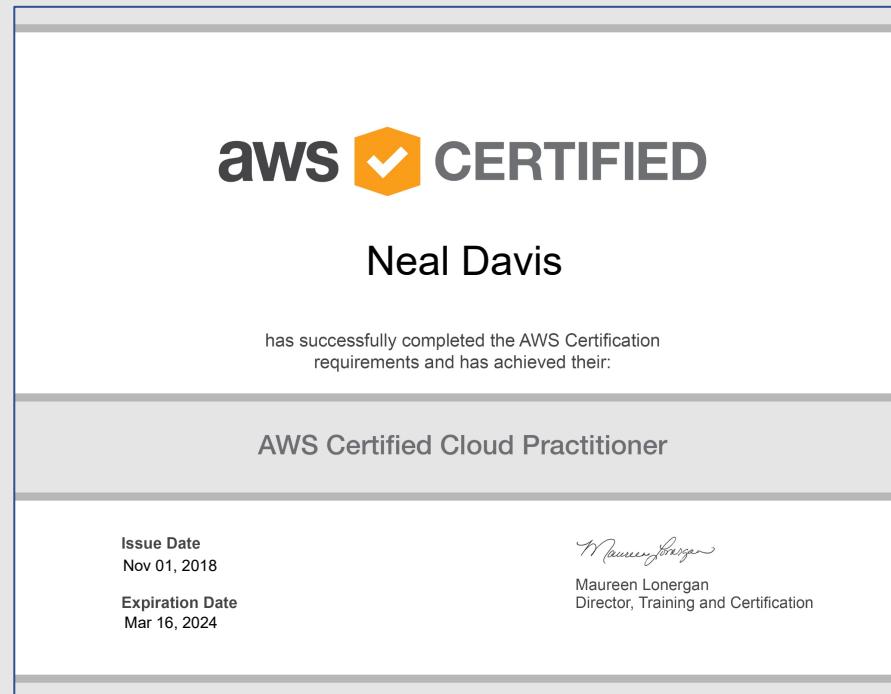


SECTION 1

Let's Get Started!

The CLF-C01 Exam





The CLF-C01 Exam

	Format 65 questions; either multiple choice or multiple response
	Type Foundational
	Delivery Method Testing center or online proctored exam
	Time 90 minutes to complete the exam
	Cost 100 USD (Practice Exam: 20 USD)
	Language Available in English, Indonesian (Bahasa), Japanese, Korean, and Simplified Chinese



The CLF-C01 Exam

Recommended knowledge:

- At least 6 months experience with AWS Cloud in any role including technical, managerial, sales, purchasing or financial
- No experience necessary for this course!



The CLF-C01 Exam

Time and Length:

- 90 minutes
- 65 questions

Scoring:

- Scaled score between 100 – 1000
- Minimum passing score of 700



The CLF-C01 Exam

Question format:

- Multiple-choice: Has one correct response and three incorrect responses
- Multiple-response: Has two or more correct responses out of five or more options



The CLF-C01 Exam

Domain 1: Cloud Concepts

- 1.1 Define the AWS Cloud and its value proposition
- 1.2 Identify aspects of AWS Cloud economics
- 1.3 List the different cloud architecture design principles

Domain 2: Security

- 2.1 Define the AWS Shared Responsibility model
- 2.2 Define AWS Cloud security and compliance concepts
- 2.3 Identify AWS access management capabilities
- 2.4 Identify resources for security support



The CLF-C01 Exam

Domain 3: Technology

- 3.1 Define methods of deploying and operating in the AWS Cloud
- 3.2 Define the AWS global infrastructure
- 3.3 Identify the core AWS services
- 3.4 Identify resources for technology support

Domain 4: Billing and Pricing

- 4.1 Compare and contrast the various pricing models for AWS
- 4.2 Recognize the various account structures in relation to AWS billing and pricing
- 4.3 Identify resources available for billing support



The CLF-C01 Exam

Domain	% of Exam
Domain 1: Cloud Concepts	26%
Domain 2: Security and Compliance	25%
Domain 3: Technology	33%
Domain 4: Billing and Pricing	16%
TOTAL	100%

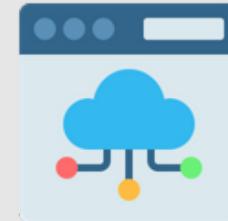
SECTION 2

Create AWS Free Tier Account

SECTION 3

Cloud Computing and AWS

Traditional IT and Cloud Computing

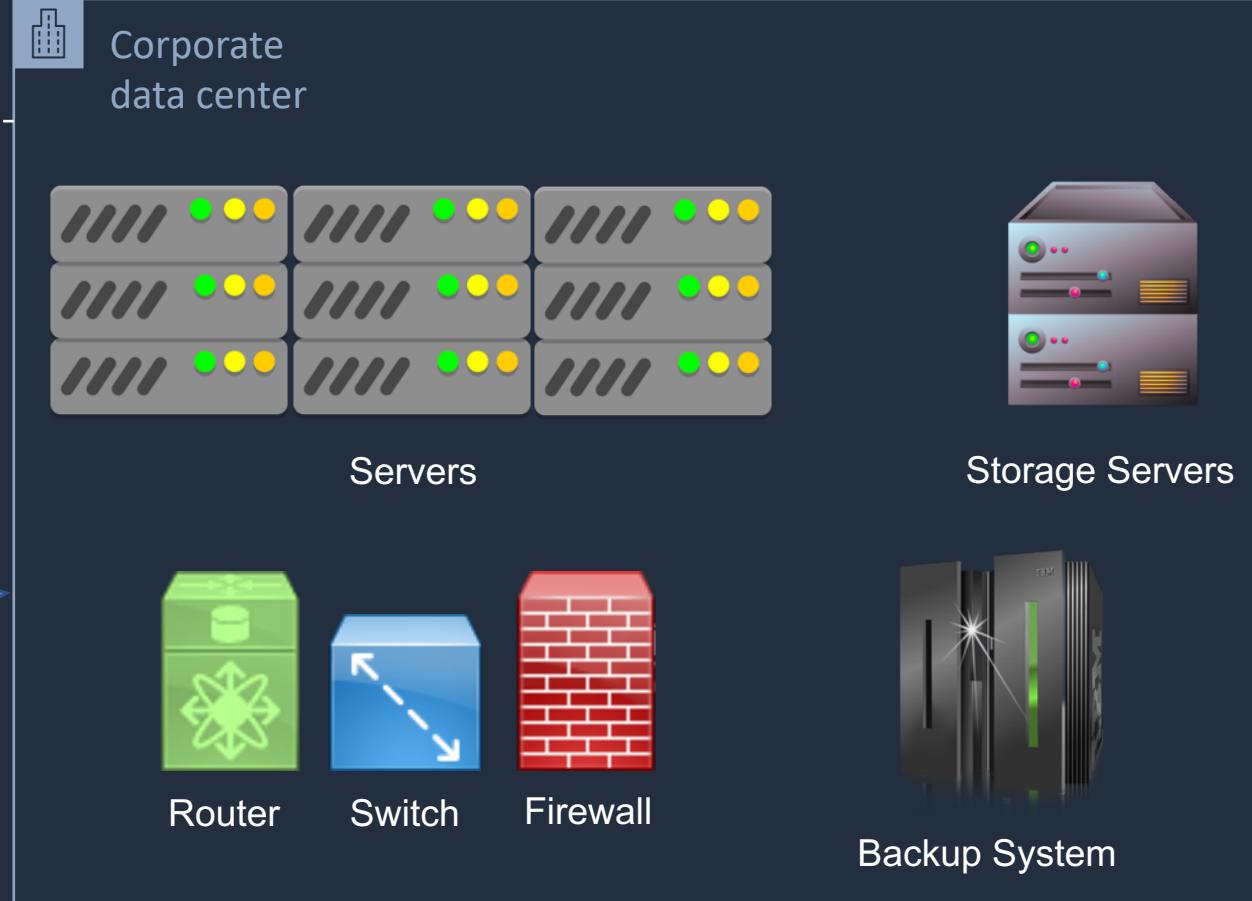




Legacy IT / Traditional IT



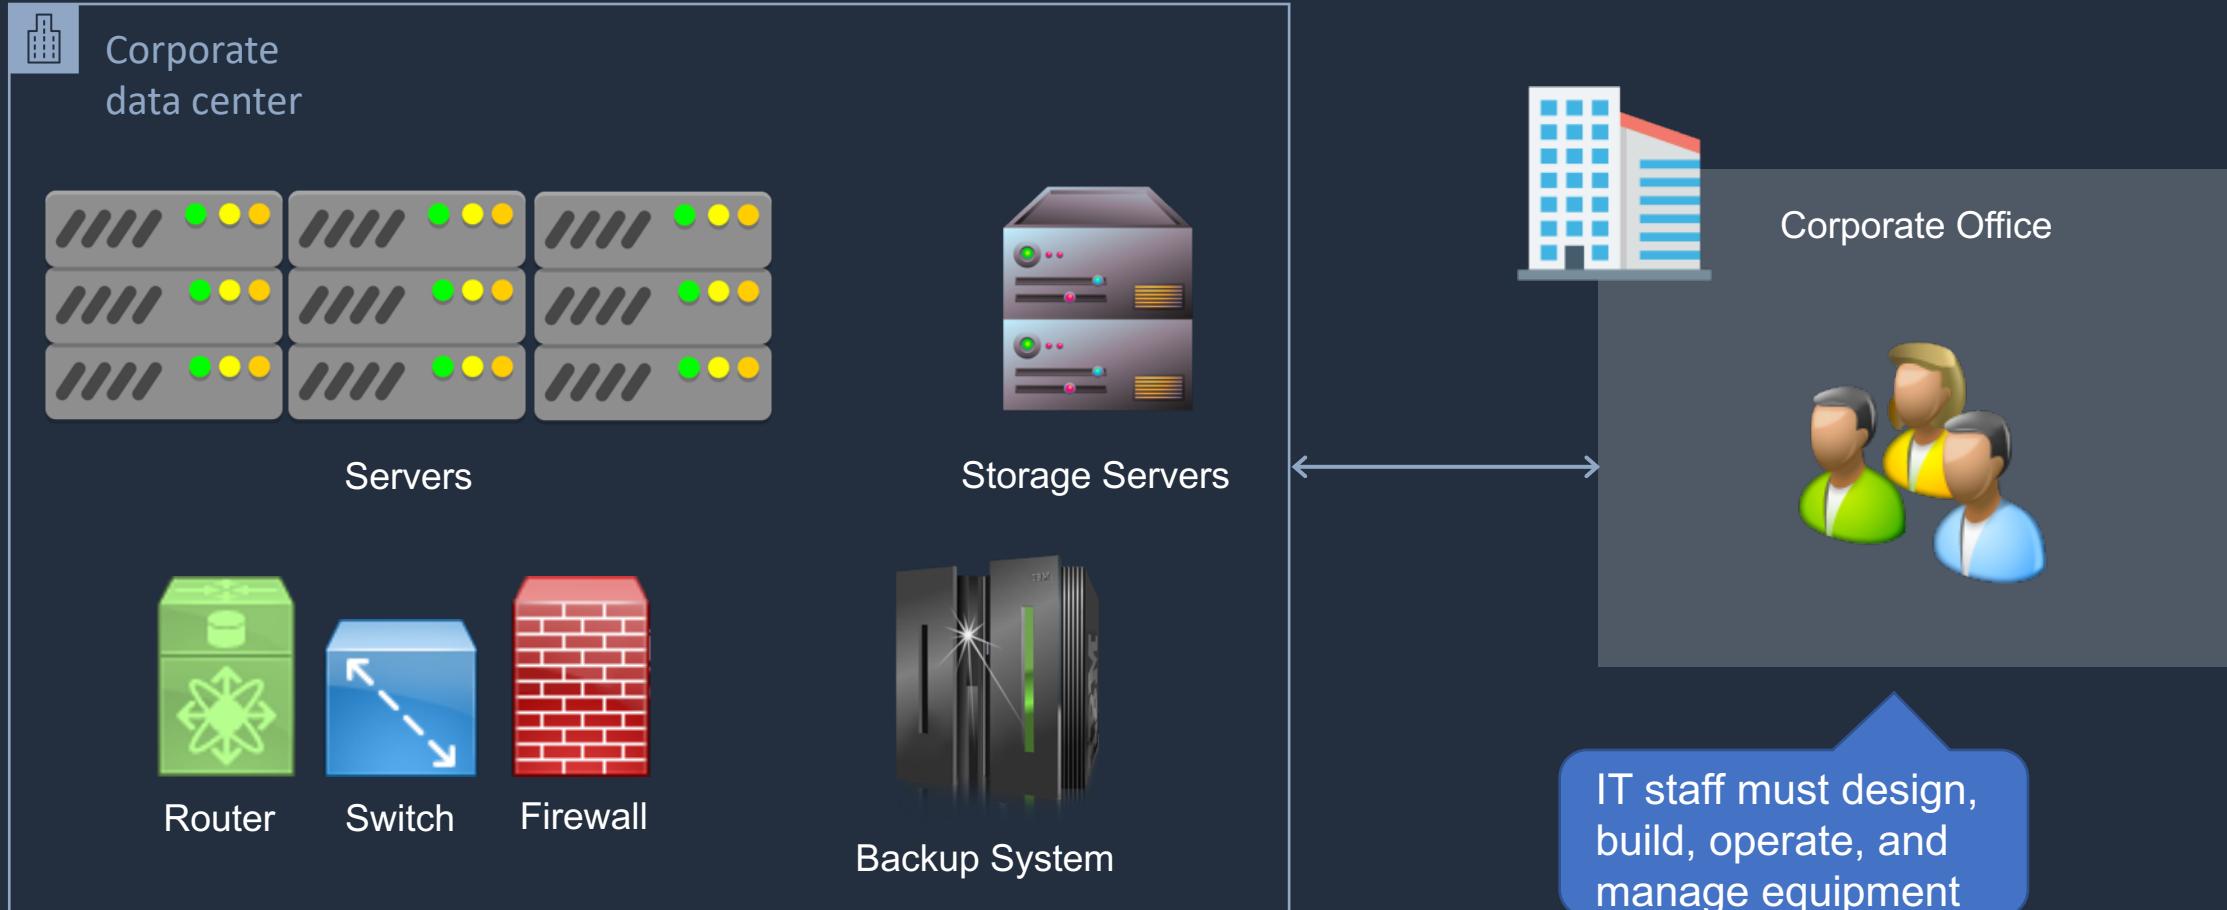
This model is very capital intensive



A company typically **leases** space in a data center, or may **own** the whole building

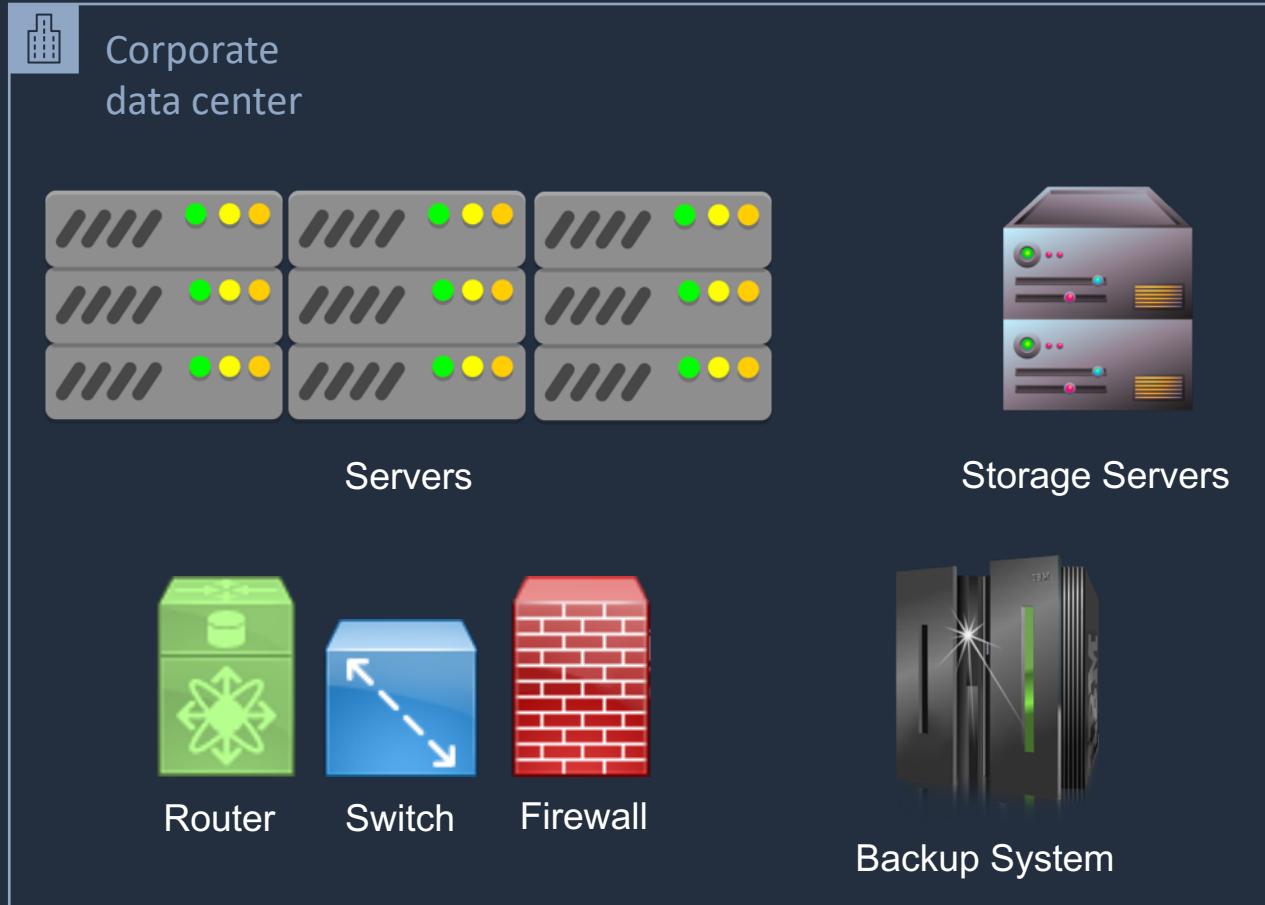


Legacy IT / Traditional IT





Legacy IT / Traditional IT



Costs:

- Data center building
- Data center security
- Physical IT hardware
- Software licensing costs
- Maintenance contracts
- Power
- Internet connectivity
- Staff wages (design, build, operations, maintenance)



Definition of Cloud Computing

On-demand,
self-service

Broad network
access

Resource
pooling



Gmail



The Internet



Dropbox



Salesforce.com

Rapid
elasticity

Measured
service



Amazon Web Services



Cloud vs Traditional IT

Cloud Computing

On-demand, self-service

Broad network access

Resource pooling

Rapid elasticity

Measured service

Traditional IT

Requires human involvement

Internal accessibility, limited public presence

Single-tenant, can be virtualized

Limited scalability

Usage is not typically measured

Examples and Benefits of Cloud Computing





Examples of Cloud Computing

Non-Cloud Services:



Email Server



File Server



Customer Relationship
Management (CRM)

Cloud Services:



Gmail



Dropbox



Salesforce

You don't **own** or **manage** the infrastructure on which the service runs

Cloud services are offered on a **subscription / consumption** model

The service **scales** as **demand** changes



Deploying a Website On-Premises

Assumes you don't have a private cloud, or don't have enough capacity

Activity:

Timeline:

1) Purchase hardware

4-12 weeks

2) Install and build

4-8 weeks

3) Acceptance testing

2-4 weeks

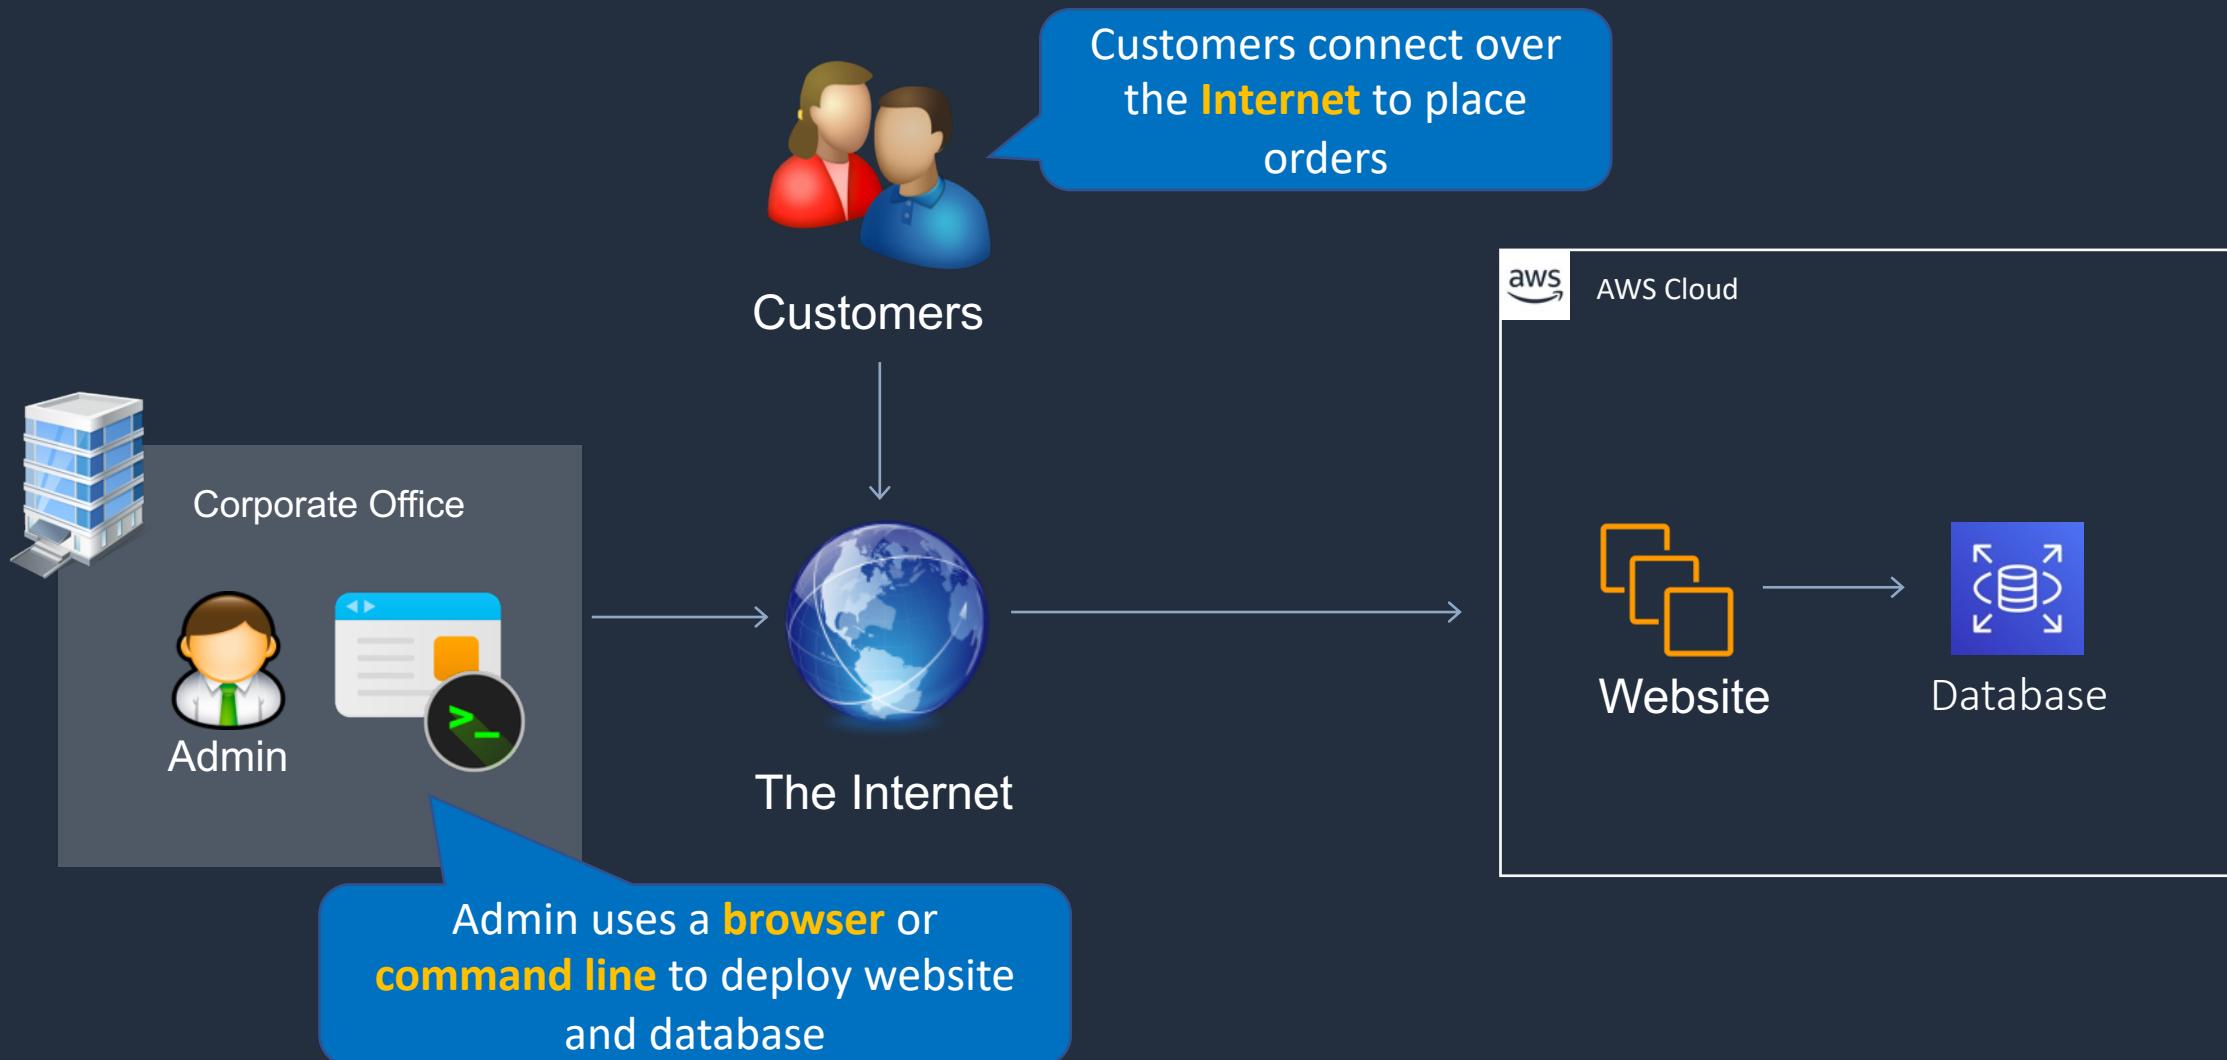
1) Handover to operations

1-2 weeks

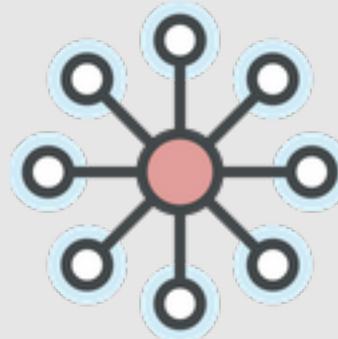
3-6 months !



Deploying a Website in the Cloud

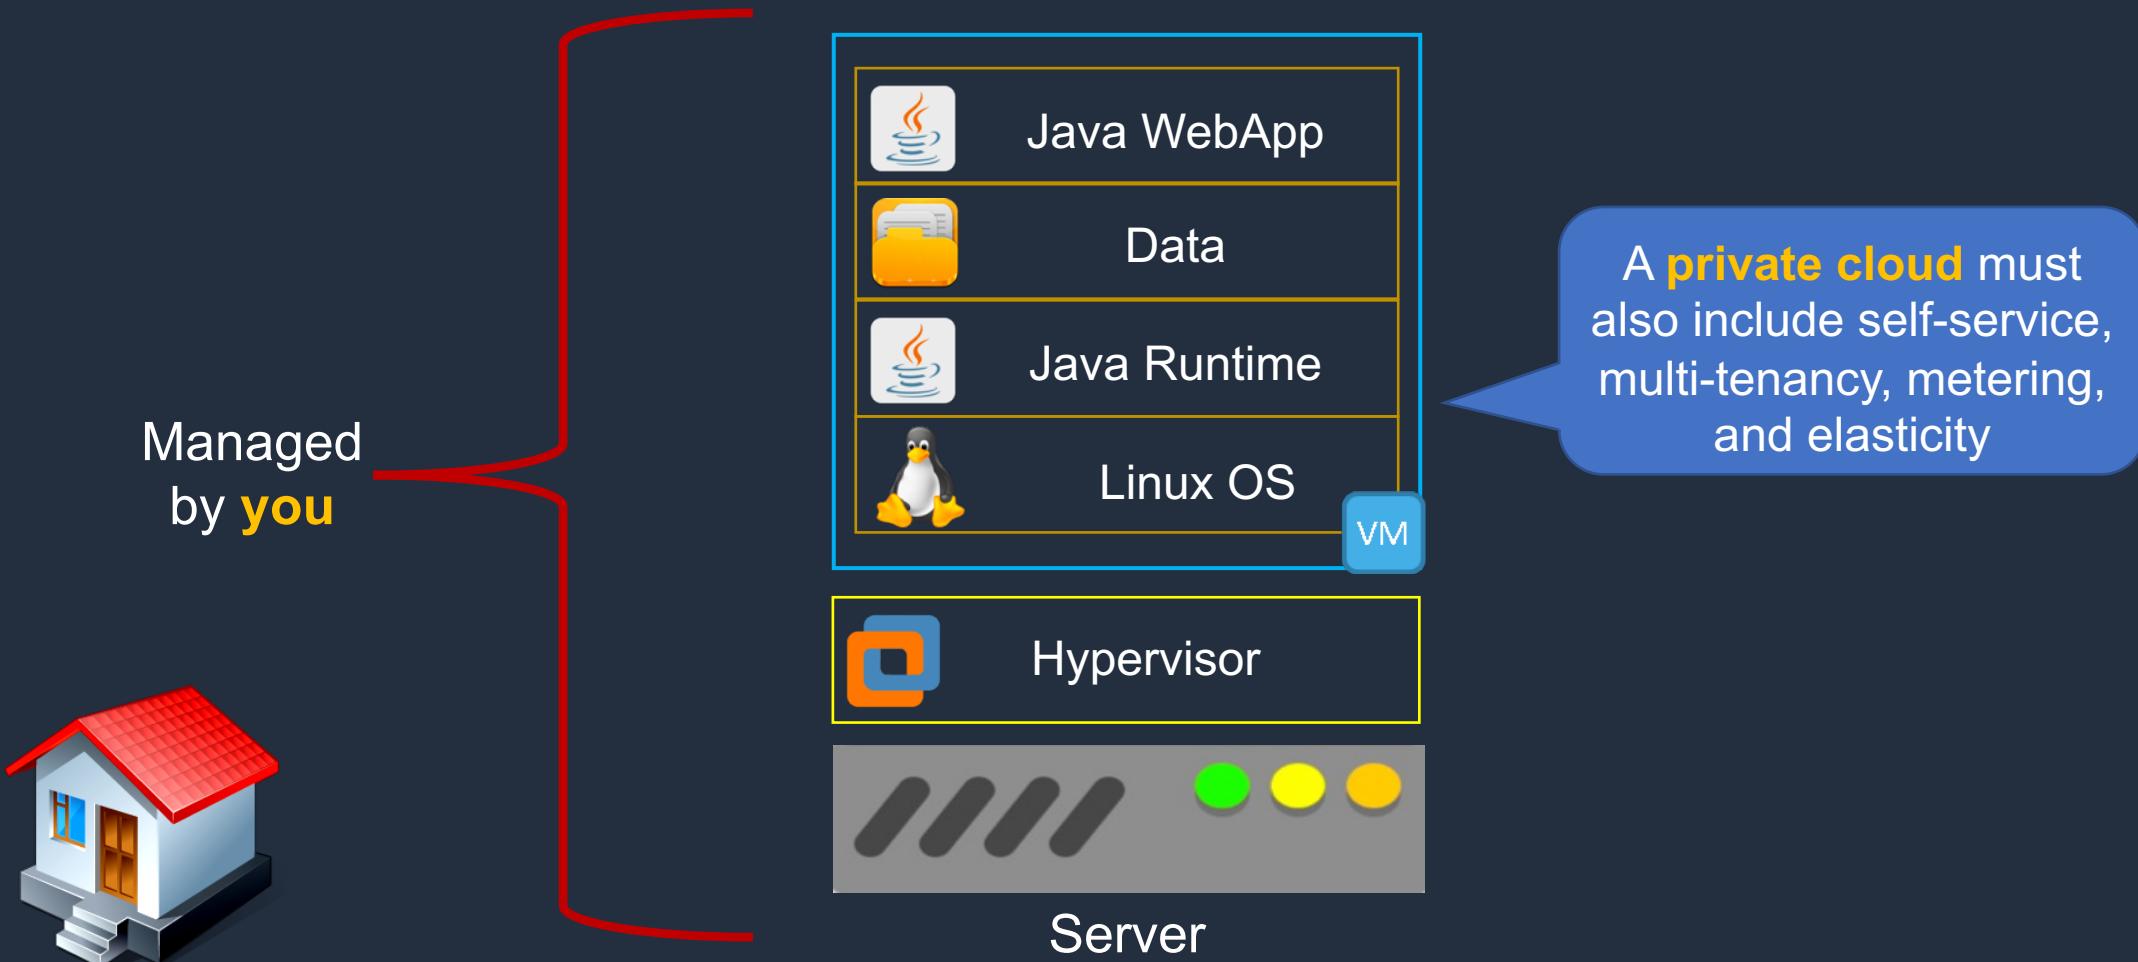


Types of Cloud Service and Deployment



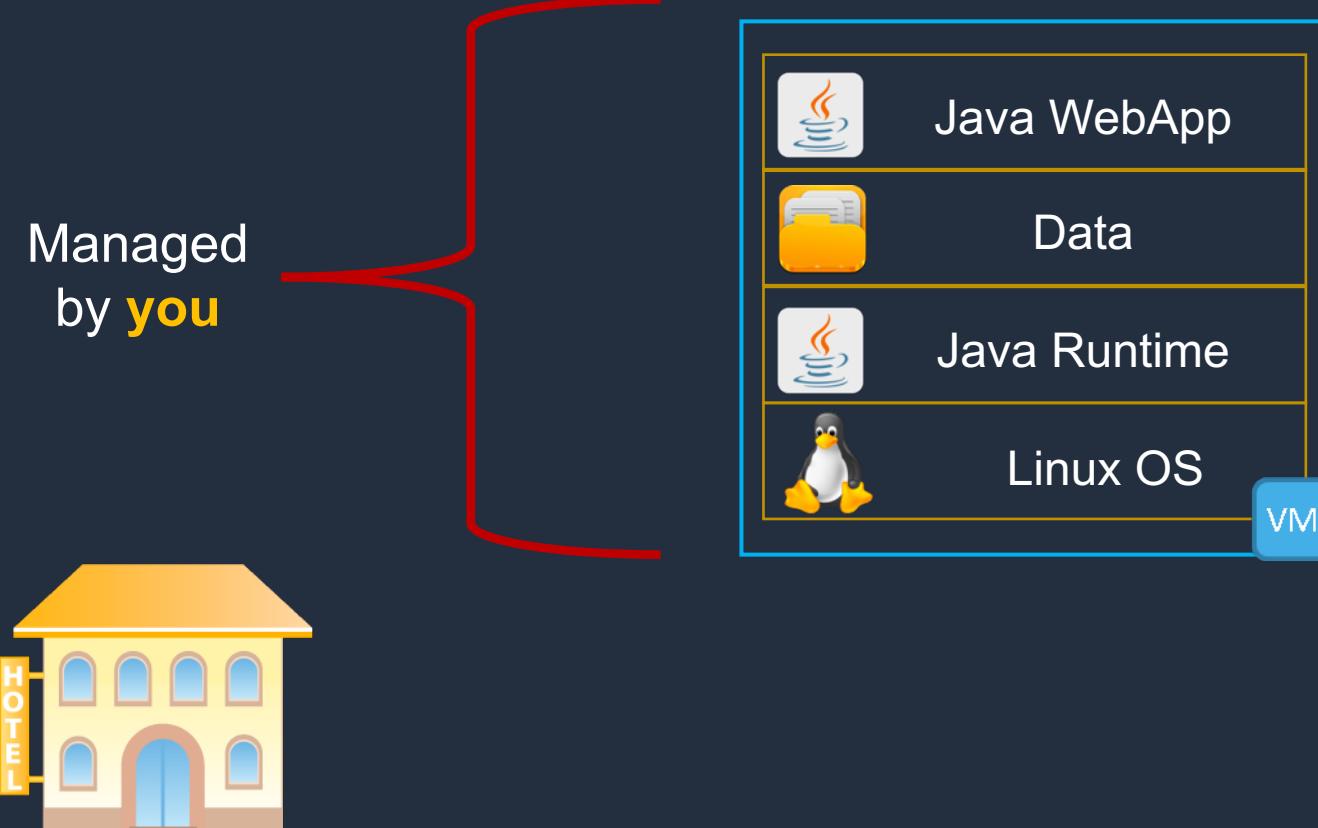


Cloud Service Models: Private Cloud





Cloud Service Models: Infrastructure as a Service (IaaS)



Examples:

- Amazon Elastic Compute Cloud (EC2)
- Azure Virtual Machines
- Google Compute Engine



Cloud Service Models: Platform as a Service (PaaS)

Managed
by **you**



Examples:

- AWS Elastic Beanstalk
- Azure WebApps
- Compute App Engine



Cloud Service Models: Software as a Service (SaaS)

Managed
by you



Java WebApp

Pure consumption
model

Examples:

- Google Apps
- Salesforce.com
- Zoom



Cloud Service Models: Comparison

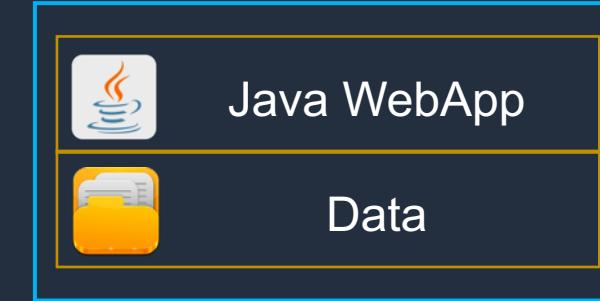
Private Cloud



IaaS



PaaS



You simply upload your **code/data** to create your application

SaaS



You simply **consume** the service - little responsibility + little control

You manage **everything** - greater responsibility + greater control

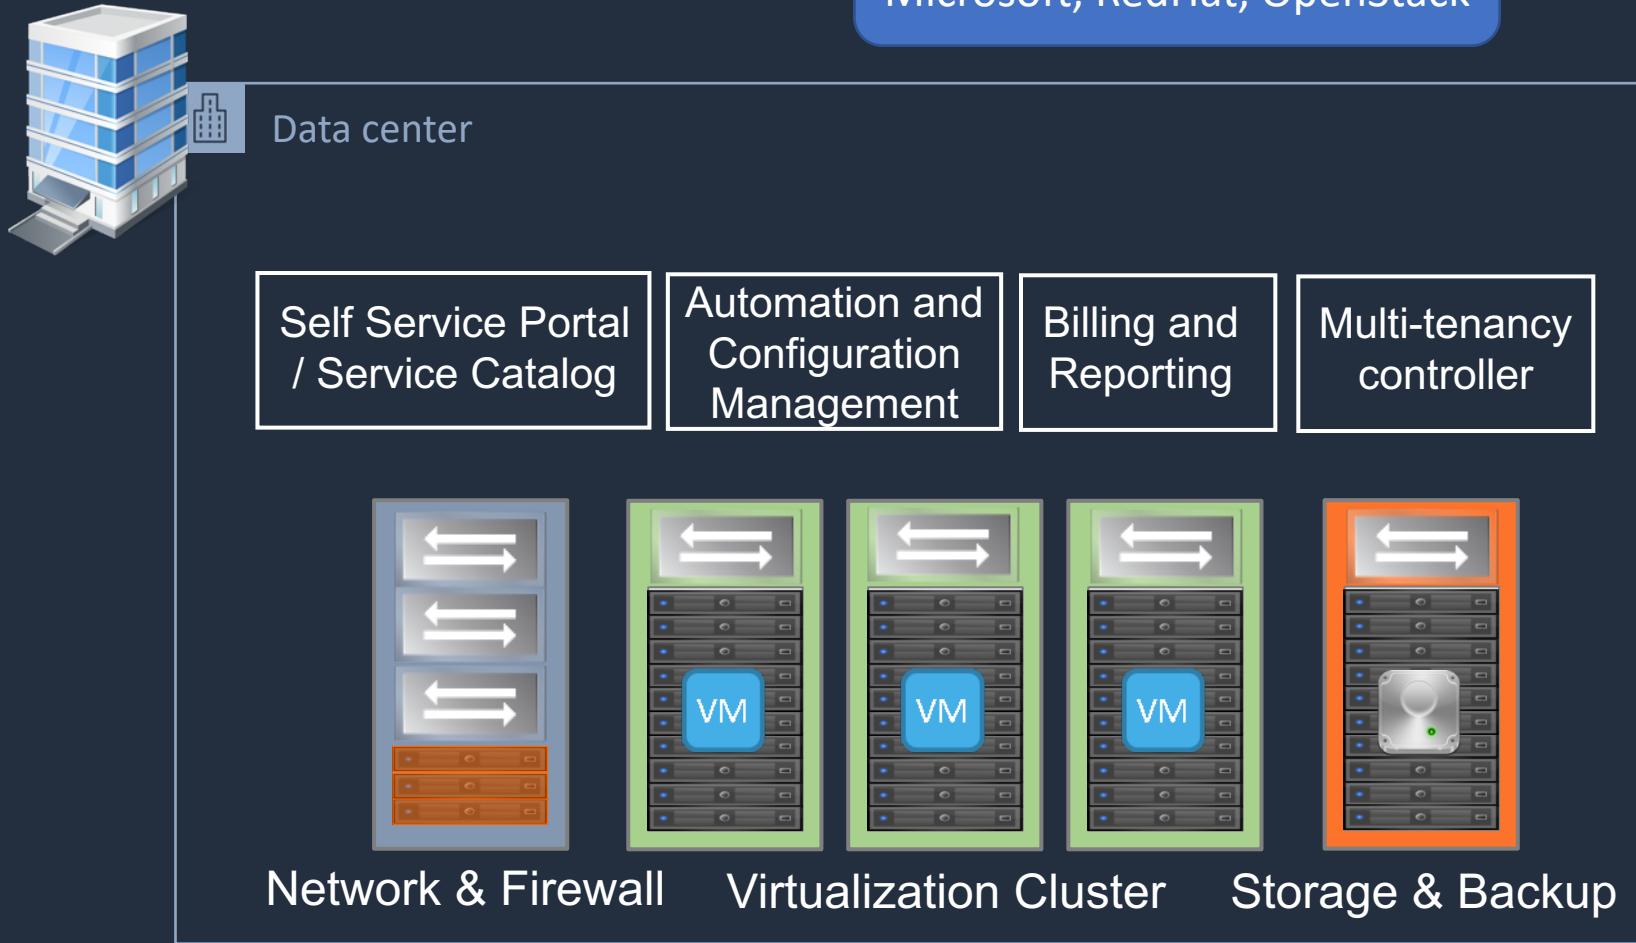


You manage from the **virtual server** upwards





Private Cloud



Cloud management software layer

Benefits

- Complete control of the entire stack
- Security – in a few cases, organizations may need to keep all or some of their applications and data in house

You **build** and **manage** the cloud deployment

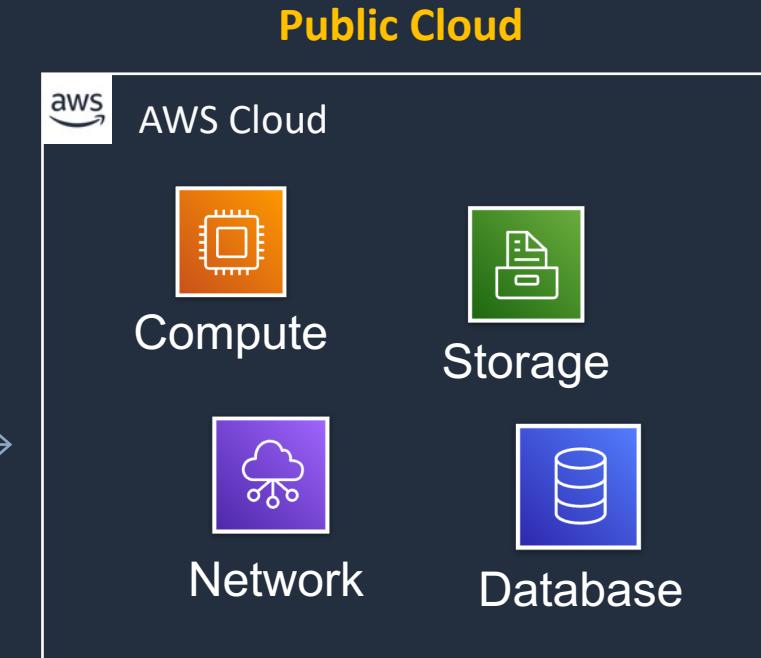


Public Cloud

Examples are AWS, Microsoft Azure, Google Cloud Platform

Benefits:

- Variable expense, instead of capital expense
- Economies of scale
- Massive elasticity



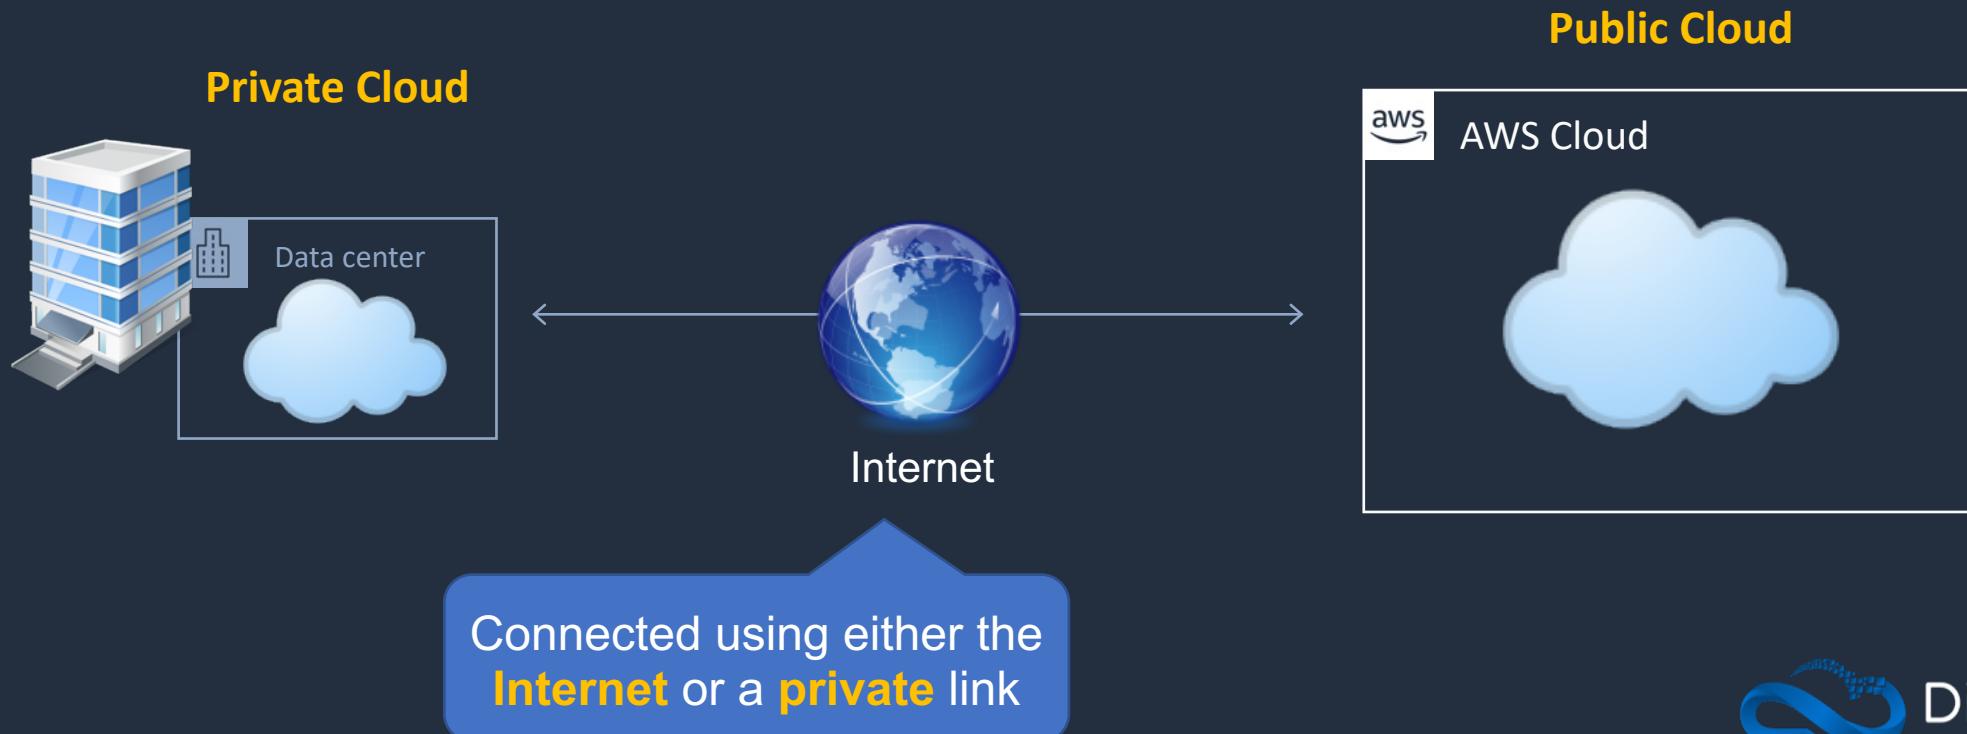
Connected using either the **Internet** or a **private link**



Hybrid Cloud

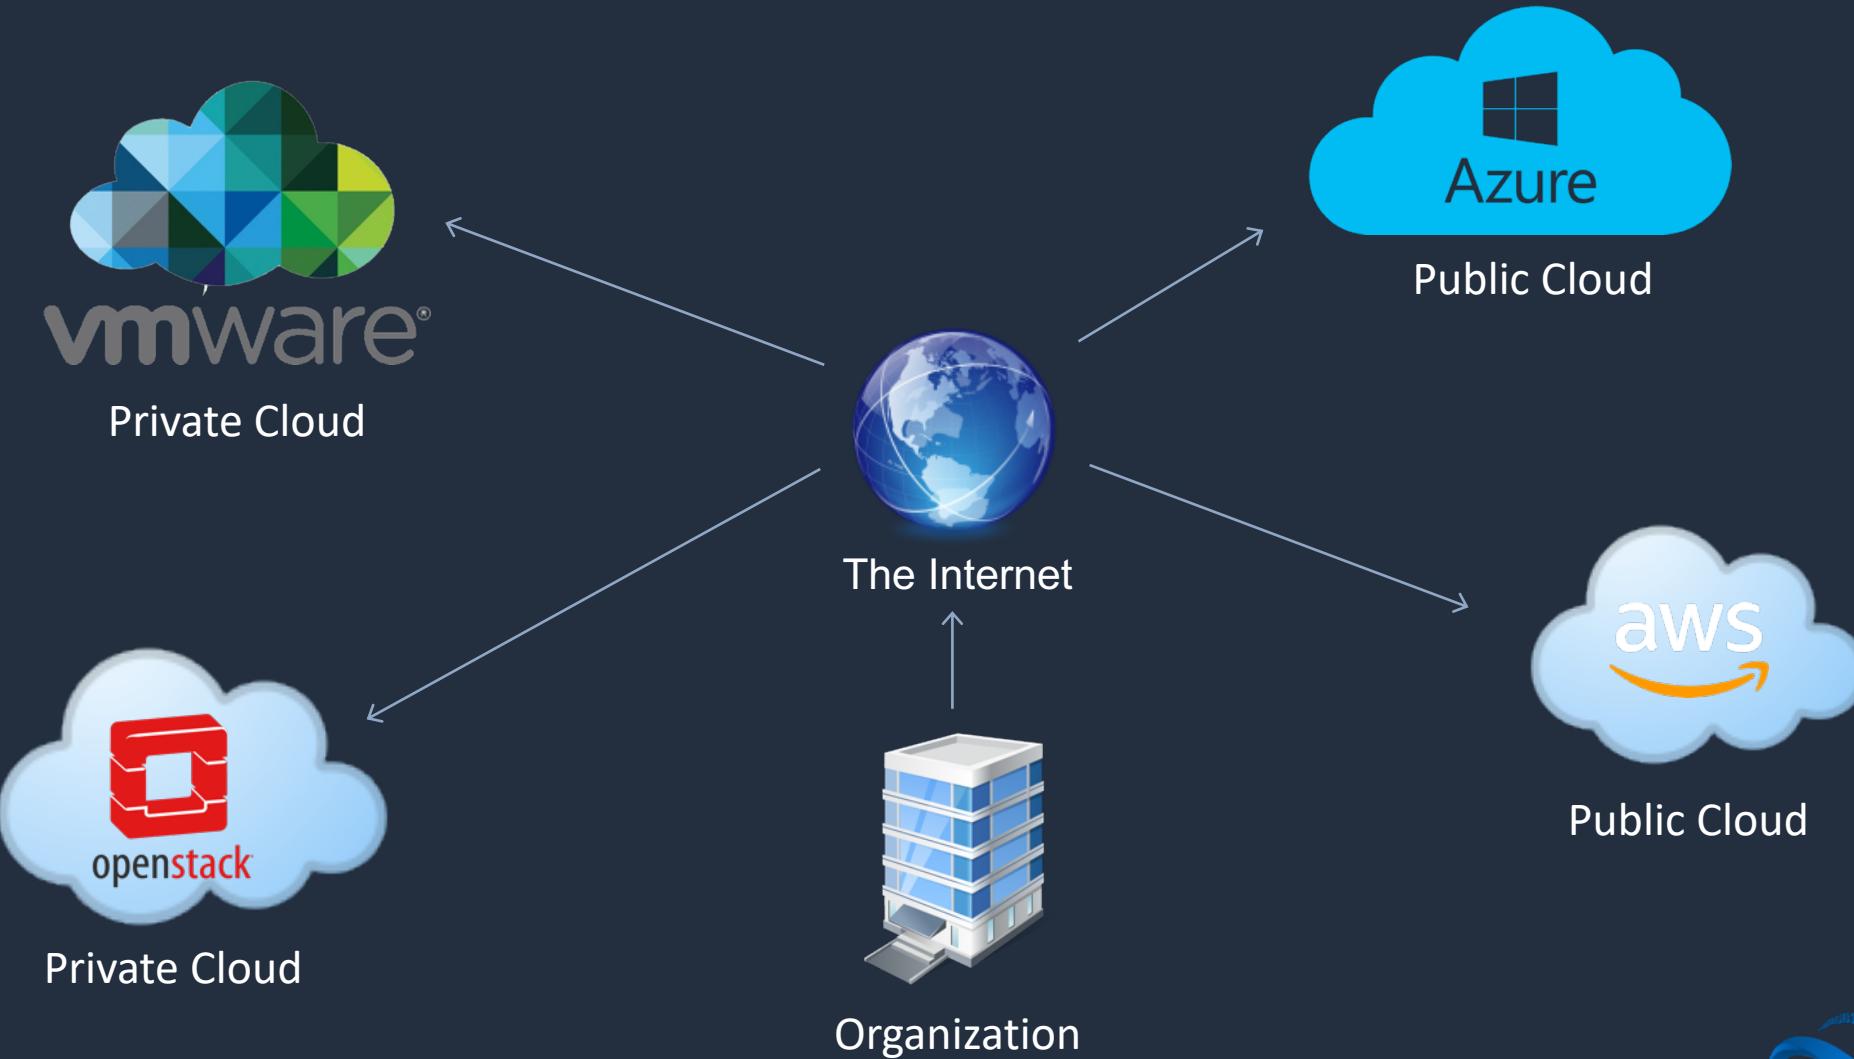
Benefits:

- Allows companies to keep the critical applications and sensitive data in a traditional data center environment or private cloud
- Take advantage of public cloud resources like SaaS, for the latest applications, and IaaS, for elastic virtual resources
- Facilitates portability of data, apps and services and more choices for deployment models



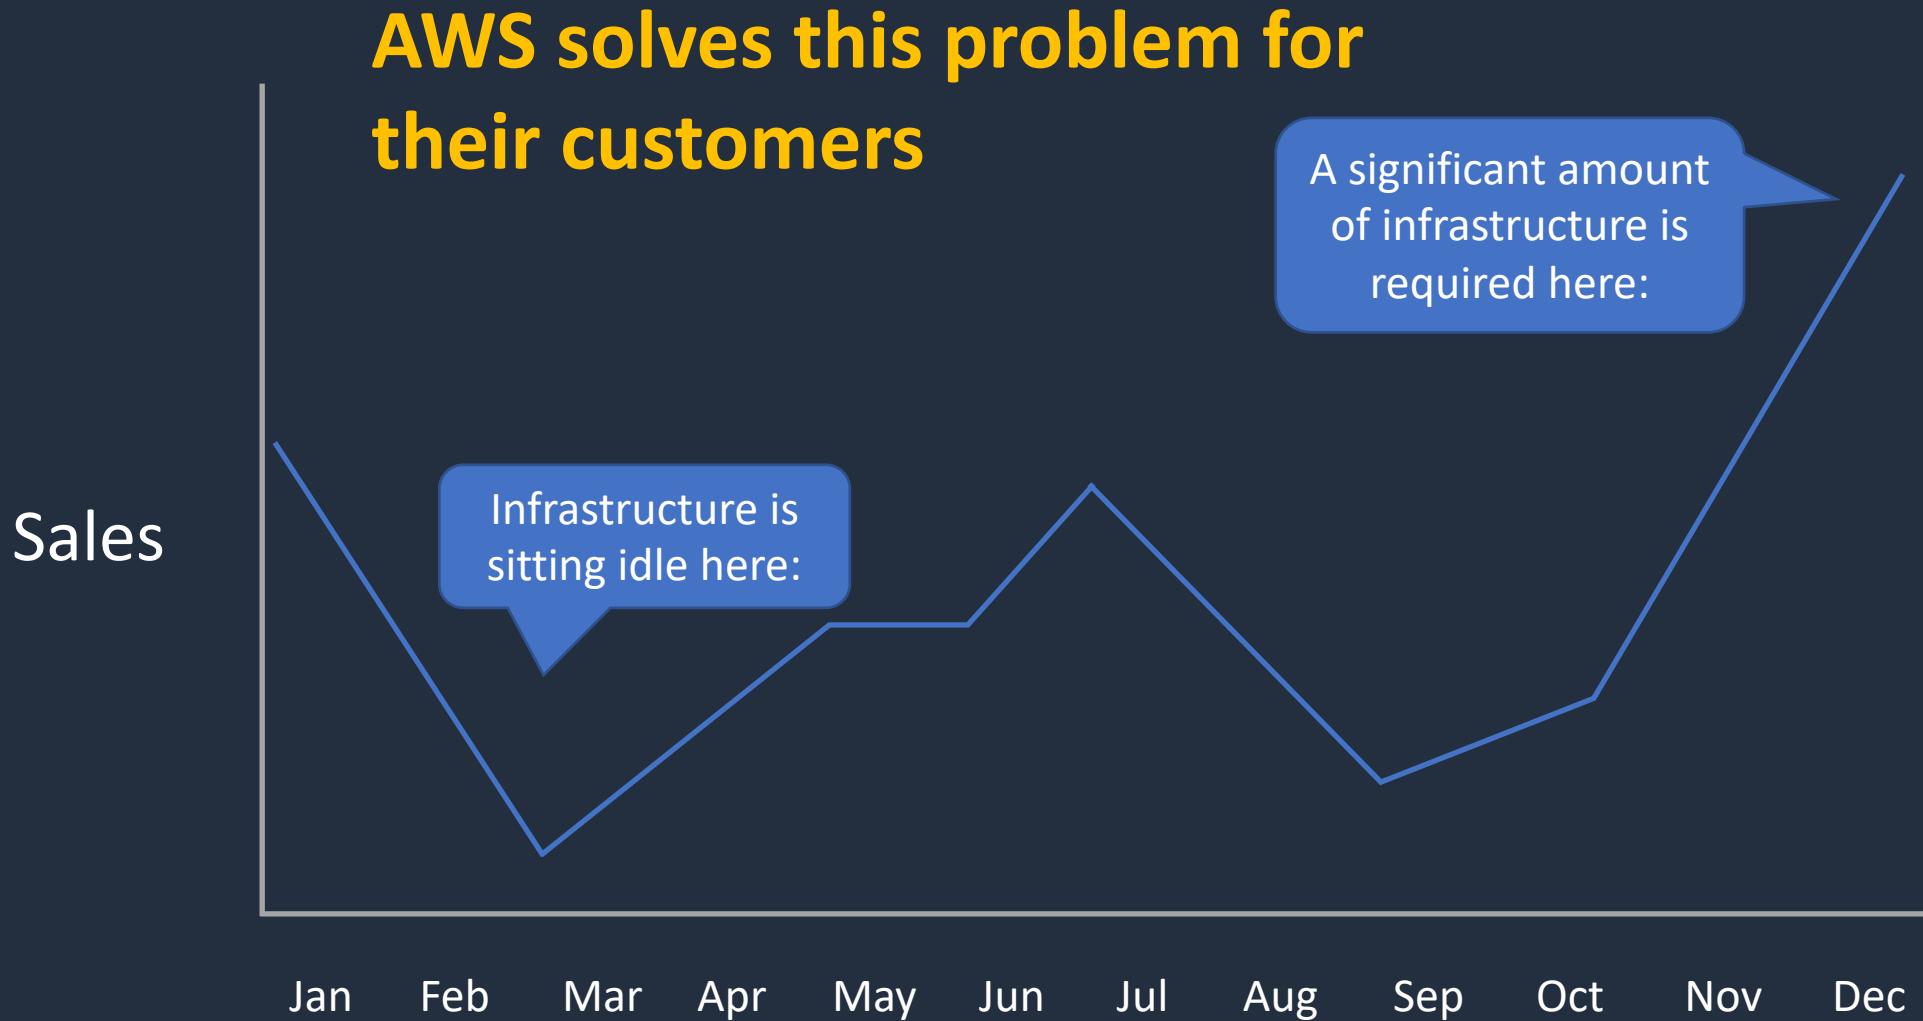


Multicloud



Overview of Amazon Web Services (AWS)





Charge for
services based
on usage

25 Regions
around the
world

Subsidiary
of Amazon



Hyperscale
Public Cloud
Provider

Services are
offered on-
demand

AWS Service Categories (a few examples)



Compute



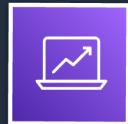
Machine Learning



Storage



Database



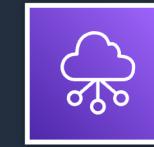
Analytics



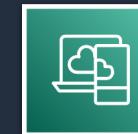
Many more categories
and over **200** services!



Media Services



Networking



End User Computing



Internet of Things

AWS Pricing Fundamentals

Compute



Amount of resources such as CPU and RAM and duration

Storage



Quantity of data stored

Outbound Data Transfer

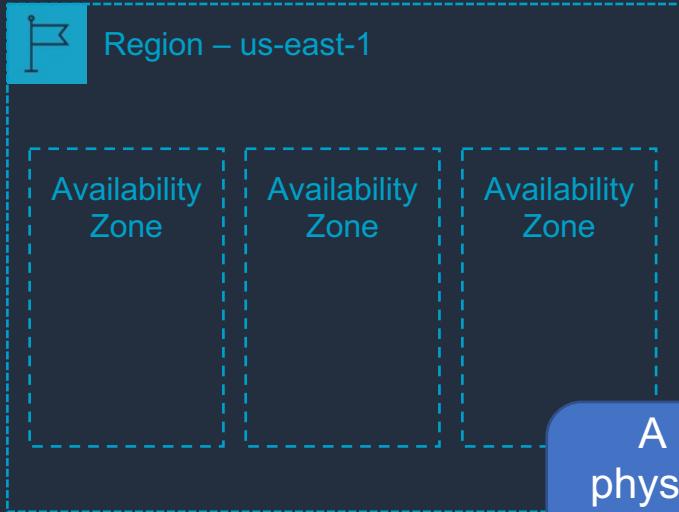


Quantity of data that is transferred out from all services

The AWS Global Infrastructure



AWS Global Infrastructure



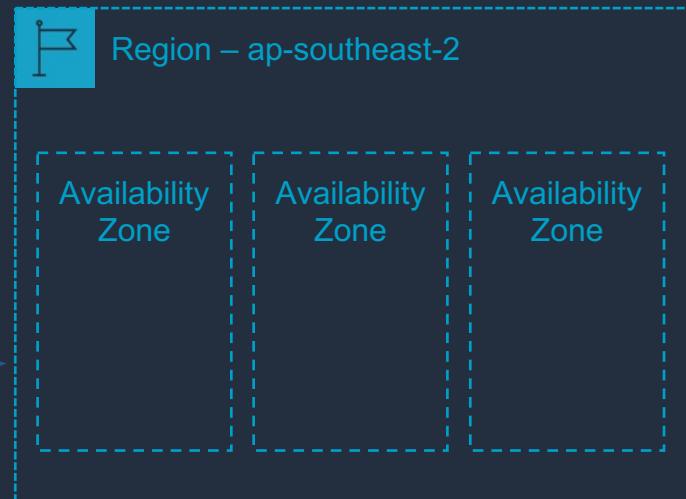
Every region is **connected** via a high bandwidth, fully redundant network

There are 25 **regions** around the world



A **Region** is a physical location in the world and is **independent**

Each region consists of **two or more** Availability Zones



An Availability Zone is composed of **one or more** data centers

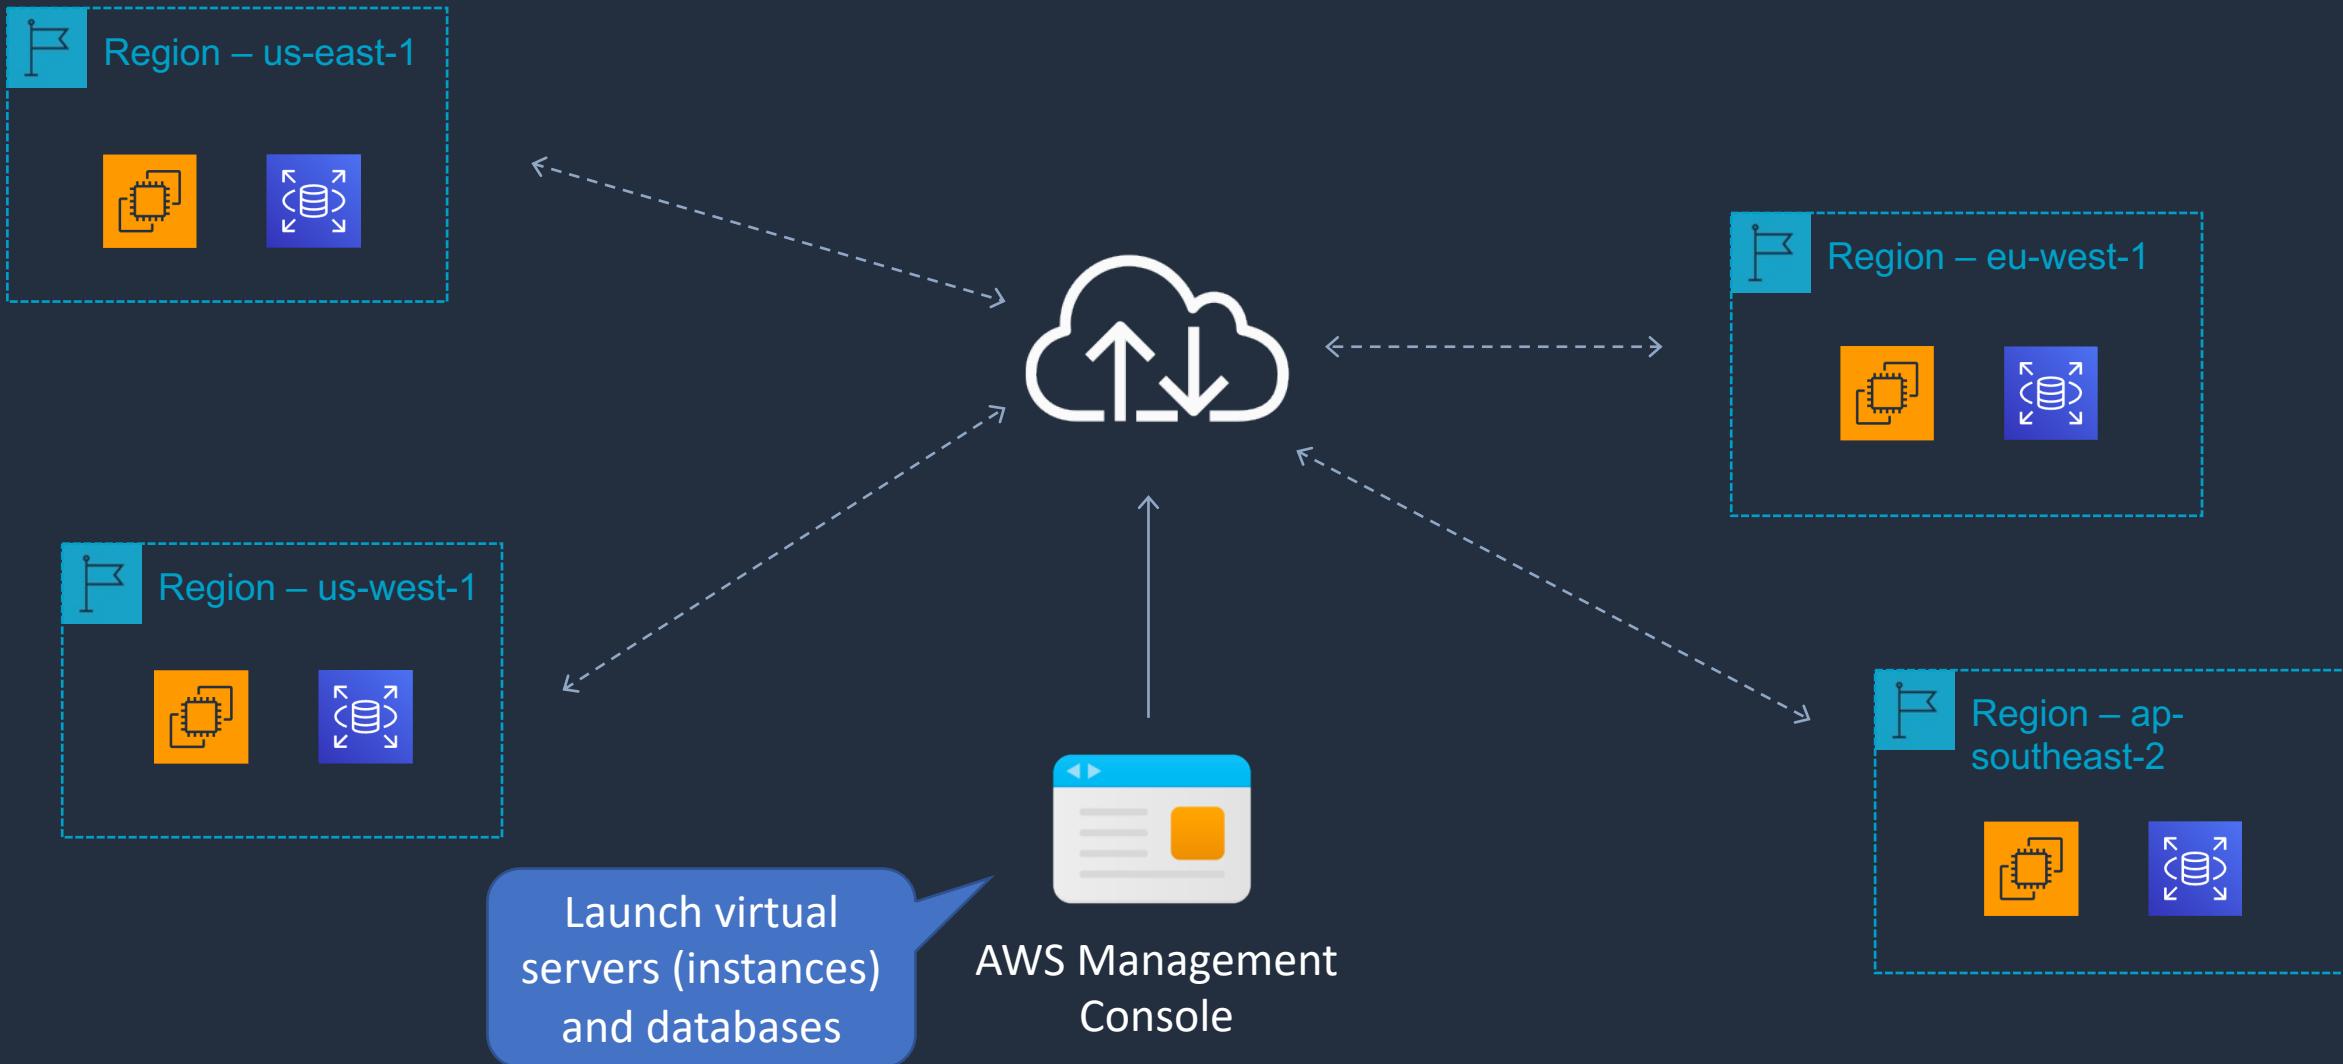
Local Zones extend regions closer to end-users



AWS Local Zones



Deploying Services Globally

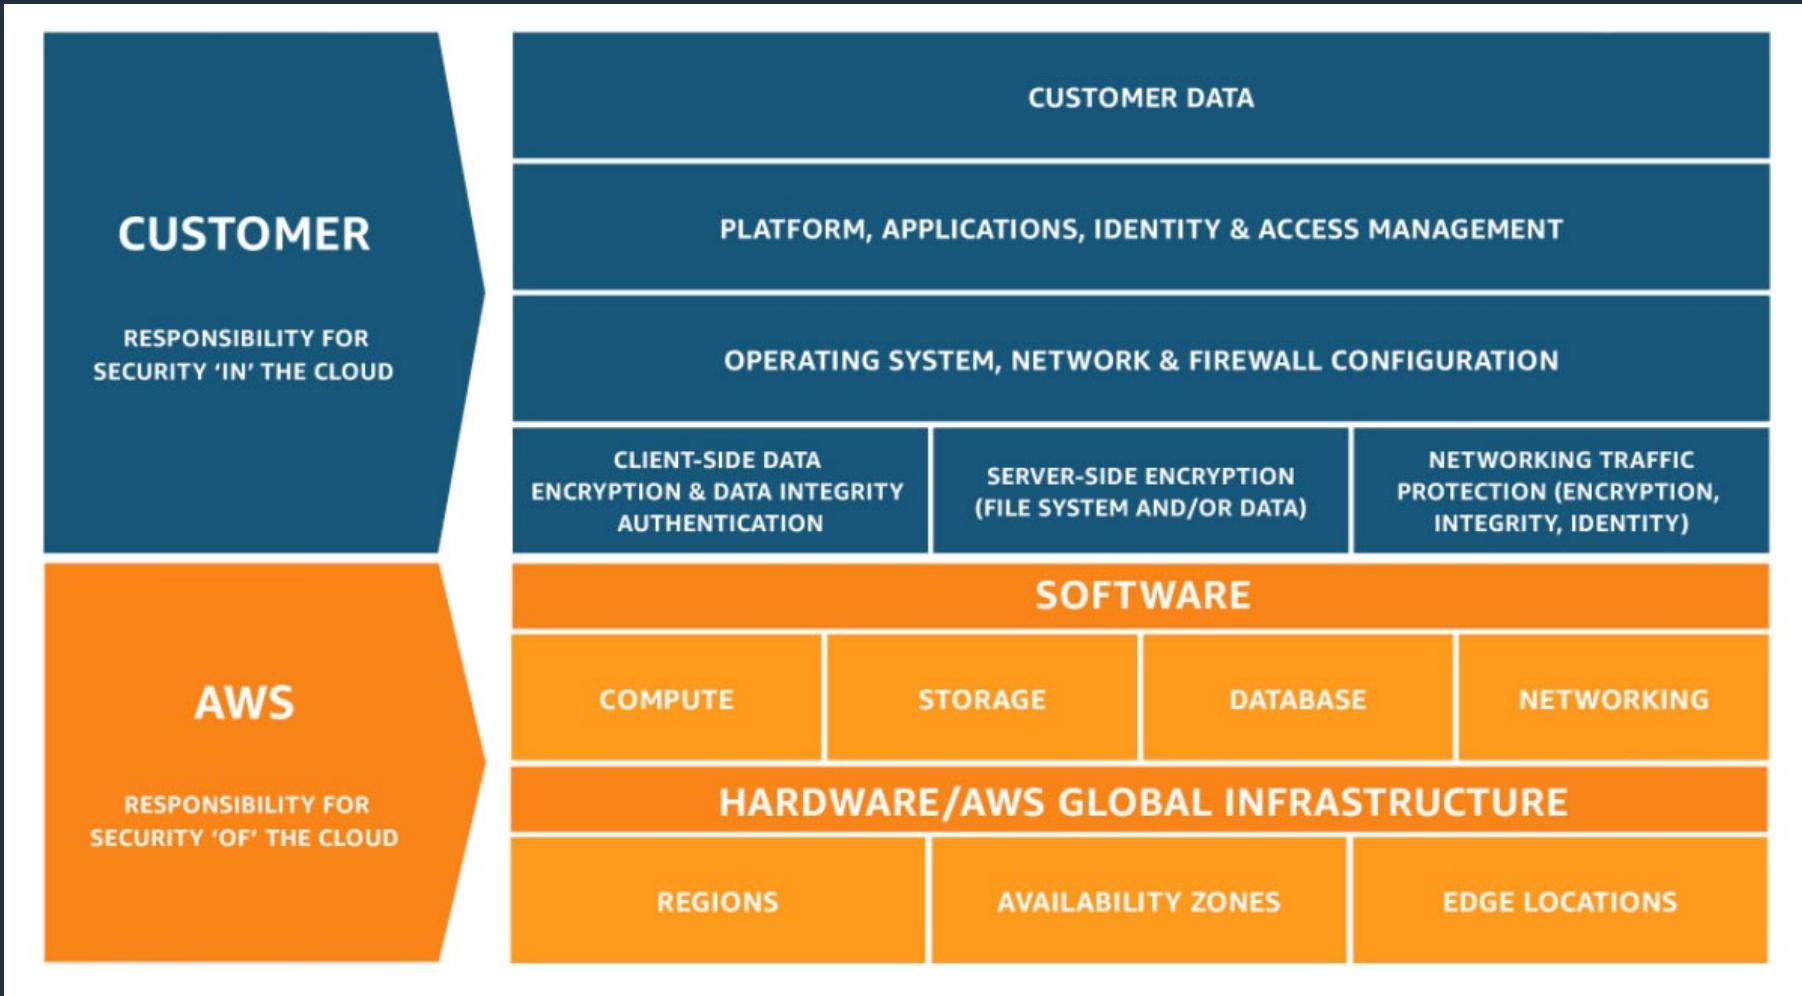


The AWS Shared Responsibility Model





The AWS Shared Responsibility Model





The AWS Shared Responsibility Model

CUSTOMER RESPONSIBILITY



Bucket with objects



Role



Multi-Factor Authentication



Security Group



Patch management



Staff training



Data encryption



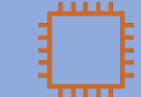
IAM User



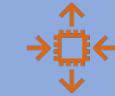
Network ACL



SSL encryption



EC2 Instance



Auto Scaling



Elastic load balancer

AWS RESPONSIBILITY



Data center



Data center security



Network router



Network switch



Server



Storage



Database Server

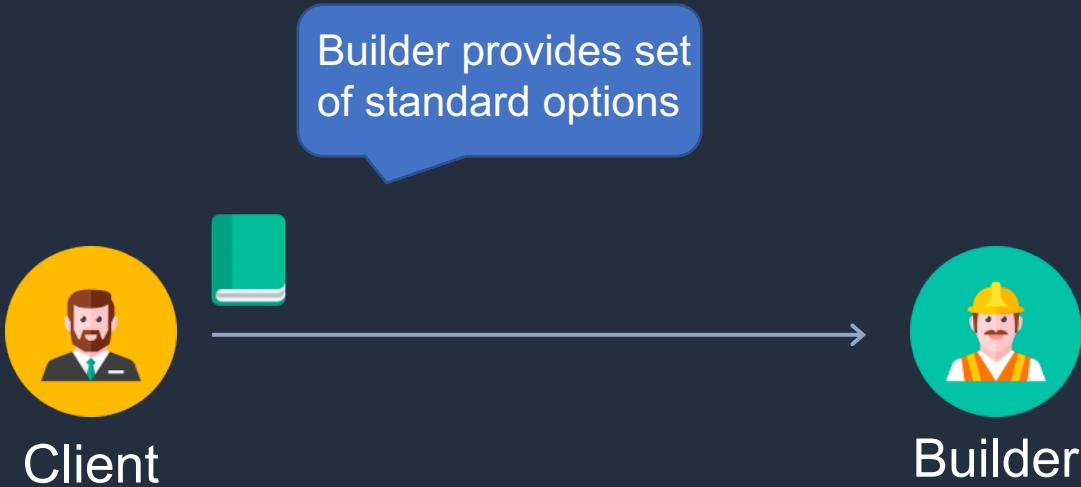


Disk drive

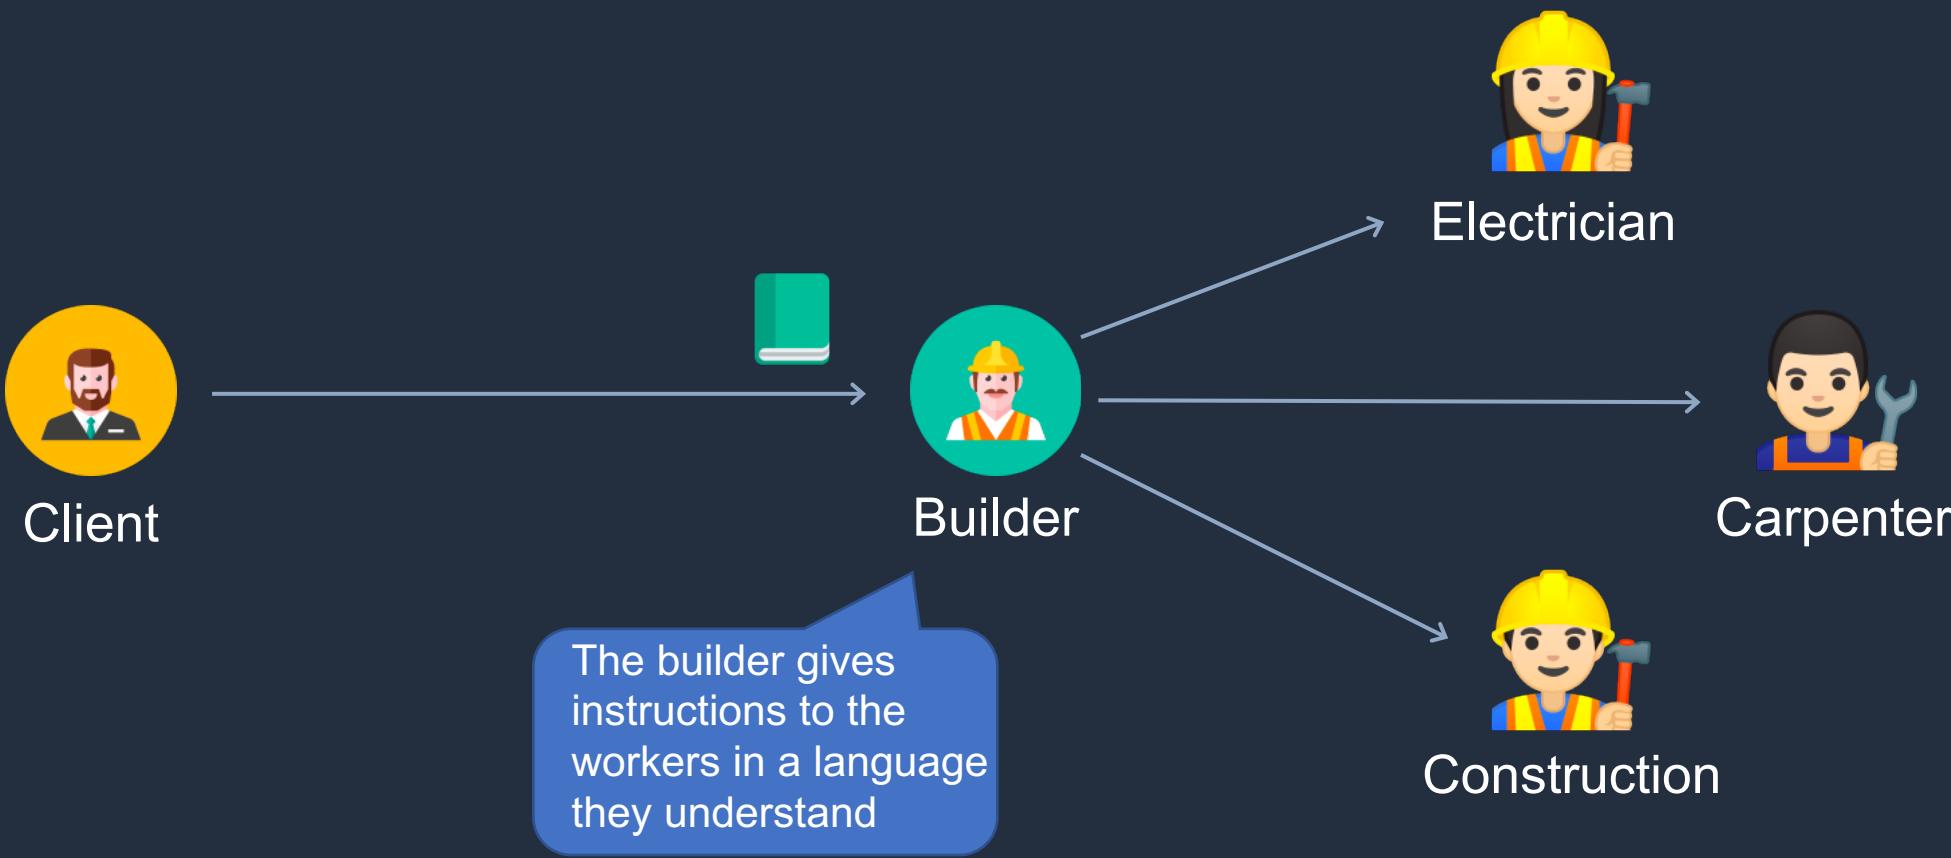
Application Programming Interfaces (APIs)



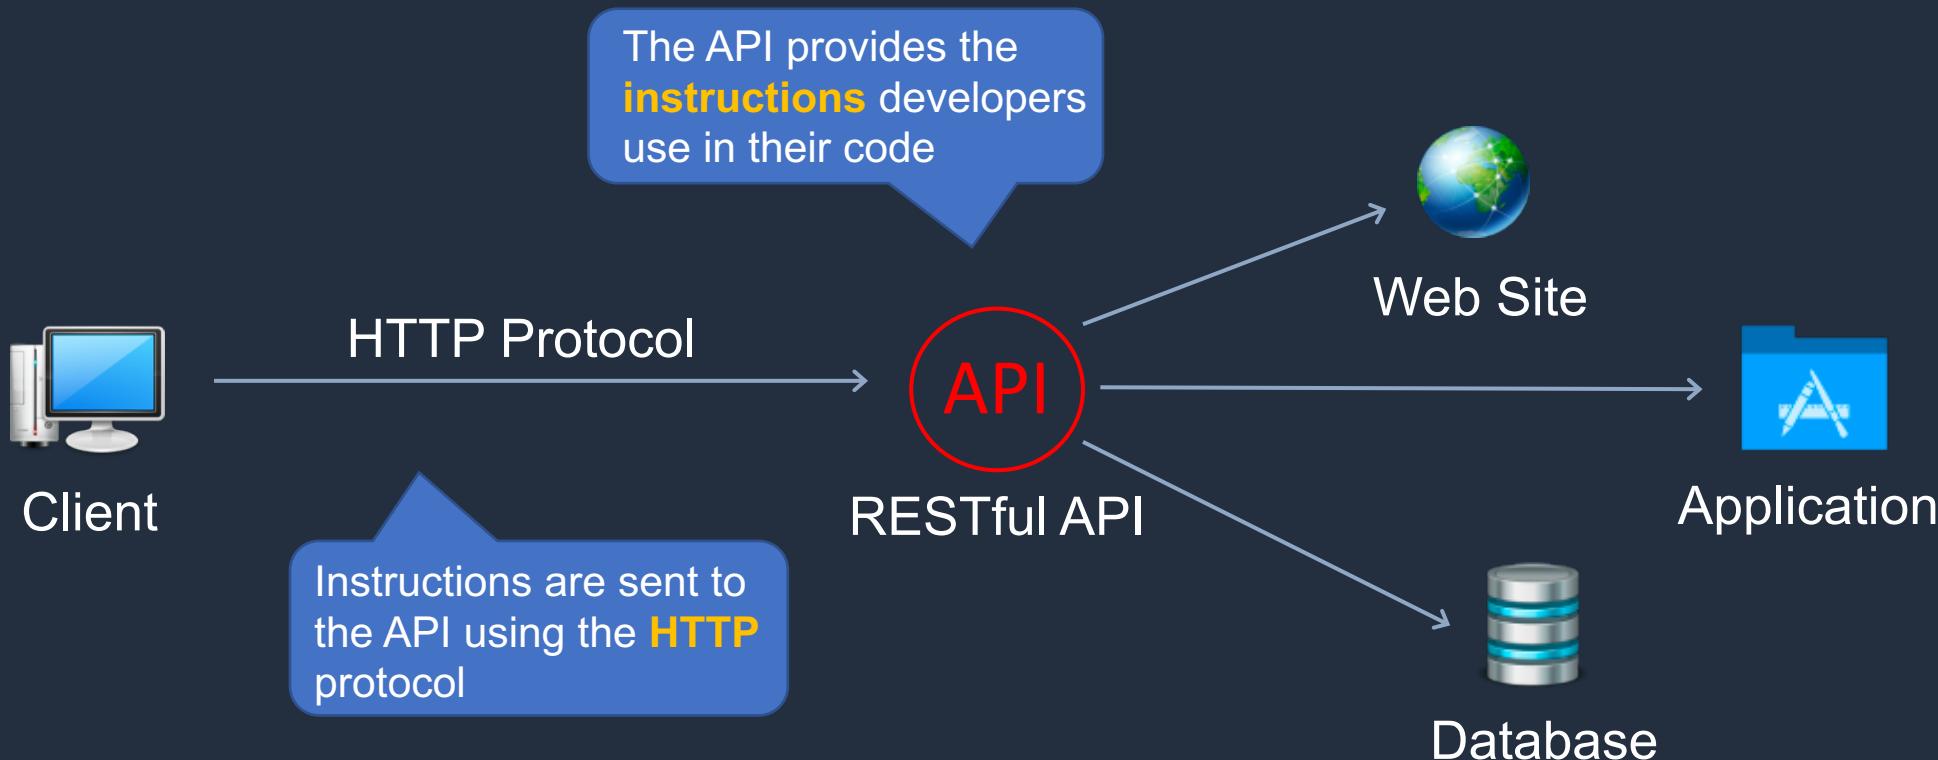
Application Programming Interfaces (APIs) – Building a house analogy



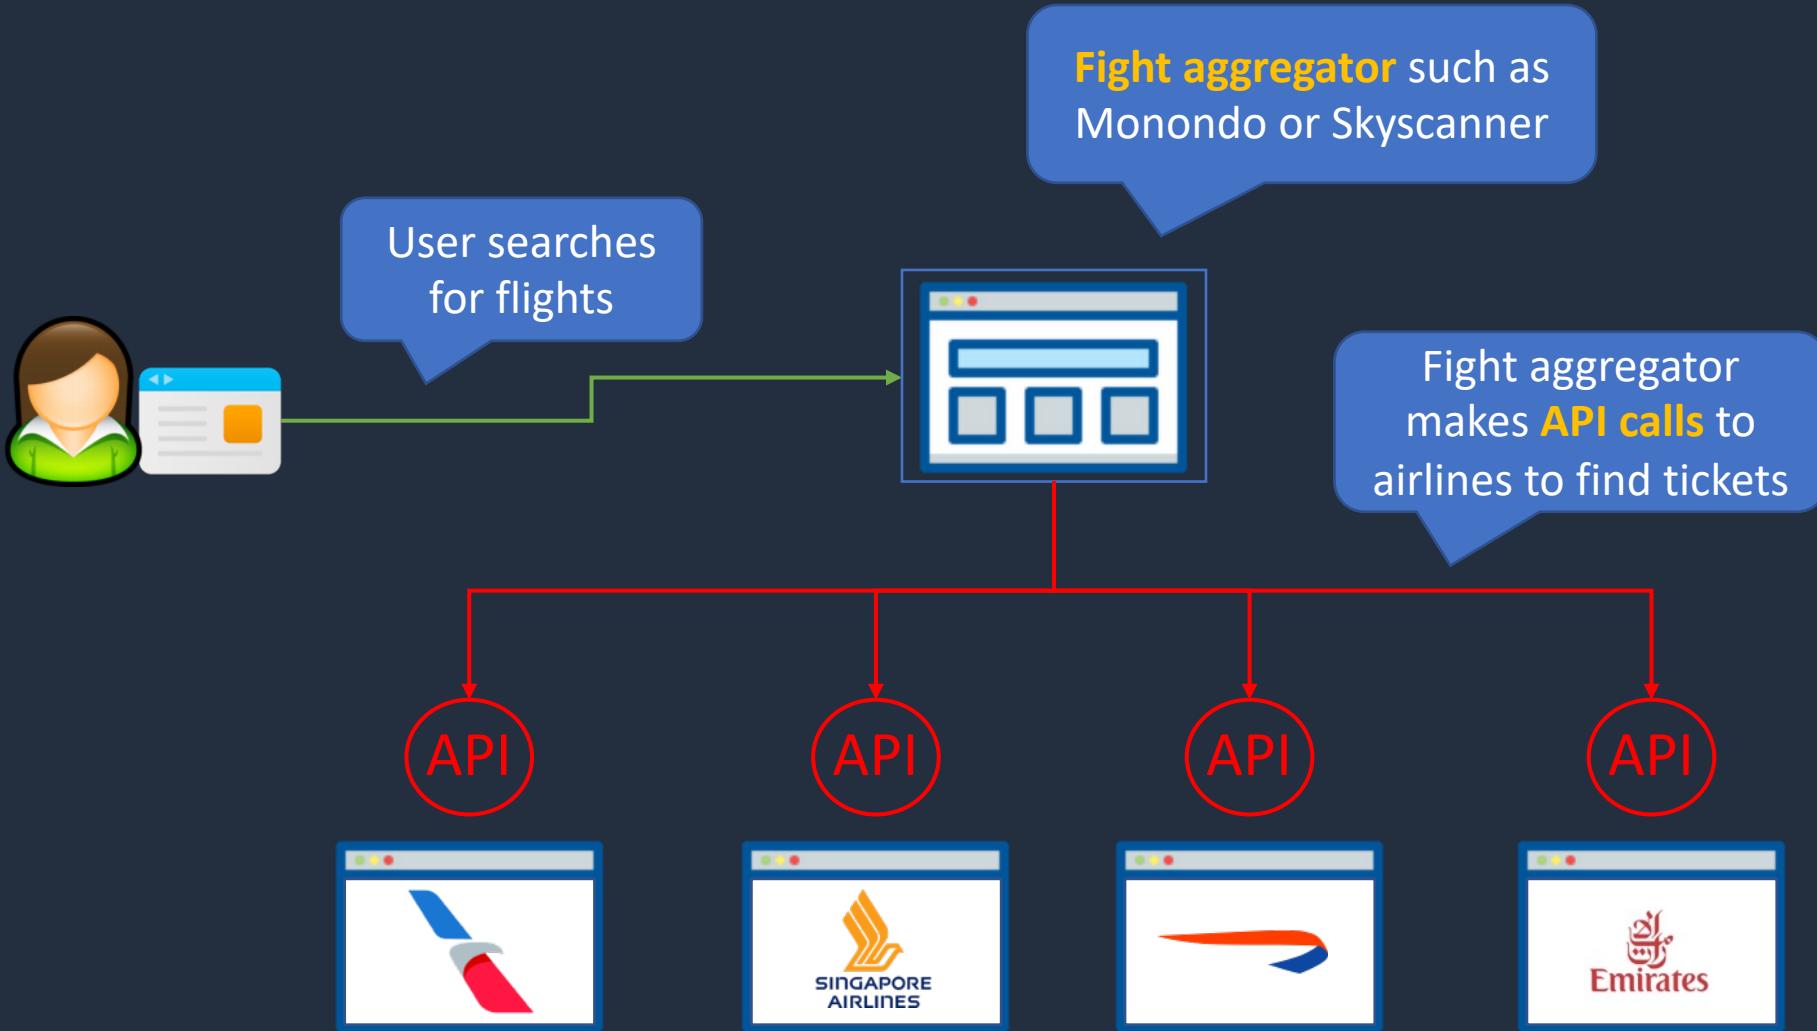
Application Programming Interfaces (APIs) – Building a house analogy



Application Programming Interfaces (APIs)



Flight Aggregator Example



Launching Cloud Services





Launching Cloud Services: Management Console

AWS Management Console

The screenshot shows the AWS Management Console navigation menu. It is organized into several sections:

- Compute**: EC2, Lightsail, Lambda, Batch, Elastic Beanstalk, Serverless Application Repository, AWS Outposts, EC2 Image Builder.
- Storage**: S3, EFS, FSx, S3 Glacier, Storage Gateway, AWS Backup.
- Database**: RDS, DynamoDB, ElastiCache, Neptune, Amazon Redshift, Amazon QLDB.
- Blockchain**: Amazon Managed Blockchain.
- Satellite**: Ground Station.
- Quantum Technologies**: Amazon Braket.
- Management & Governance**: AWS Organizations, CloudWatch, AWS Auto Scaling, CloudFormation, CloudTrail, Config, OpsWorks, Service Catalog, Systems Manager, AWS AppConfig, Trusted Advisor, Control Tower, AWS License Manager, AWS Well-Architected Tool.
- Analytics**: Athena, EMR, CloudSearch, Elasticsearch Service, Kinesis, QuickSight, Data Pipeline, AWS Data Exchange, AWS Glue, AWS Lake Formation, MSK.
- Business Applications**: Alexa for Business, Amazon Chime, WorkMail, Amazon Honeycode.
- End User Computing**: WorkSpaces, AppStream 2.0, WorkDocs, WorkLink.
- Security, Identity, & Compliance**: IAM, Resource Access Manager, Cognito, Secrets Manager, GuardDuty, Inspector, Amazon Macie, AWS Single Sign-On, Certificate Manager, Key Management Service, CloudHSM, Directory Service.
- Internet Of Things**: IoT Core, FreeRTOS, IoT 1-Click, IoT Analytics, IoT Device Defender, IoT Device Management, IoT Events, IoT Greengrass, IoT SiteWise, IoT Things Graph.

A web-based console accessed through a standard web browser



Launching Cloud Services: Command Line

Command Line

This command launches a virtual server (instance) on AWS



```
aws ec2 run-instances --image-id ami-xxxxxxxx --count 1 --instance-type t2.micro
```



```
aws s3 ls s3://mys3databucket
```

This command lists the contents of a storage container (bucket) on Amazon S3



Launching Cloud Services: Software Development Kit

A developer writes code in an **integrated development environment** (IDE)

The screenshot shows a Visual Studio Code interface with the following details:

- File Explorer:** Shows the project structure under "MY-SAM-APP (WORKSPACE)".
- Editor:** The file "app.js" is open, showing a Node.js Lambda function handler. The code defines a "lambdaHandler" function that makes an axios request to "url" and returns a JSON response with "hello world". It includes error handling for logging errors.
- Status Bar:** Shows the current line (Ln 29, Col 1), character count (27 selected), and file encoding (UTF-8 CRLF JavaScript AWS:default).

```
15 *  
16 */  
17 exports.lambdaHandler = async (event, context) => {  
18   try {  
19     const ret = await axios(url);  
20     response = {  
21       statusCode: 200,  
22       body: JSON.stringify({  
23         message: 'hello world',  
24         // location: ret.data.trim()  
25       })  
26     }  
27   } catch (err) {  
28     console.log(err);  
29     return err;  
30   }  
31  
32   return response;  
33 };  
34 };
```

The code leverages the **SDK** to work with cloud services

Create your AWS Free Tier Account



Configure Account and Create a Billing Alarm



AWS CLI and CloudShell

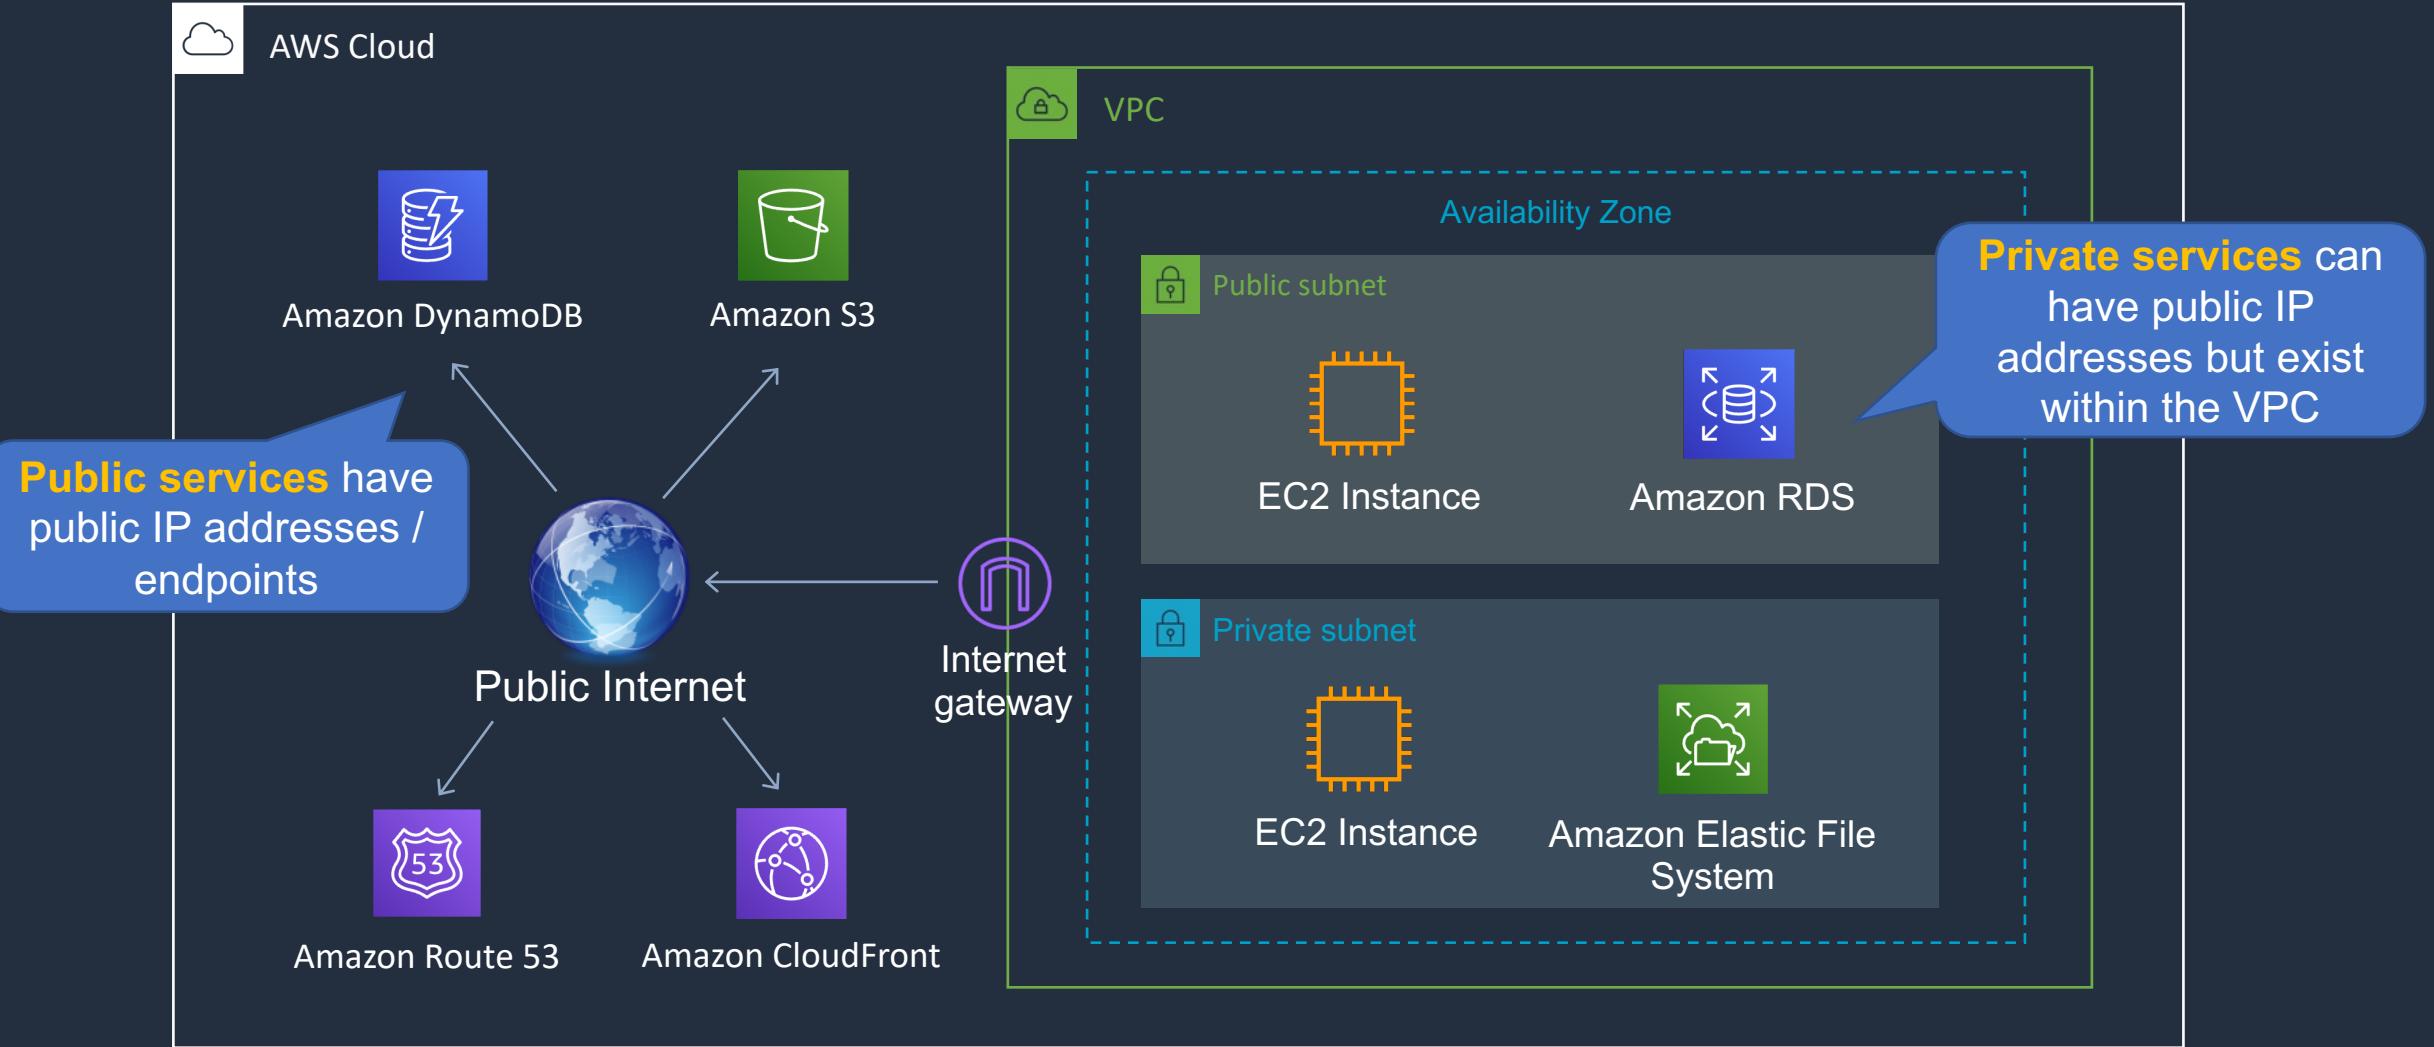


AWS Public and Private Services





AWS Public and Private Services



The 6 Advantages of Cloud Computing





The 6 Advantages of Cloud Computing

1. Trade capital expense for variable expense

CAPEX



Purchase servers



Tax deductible over depreciation lifetime

OPEX



Pay as you go



Tax deductible in same year



The 6 Advantages of Cloud Computing

2. Benefit from massive economies of scale





The 6 Advantages of Cloud Computing

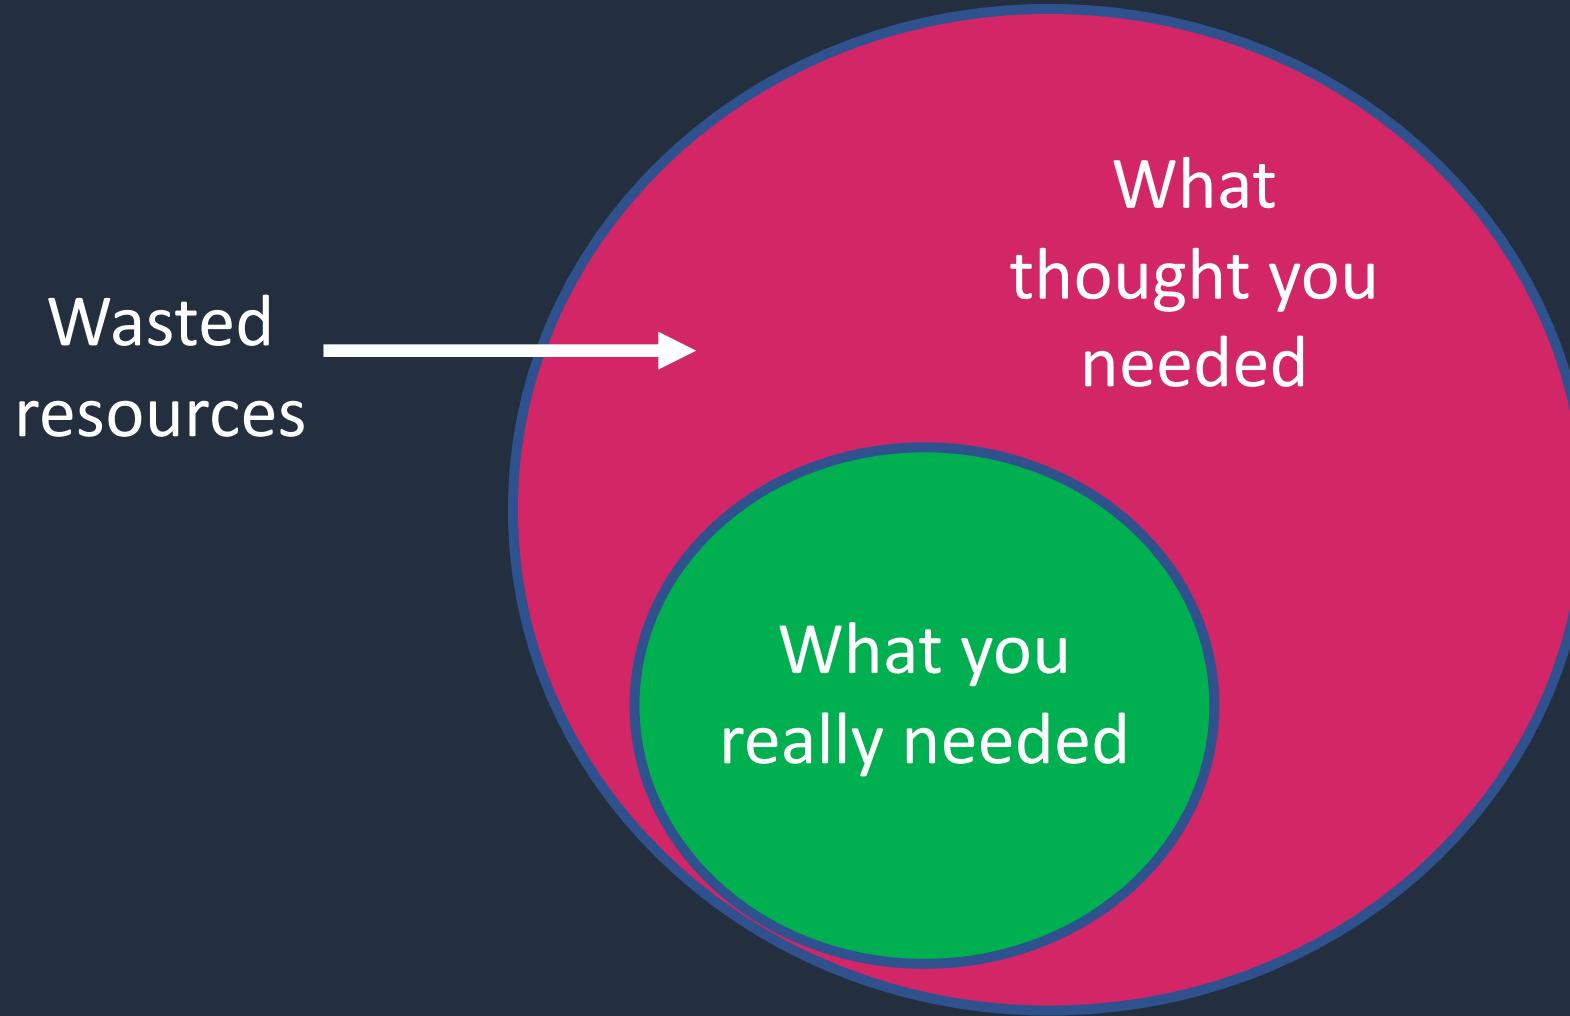
2. Benefit from massive economies of scale

- Aggregated usage across hundreds of thousands of customers = lower variable costs for customers



The 6 Advantages of Cloud Computing

3. Stop guessing capacity





The 6 Advantages of Cloud Computing

4. Increase speed and agility



Speed = deploy resources easily and quickly

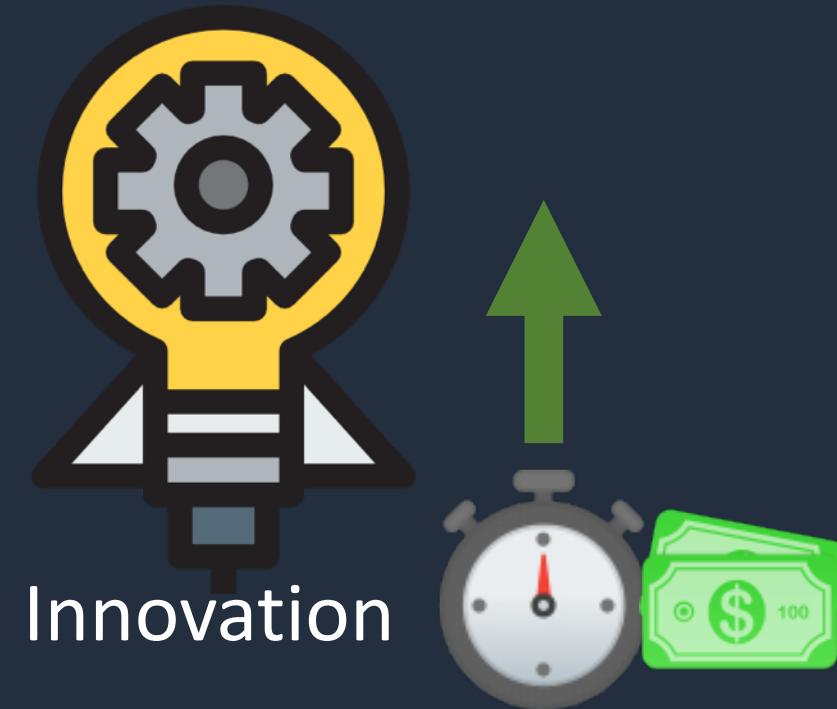


Agility = react to change ; speed to market



The 6 Advantages of Cloud Computing

5. Stop spending money running and maintaining data centers





The 6 Advantages of Cloud Computing

6. Go global in minutes



SECTION 4

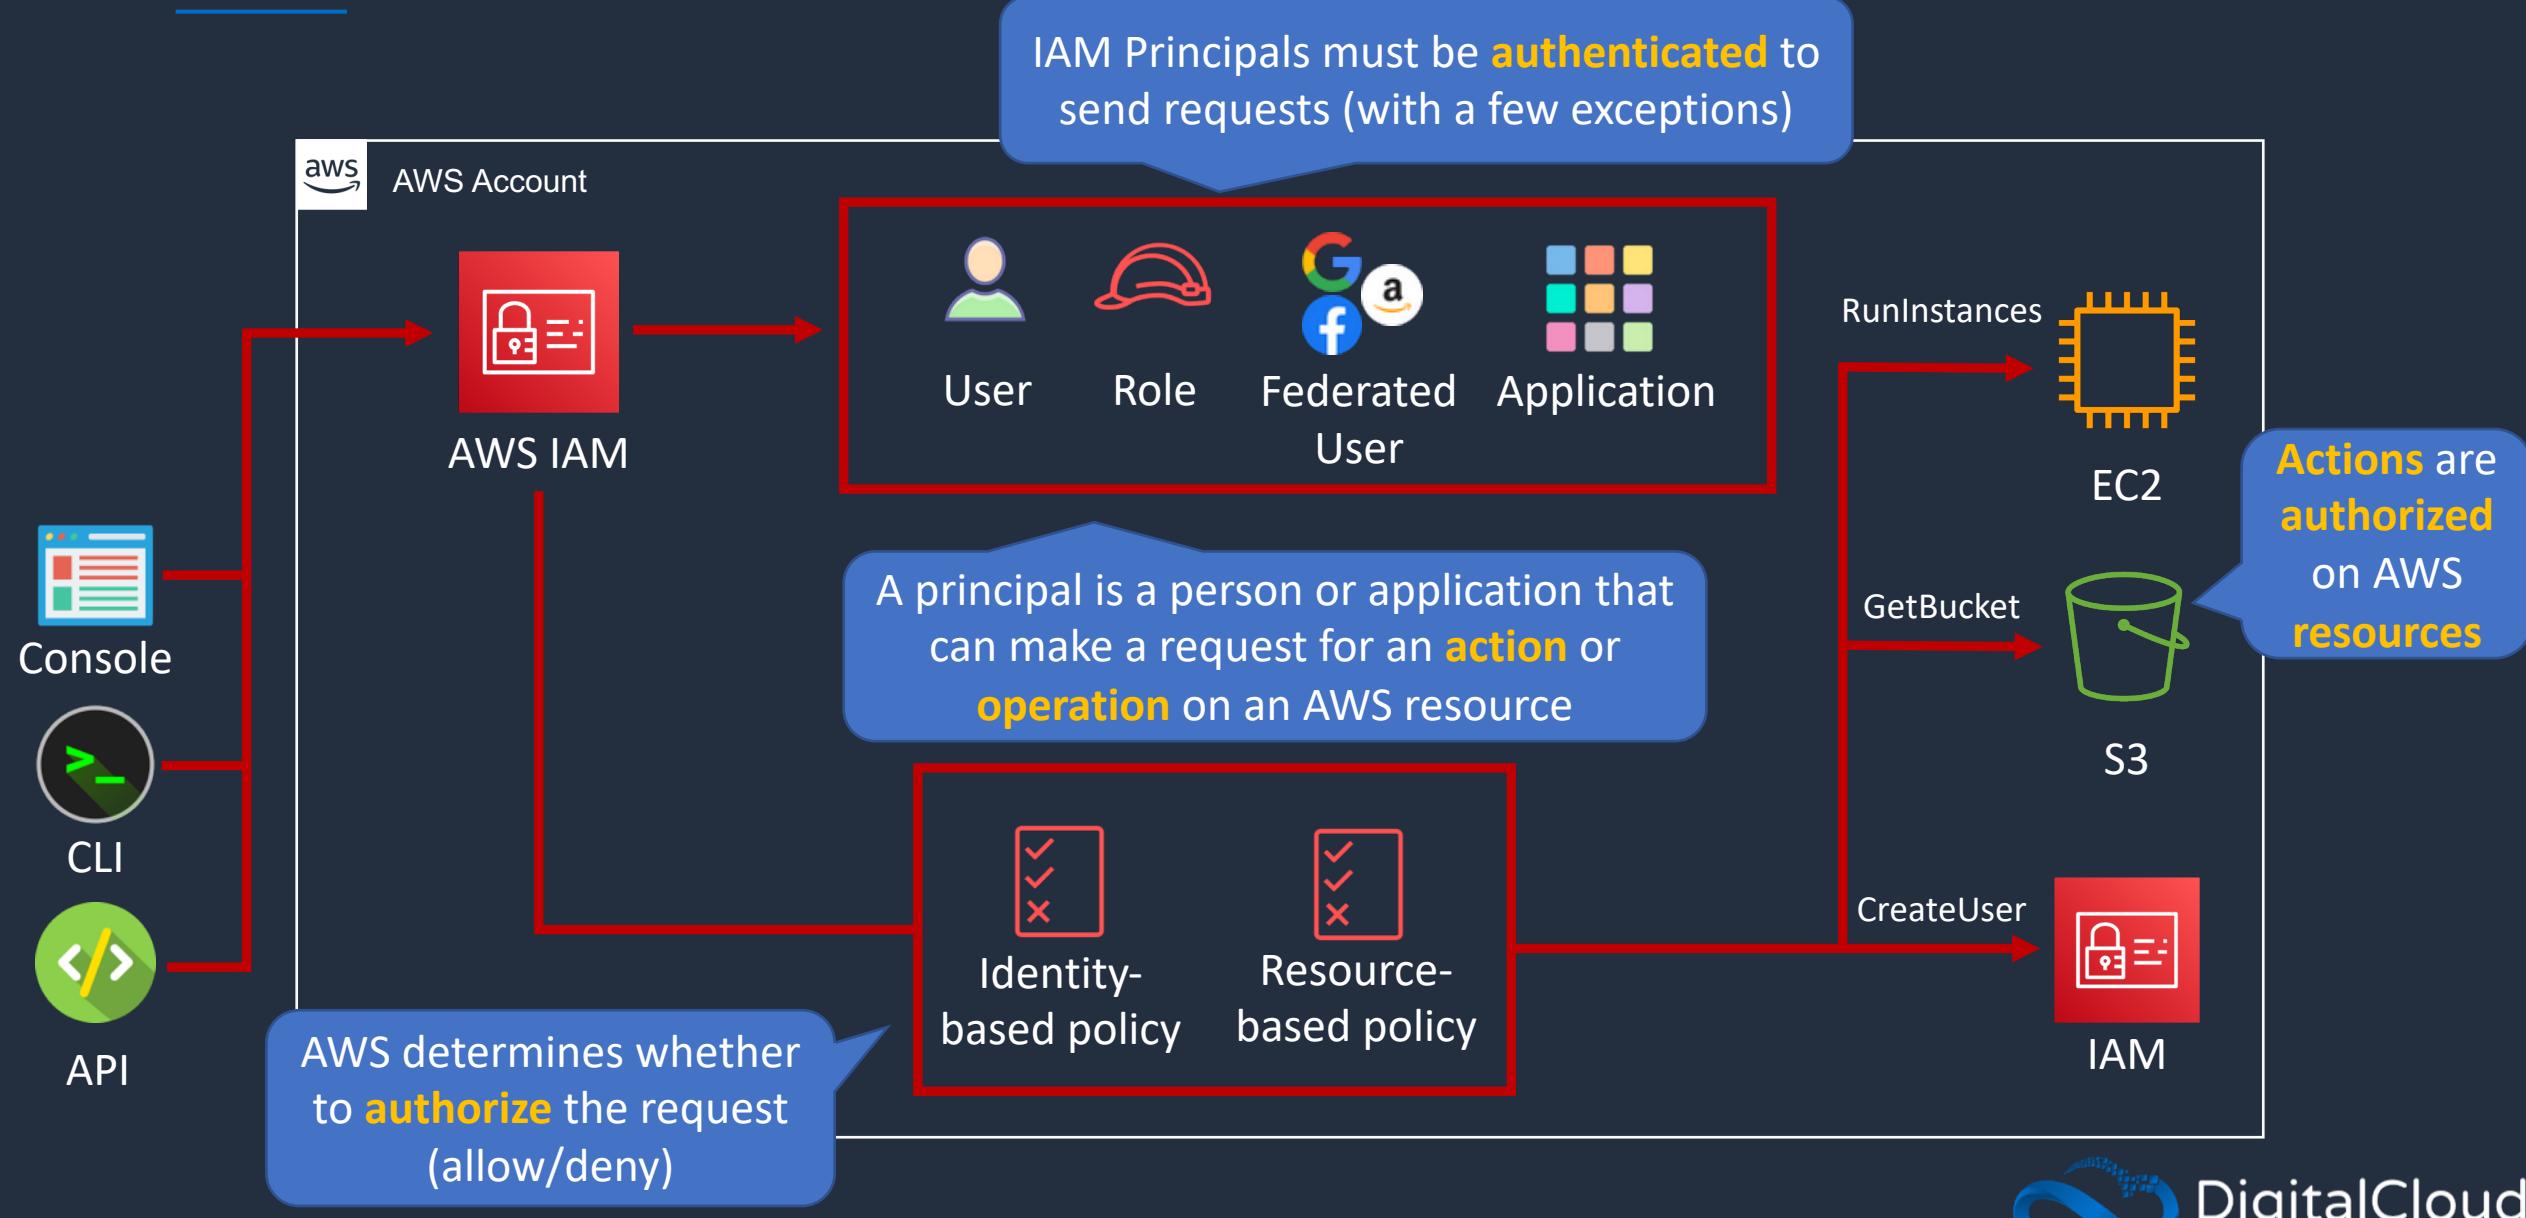
Identity and Access Management (AWS IAM)

AWS IAM Overview





AWS Identity and Access Management (IAM)

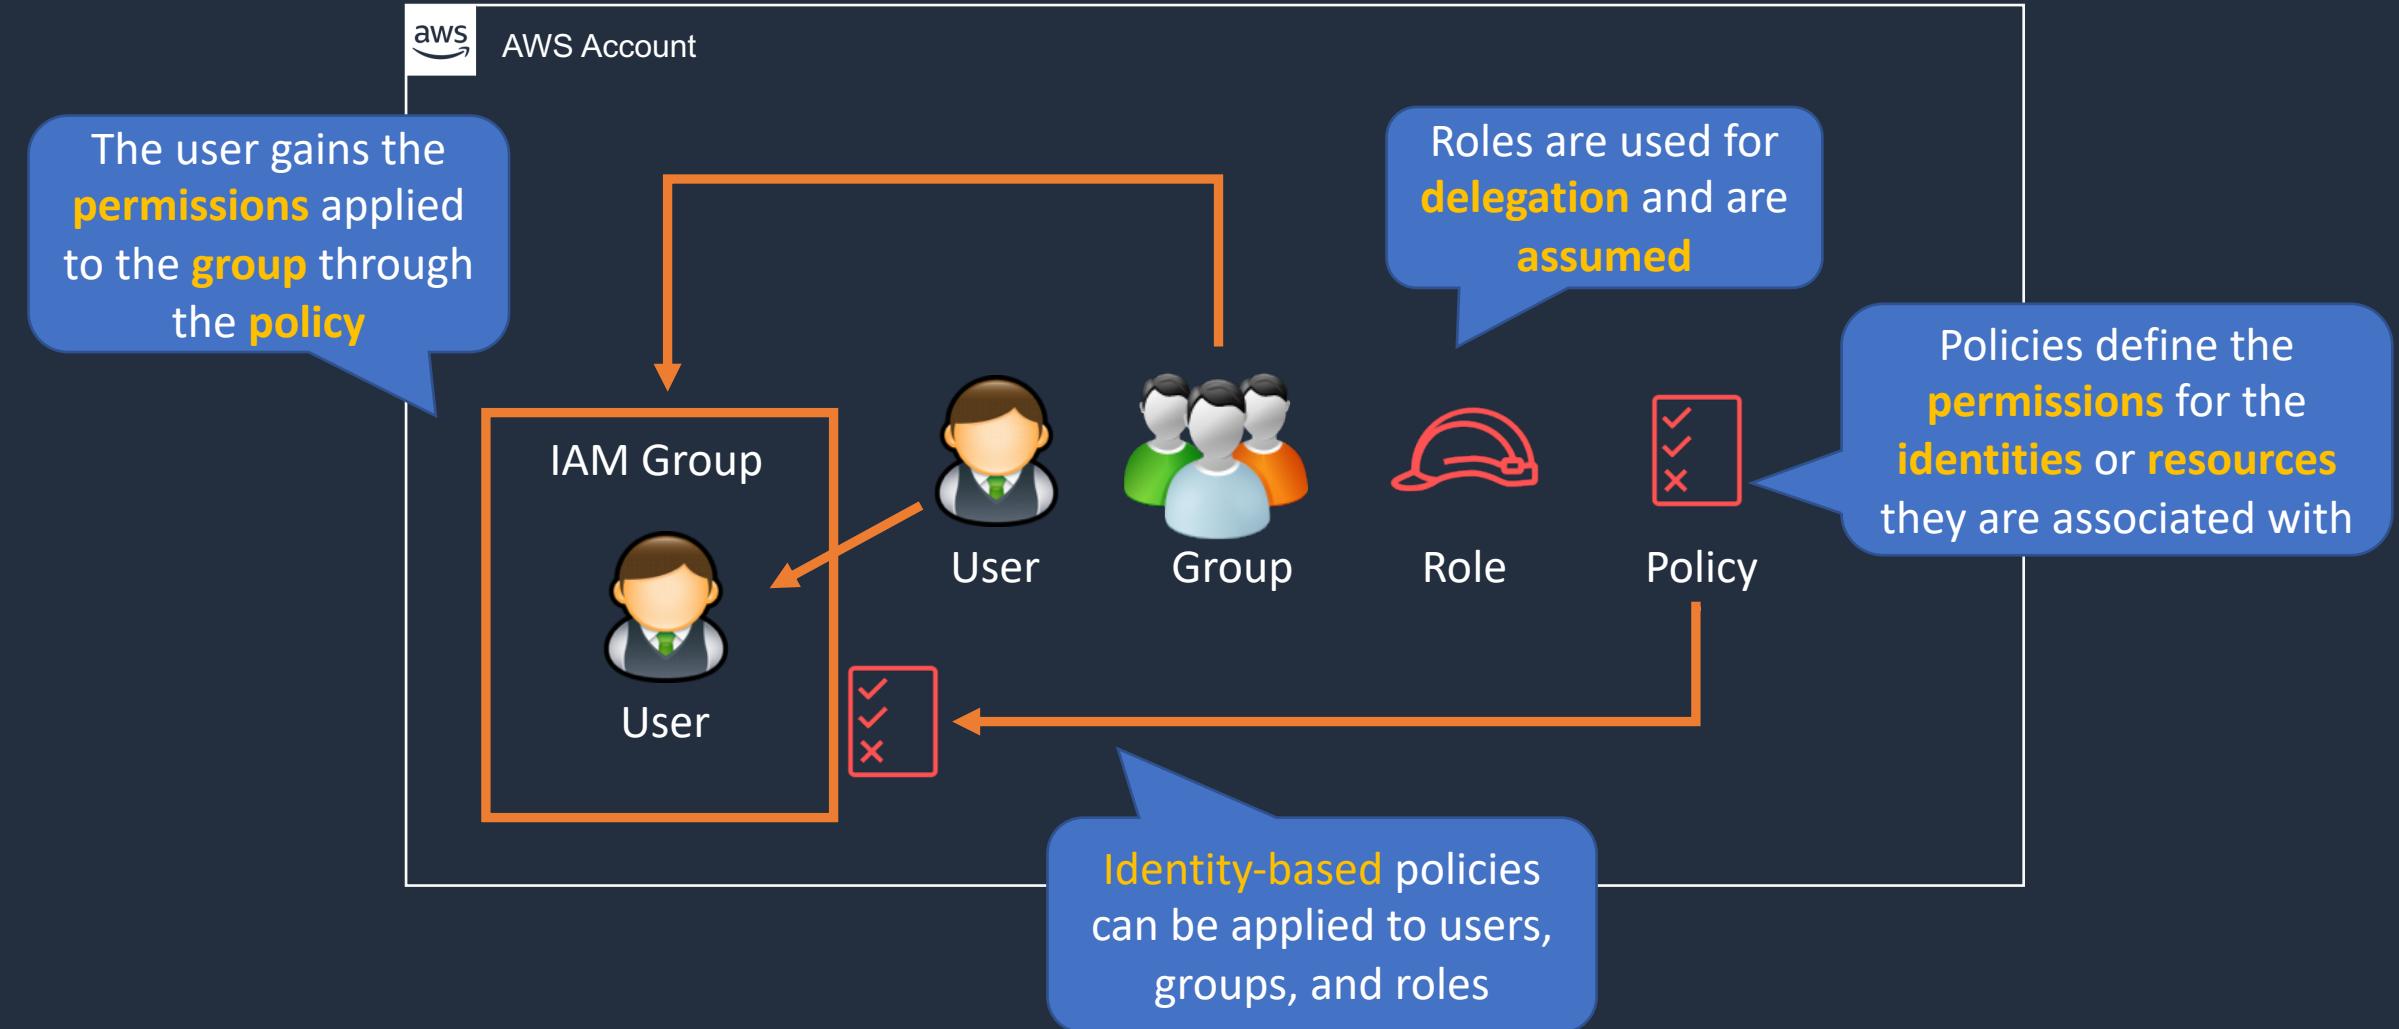


IAM Users, Groups, Roles, and Policies



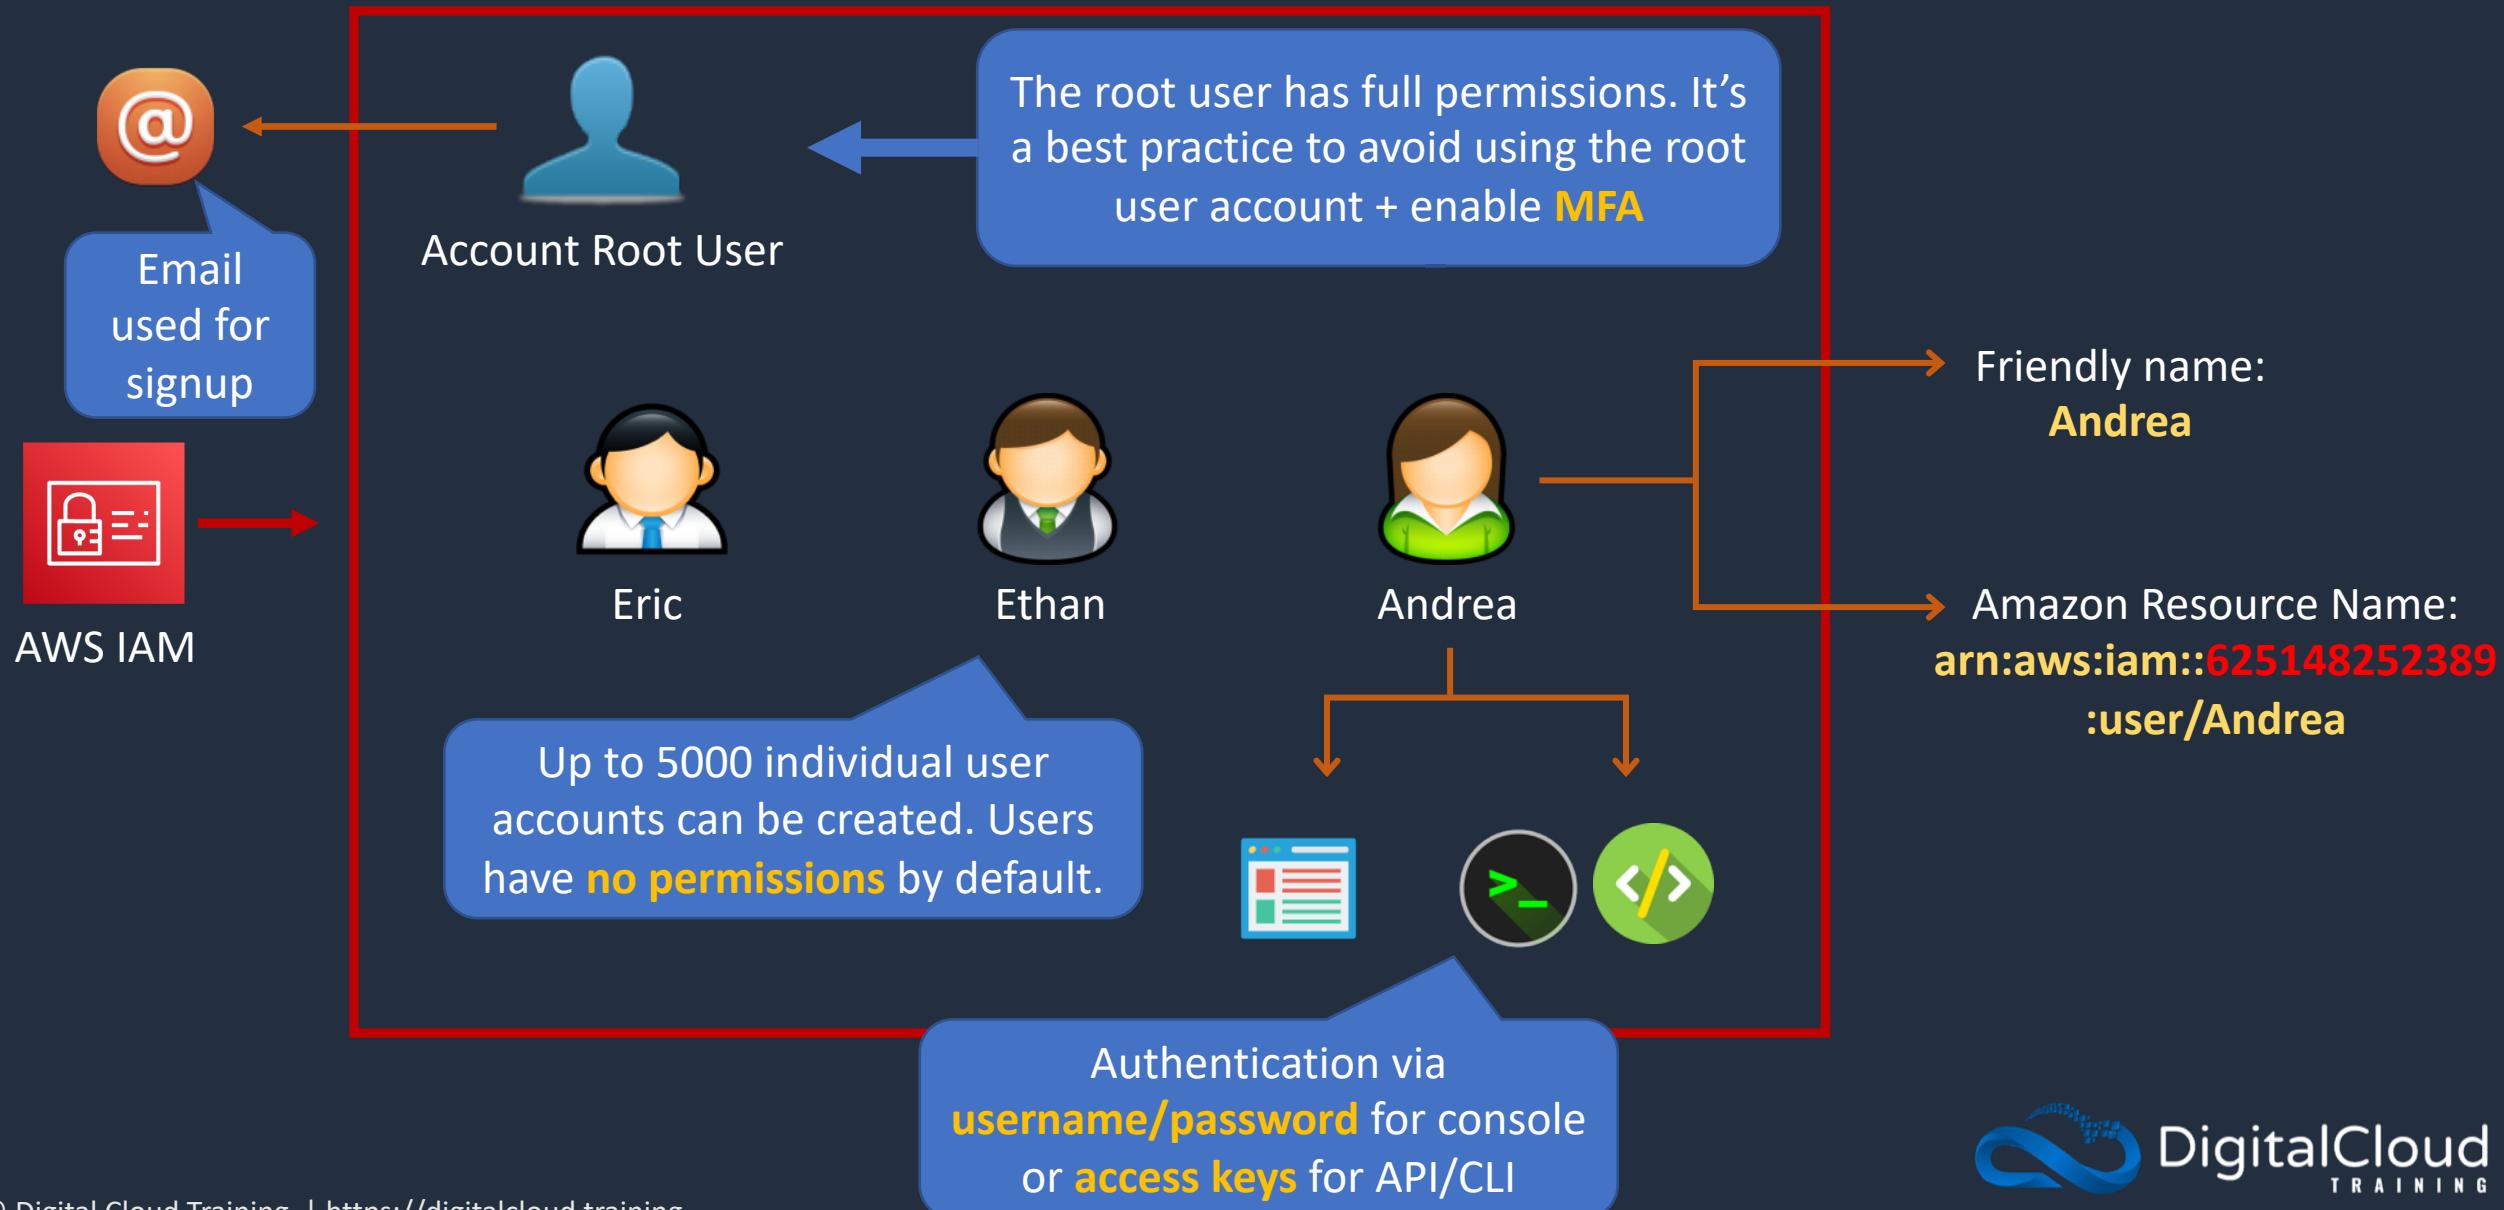


Users, Groups, Roles and Policies



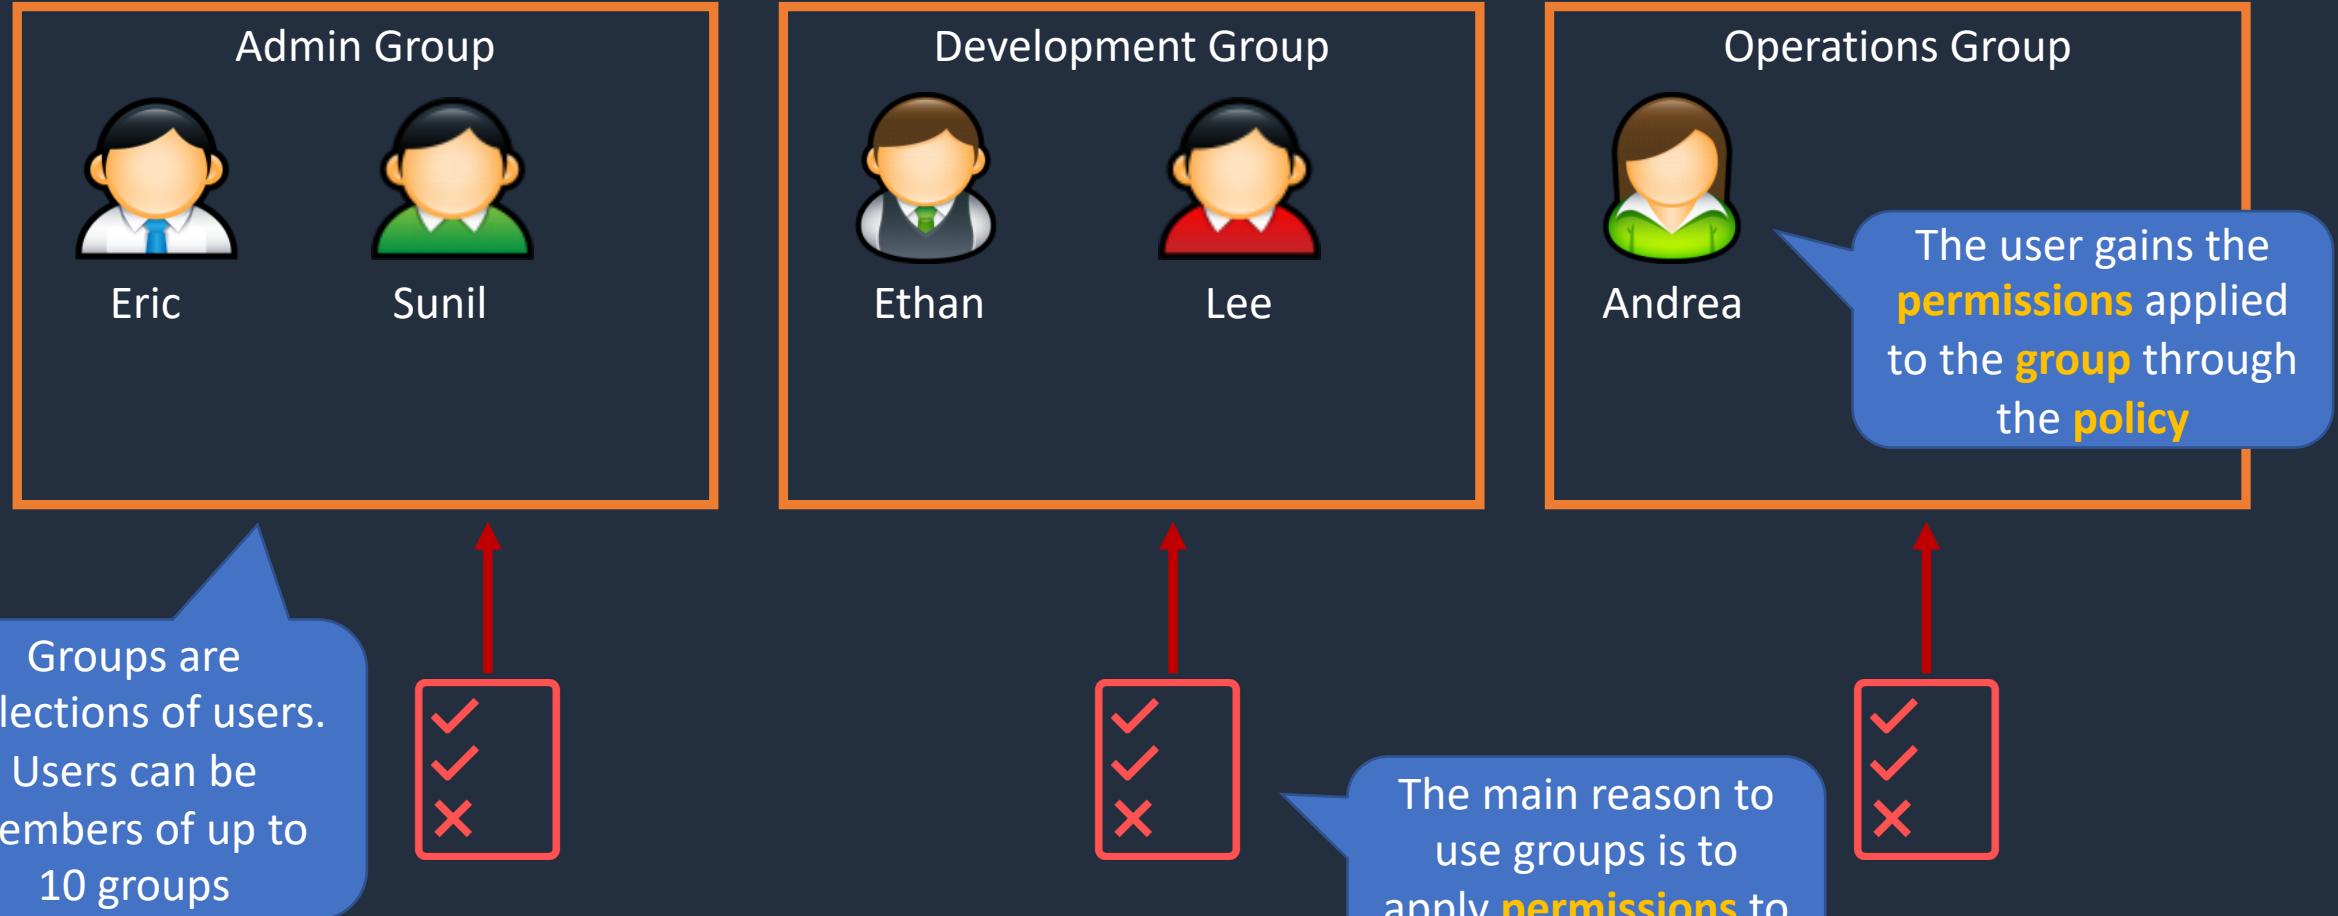


IAM Users





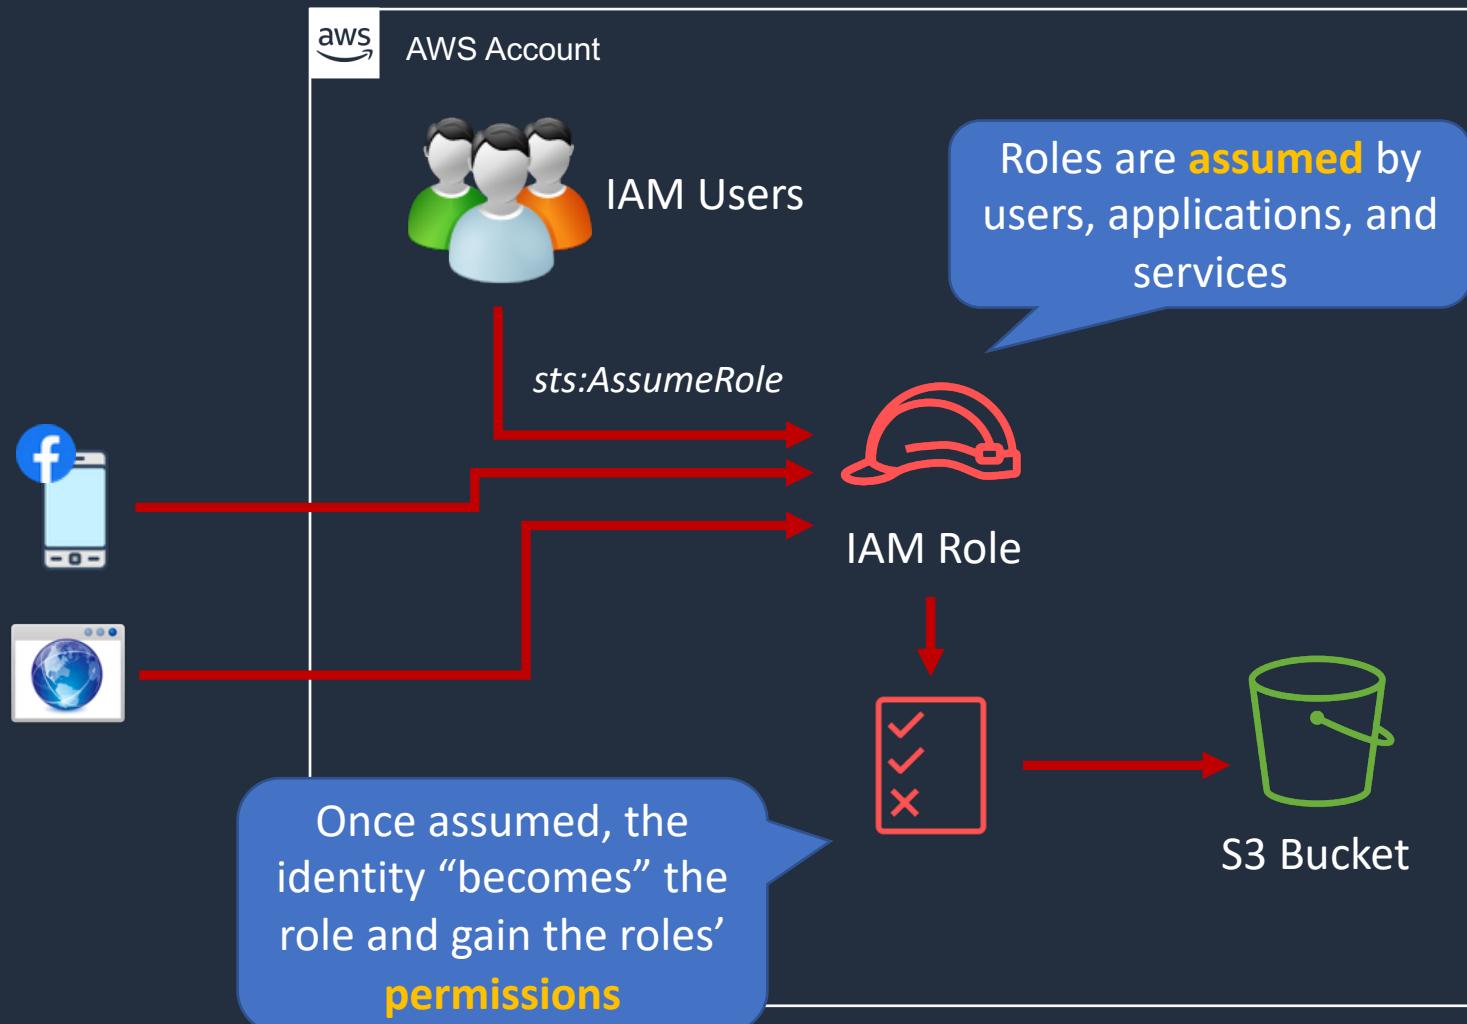
IAM Groups





IAM Roles

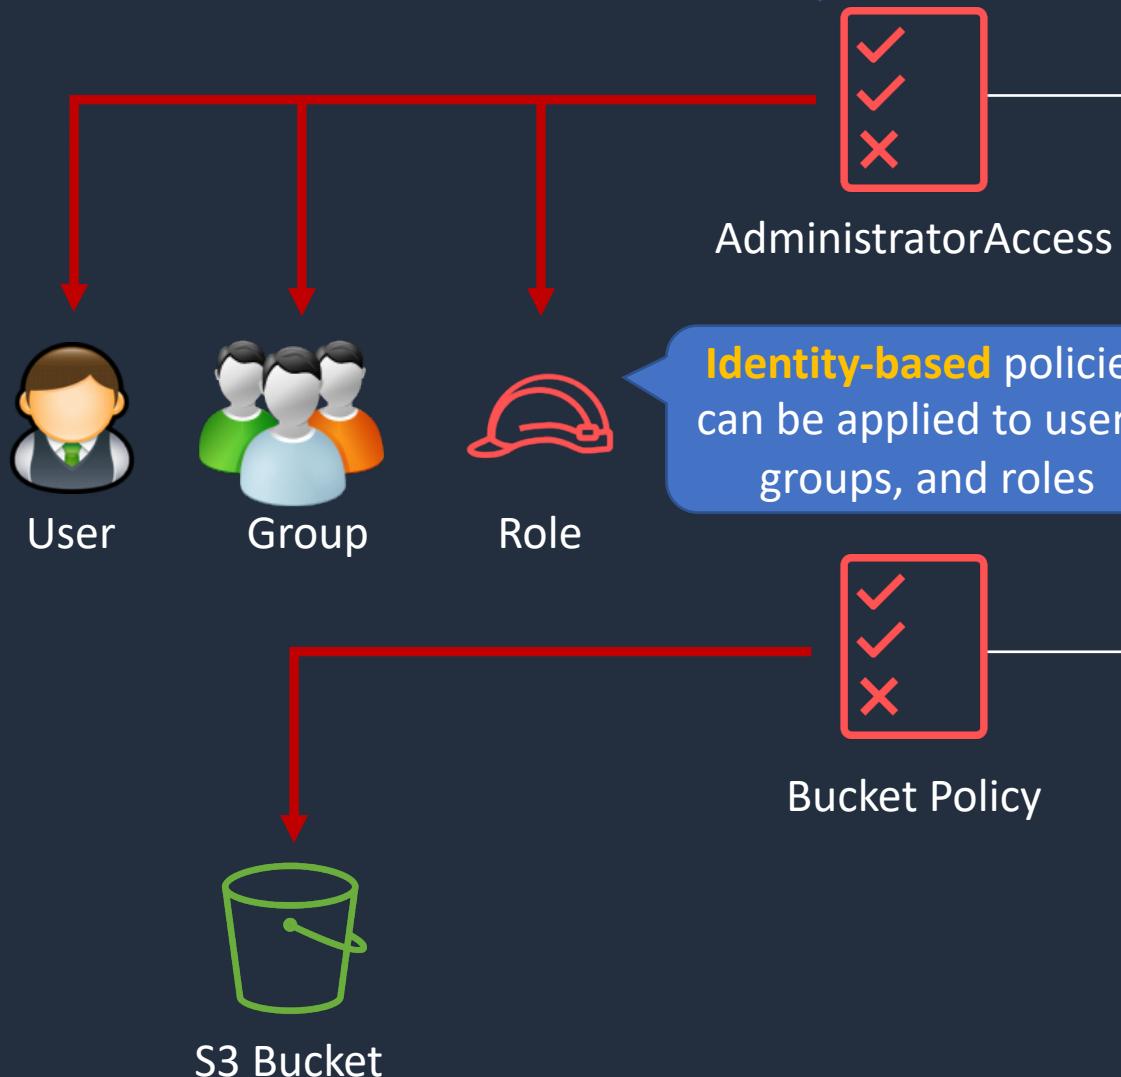
An **IAM role** is an IAM **identity** that has specific **permissions**





IAM Policies

Policies are **documents** that define **permissions** and are written in **JSON**



```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Action": "*",  
      "Resource": "*"  
    }  
  ]  
}
```

All permissions are **implicitly denied** by default

```
{  
  "Version": "2012-10-17",  
  "Id": "Policy1561964929358",  
  "Statement": [  
    {  
      "Sid": "Stmt1561964454052",  
      "Effect": "Allow",  
      "Principal": {  
        "AWS": "arn:aws:iam::515148227241:user/Paul"  
      },  
      "Action": "s3:*",  
      "Resource": "arn:aws:s3:::dctcompany",  
      "Condition": {  
        "StringLike": {  
          "s3:prefix": "Confidential/*"  
        }  
      }  
    }  
  ]  
}
```

Resource-based policies apply to **resources** such as S3 buckets or DynamoDB tables

Setup Individual User Account

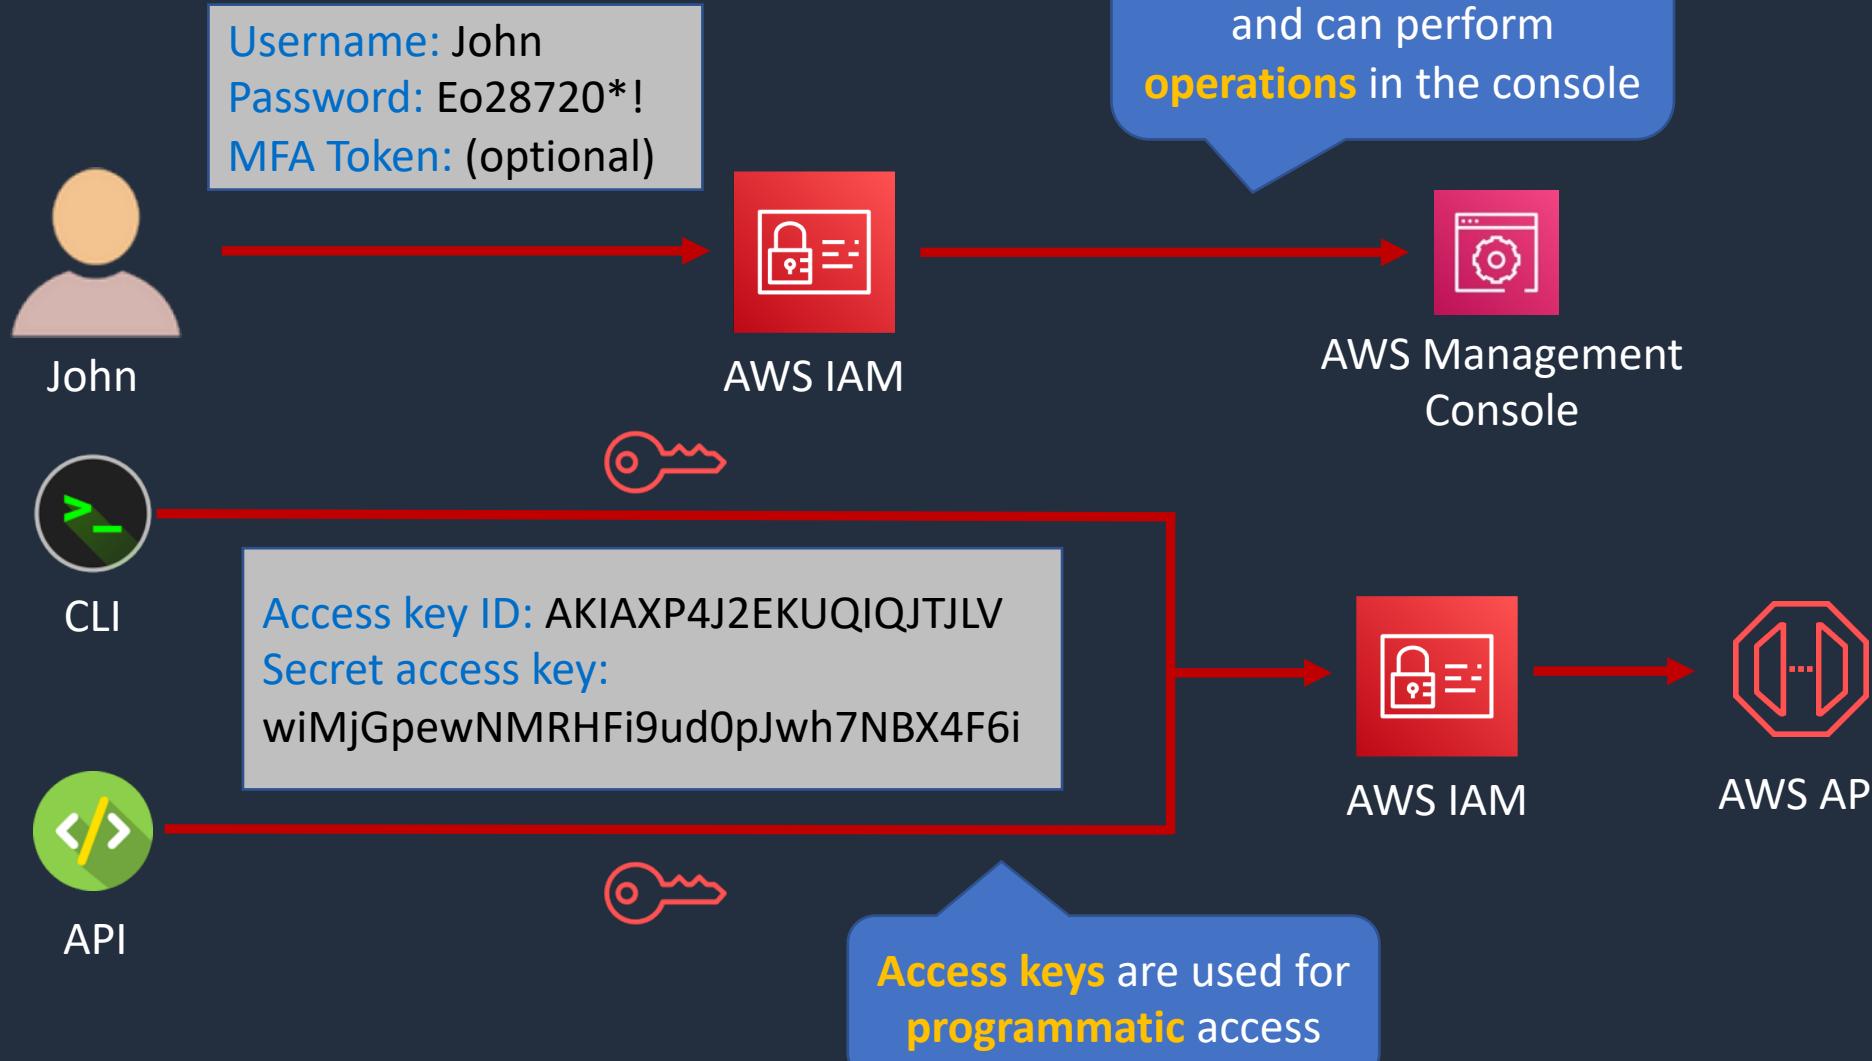


IAM Authentication and MFA





IAM Authentication Methods





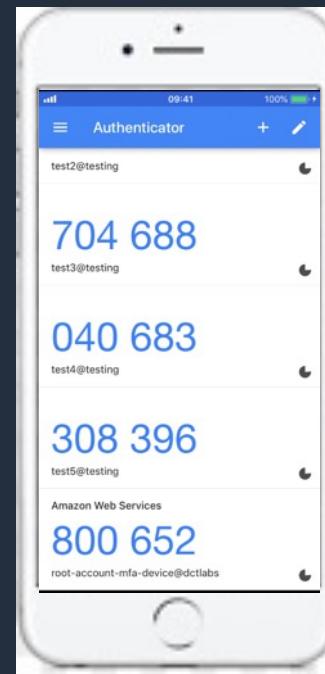
Multi-Factor Authentication

Something you **know**:

EJPx!*21p9%

Password

Something you **have**:



Something you **are**:





Multi-Factor Authentication

Something you **know**:



IAM User

EJPx!*21p9%

Password

Something you **have**:



Virtual MFA

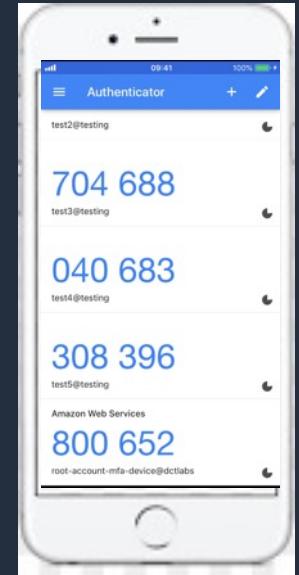


Physical MFA

e.g. Google Authenticator on
your smart phone



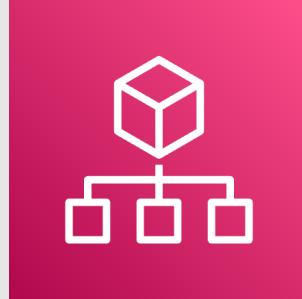
Physical tokens can
be purchased from
third parties



Setup Multi-Factor Authentication (MFA)

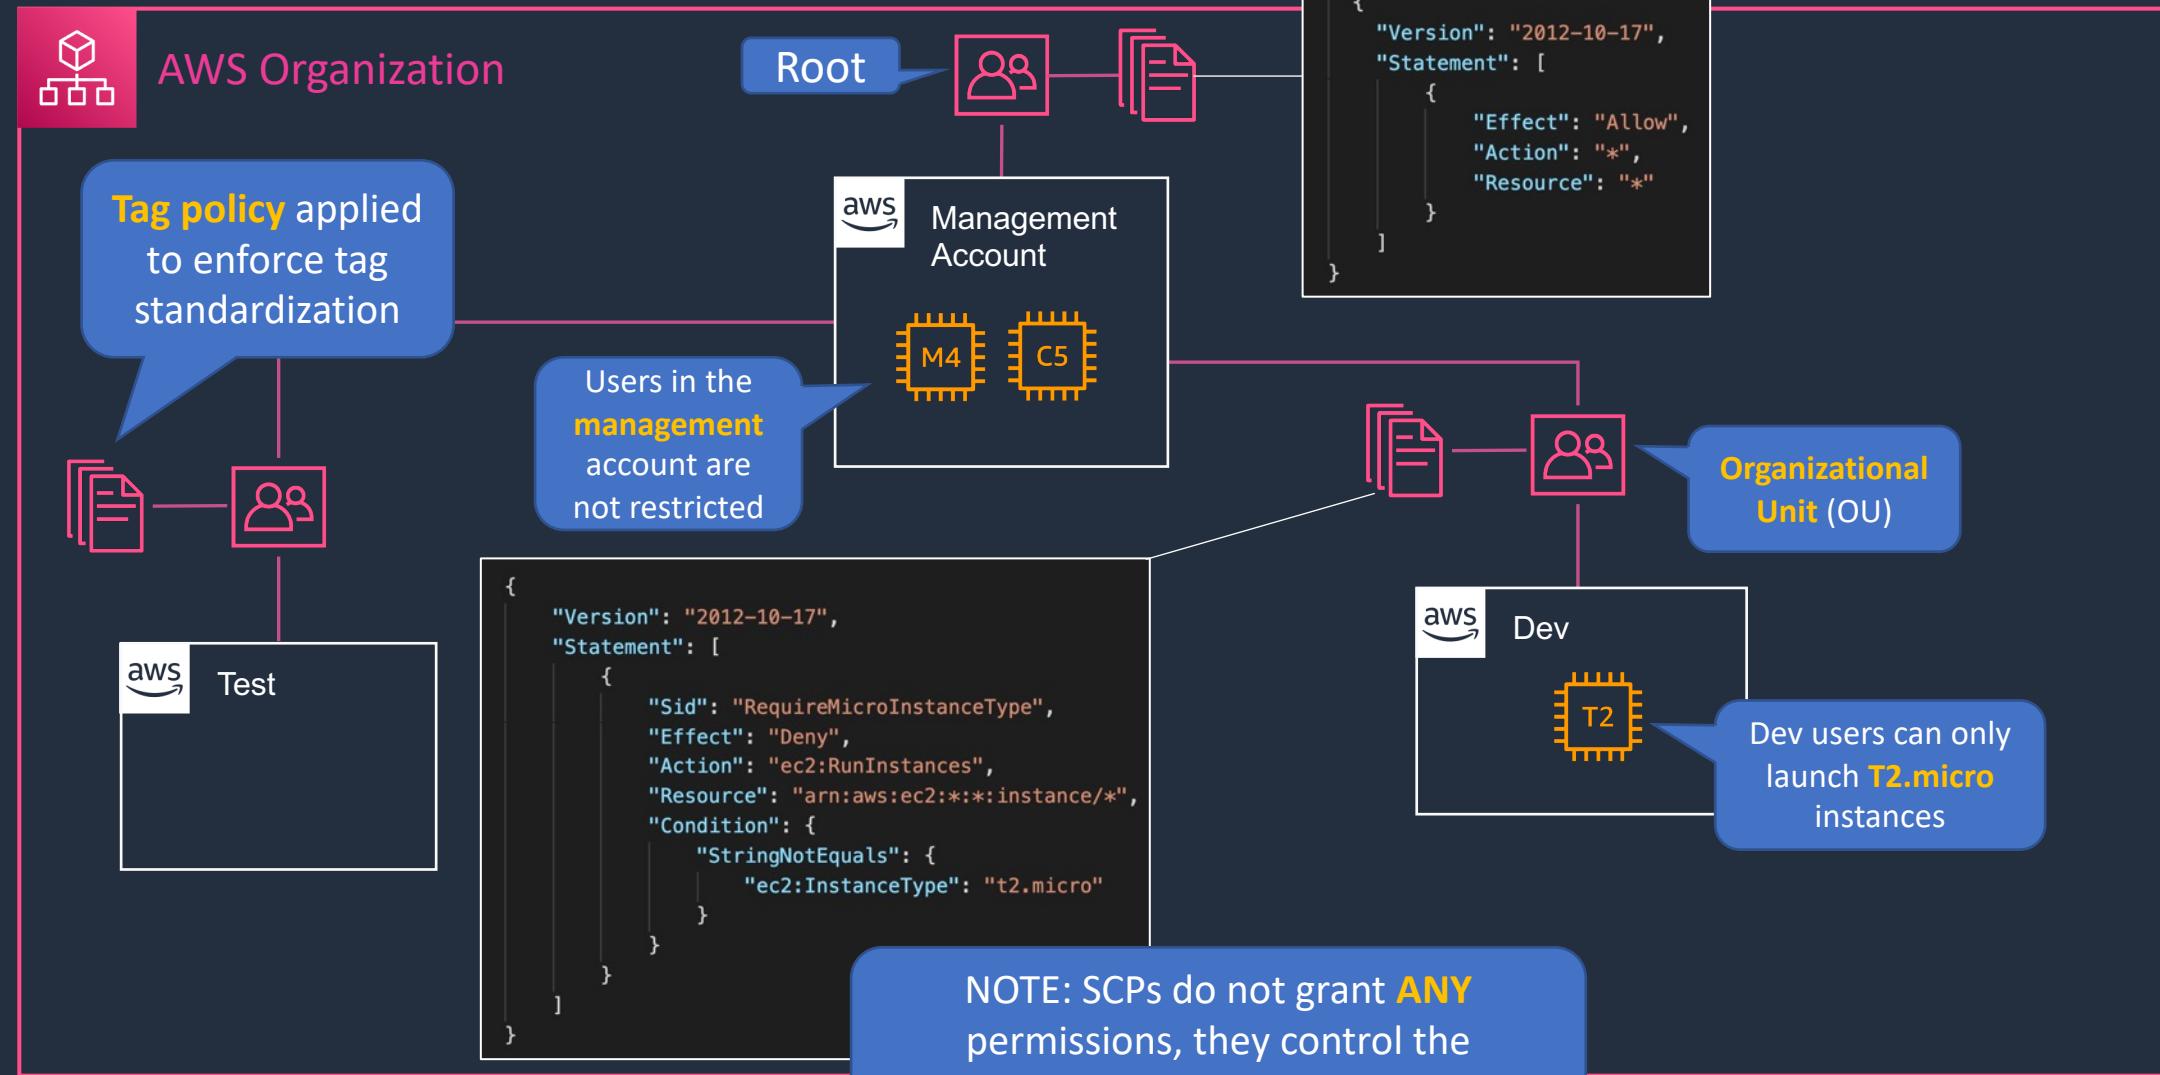


Service Control Policies (SCPs)



Service Control Policies

SCPs control the maximum available permissions



IAM Password Policy



IAM Best Practices





AWS IAM Best Practices

- Lock away your AWS account root user access keys
- Create individual IAM users
- Use groups to assign permissions to IAM users
- Grant least privilege
- Get started using permissions with AWS managed policies
- Use customer managed policies instead of inline policies
- Use access levels to review IAM permissions
- Configure a strong password policy for your users
- Enable MFA



AWS IAM Best Practices

- Use roles for applications that run on Amazon EC2 instances
- Use roles to delegate permissions
- Do not share access keys
- Rotate credentials regularly
- Remove unnecessary credentials
- Use policy conditions for extra security
- Monitor activity in your AWS account

SECTION 5

AWS Compute Services

Computing Basics





Computing Basics

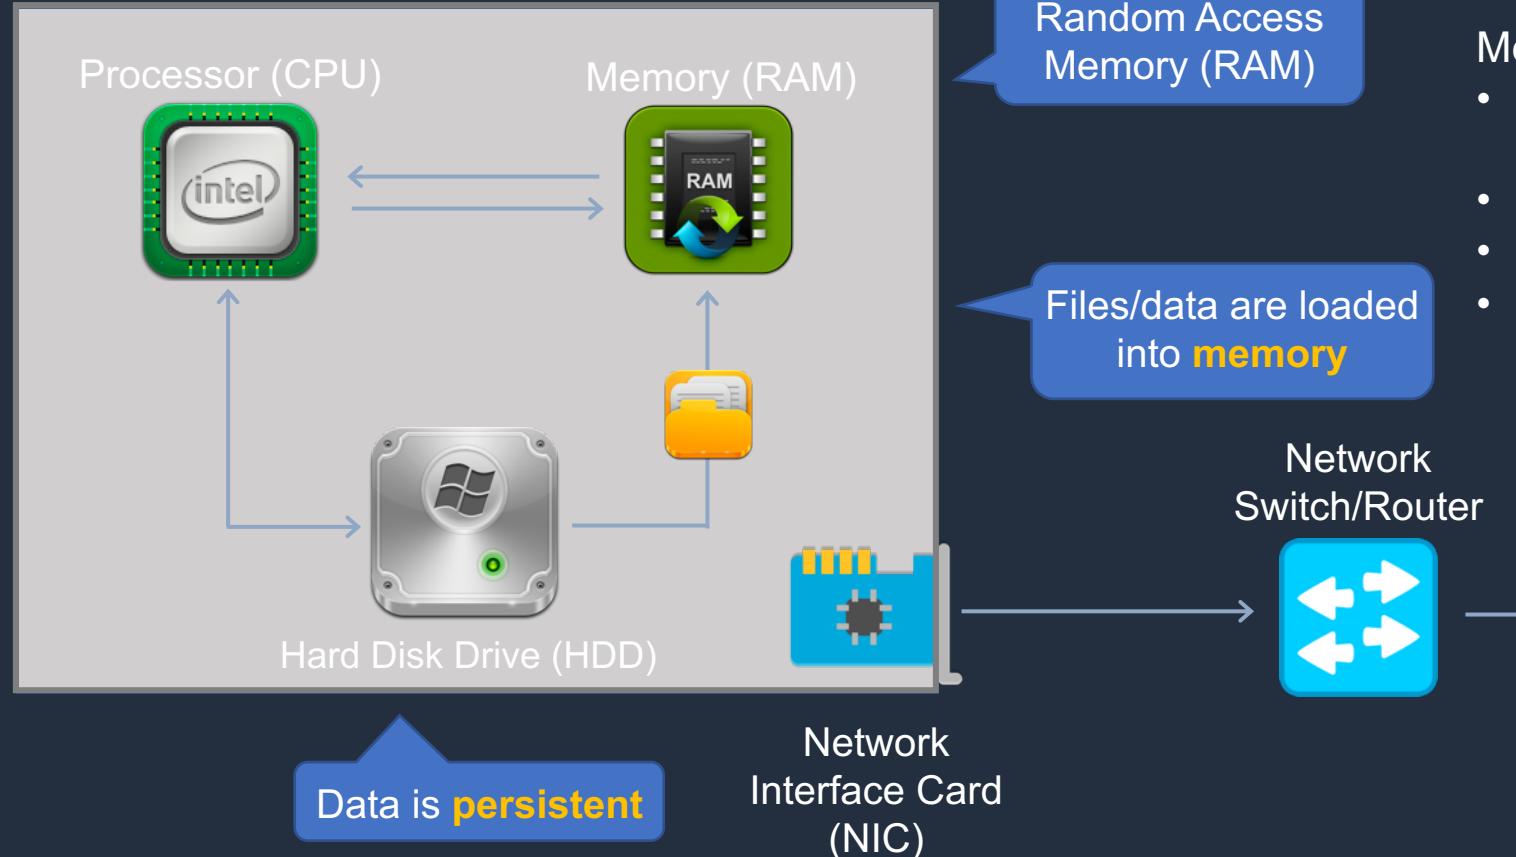




Computing Basics

Central Processing Unit (CPU)

RAM is non-persistent storage

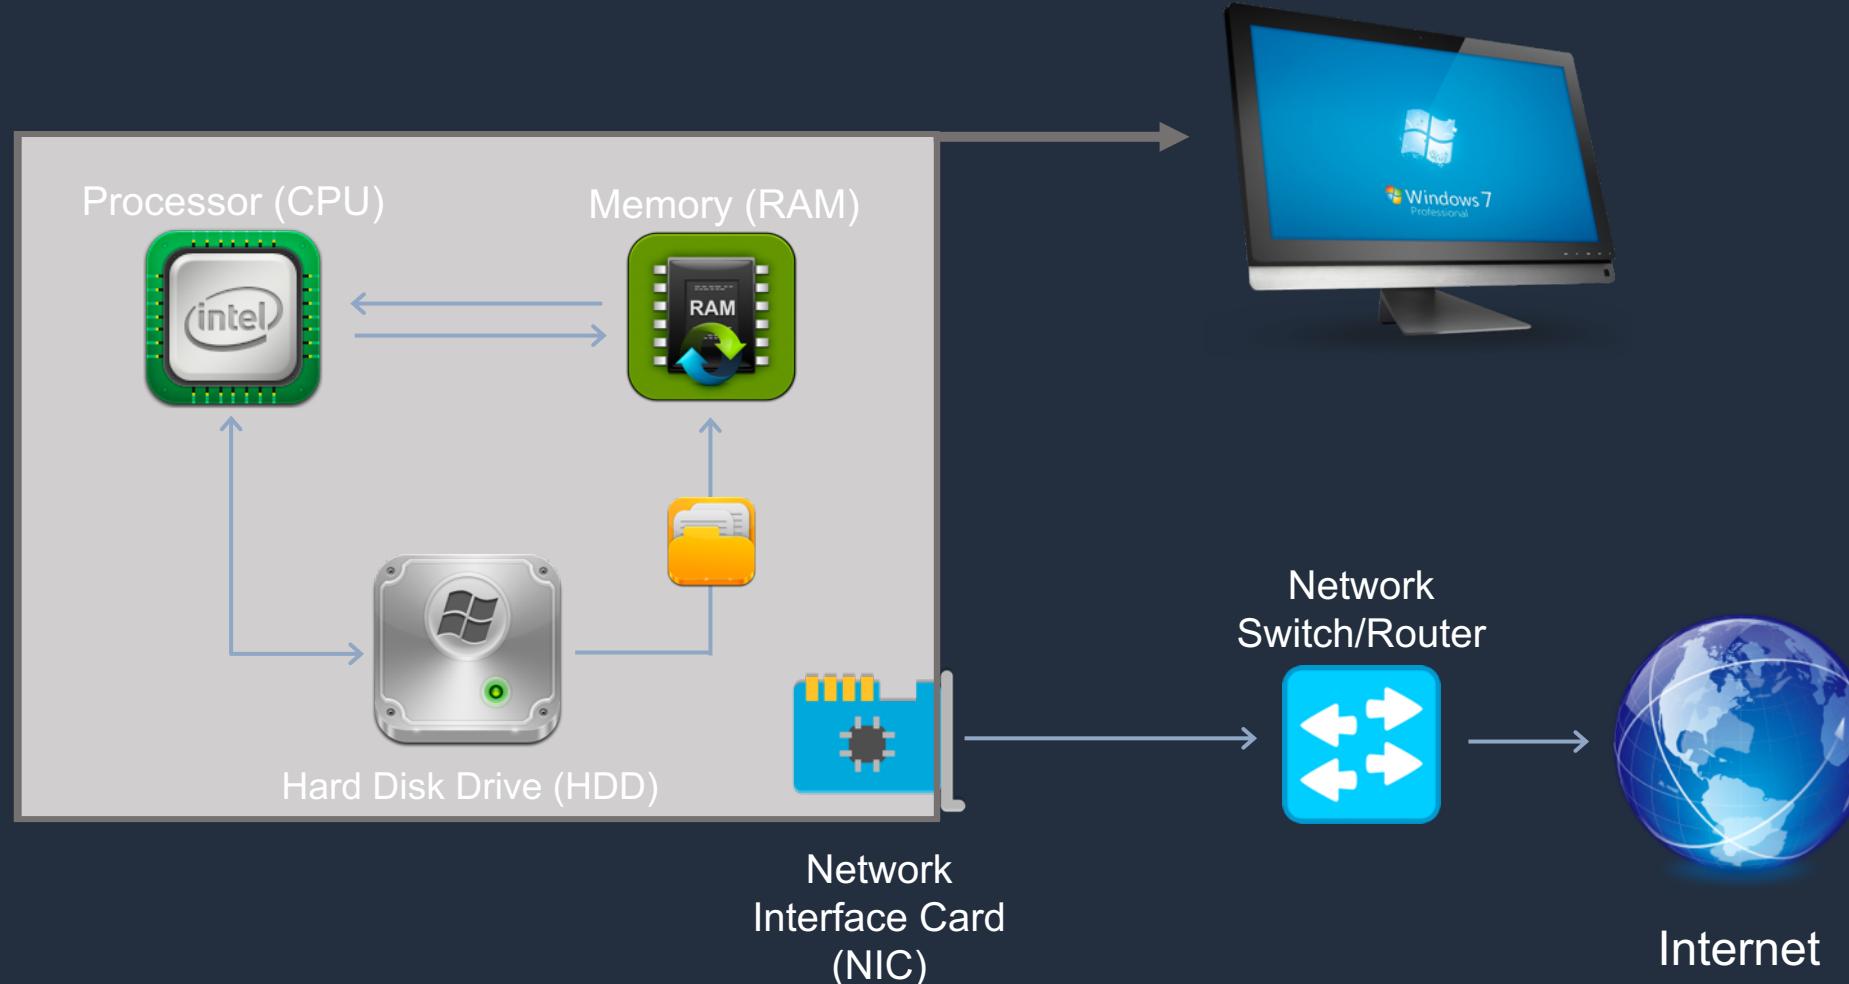


Measurements:

- CPU is measured in Gigahertz (Ghz)
- RAM is measured in Gigabyte (GB)
- HDD is measured in Gigabyte (GB)
- NIC is measured in Megabits per second (Mbps) or Gigabits per second (Gbps)

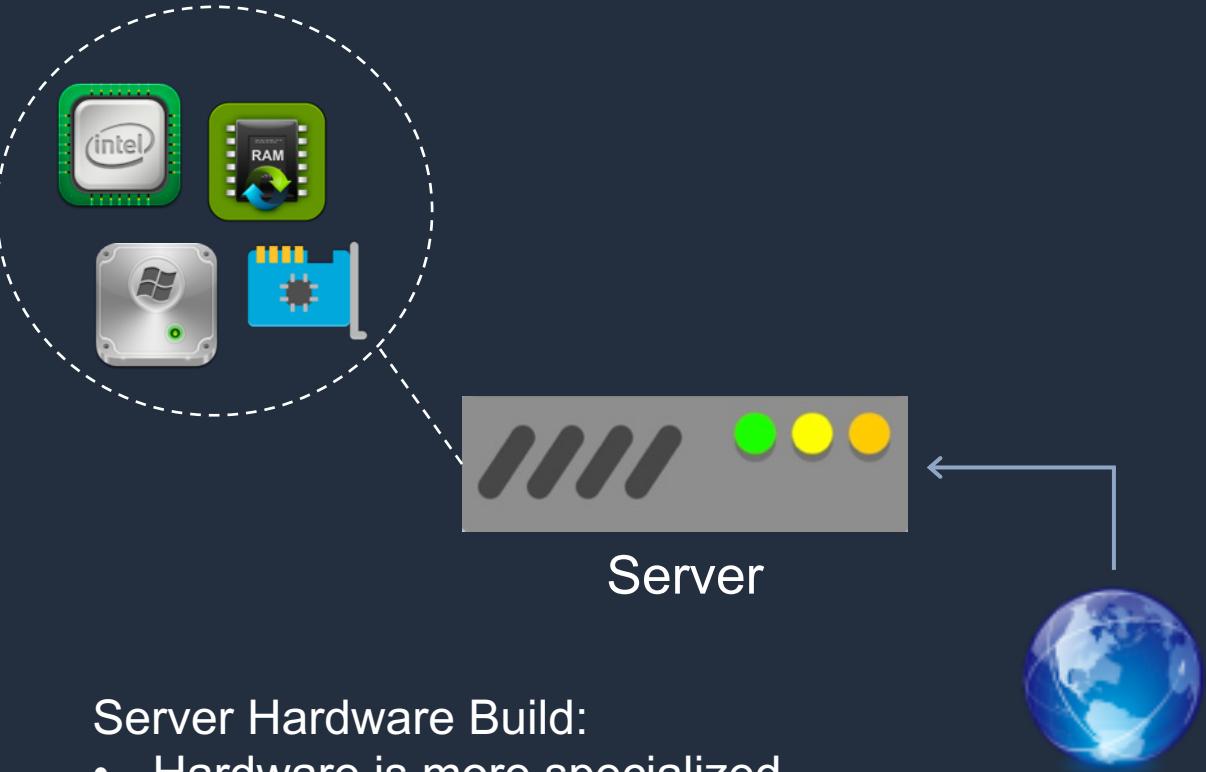


Computing Basics





Servers vs Desktops/Laptops



Server Hardware Build:

- Hardware is more specialized
- Much higher prices compared to desktops / laptops
- Includes redundancy

Servers can be used by many users over a network

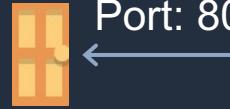




Client / Server Computing



Web Server



Port: 80

Protocol: HTTP



File Server



Port: 445

Protocol: SMB



Email Server



Port: 25

Protocol: SMTP



The client application finds the server by **IP address**

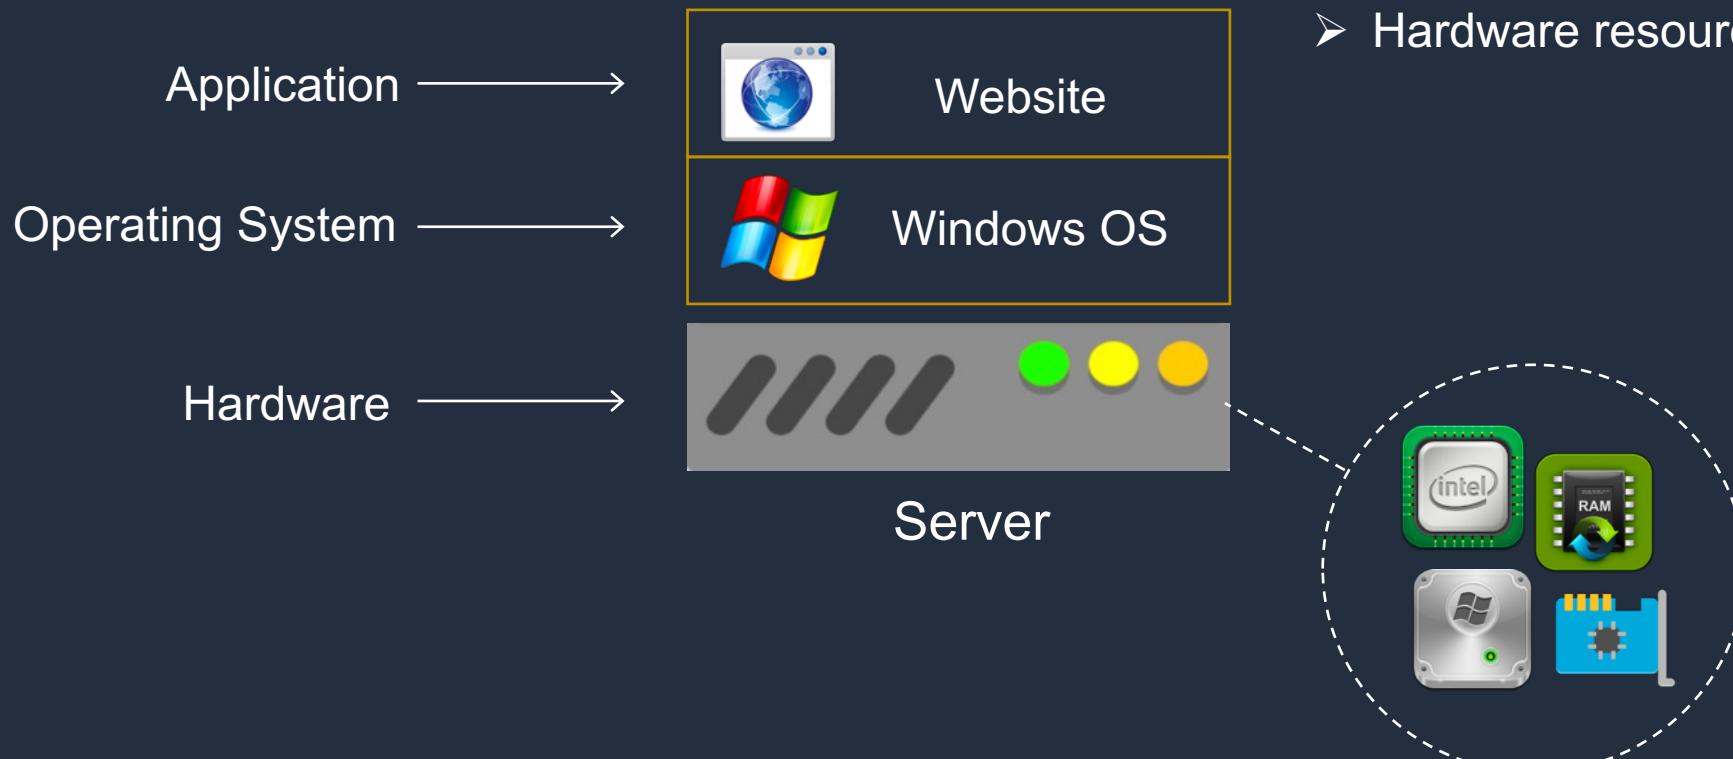
A **port** is like a door into the server

Server Virtualization





Without Server Virtualization

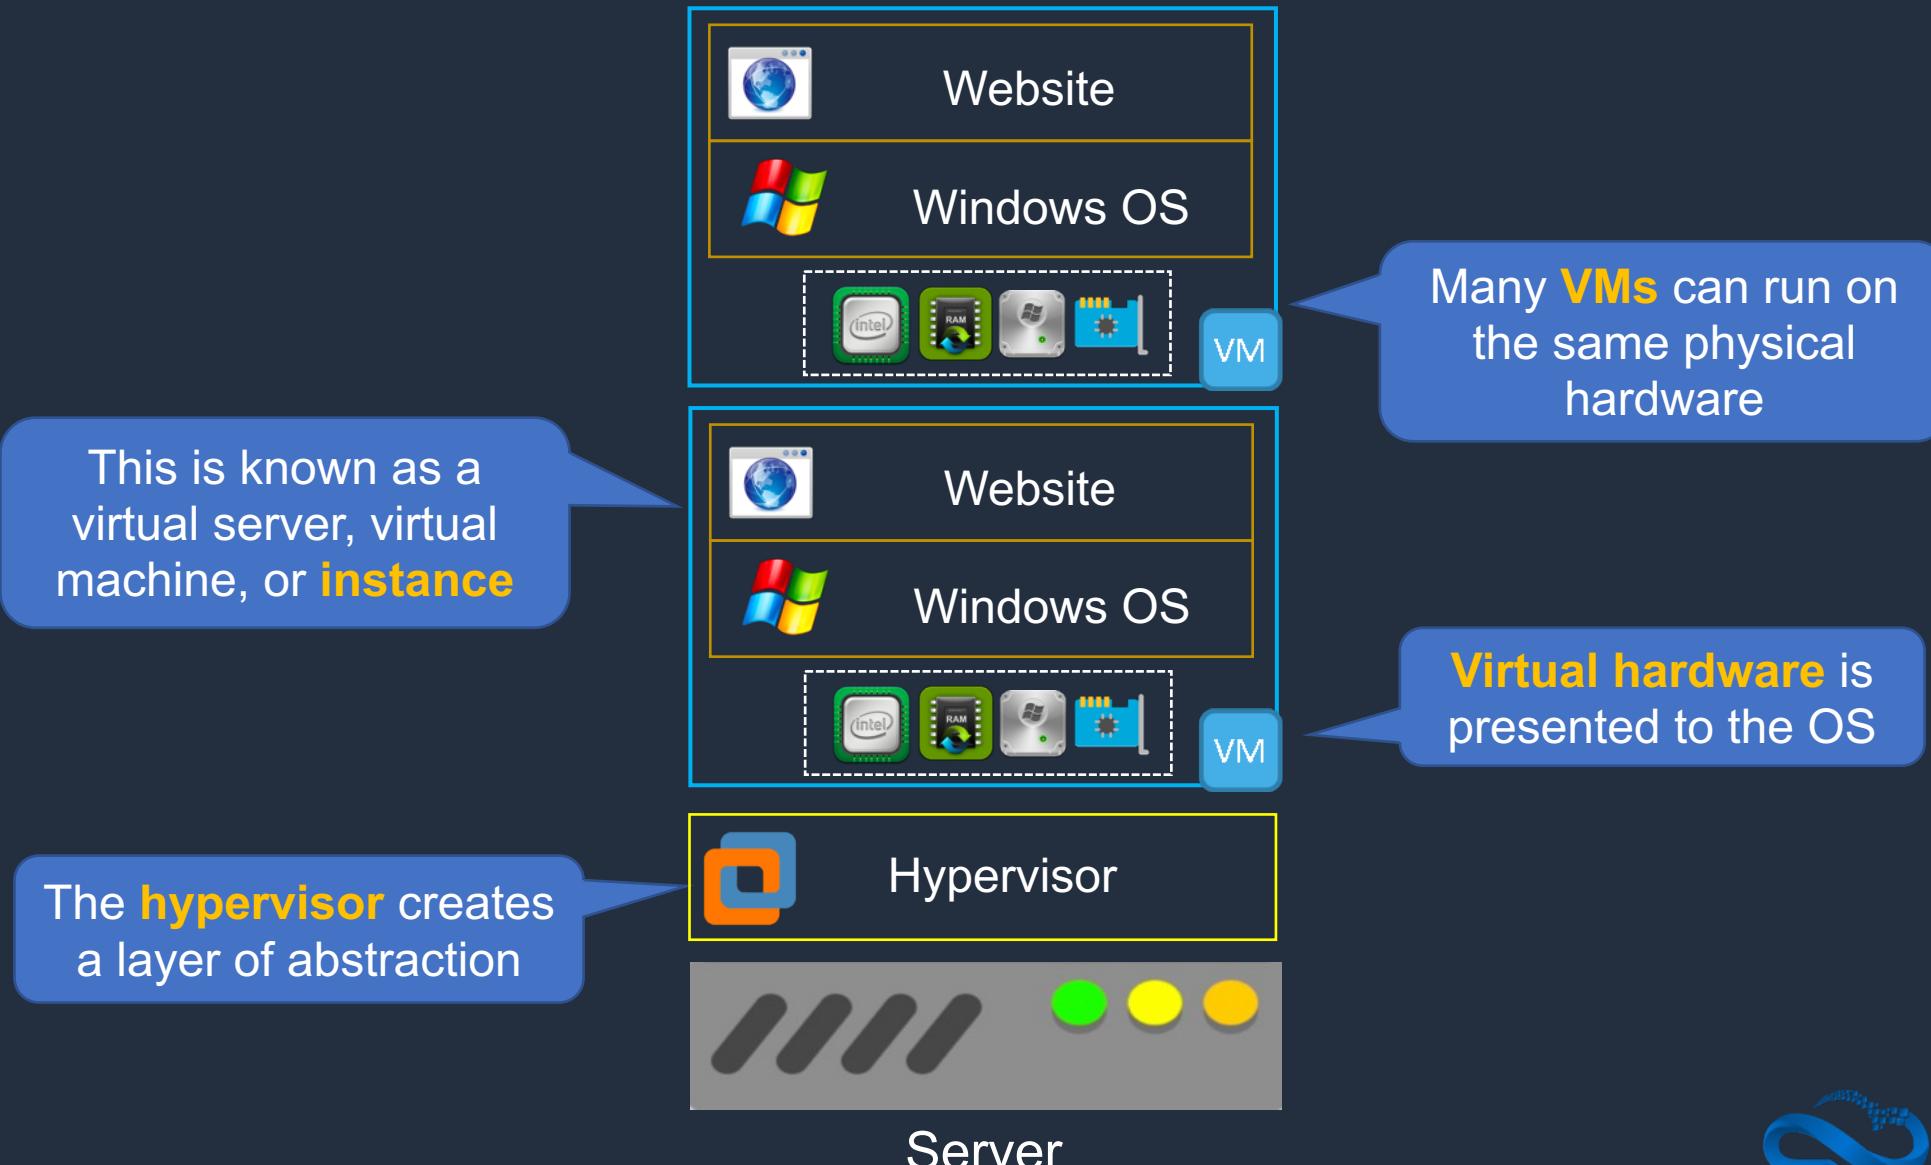


Limitations:

- OS is tied to hardware (no portability)
- Hardware resources may be underutilized

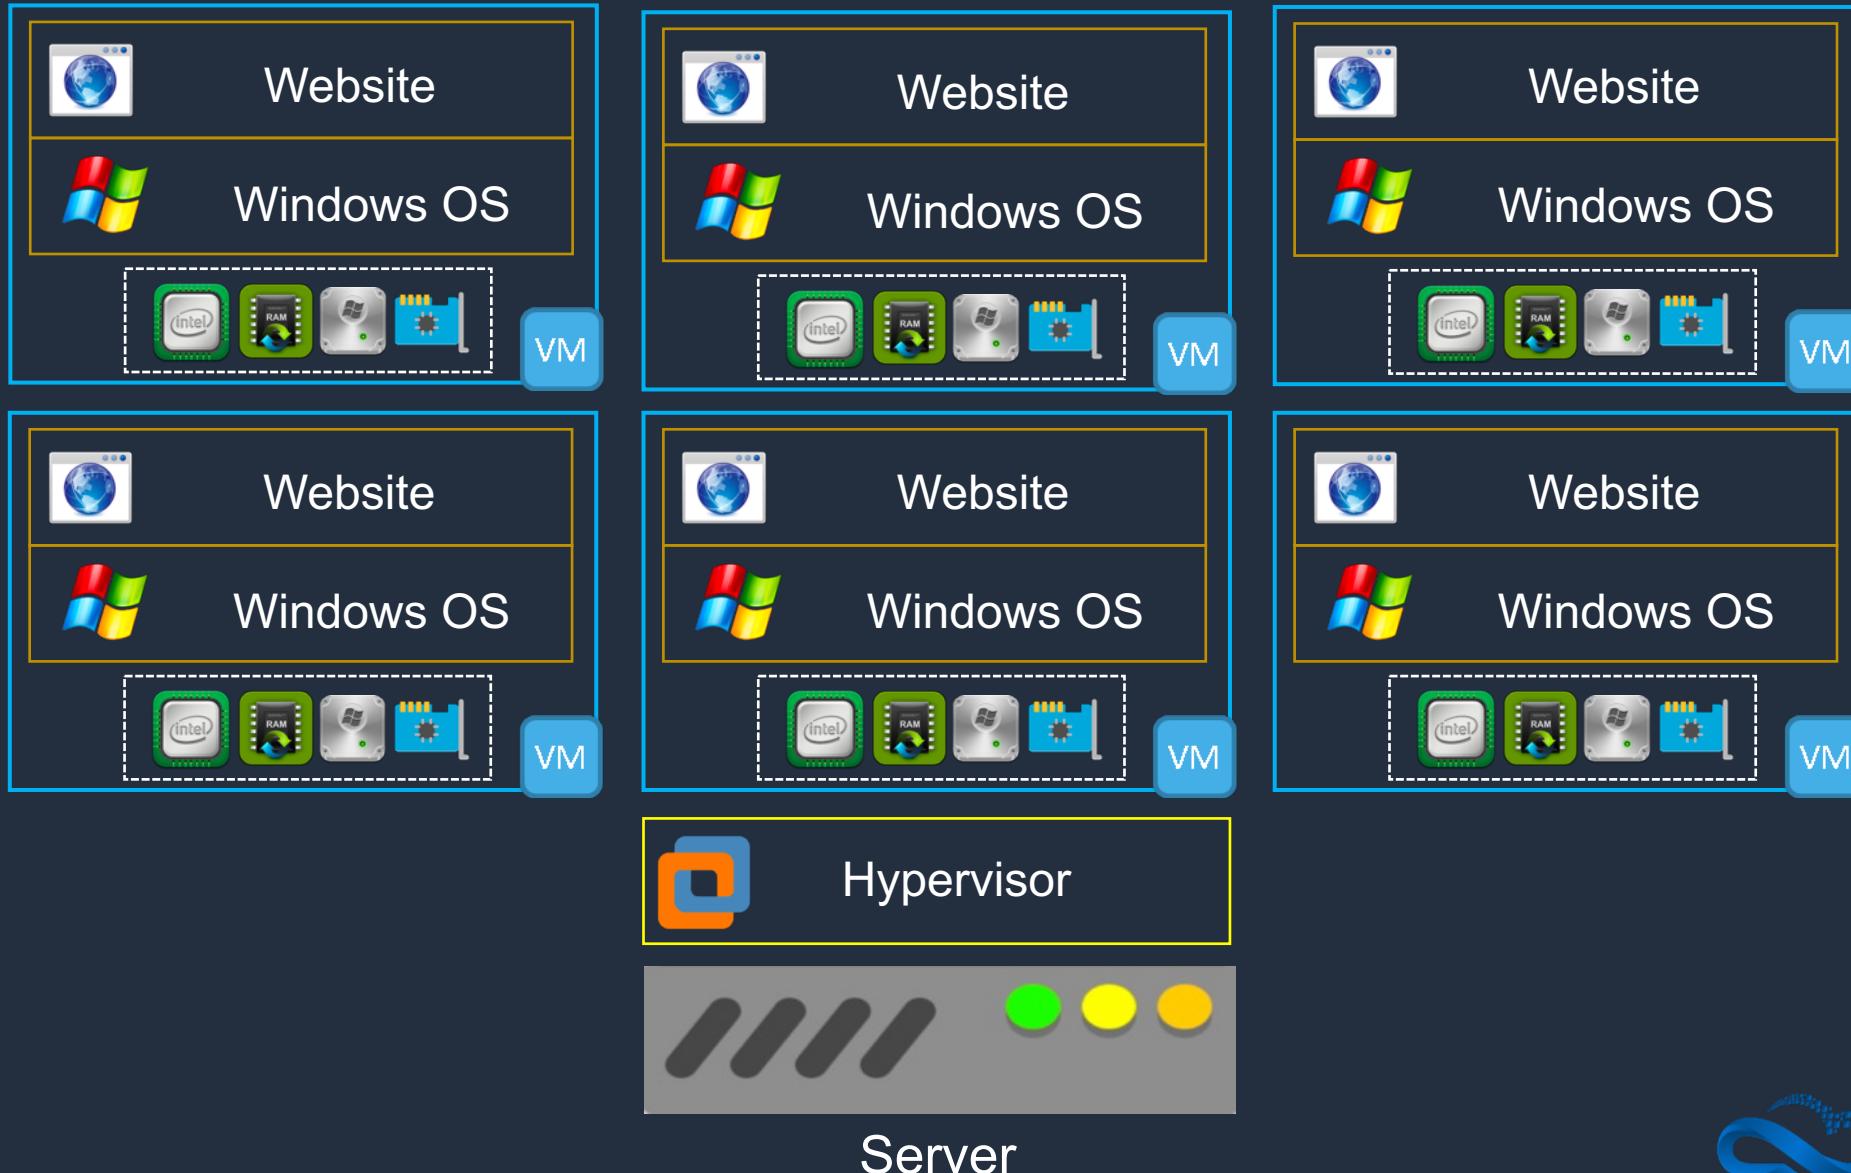


Server Virtualization





Server Virtualization

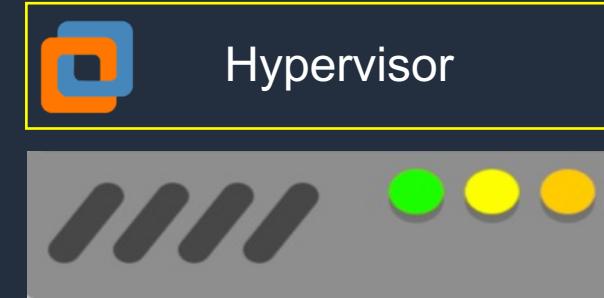




Server Virtualization



Server



Server



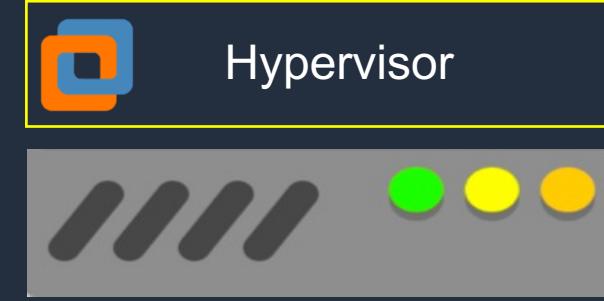
Server Virtualization



Server

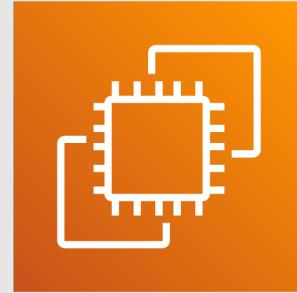


Hypervisor



Server

Amazon Elastic Compute Cloud (EC2)



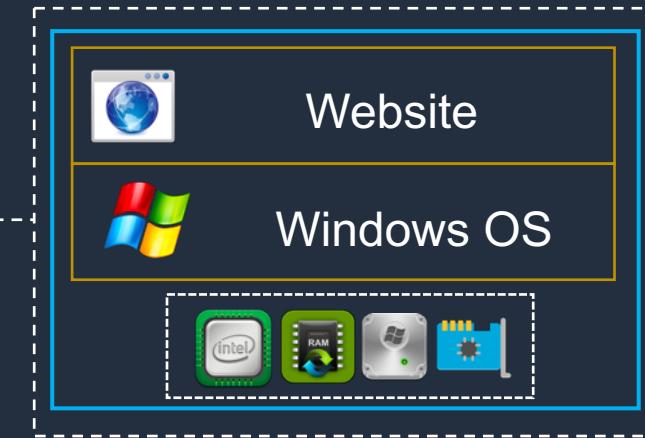
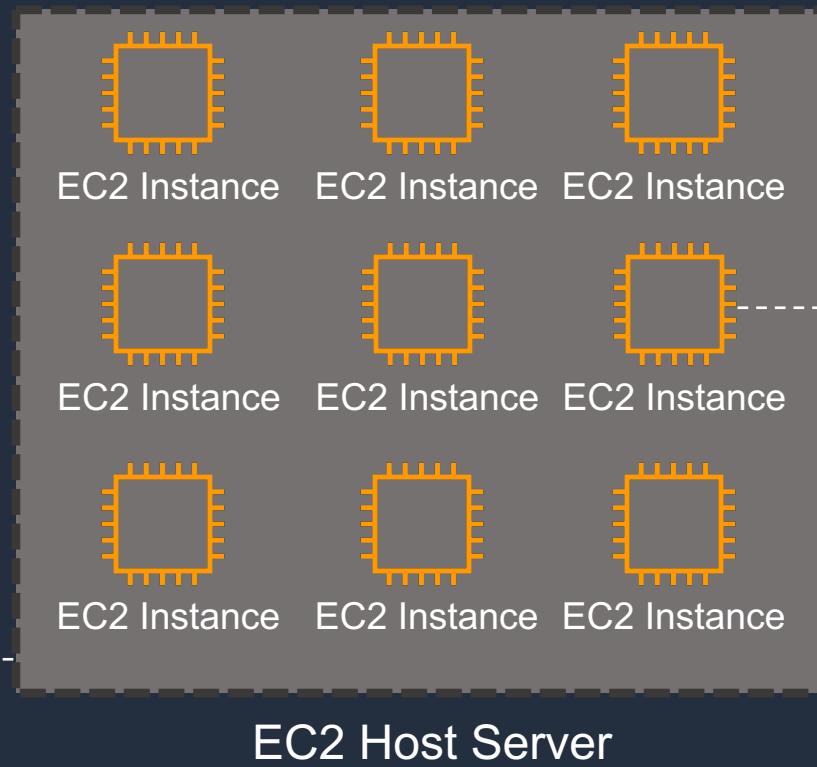


Amazon EC2

EC2 instances run
Windows, Linux, or
MacOS

An **EC2 instance**
is a virtual server

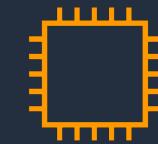
EC2 hosts are
managed by AWS



A selection of **instance types**
come with varying combinations
of CPU, memory, storage and
networking



Launching an EC2 Instance



Select an **instance type**

An **AMI** defines the configuration of the instance

EC2 Instance

A **snapshot** is a point-in-time backup of an instance



Amazon Machine Image (AMI)



EBS Snapshot



Linux



Microsoft Windows

You can customize your instance and create a **custom AMI**



Customized AMI

Family	Type	vCPUs	Memory (GiB)
General purpose	t2.micro	1	1
Compute optimized	c5n.large	2	5.25
Memory optimized	r5ad.large	2	16
Storage optimized	d2.xlarge	4	30.5
GPU instances	g2.2xlarge	8	15

The **instance type** defines the hardware profile (and cost)



Benefits of Amazon EC2

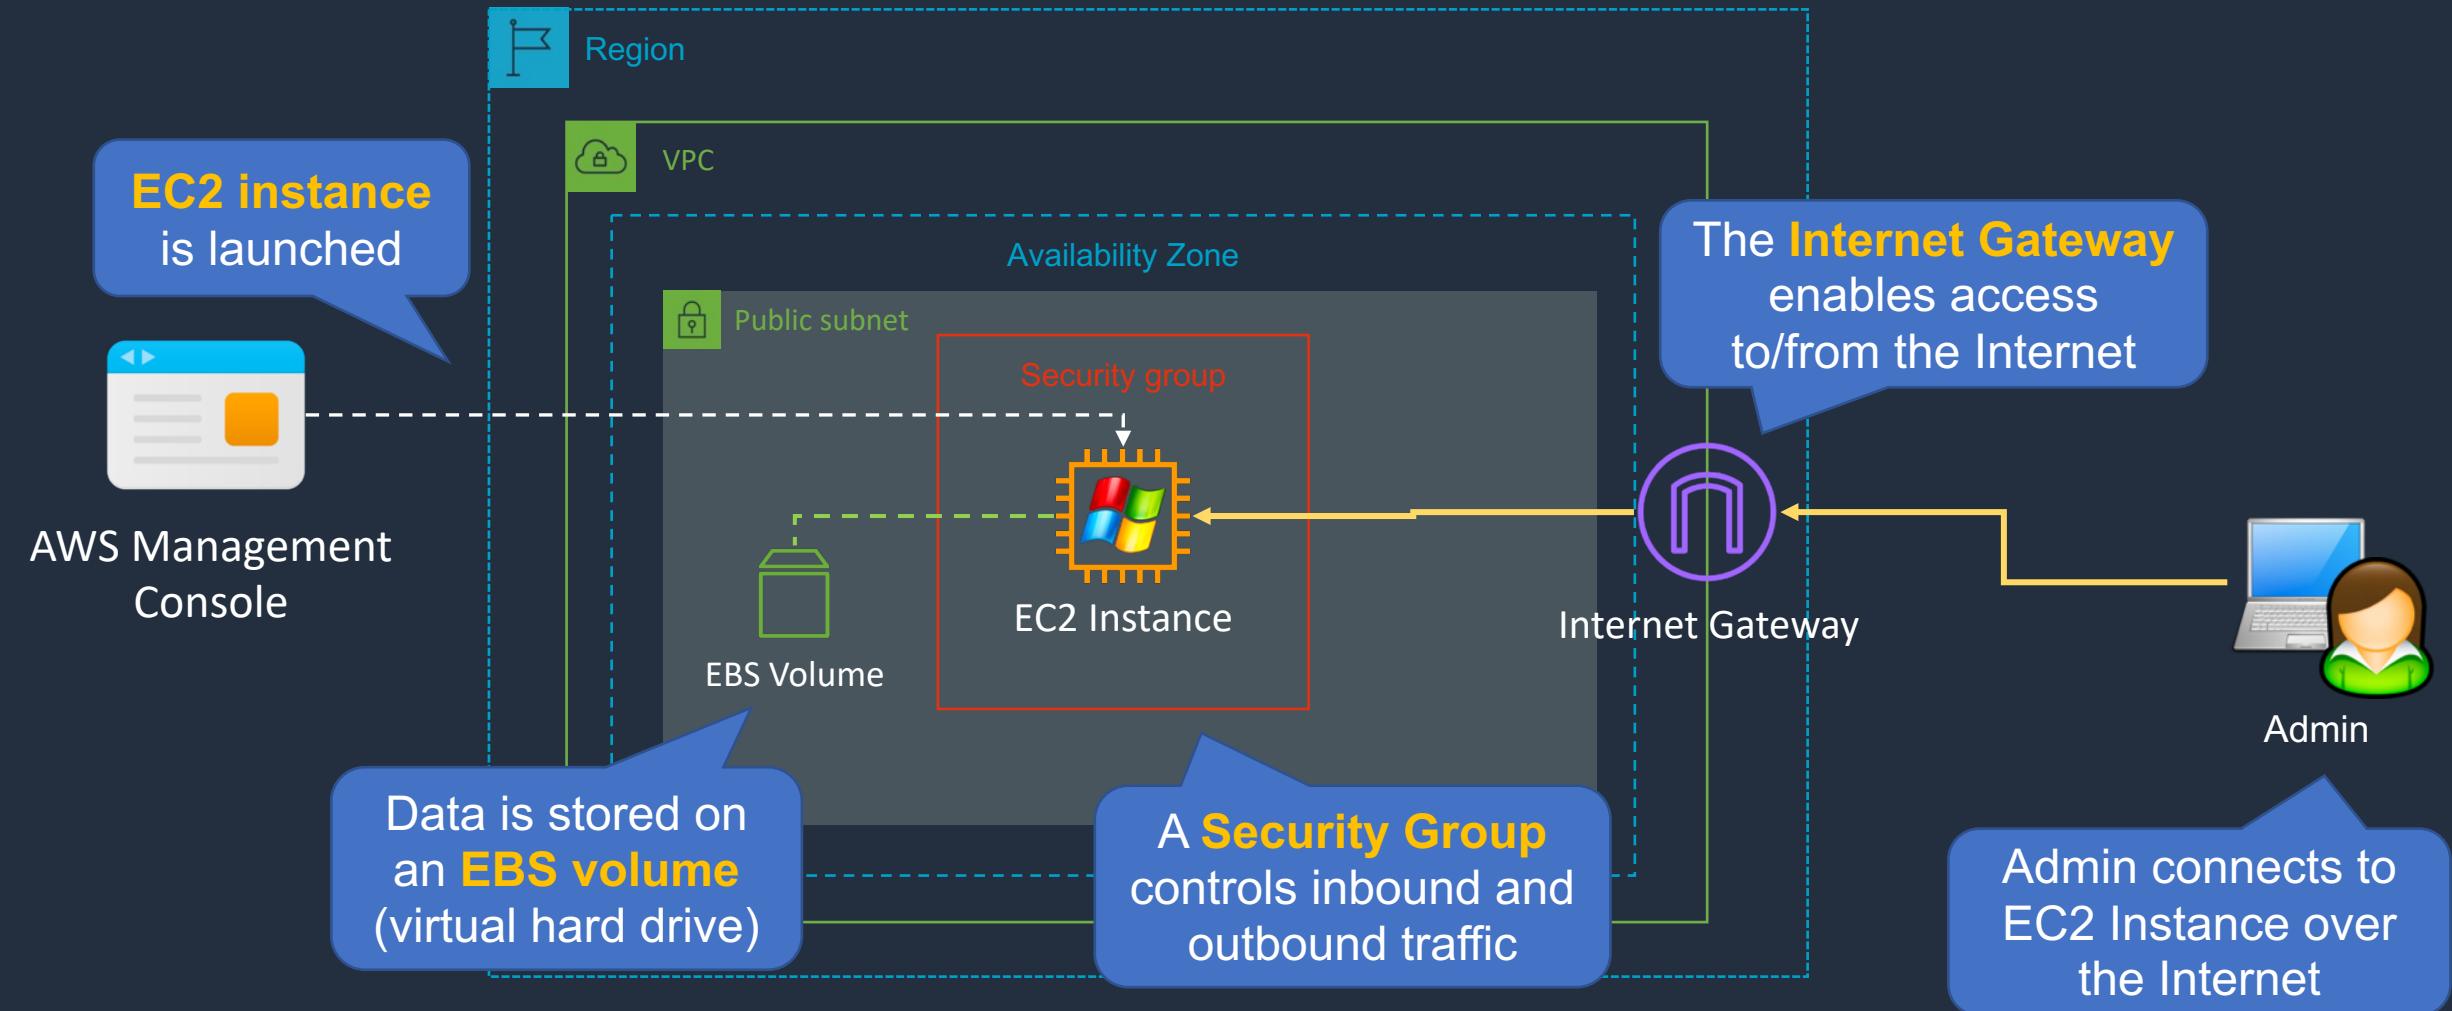
- **Elastic computing** – easily launch hundreds to thousands of EC2 instances within minutes
- **Complete control** – you control the EC2 instances with full root/administrative access
- **Flexible** – Choice of instance types, operating systems, and software packages
- **Reliable** – EC2 offers very high levels of availability and instances can be rapidly commissioned and replaced
- **Secure** – Fully integrated with Amazon VPC and security features
- **Inexpensive** – Low cost, pay for what you use

Launch EC2 Instances (Windows + Linux)





Amazon EC2 Instance in a Public Subnet



EC2 Instance Connect and SSH



RDP to Windows Instance



Amazon EC2 User Data and Metadata





Amazon EC2 User Data

The code is run when the instance starts for the **first time**



AWS Management Console

Batch and **PowerShell** scripts can be run on Windows

User data i As text As file Input is already base64 encoded

```
#!/bin/bash
yum update -y
yum install -y httpd
systemctl start httpd
systemctl enable httpd
```

Limited to
16 KB



EC2 Instance

EC2 Instance with a **web service** is launched



Amazon EC2 Metadata

- Instance metadata is data about your EC2 instance
- Instance metadata is available at <http://169.254.169.254/latest/meta-data>
- Examples:



```
[ec2-user@ip-172-31-42-248 ~]$ curl http://169.254.169.254/latest/meta-data
ami-id
ami-launch-index
ami-manifest-path
block-device-mapping/
events/
hibernation/
hostname
identity-credentials/
instance-action
instance-id
instance-life-cycle
instance-type
local-hostname
local-ipv4
```



Amazon EC2 Metadata

- Examples ctd.:

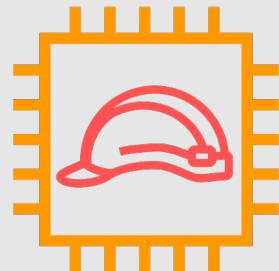
```
[ec2-user@ip-172-31-42-248 ~]$ curl http://169.254.169.254/latest/meta-data/local-ipv4  
172.31.42.248[ec2-user@ip-172-31-42-248 ~]$
```

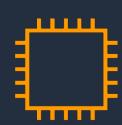
```
[ec2-user@ip-172-31-42-248 ~]$ curl http://169.254.169.254/latest/meta-data/public-ipv4  
3.26.54.18[ec2-user@ip-172-31-42-248 ~]$
```

[HOL] Launch Instance with User Data and Metadata

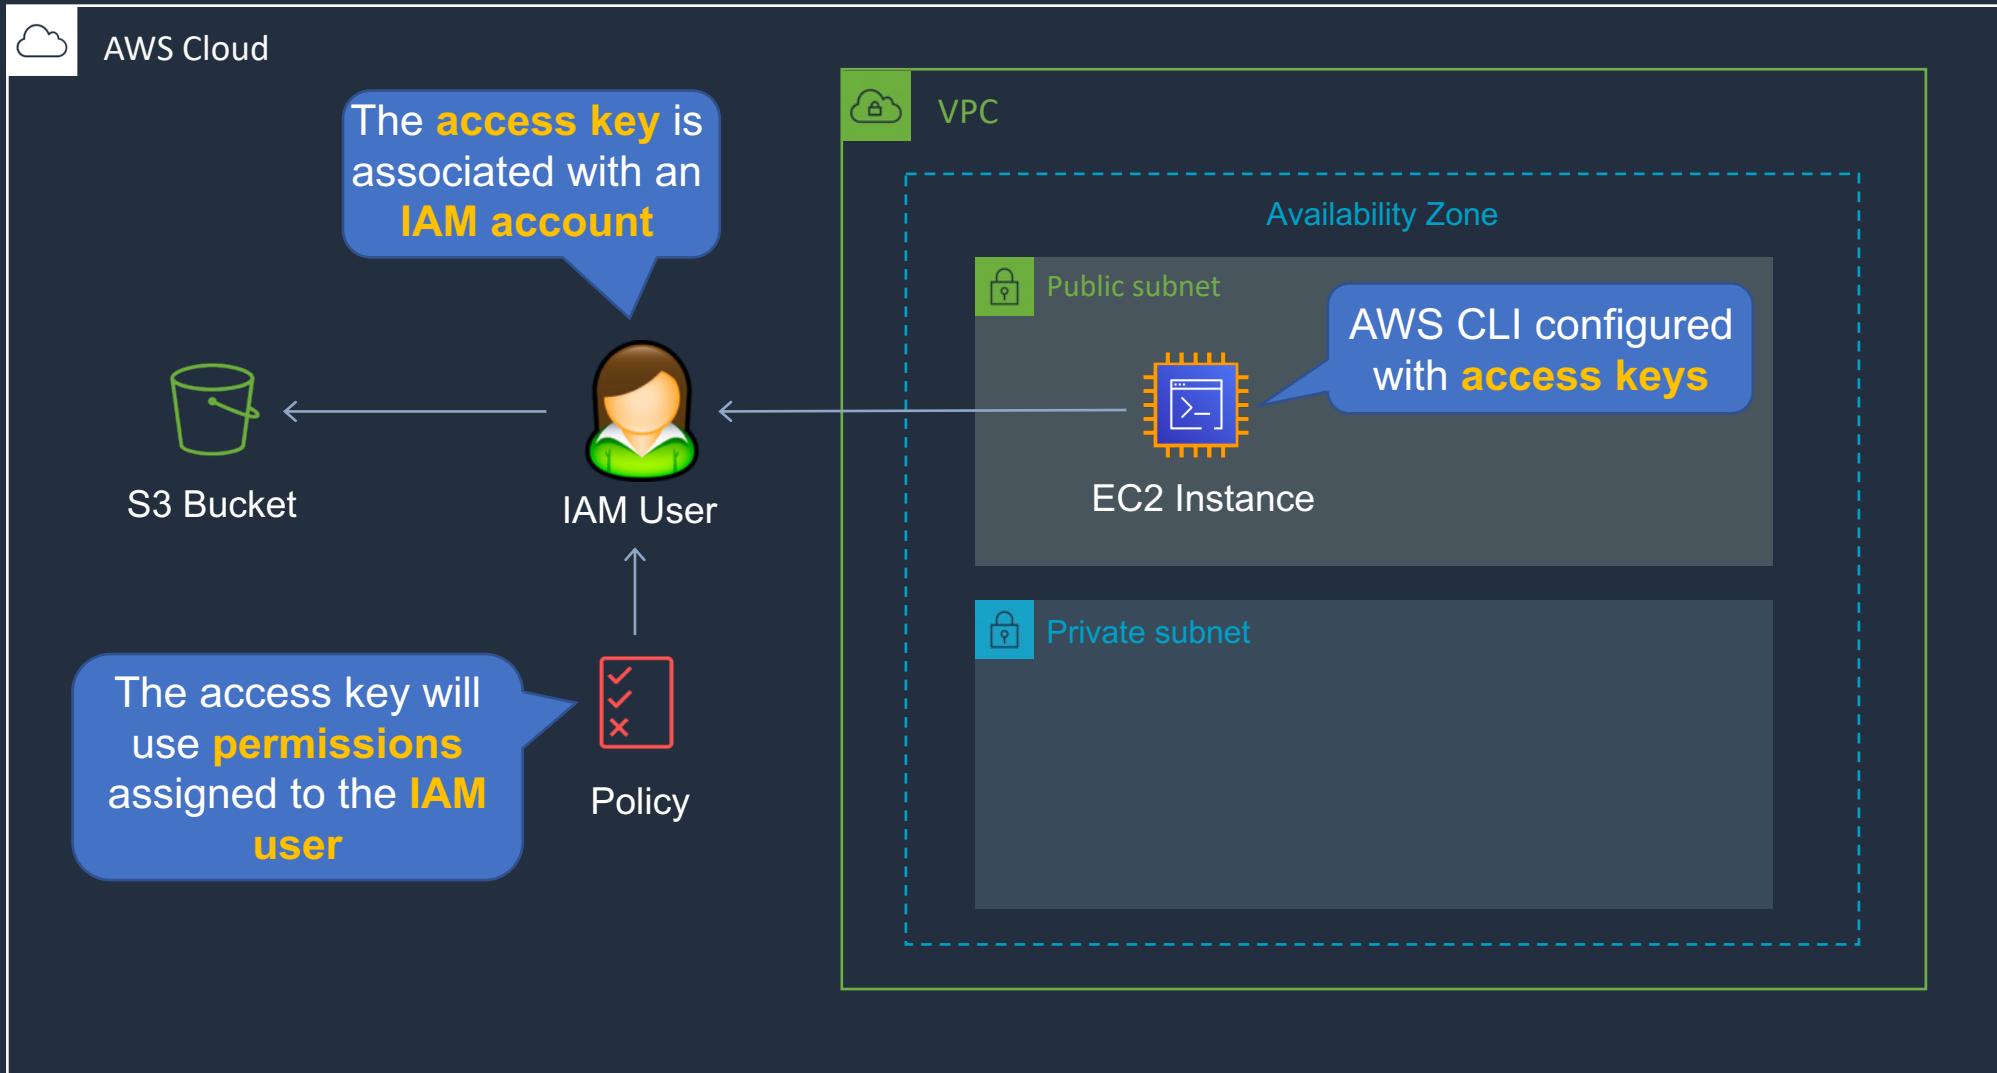


Accessing Services – Access Keys and IAM Roles



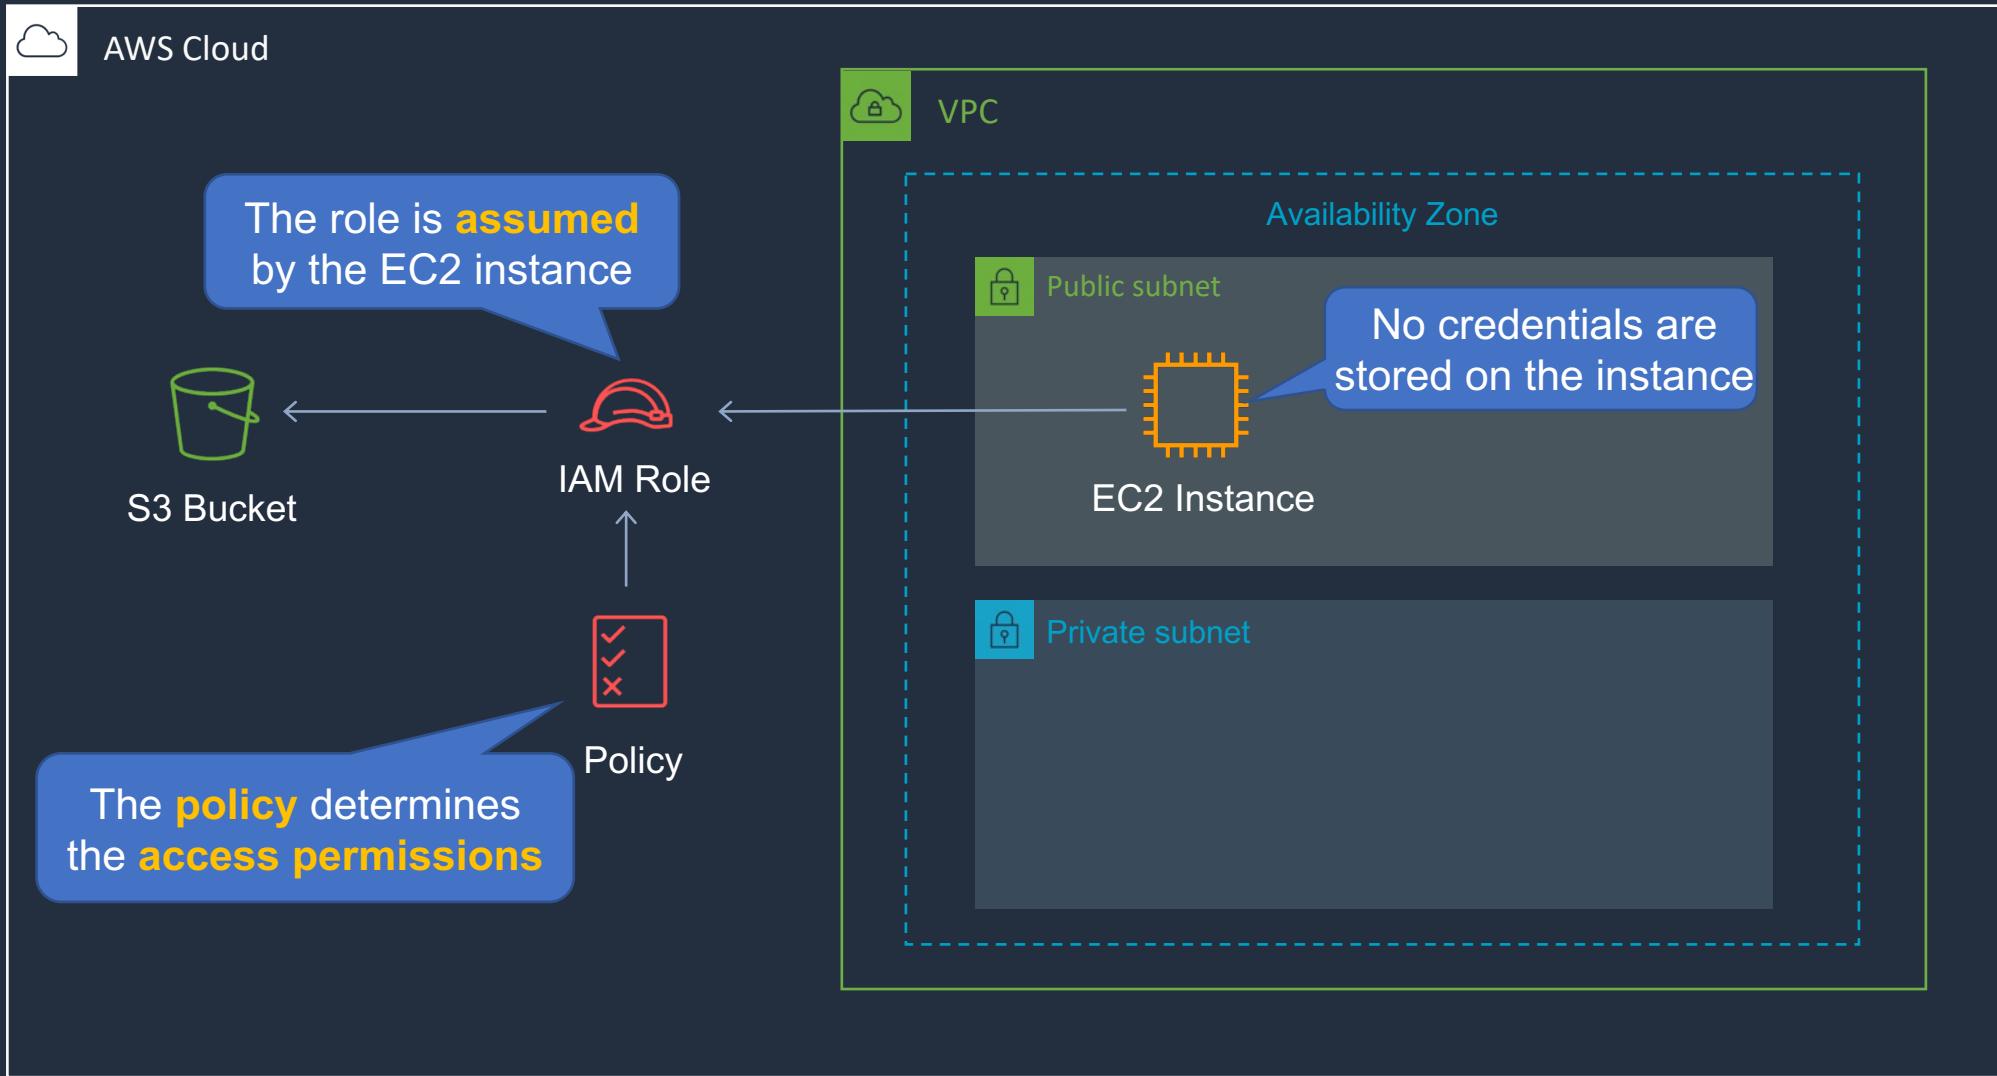


Access Keys





Amazon EC2 Instance Profiles (IAM Roles for EC2)



Access Keys and IAM Roles



AWS Batch





AWS Batch



Launch a **Batch Job**



Job **Definition**

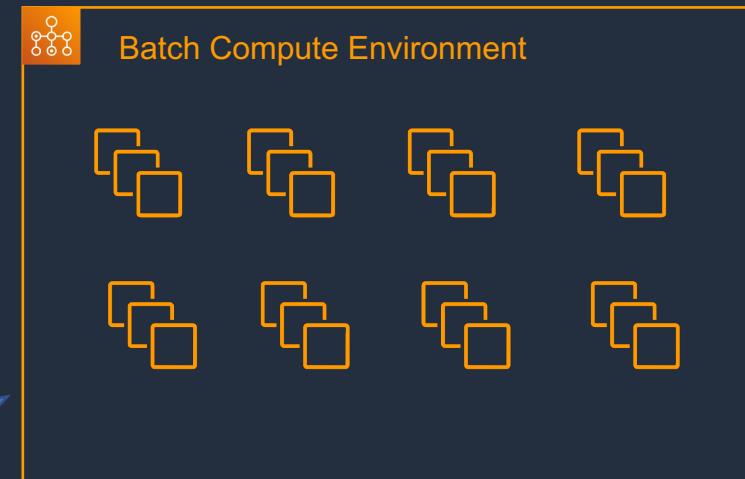


Job **Queue**

A job is submitted to a **queue** until **scheduled** onto a compute environment

A job is a unit of work such as a **shell script**, **executable** or **Docker container image**

Batch **launches**, **manages**, and **terminates** resources as required (EC2 and ECS/Fargate)



Managed or **unmanaged** resources used to run the job

Amazon LightSail





Amazon LightSail

- Low cost and ideal for users with less technical expertise
- Compute, storage, and network
- Preconfigured virtual servers

Amazon's Simple Cloud Server

Amazon Lightsail - Powerful virtual servers built for reliability & performance

FREE TIER OFFER

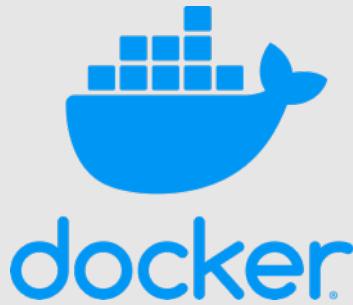
1 MONTH FREE*

1 Month Free Trial Starting at \$3.50/month Starting at \$8/month

- Virtual servers, databases and load balancers
- SSH and RDP access
- Can access Amazon VPC

Exam tip: typically comes up in use cases where an easy method of deploying a virtual server is required by a user with little or no AWS expertise

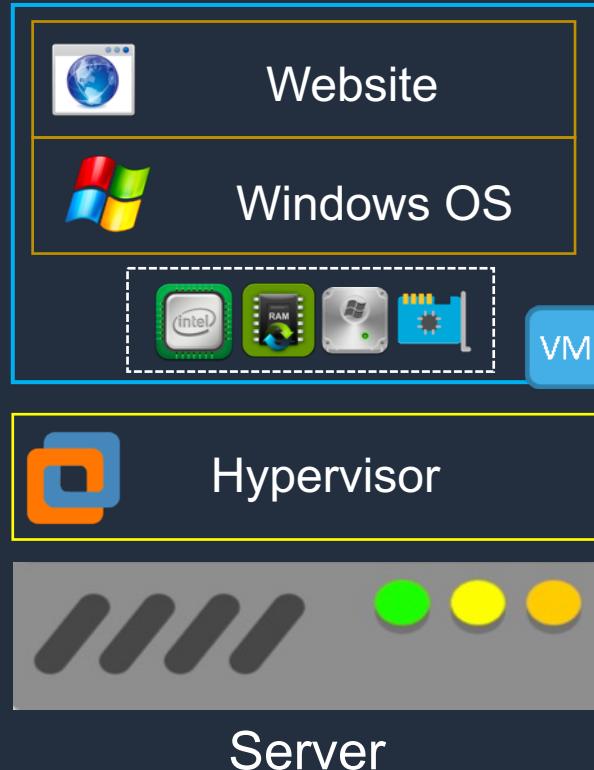
Docker Containers and Microservices





Server Virtualization vs Containers

Every VM-instance needs an **operating system** which uses significant resources





Docker Containers

Containers **start up** very **quickly**

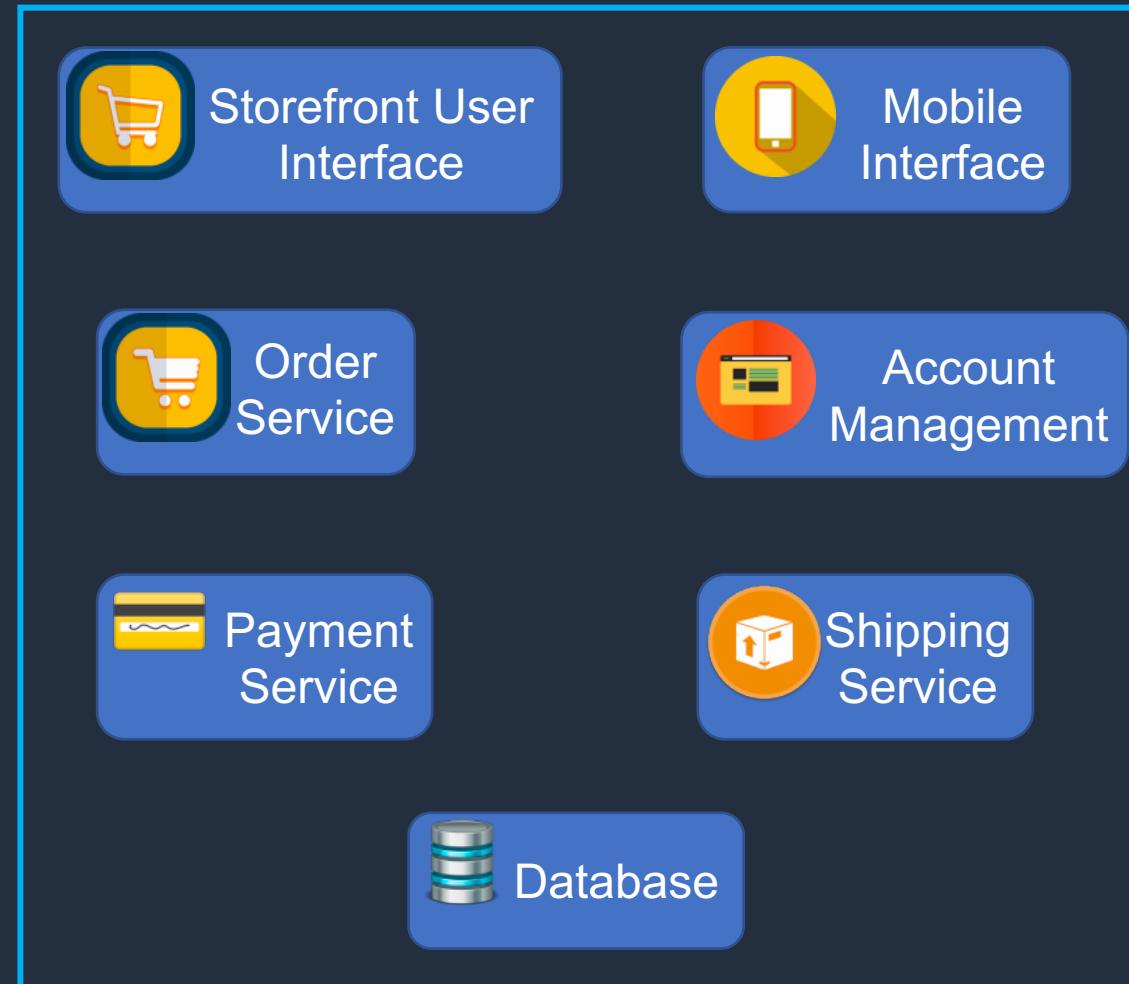
A **container** includes all the code, settings, and dependencies for running the application



Containers are very resource **efficient**

Each container is **isolated** from other containers

Monolithic Application



Monolithic Application

Updates to, or failures of, any single component can take down the whole application



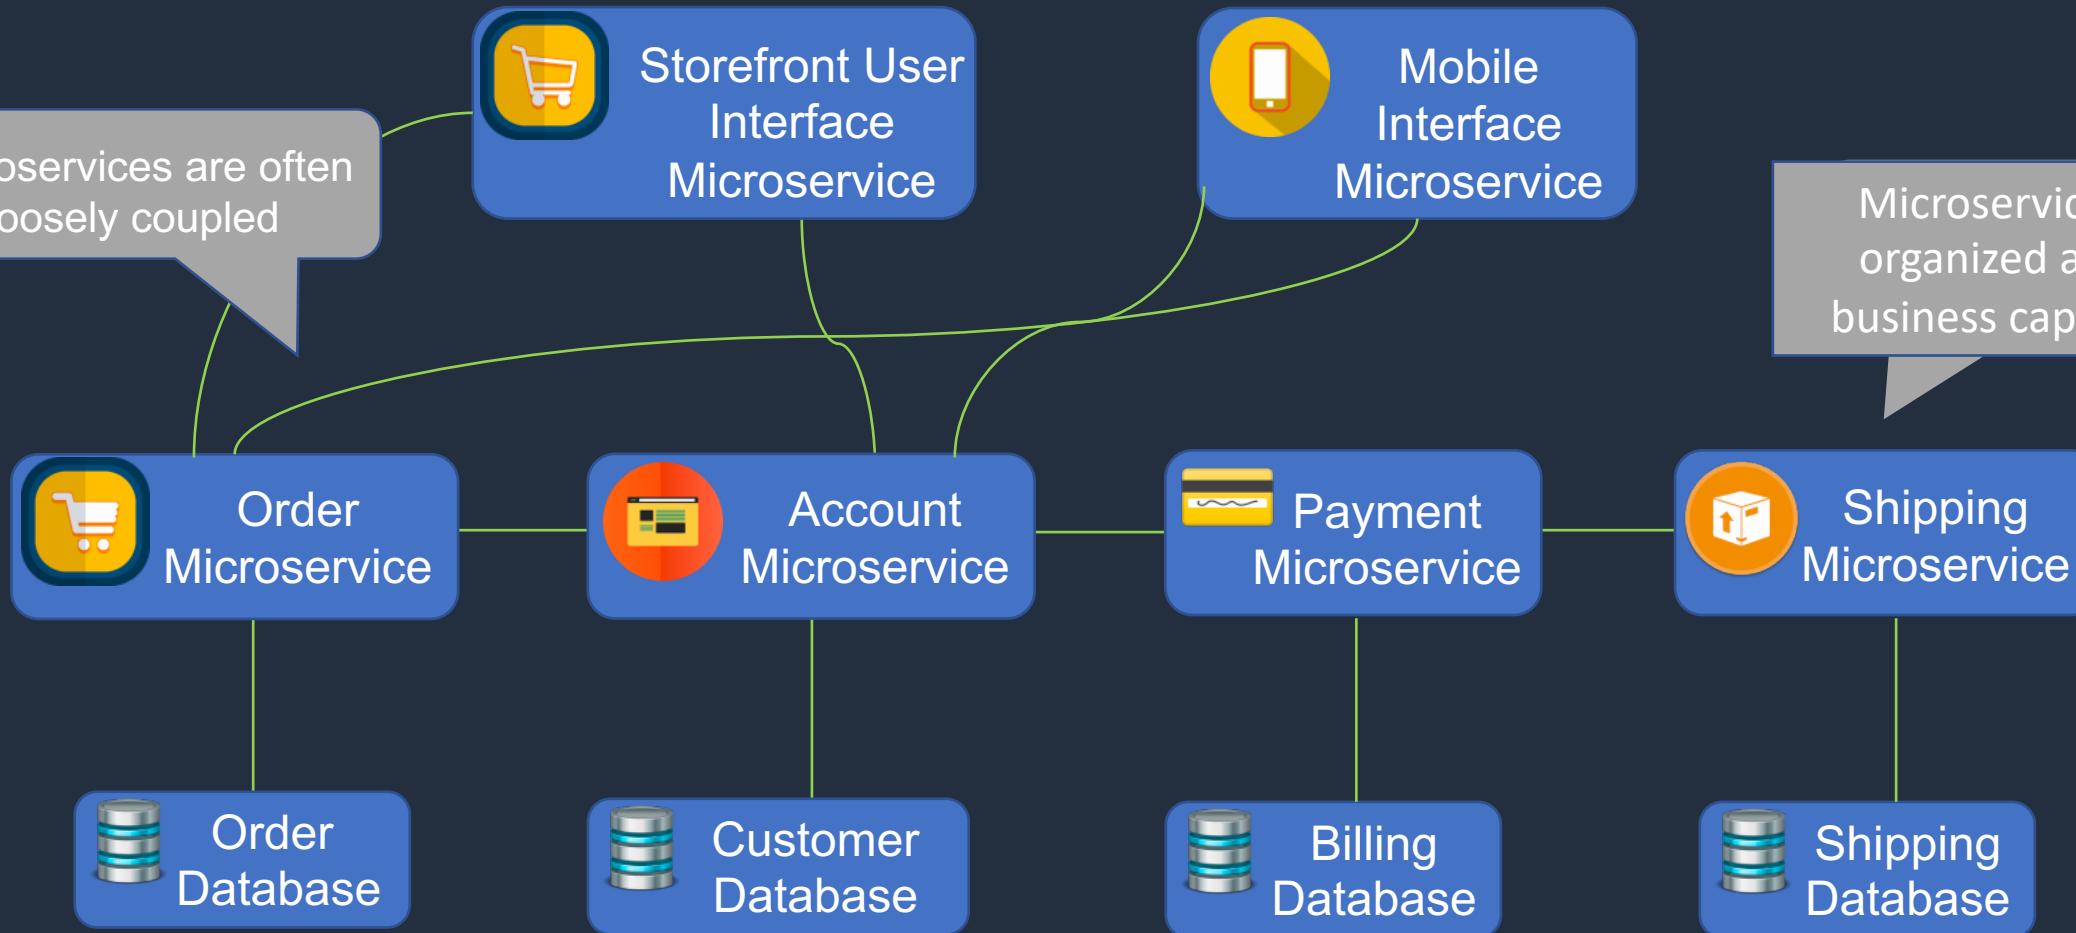
The user interface, business logic, and data access layer are combined on a single platform



Microservices Application

A microservice is an independently deployable unit of code

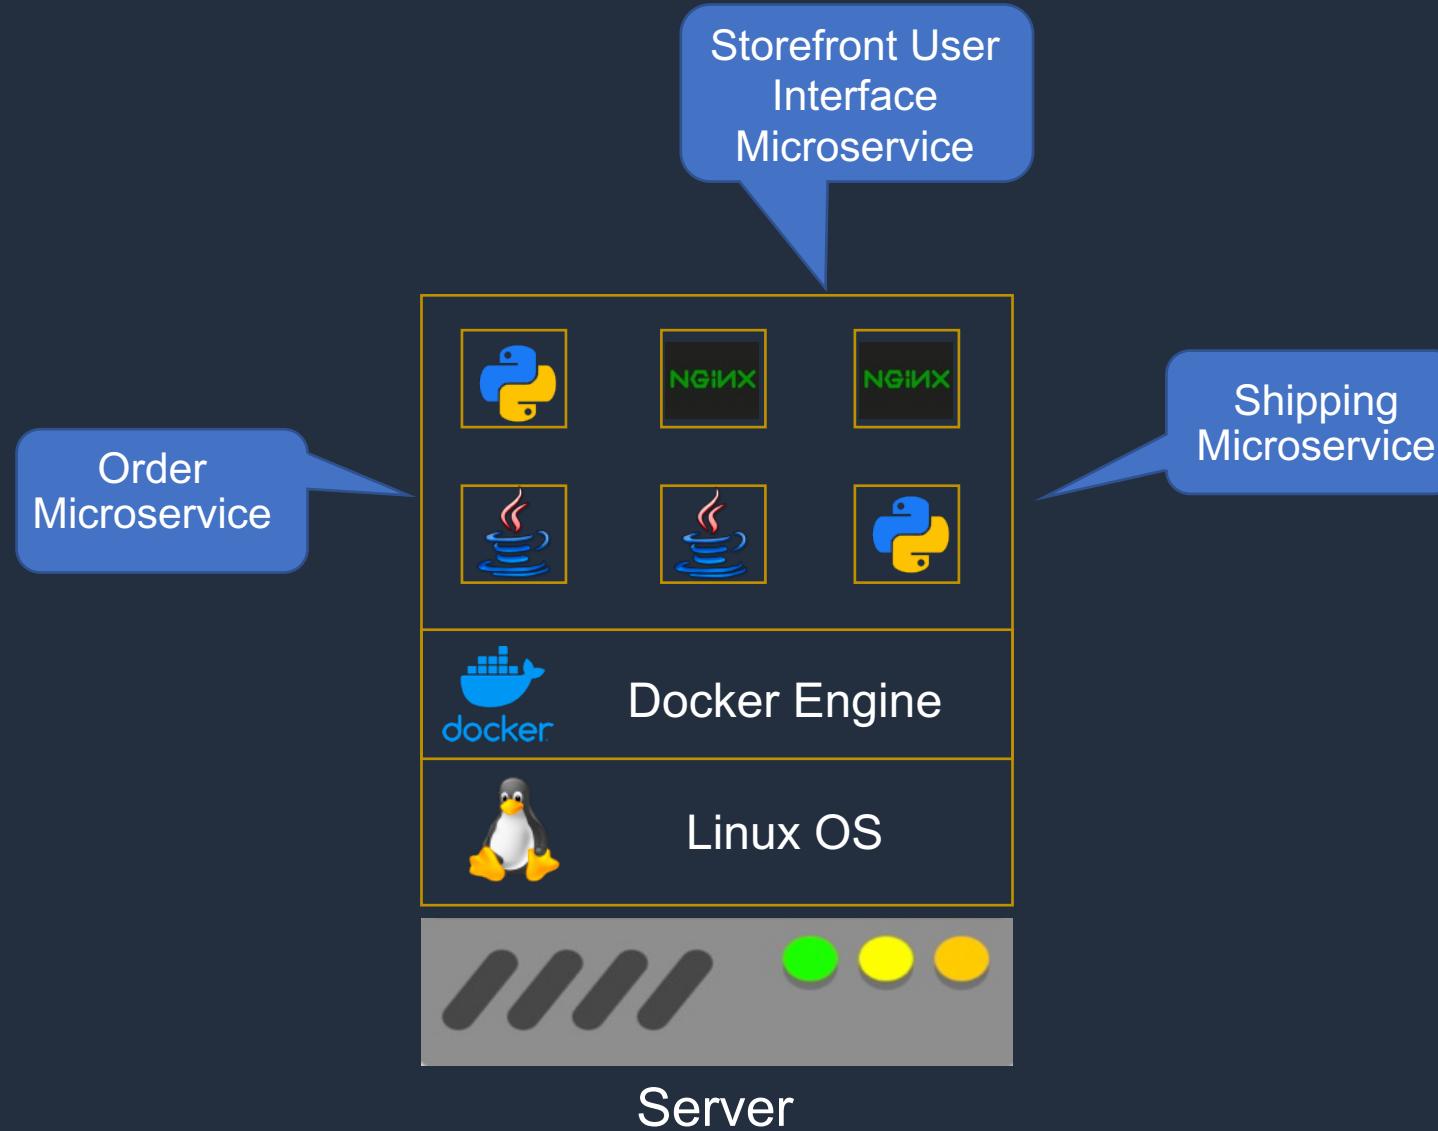
Microservices are often loosely coupled



Microservices are organized around business capabilities



Microservices Application





Microservices Application

Microservices can also be
spread across hosts

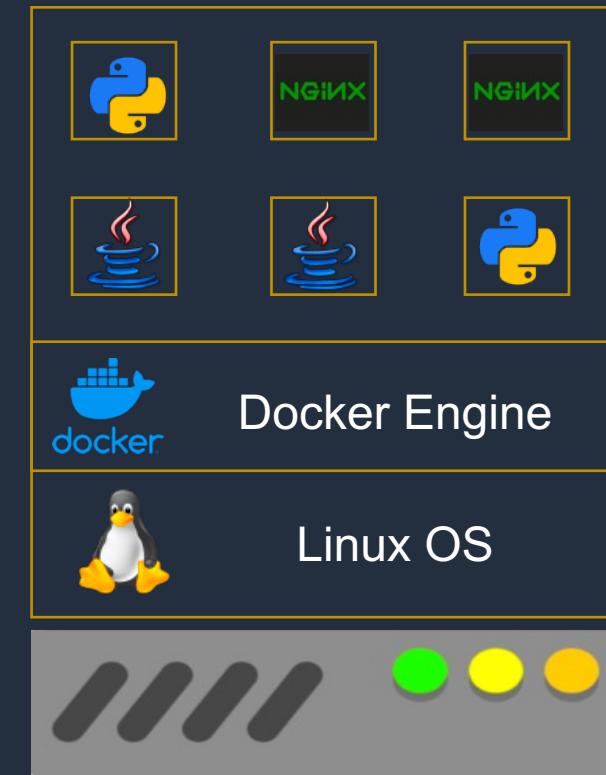
Many instances of each microservice
can run on each host



Server



Server



Server

Amazon Elastic Container Service (ECS)





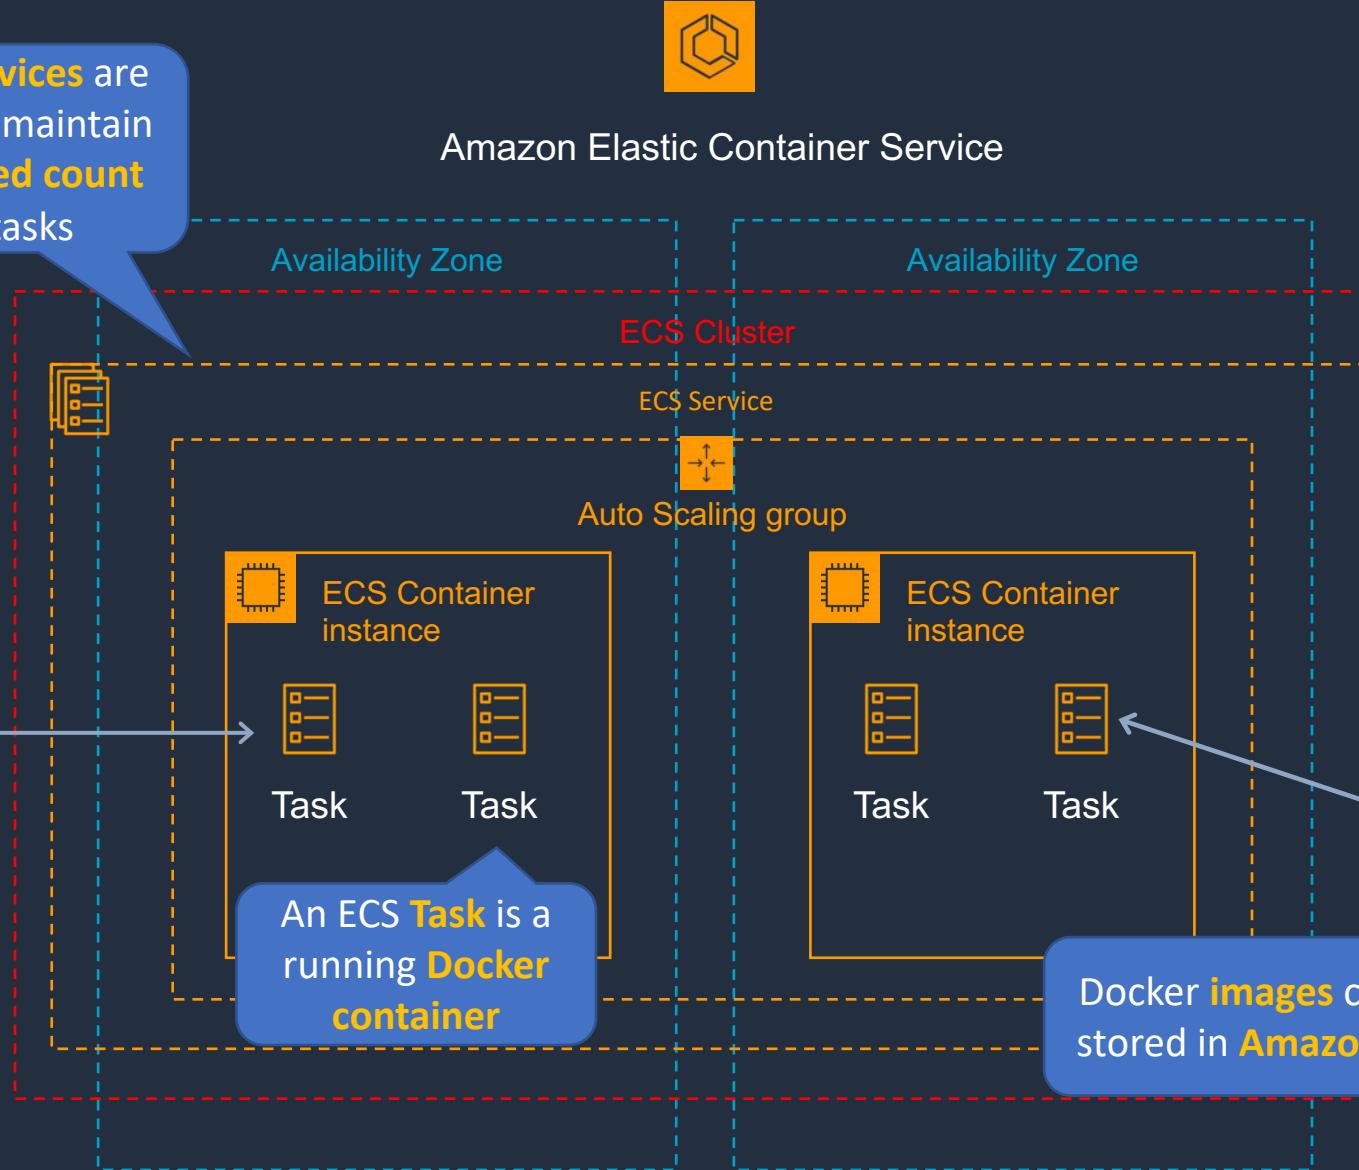
Amazon ECS

An ECS **Task** is created from a **Task Definition**

Task Definition

```
{  
  "containerDefinitions": [  
    {  
      "name": "wordpress",  
      "links": [  
        "mysql"  
      ],  
      "image": "wordpress",  
      "essential": true,  
      "portMappings": [  
        {  
          "containerPort": 80,  
          "hostPort": 80  
        }  
      ],  
      "memory": 500,  
      "cpu": 10  
    }  
  ]  
}
```

ECS **Services** are used to maintain a **desired count** of tasks

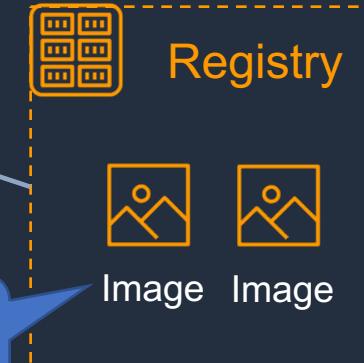


Amazon Elastic Container Service

An Amazon **ECS Cluster** is a logical grouping of **tasks** or **services**

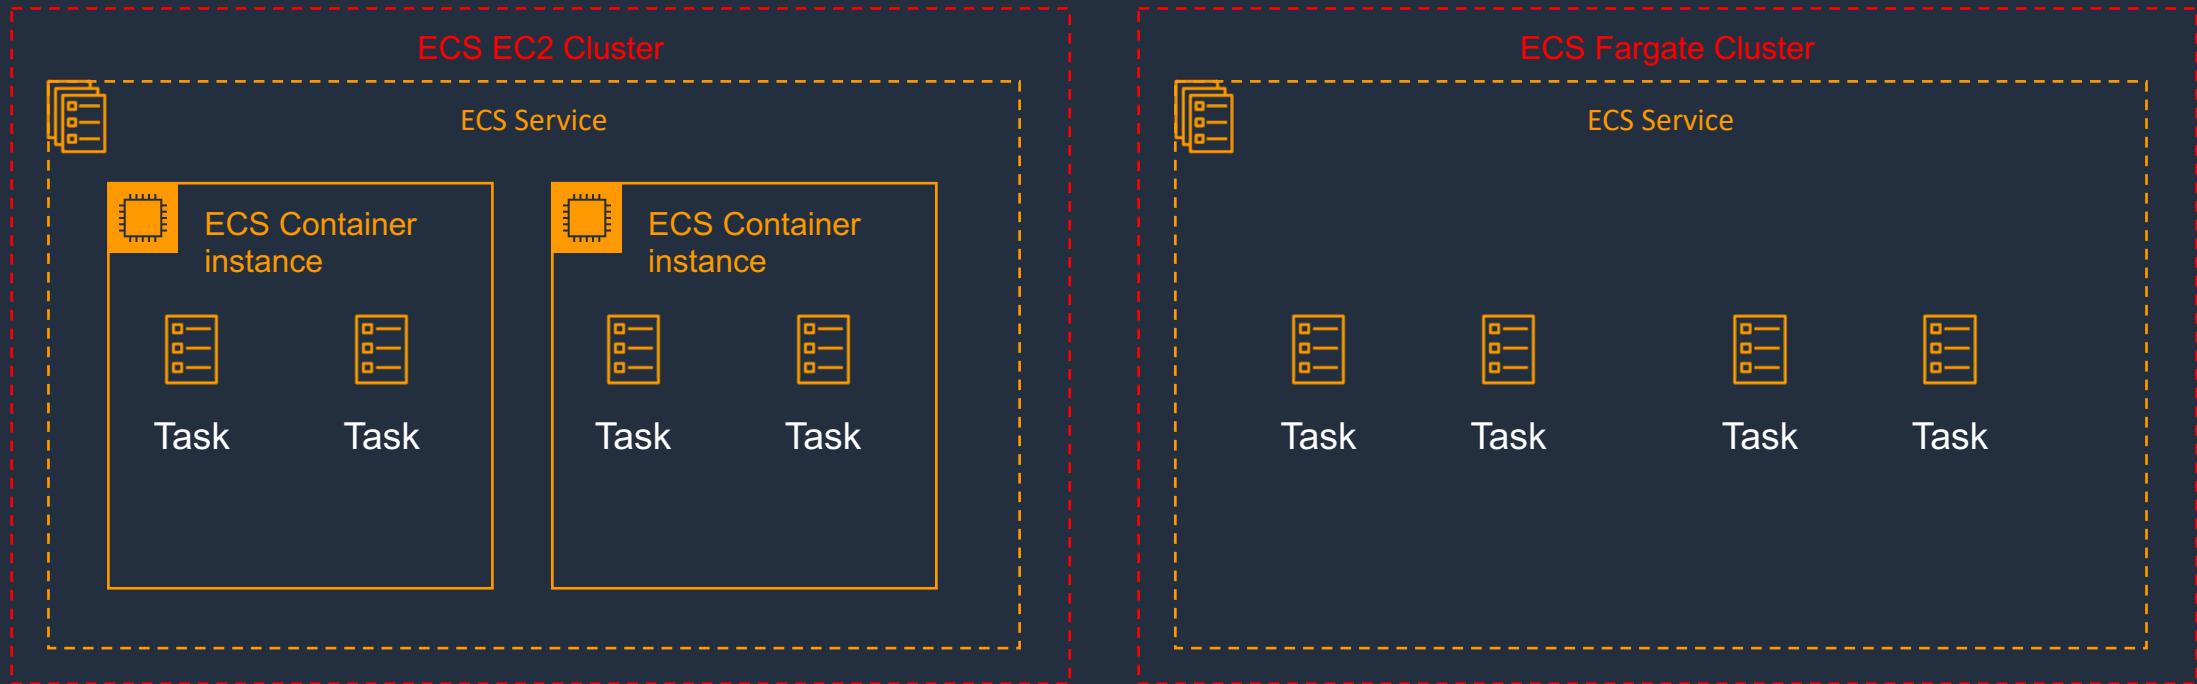


Amazon Elastic Container Registry





Amazon ECS



EC2 Launch Type

- You explicitly provision EC2 instances
- You're responsible for managing EC2 instances
- Charged per running EC2 instance
- EFS and EBS integration
- You handle cluster optimization
- More granular control over infrastructure

Fargate Launch Type

- Fargate automatically provisions resources
- Fargate provisions and manages compute
- Charged for running tasks
- No EFS and EBS integration
- Fargate handles cluster optimization
- Limited control, infrastructure is automated

Launch Docker Container on ECS



SECTION 6

AWS Storage Services

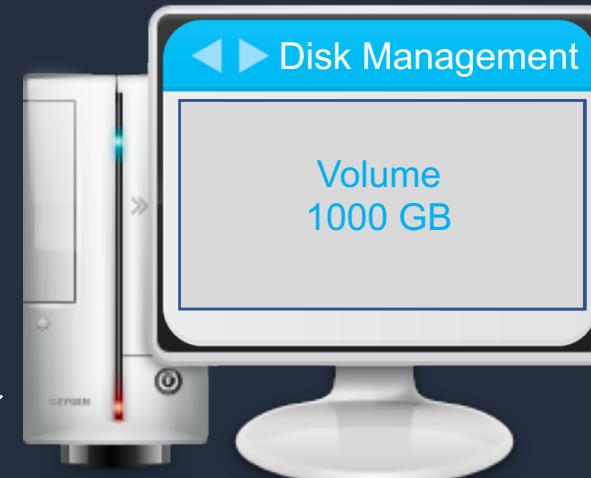
Block vs File vs Object Storage





Hard Drives

Hard drives are **block-based** storage systems

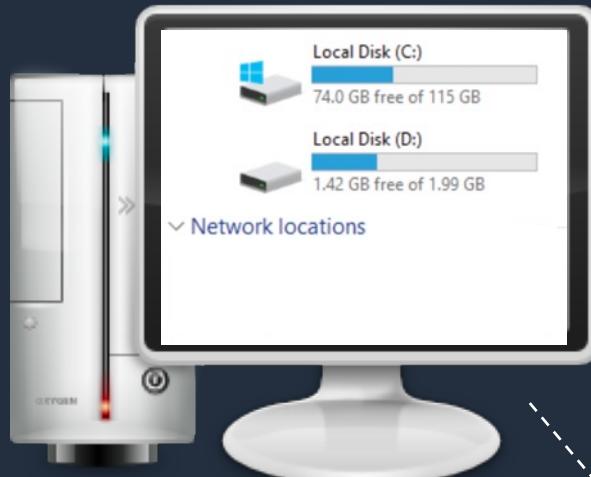


The Operating System (OS) can be used to create **volumes**. A volume can be partitioned and formatted

Hard drives are block-based storage systems



Network Attached Storage



The Operating System (OS) sees a **filesystem** that is mapped to a local drive letter

The NAS “shares” **filesystems** over the network



Network Switch



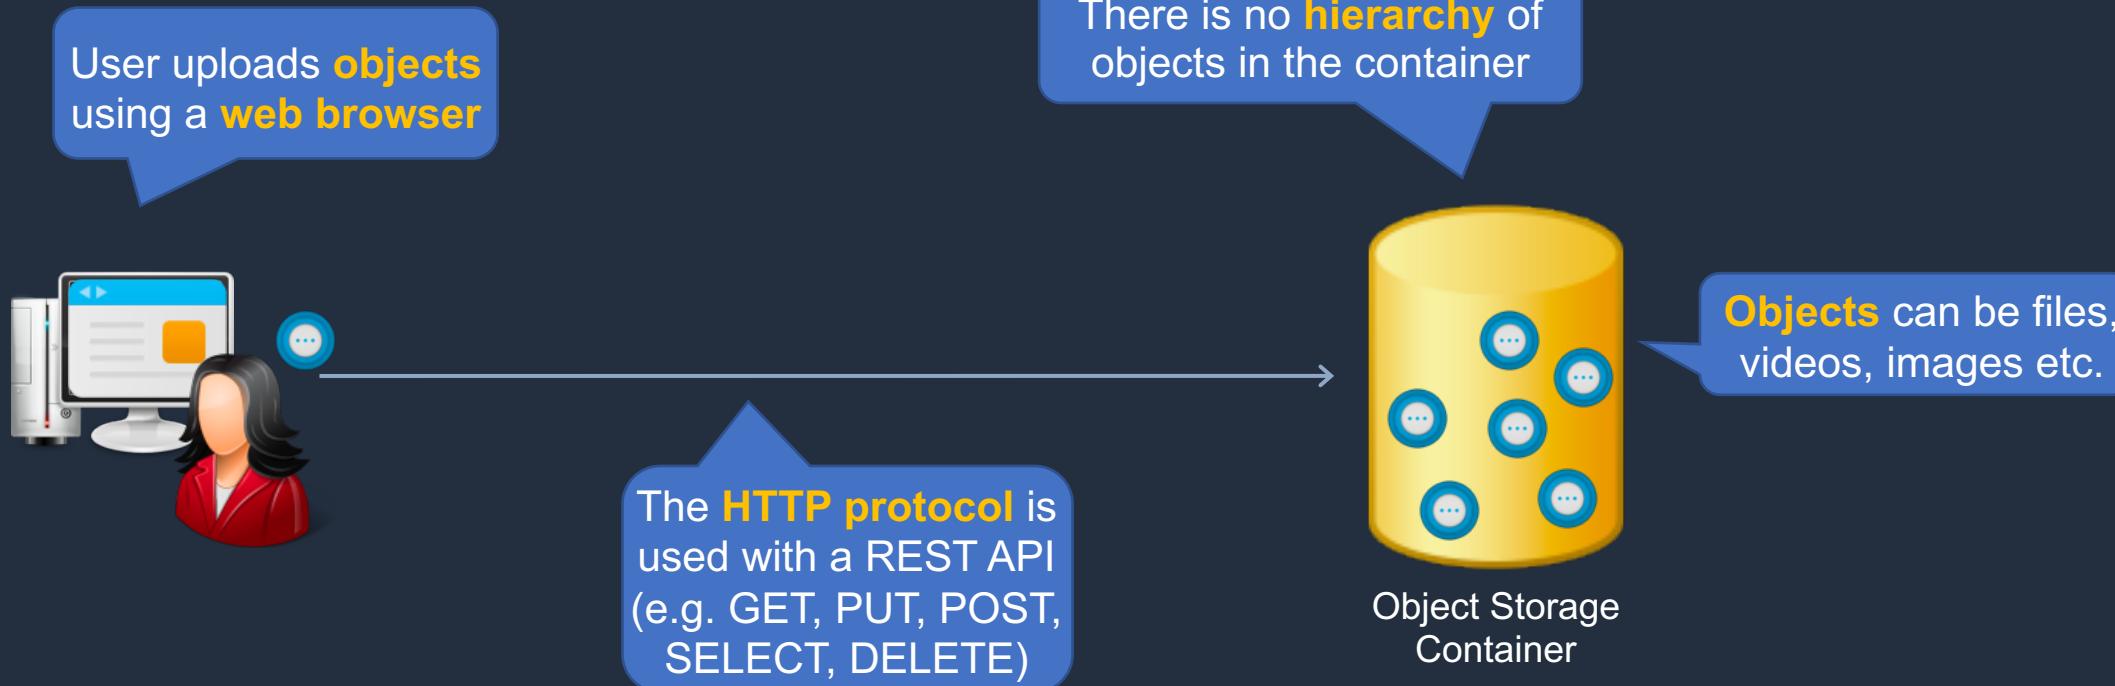
Network Attached Storage Server (NAS)



NAS devices are file-based storage systems



Object Storage Systems





Block, File, and Object Storage

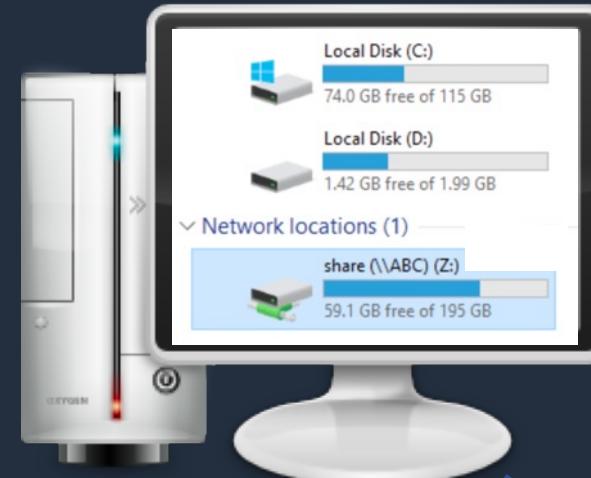
The OS sees **volumes** that can be partitioned and formatted

Block Storage



A **filesystem** can be shared by many users/computers

File Storage



Massively scalable, low cost

Object Storage



There is **no hierarchy** of objects in the container

Uses a **REST API**

The OS reads/writes at the **block level**. Disks can be internal, or network attached

A **filesystem** is “mounted” to the OS using a **network share**



AWS Storage Services

Block Storage



Amazon Elastic Block
Store

File Storage



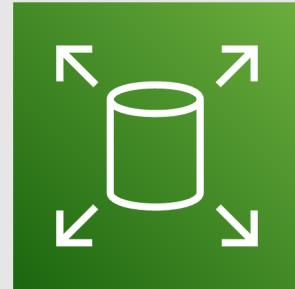
Amazon Elastic
File System

Object Storage



Amazon Simple
Storage Service (S3)

Amazon Elastic Block Store (EBS)

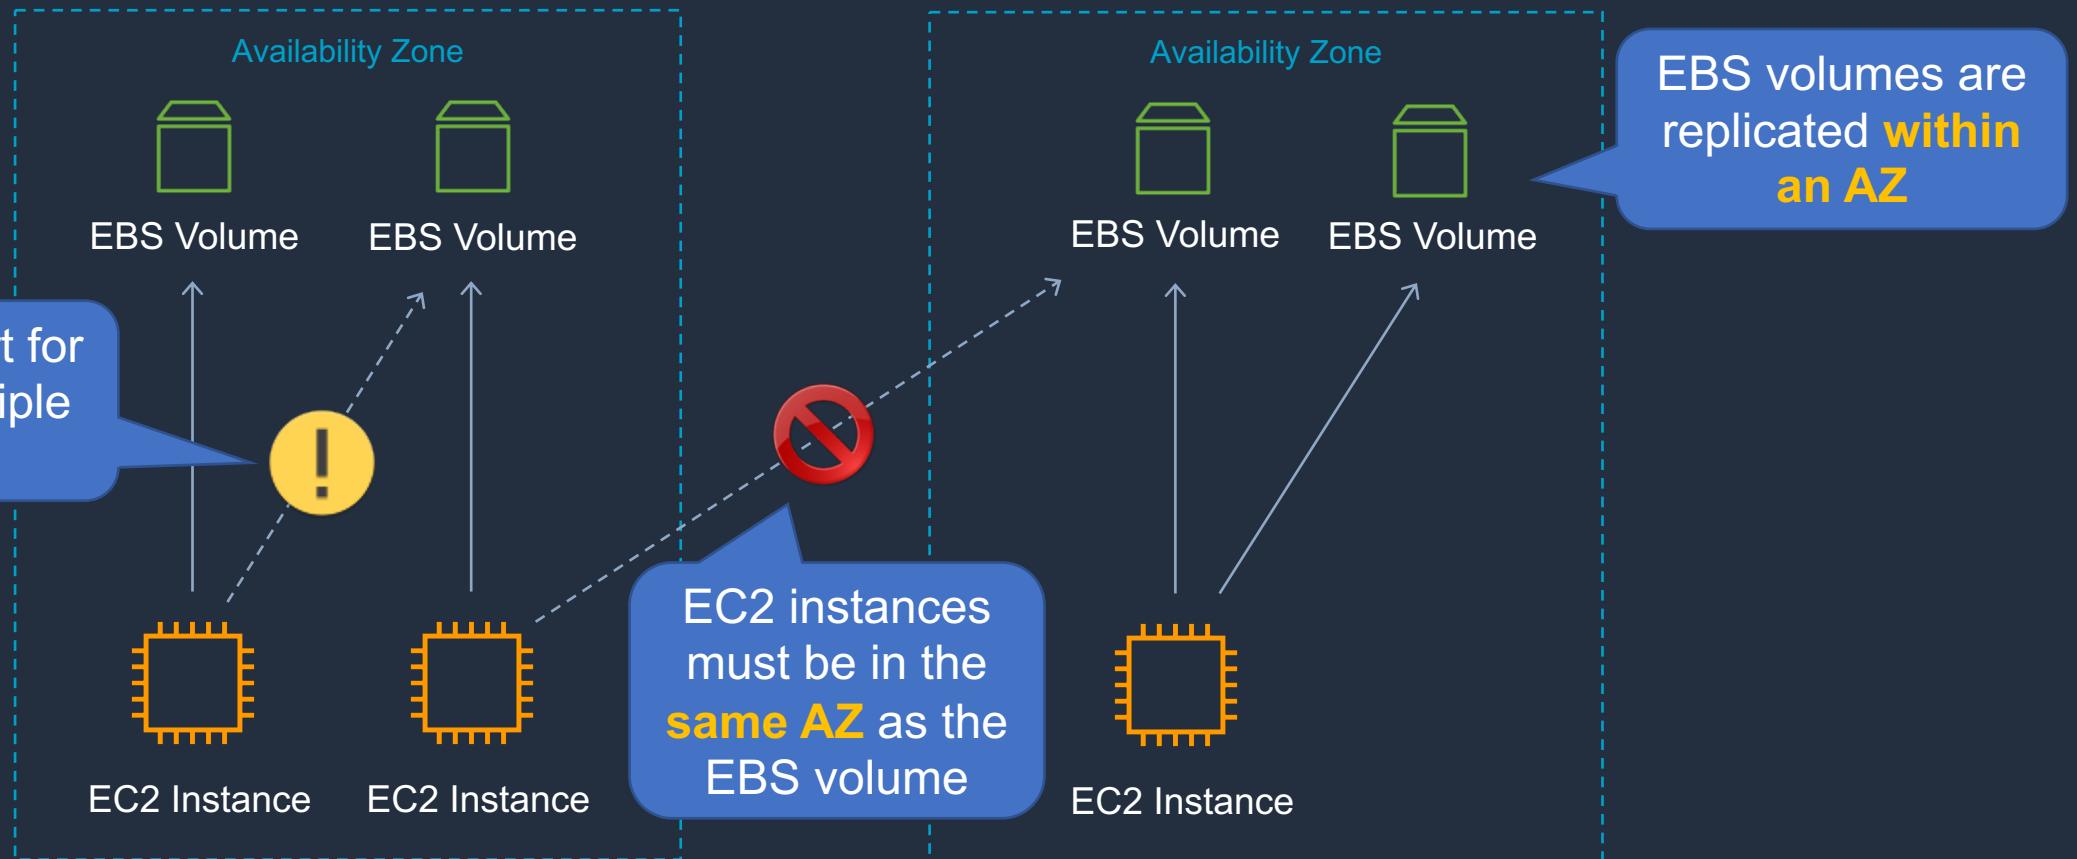




Amazon EBS



Amazon Elastic Block Store (EBS)





Amazon EBS SSD-Backed Volumes

New and **not** on
the exam yet

New and **not** on
the exam yet

	General Purpose SSD		Provisioned IOPS SSD		
Volume type	gp3	gp2	io2 Block Express ‡	io2	io1
Durability	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)	99.999% durability (0.001% annual failure rate)		99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)
Use cases	<ul style="list-style-type: none">Low-latency interactive appsDevelopment and test environments		Workloads that require sub-millisecond latency, and sustained IOPS performance or more than 64,000 IOPS or 1,000 MiB/s of throughput		<ul style="list-style-type: none">Workloads that require sustained IOPS performance or more than 16,000 IOPSI/O-intensive database workloads
Volume size	1 GiB - 16 TiB		4 GiB - 64 TiB	4 GiB - 16 TiB	
Max IOPS per volume (16 KiB I/O)	16,000		256,000	64,000 †	
Max throughput per volume	1,000 MiB/s	250 MiB/s *	4,000 MiB/s	1,000 MiB/s †	
Amazon EBS Multi-attach	Not supported		Not supported	Supported	
Boot volume	Supported				



Amazon EBS HDD-Backed Volumes

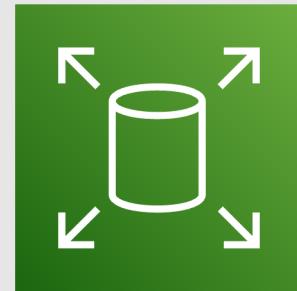
	Throughput Optimized HDD	Cold HDD
Volume type	st1	sc1
Durability	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)
Use cases	<ul style="list-style-type: none">• Big data• Data warehouses• Log processing	<ul style="list-style-type: none">• Throughput-oriented storage for data that is infrequently accessed• Scenarios where the lowest storage cost is important
Volume size	125 GiB - 16 TiB	125 GiB - 16 TiB
Max IOPS per volume (1 MiB I/O)	500	250
Max throughput per volume	500 MiB/s	250 MiB/s
Amazon EBS Multi-attach	Not supported	Not supported
Boot volume	Not supported	Not supported



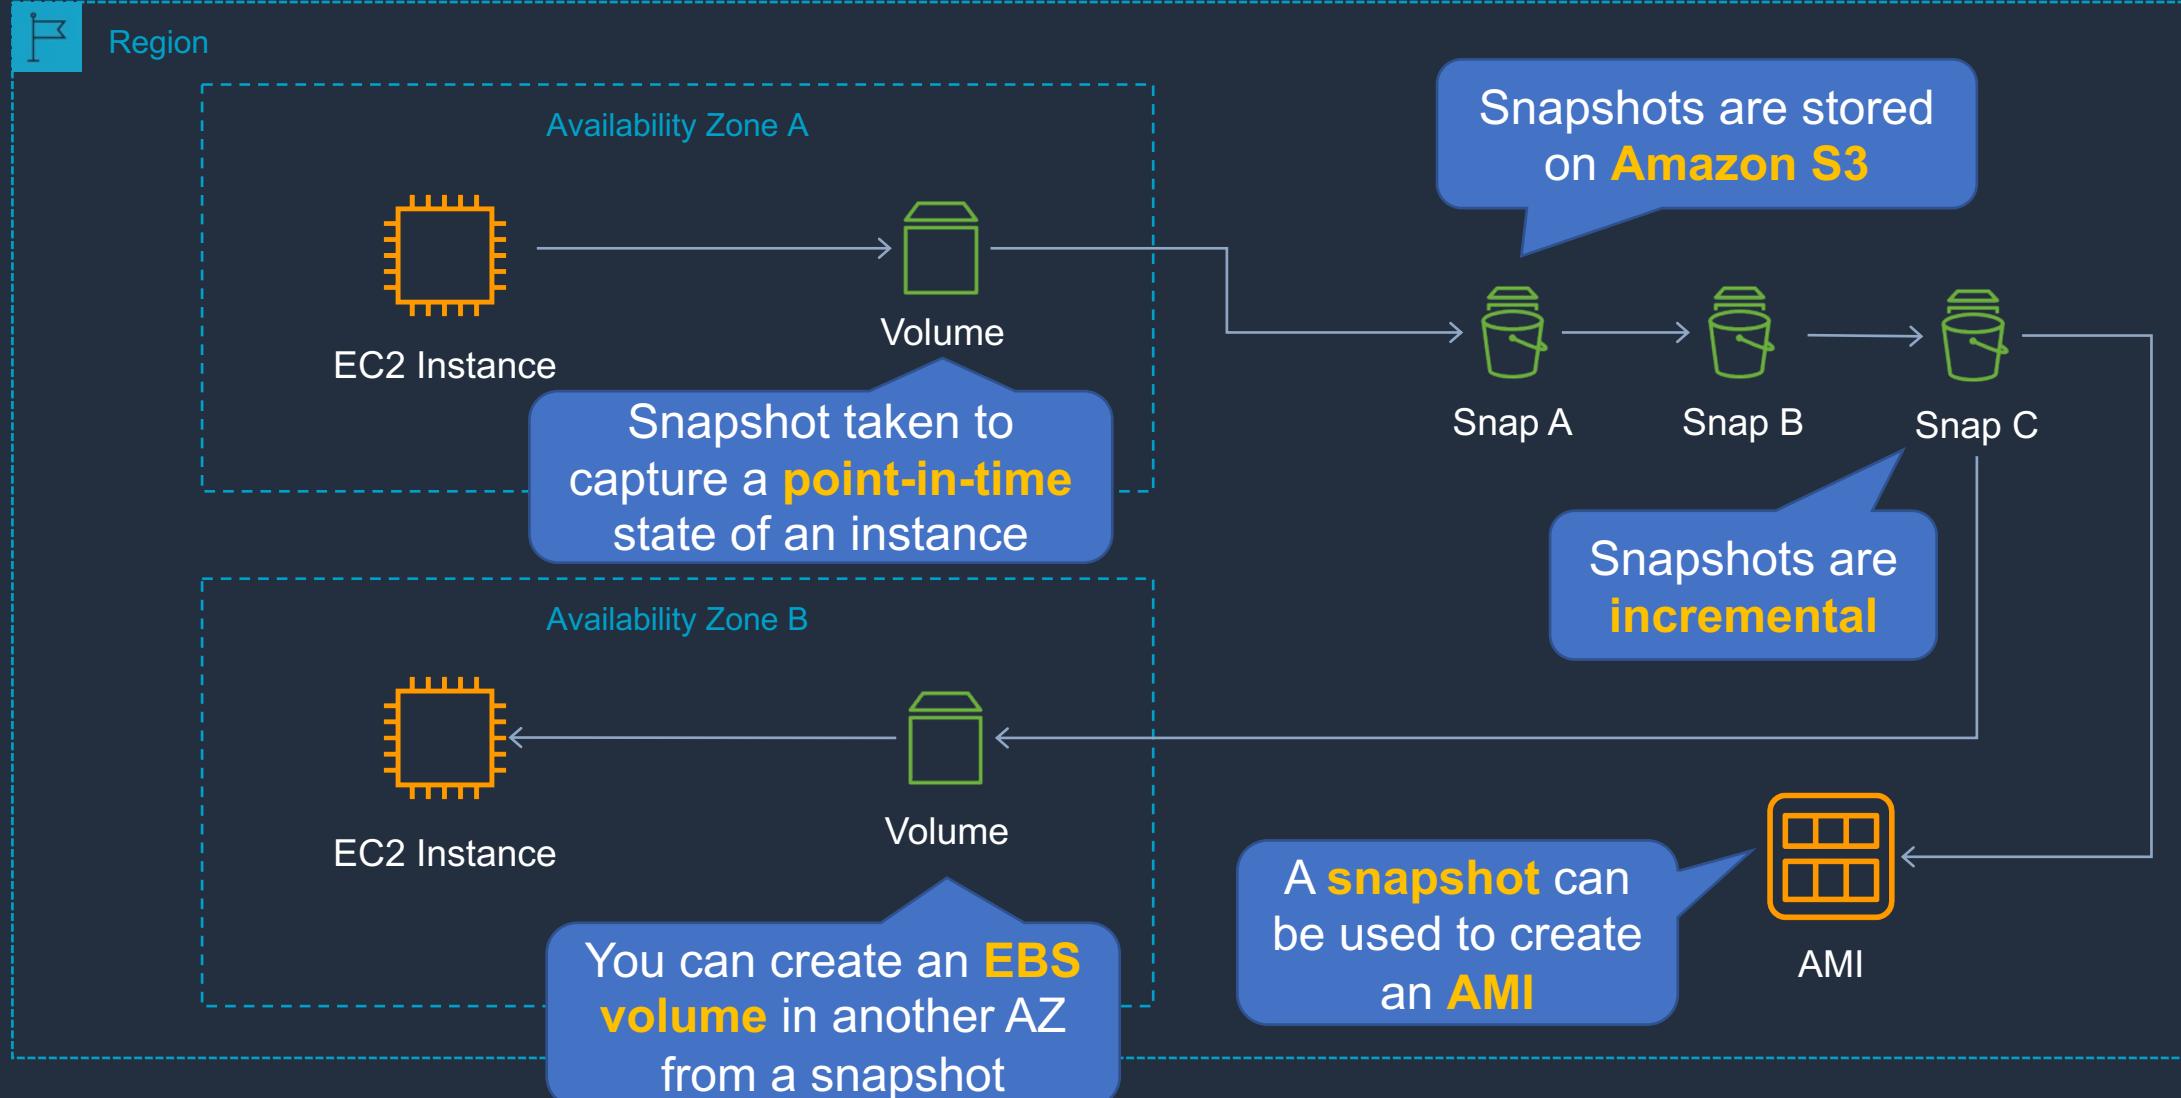
Amazon EBS

- EBS volume data persists **independently** of the life of the instance
- EBS volumes do not need to be attached to an instance
- You can attach multiple EBS volumes to an instance
- You can use multi-attach to attach a volume to multiple instances but with some constraints
- EBS volumes must be in the **same AZ** as the instances they are attached to
- Root EBS volumes **are deleted** on termination by default
- Extra non-boot volumes **are not deleted** on termination by default

Amazon EBS Snapshots and DLM



Amazon EBS Snapshots

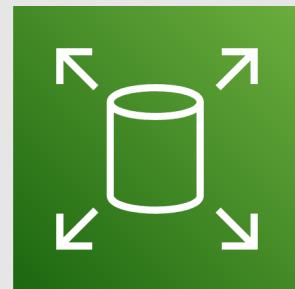




Amazon Data Lifecycle Manager (DLM)

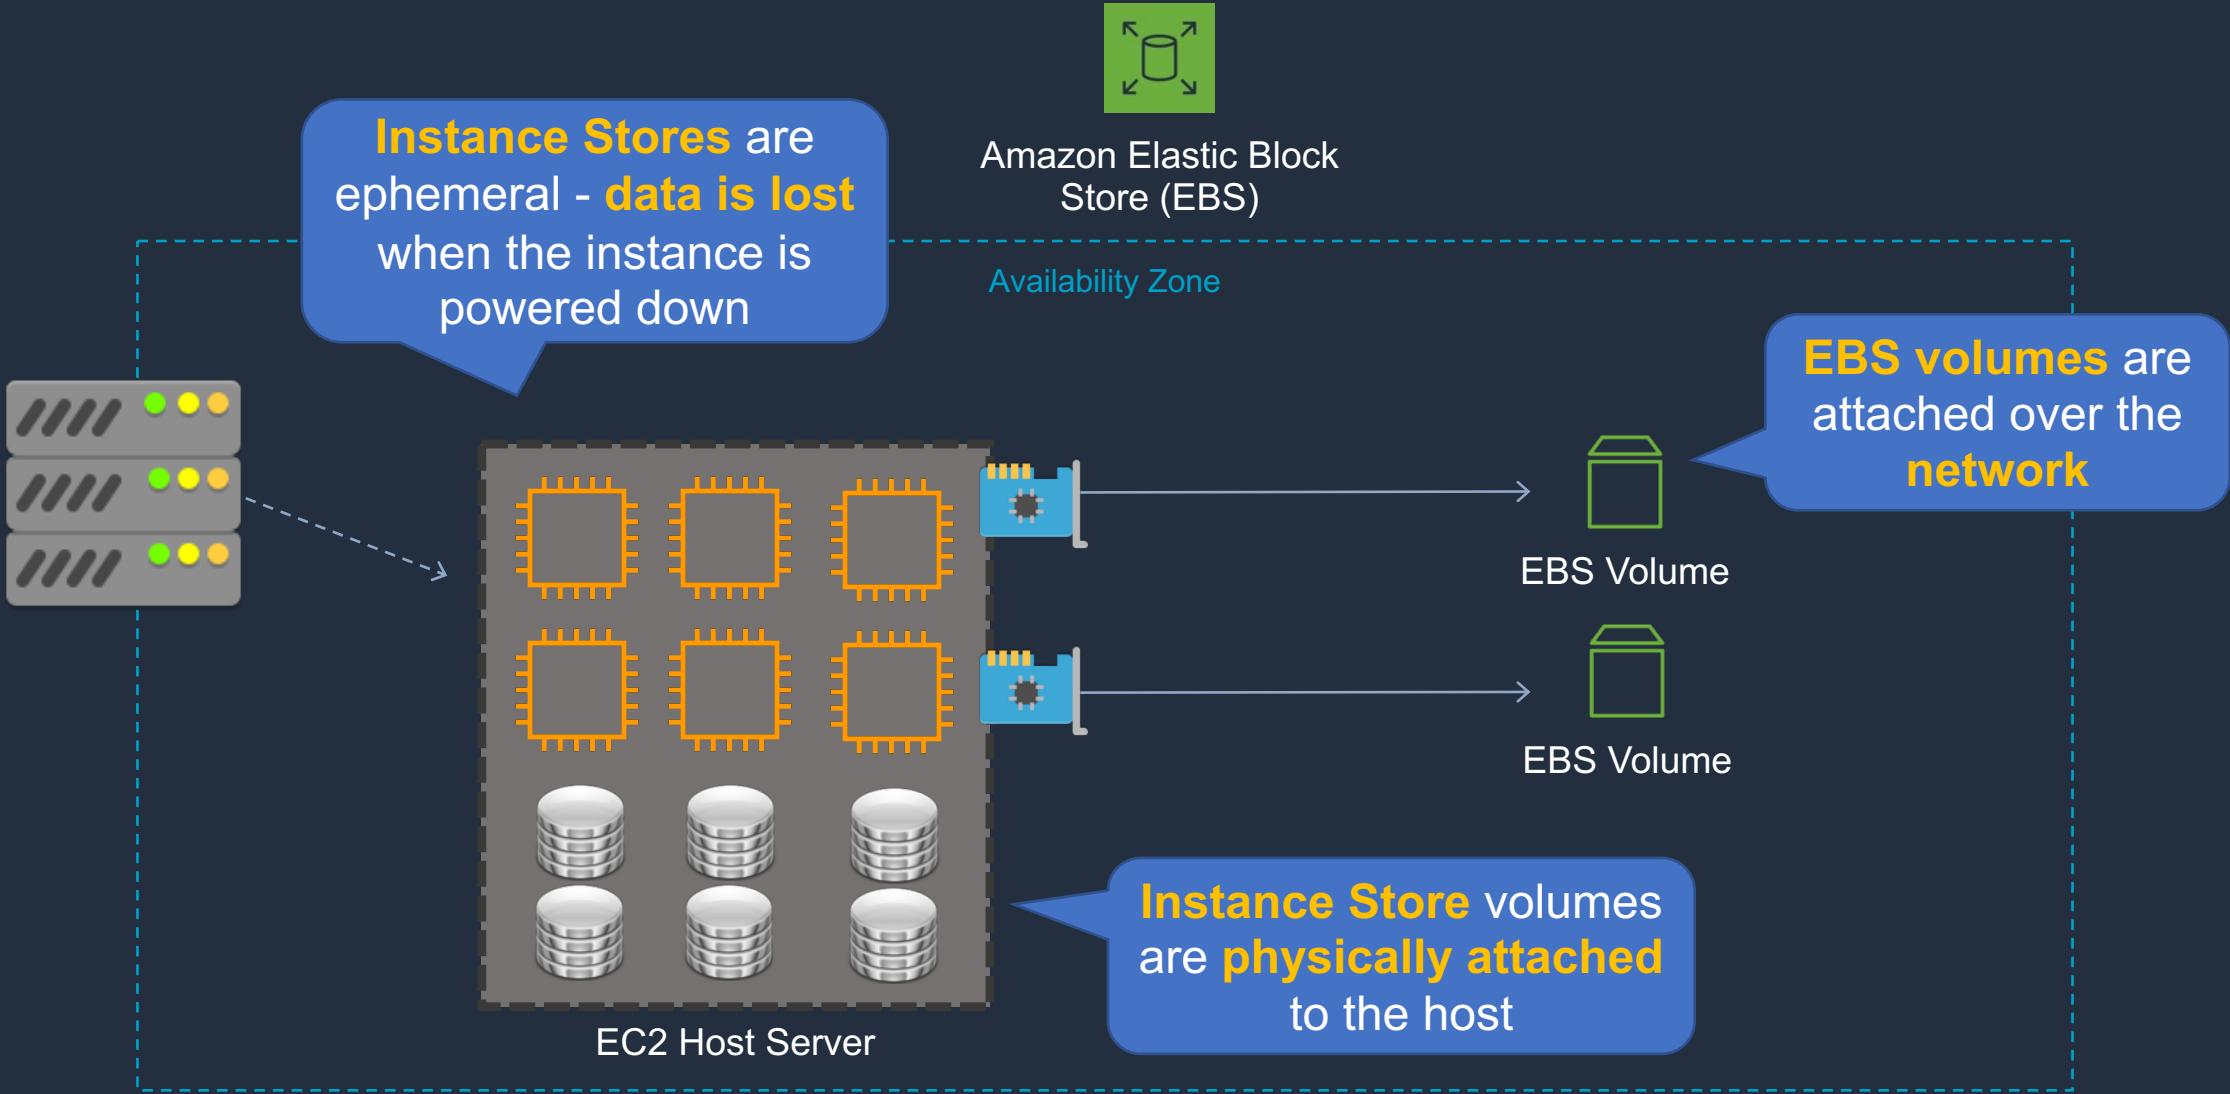
- DLM automates the creation, retention, and deletion of EBS snapshots and EBS-backed AMIs
- DLM helps with the following:
 - Protects valuable data by enforcing a regular backup schedule
 - Create standardized AMIs that can be refreshed at regular intervals
 - Retain backups as required by auditors or internal compliance
 - Reduce storage costs by deleting outdated backups
 - Create disaster recovery backup policies that back up data to isolated accounts

EC2 Instance Store Volumes





EBS vs instance store



EBS Volumes and Snapshots



Amazon Machine Images (AMI)





Amazon Machine Images (AMIs)

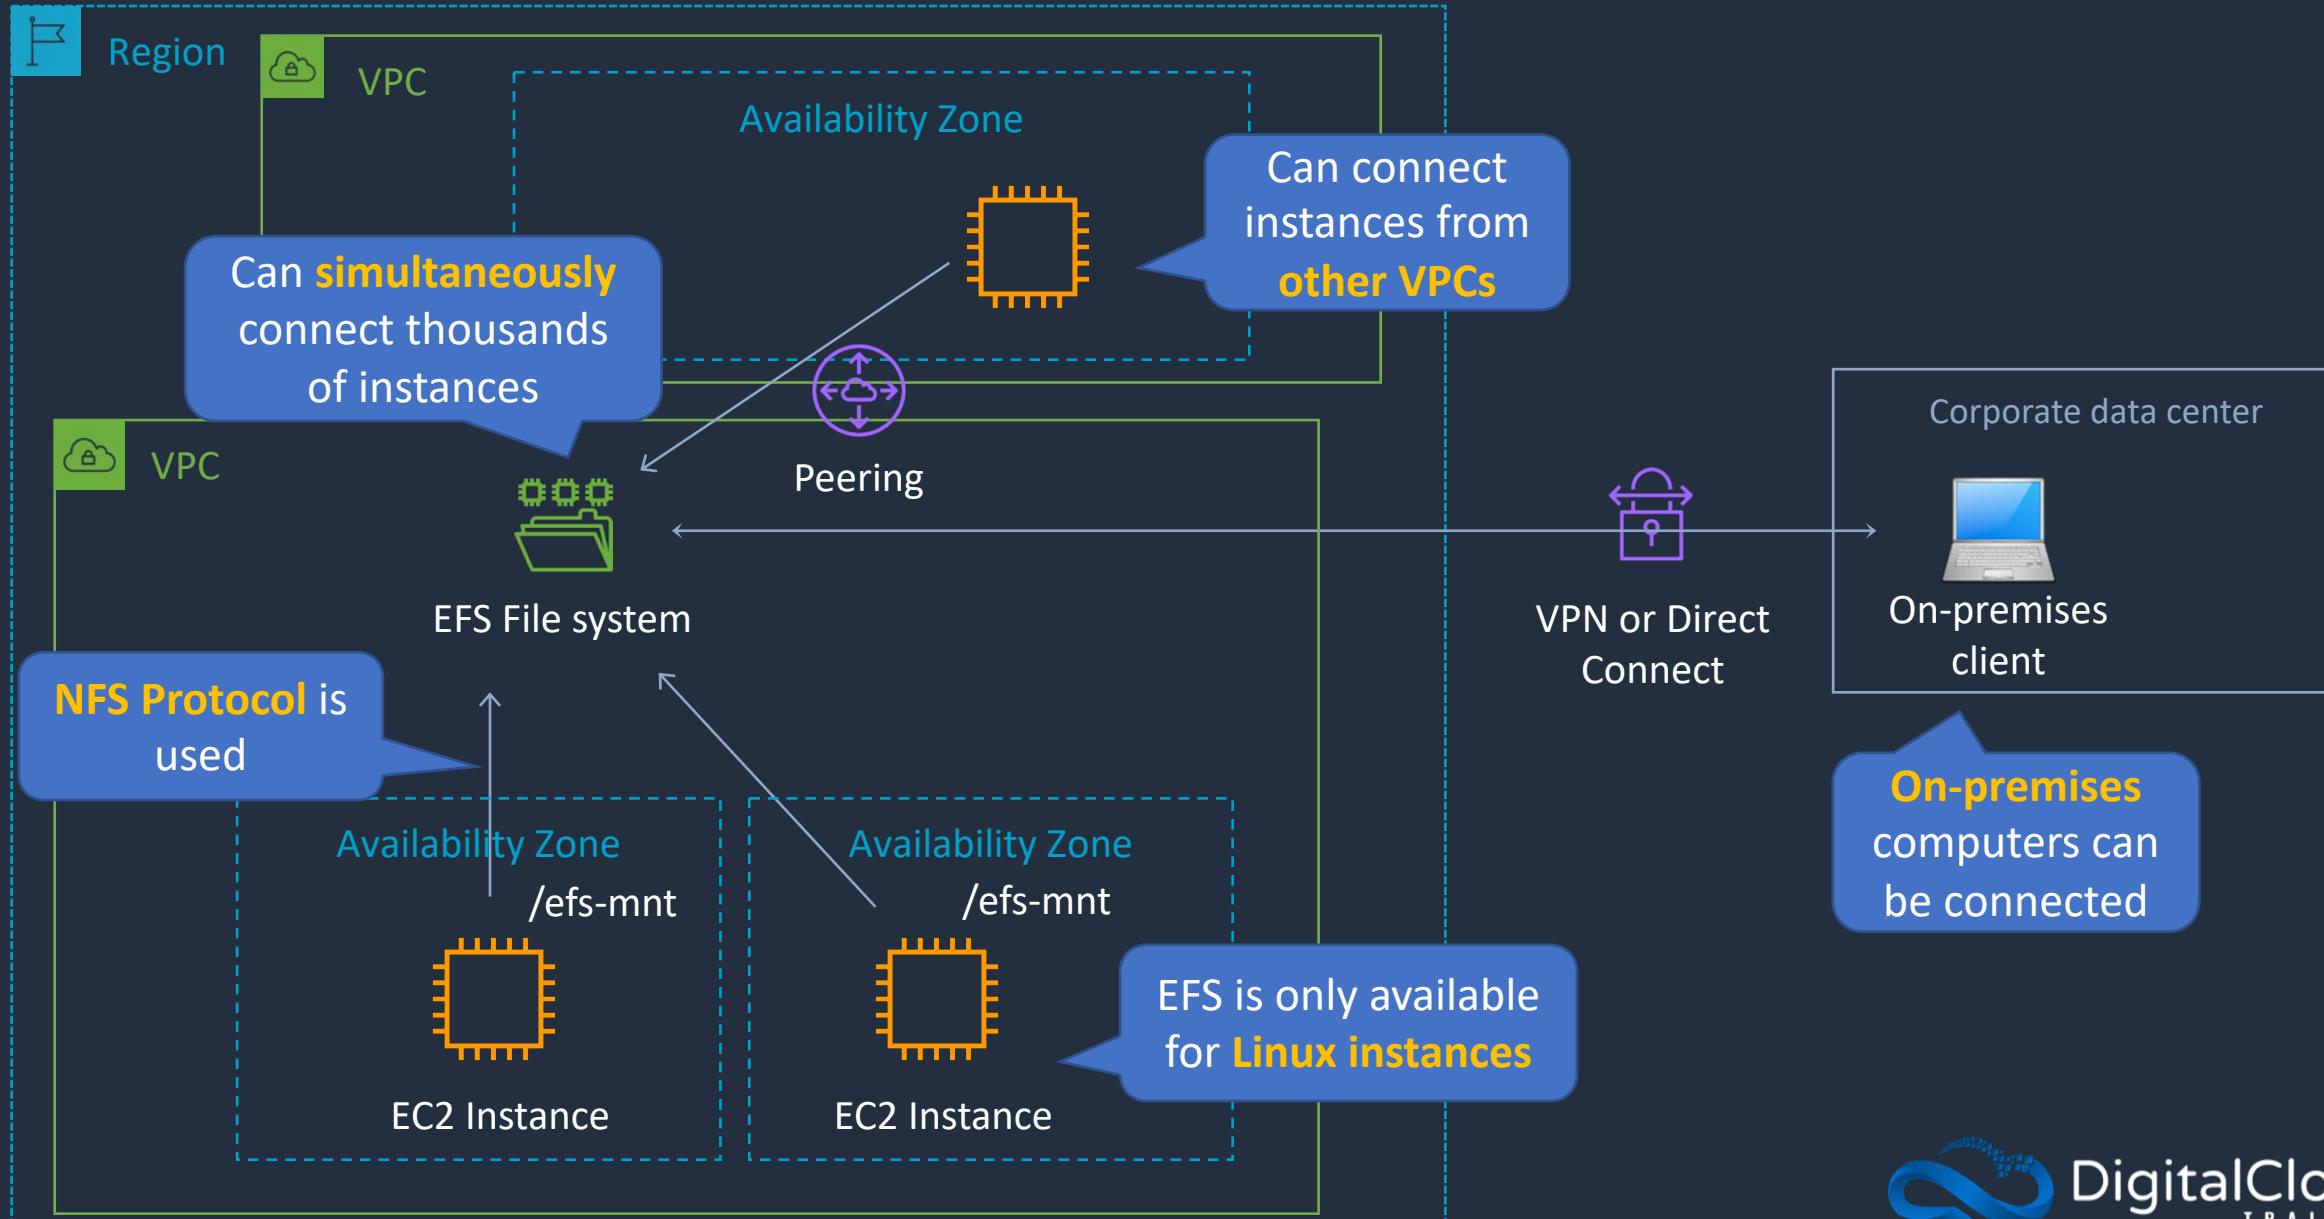
- An **Amazon Machine Image** (AMI) provides the information required to launch an instance
- An AMI includes the following:
 - One or more EBS snapshots, or, for instance-store-backed AMIs, a template for the root volume of the instance (for example, an operating system, an application server, and applications)
 - Launch permissions that control which AWS accounts can use the AMI to launch instances
 - A block device mapping that specifies the volumes to attach to the instance when it's launched
- AMIs come in three main categories:
 - **Community AMIs** - free to use, generally you just select the operating system you want
 - **AWS Marketplace AMIs** - pay to use, generally come packaged with additional, licensed software
 - **My AMIs** - AMIs that you create yourself

Amazon Elastic File System (EFS)

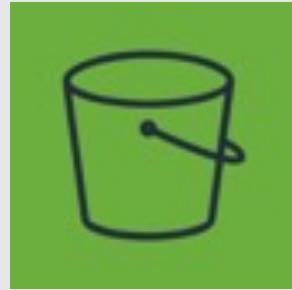




Amazon EFS



Amazon Simple Storage Service (S3)





Amazon S3



Bucket

A **bucket** is a container for objects

[http://*bucket*.s3.*aws-region*.amazonaws.com](http://bucket.s3.amazonaws.com)

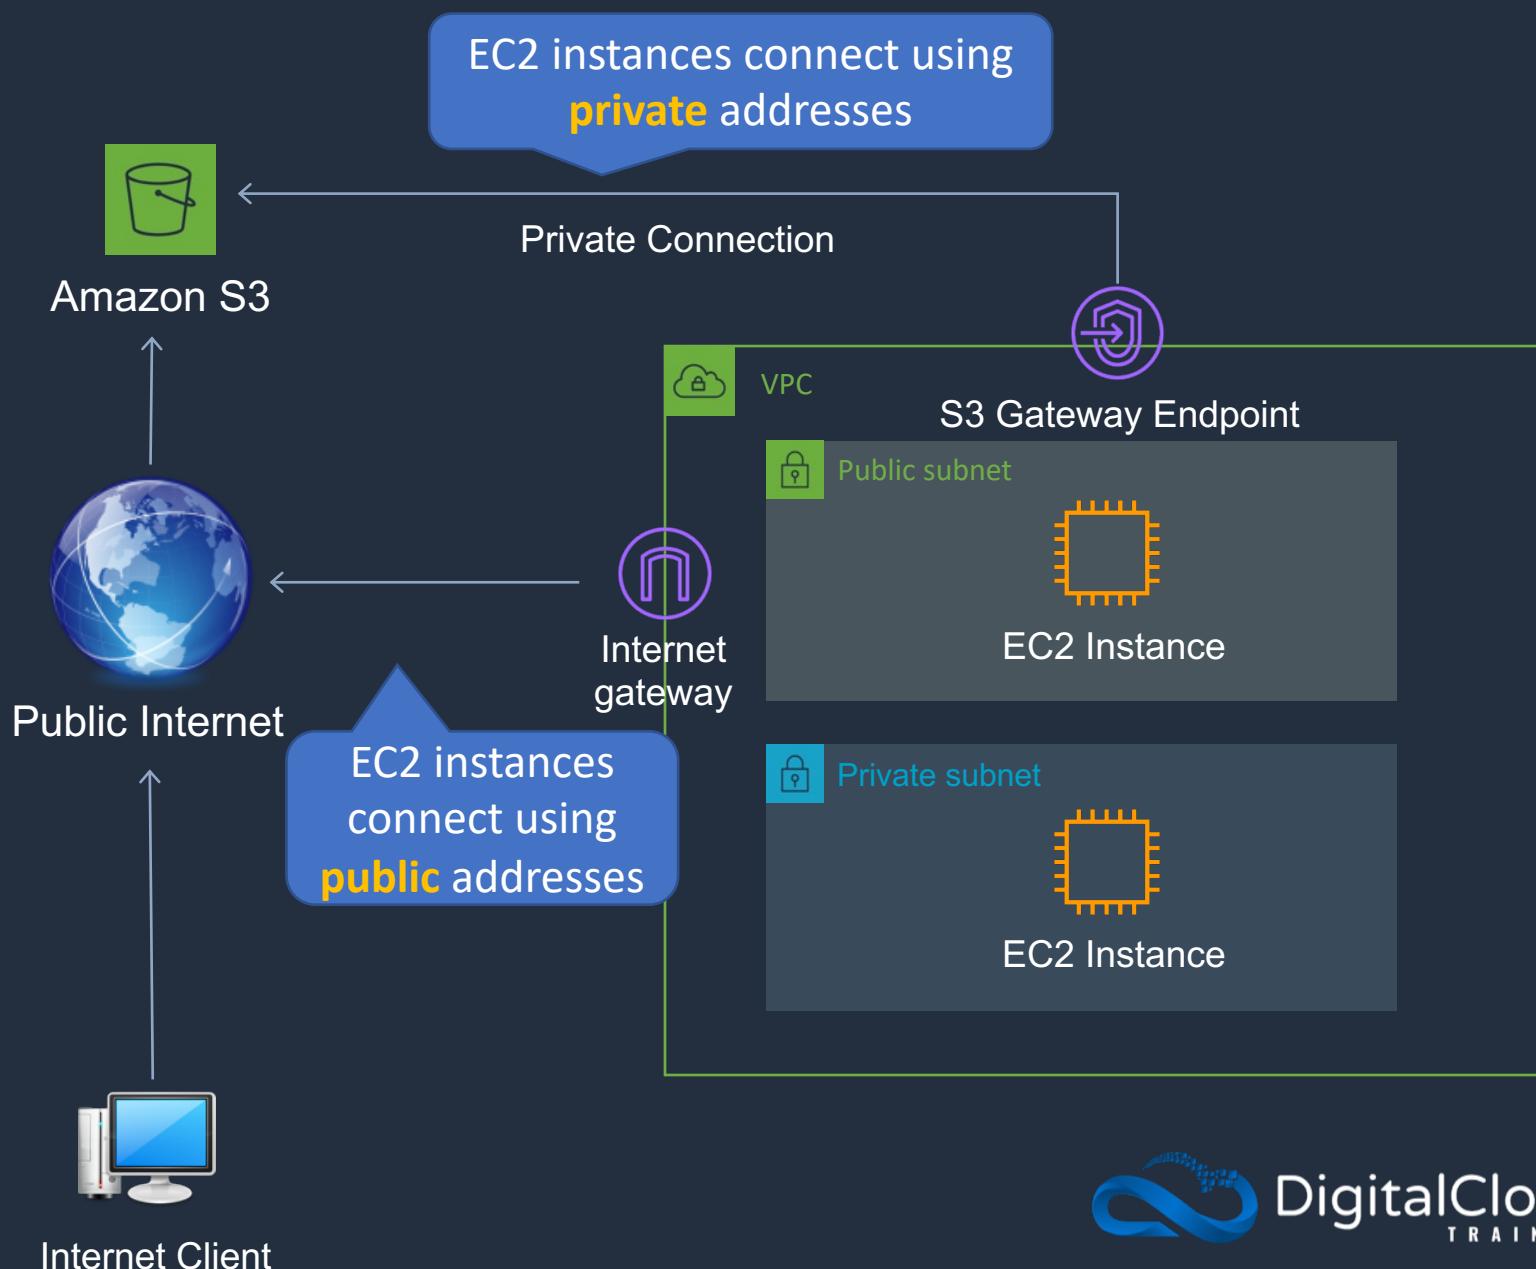
[http://s3.*aws-region*.amazonaws.com/*bucket*](http://s3.<i>aws-region</i>.amazonaws.com/<i>bucket</i>)



Object

An object consists of:

- Key (name of objects)
- Version ID
- Value (actual data)
- Metadata
- Subresources
- Access control information





Amazon S3

- You can store any type of file in S3
- Files can be anywhere from 0 bytes to 5 TB
- There is unlimited storage available
- S3 is a universal namespace so bucket names must be unique globally
- However, you create your buckets within a REGION
- It is a best practice to create buckets in regions that are physically closest to your users to reduce latency



Amazon S3 – Additional Features

S3 Capability	What it Does
Transfer Acceleration	Speed up data uploads using CloudFront in reverse
Requester Pays	The requester rather than the bucket owner pays for requests and data transfer
Events	Trigger notifications to SNS, SQS, or Lambda when certain events happen in your bucket
Static Web Hosting	Simple and massively scalable static website hosting
Versioning and Replication	Retain versions of objects and replicate objects within and across AWS Regions

Amazon S3 Storage Classes





Amazon S3 Availability and Durability

Availability

- Measures how readily available the service is
- Measured as a percentage
- S3 availability SLA varies between storage classes

Durability

- Measures the likelihood of data loss
- All storage classes offer 99.99999999% durability
- This means that if you store 100 billion objects in S3, you will lose one object at most



Amazon S3 Storage Classes

	S3 Standard	S3 Intelligent Tiering	S3 Standard-IA	S3 One Zone-IA	S3 Glacier	S3 Glacier Deep Archive
Designed for durability	99.99999999%	99.99999999%	99.99999999%	99.99999999%	99.99999999%	99.99999999%
Designed for availability	99.99%	99.9%	99.9%	99.5%	99.99%	99.99%
Availability SLA	99.9%	99%	99%	99%	99.9%	99.9%
Availability Zones	≥3	≥3	≥3	1	≥3	≥3
Minimum capacity charge per object	N/A	N/A	128KB	128KB	40KB	40KB
Minimum storage duration charge	N/A	30 days	30 days	30 days	90 days	180 days
Retrieval fee	N/A	N/A	Per GB retrieved	Per GB retrieved	Per GB retrieved	Per GB retrieved
First byte latency	milliseconds	milliseconds	milliseconds	milliseconds	select minutes or hours	select hours
Storage type	Object	Object	Object	Object	Object	Object
Lifecycle transitions	Yes	Yes	Yes	Yes	Yes	Yes

Create Amazon S3 Bucket



S3 Versioning, Replication and Lifecycle Rules





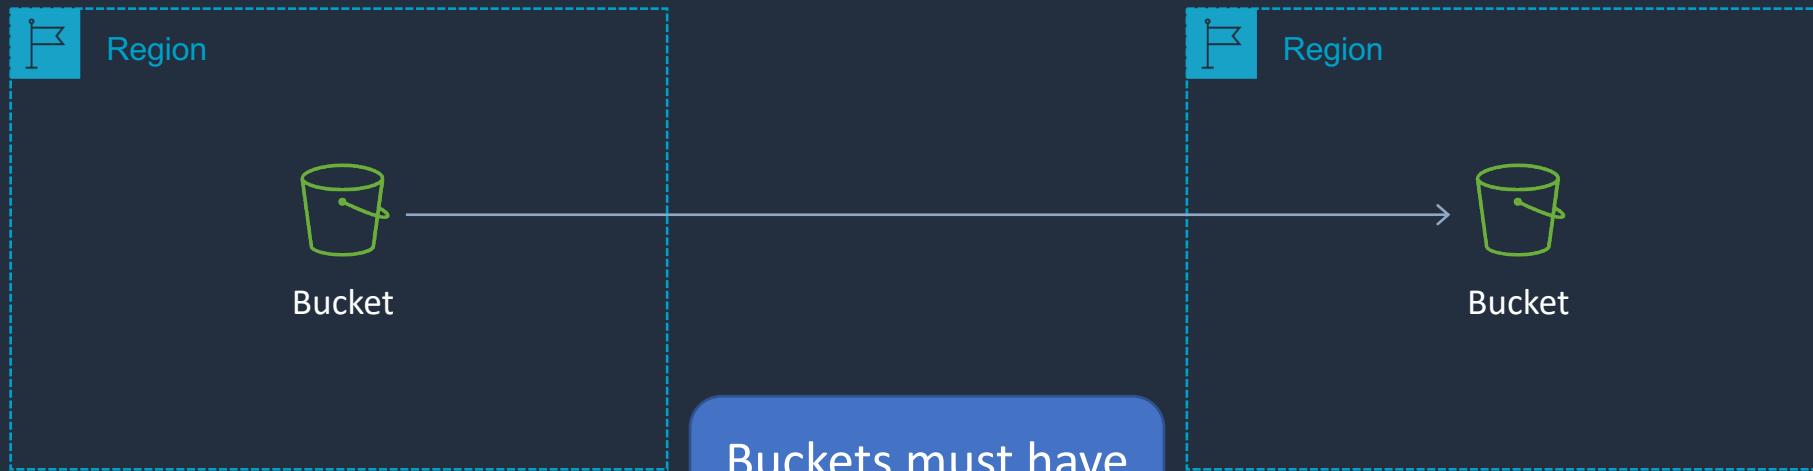
Amazon S3 Versioning

- Versioning is a means of keeping **multiple variants** of an **object** in the same bucket
- Use versioning to preserve, retrieve, and restore every version of every object stored in your Amazon S3 bucket
- Versioning-enabled buckets enable you to recover objects from accidental deletion or overwrite



Amazon S3 Replication

Cross-Region Replication (CRR)



Buckets must have
versioning enabled



Configure Replication and Lifecycle



Configure S3 Static Website



S3 Permissions and Bucket Policies



Archiving with S3 Glacier





Amazon S3 Glacier

- Extremely low cost and you pay only for what you need with no commitments of upfront fees
- Two classes **Glacier** and **Glacier Deep Archive**
- Three options for access to archives, listed in the table below:

	Expedited	Standard	Bulk
Data access time (Glacier)	1-5 minutes	3-5 hours	5-12 hours
Data access time (Deep Archive)	N/A	12 hours	48 hours



Object Lock and Glacier Vault Lock

S3 Object Lock

- Store objects using a write-once-read-many (WORM) model
- Prevent objects from being deleted or overwritten for a fixed time or indefinitely

S3 Glacier Vault Lock

- Also used to enforce a WORM model
- Can apply a policy and lock the policy from future edits
- Use for compliance objectives and data retention

AWS Storage Gateway



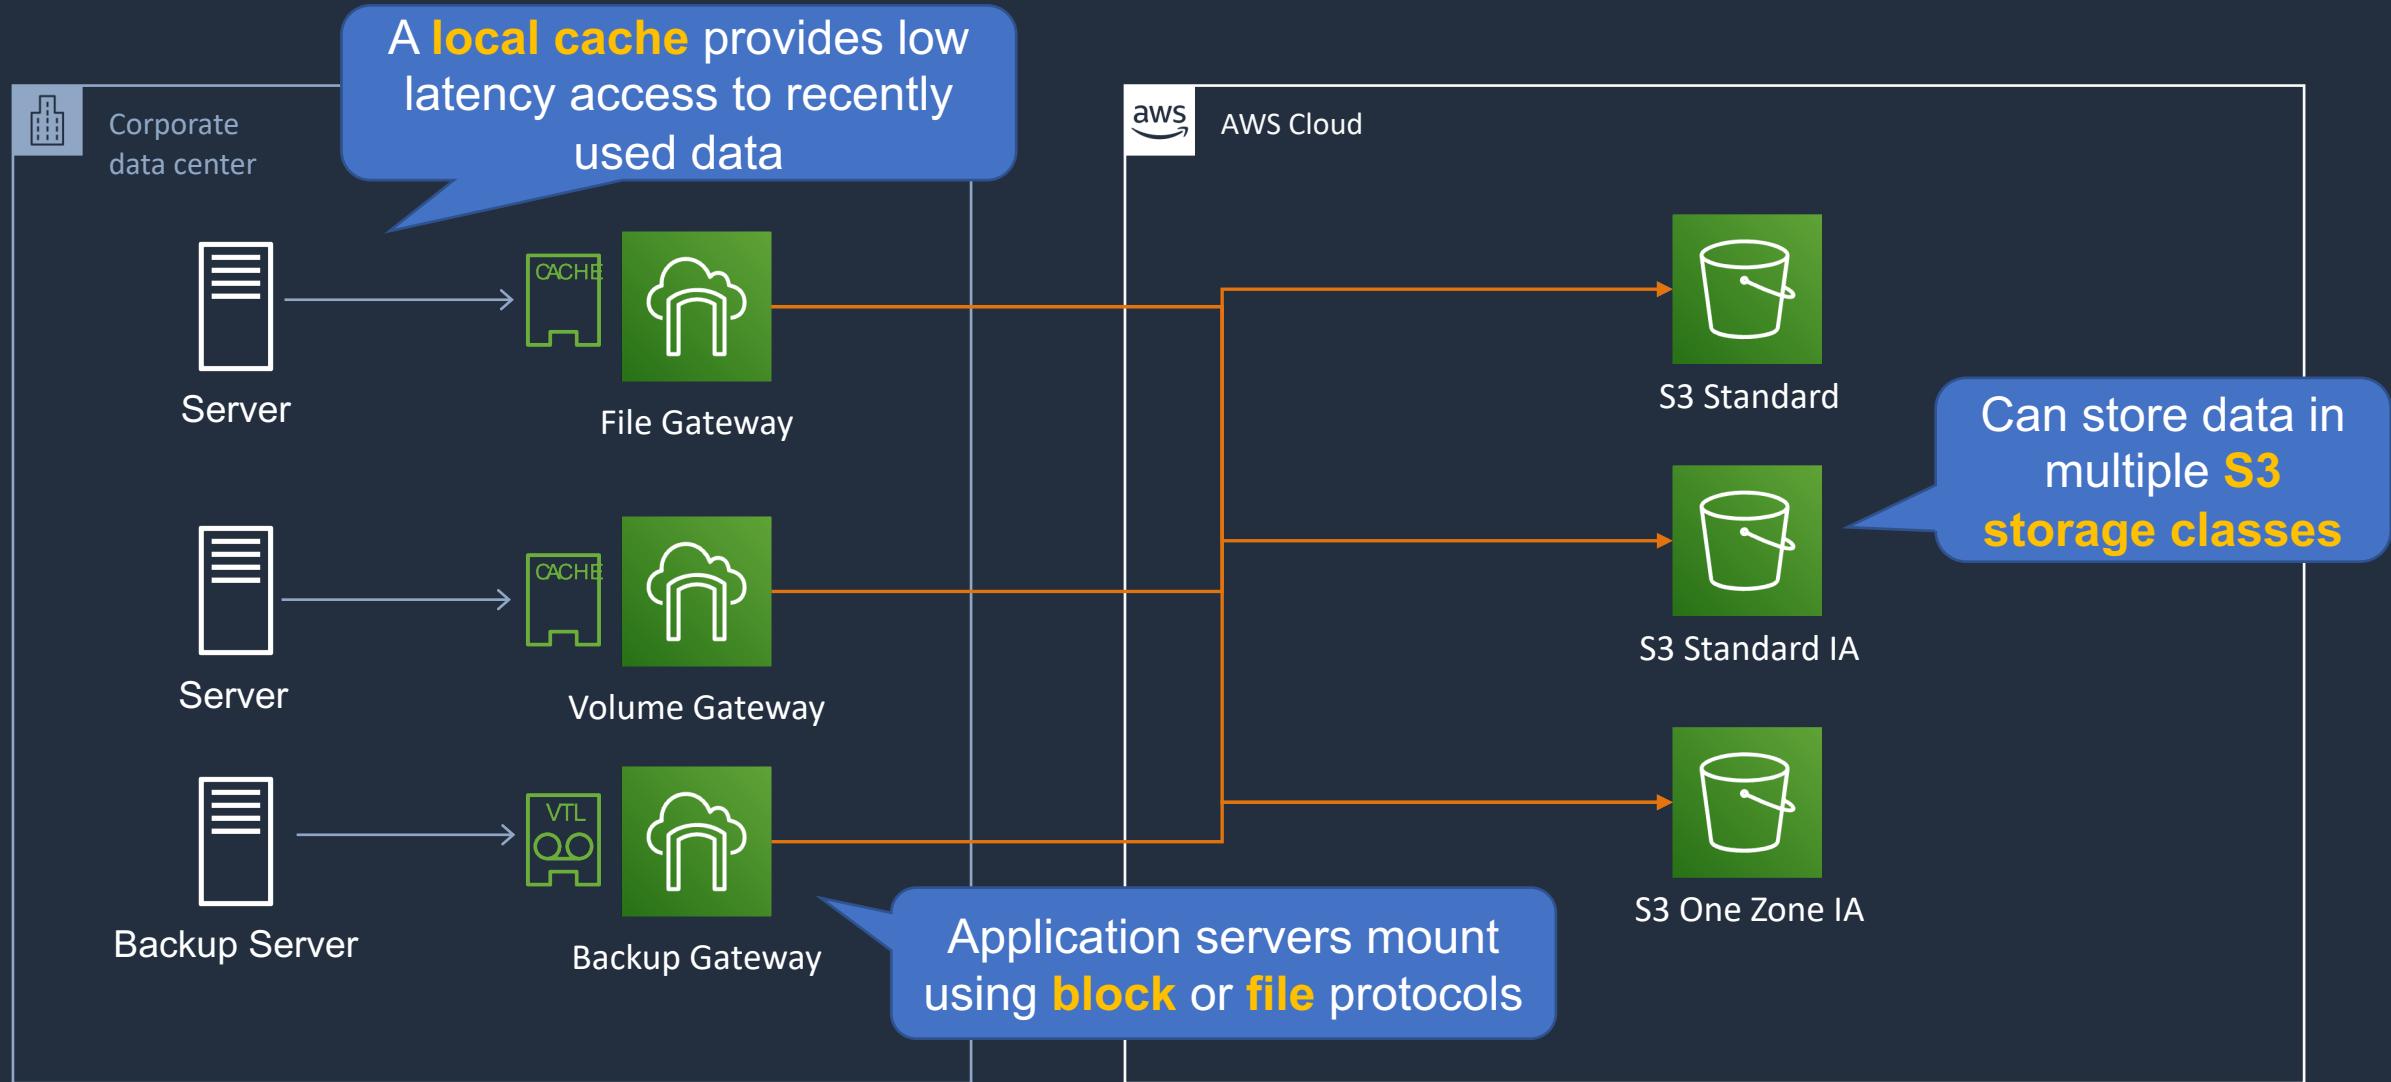


AWS Storage Gateway

- Hybrid cloud storage service
- Access cloud storage from on-premises applications
- Enables access to proprietary object storage (S3) using standard protocols
- Use cases:
 - Moving backups to the cloud
 - Using on-premises file shares backed by cloud storage
 - Low latency access to data in AWS for on-premises applications
 - Disaster recovery



AWS Storage Gateway



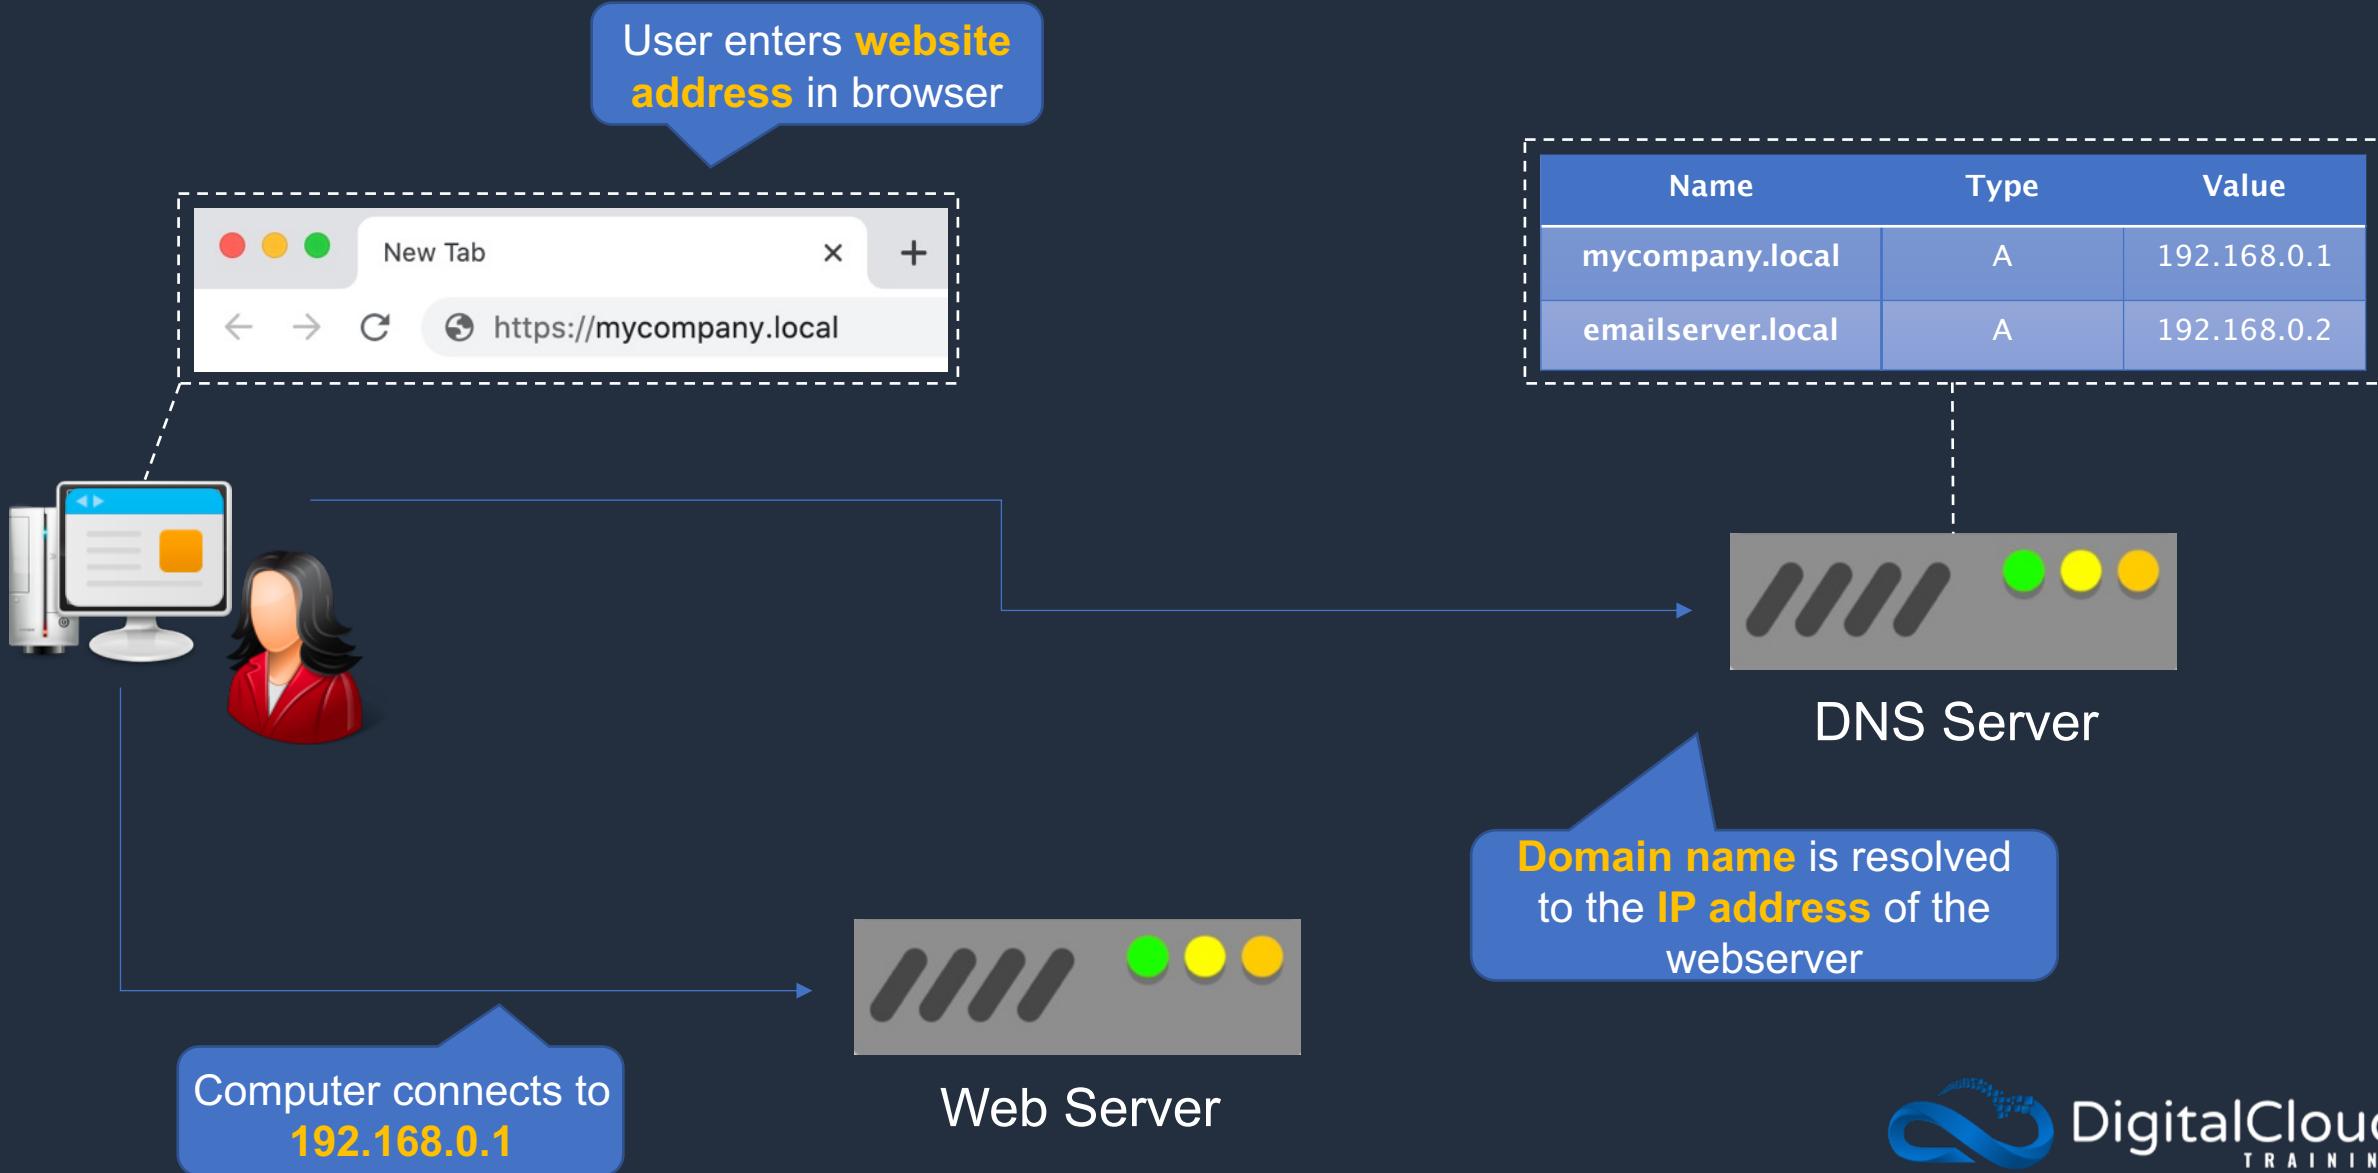
SECTION 7

DNS, Elastic Load Balancing, and Auto Scaling

DNS and Amazon Route 53

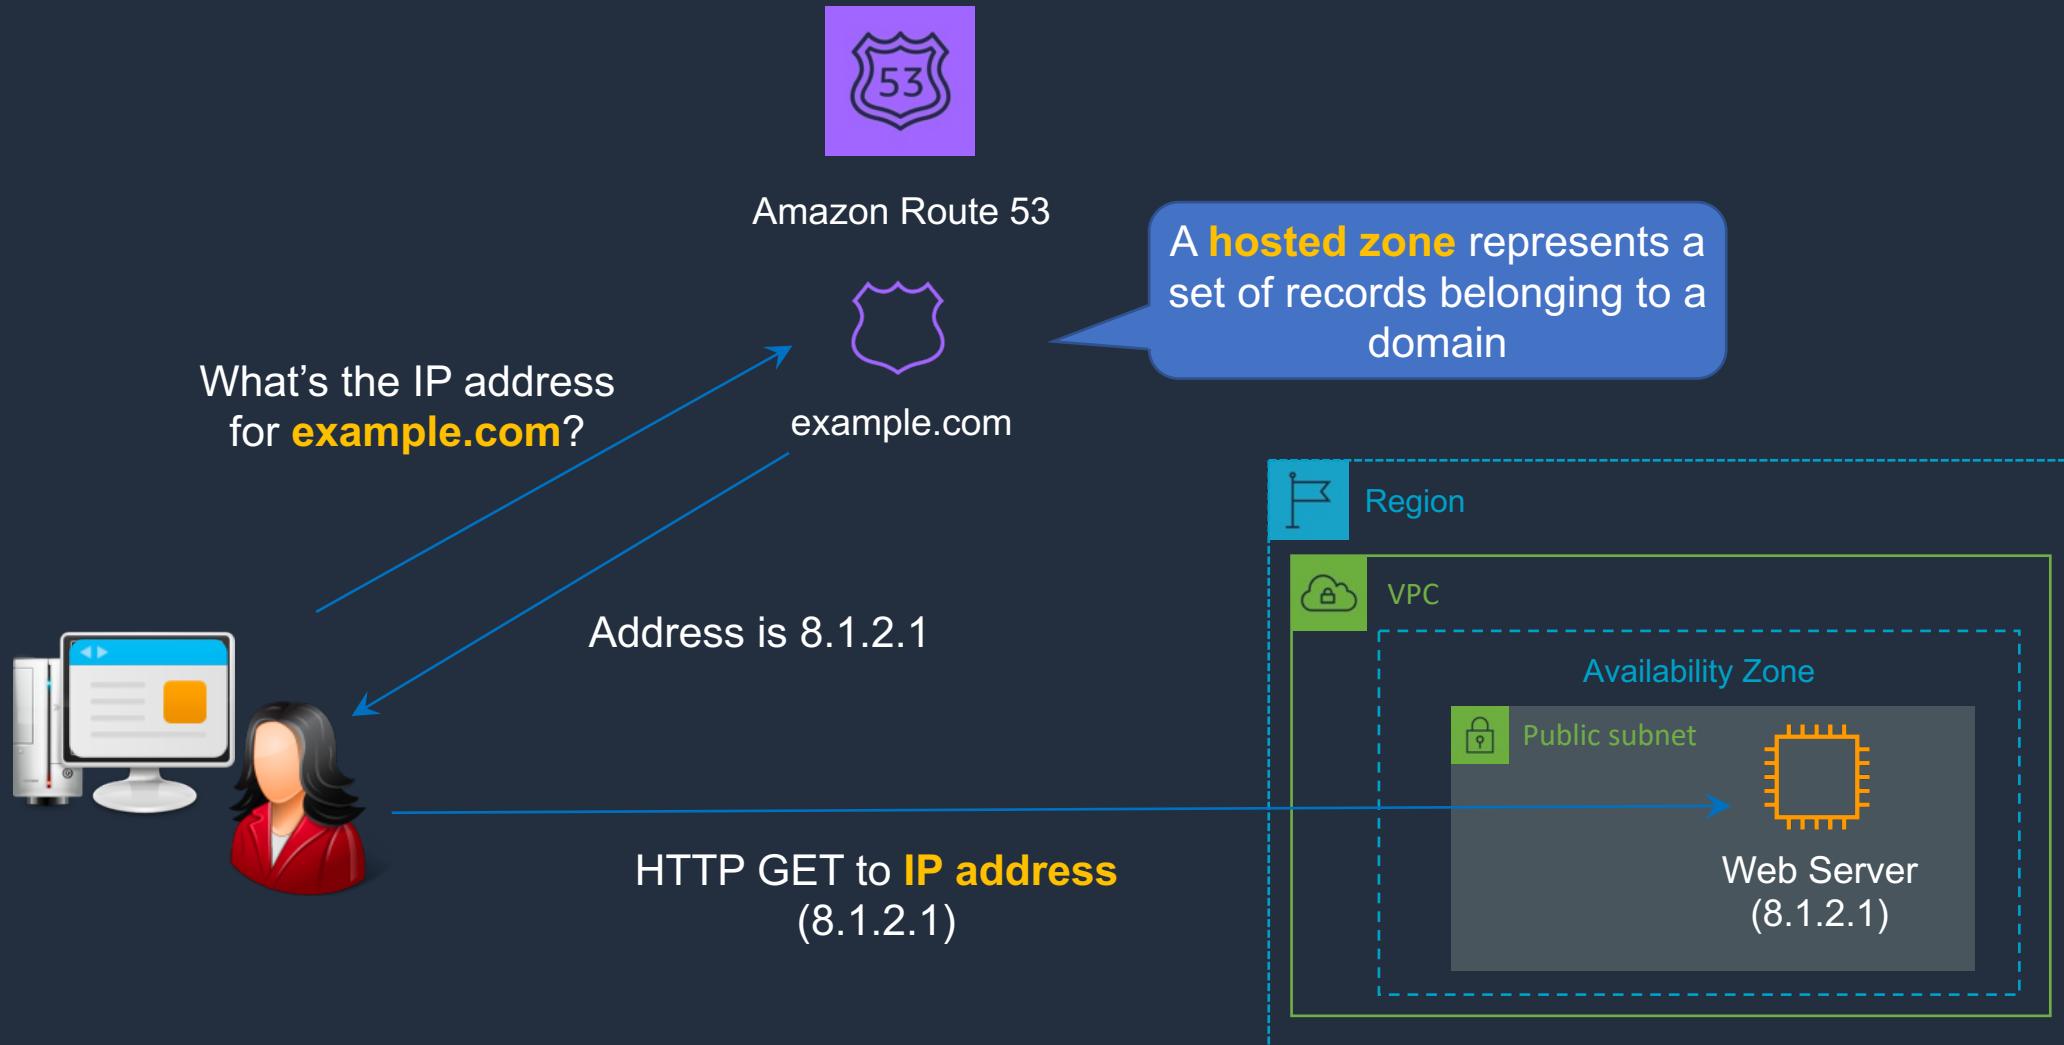


The Domain Name System (DNS)





Amazon Route 53



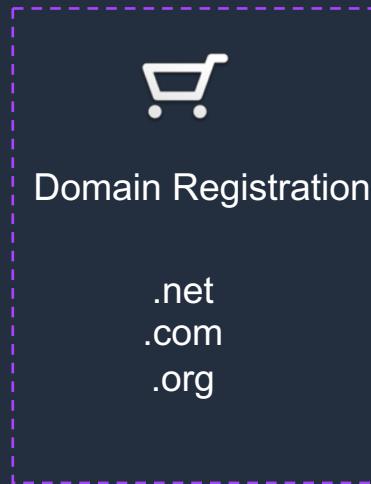
Amazon Route 53 Routing Policies

Routing Policy	What it does
Simple	Simple DNS response providing the IP address associated with a name
Failover	If primary is down (based on health checks), routes to secondary destination
Geolocation	Uses geographic location you're in (e.g. Europe) to route you to the closest region
Geoproximity	Routes you to the closest region within a geographic area
Latency	Directs you based on the lowest latency route to resources
Multivalue answer	Returns several IP addresses and functions as a basic load balancer
Weighted	Uses the relative weights assigned to resources to determine which to route to

Amazon Route Features



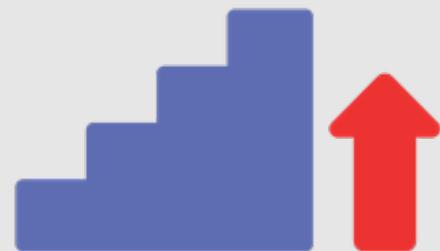
Amazon Route 53



Register Domain with Route 53 (Optional)



Elasticity: Scaling Up vs Out





Scaling Up (vertical scaling)





Scaling Up (vertical scaling)

Scaling up means
adding resources
to the instance



Limitation is that you
have a **single point of
failure** (SPOF)



Scaling Out (horizontal scaling)

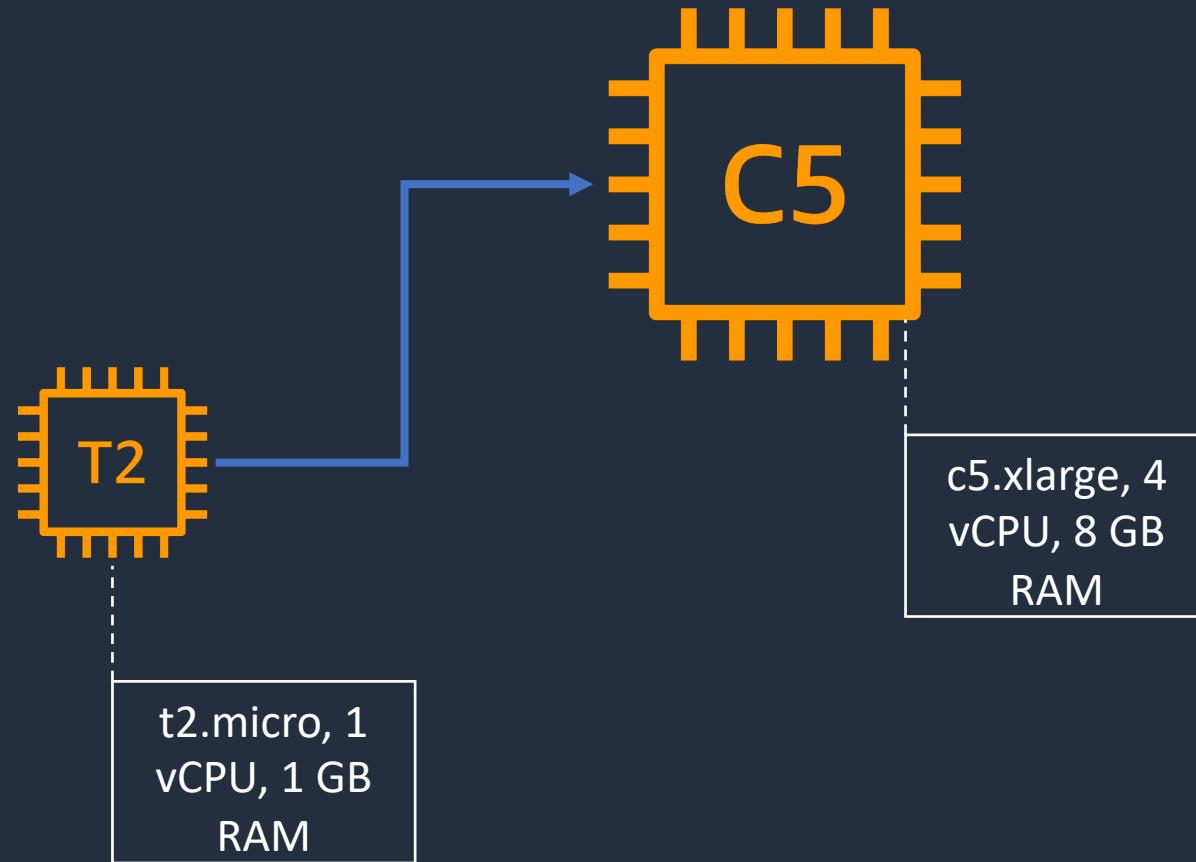
Scaling out provides greater **resiliency**



Scaling out can be used to add almost unlimited capacity

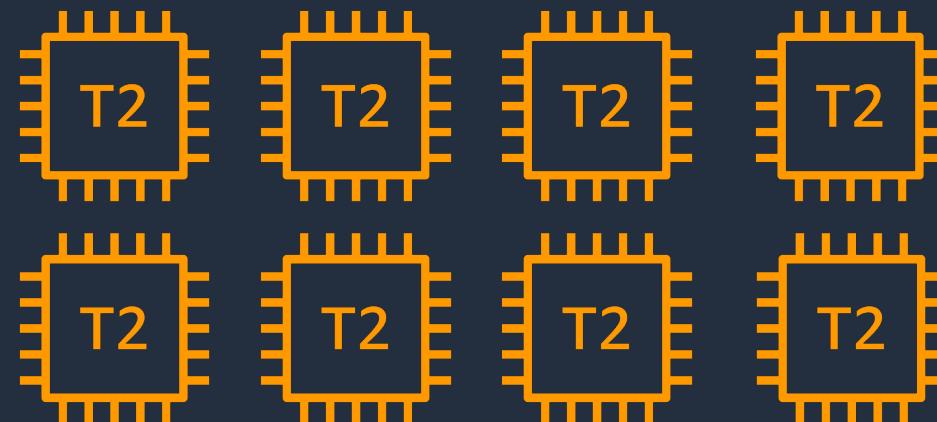


Scaling Up (vertical scaling)

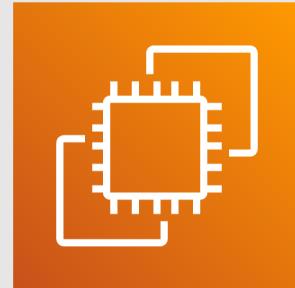




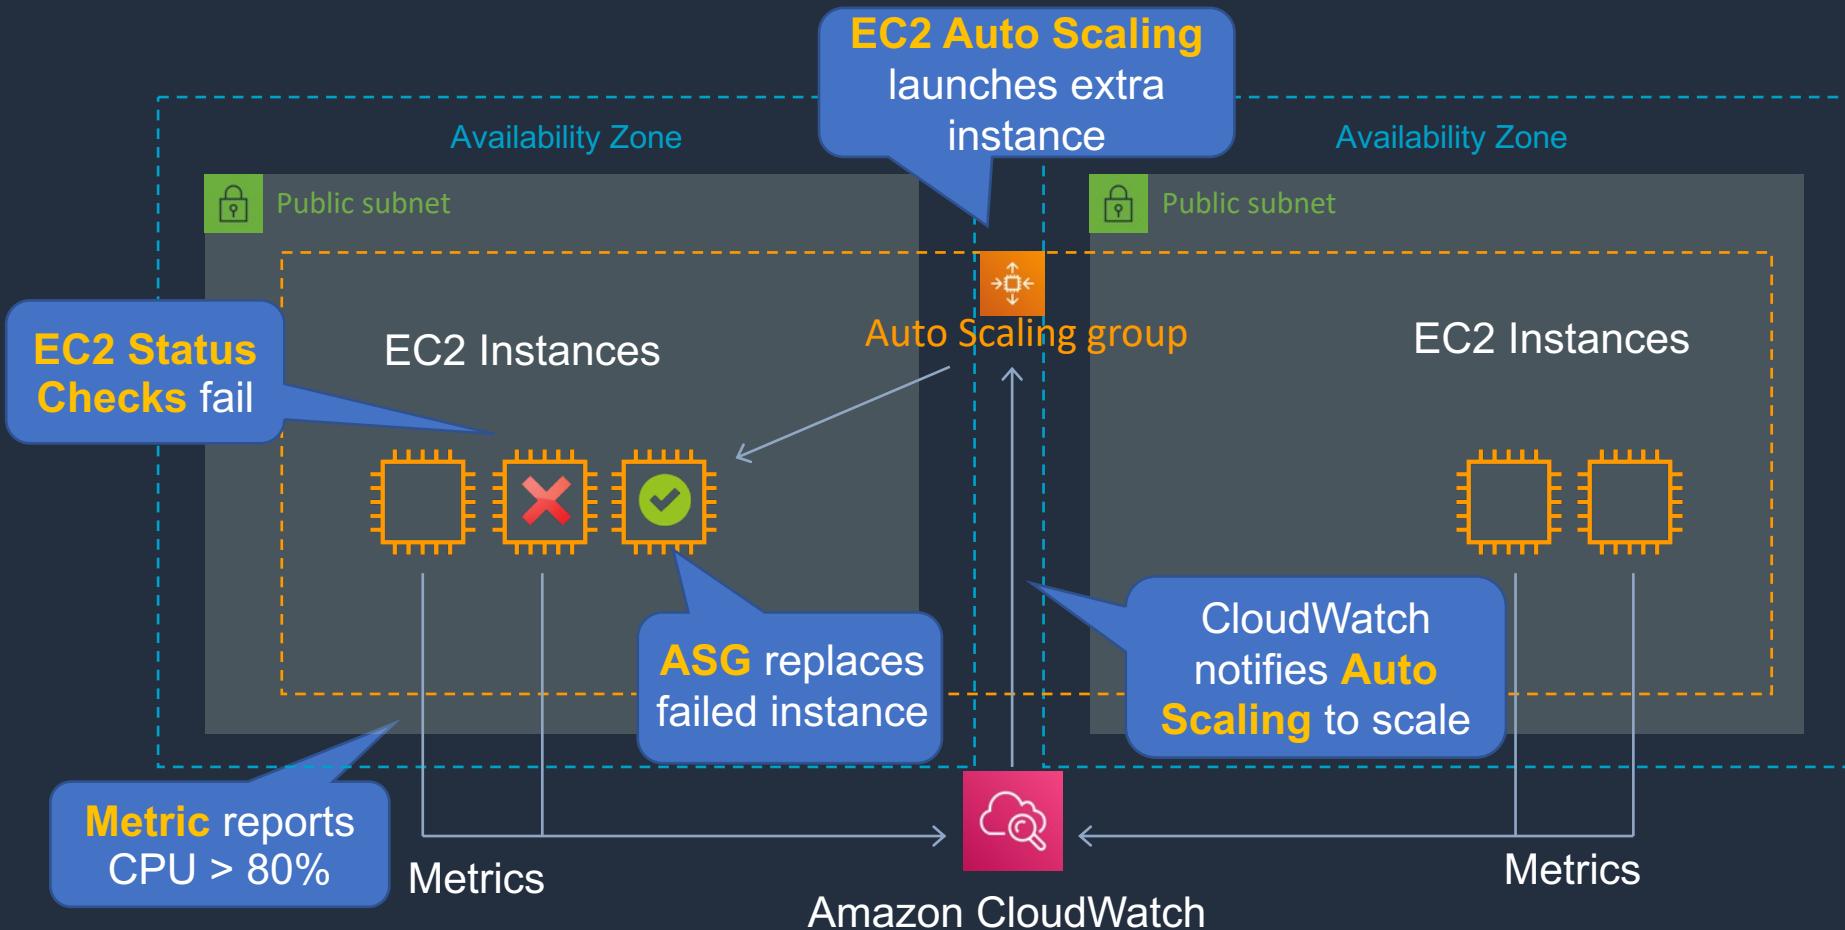
Scaling Out (horizontal scaling)



Amazon EC2 Auto Scaling



Amazon EC2 Auto Scaling





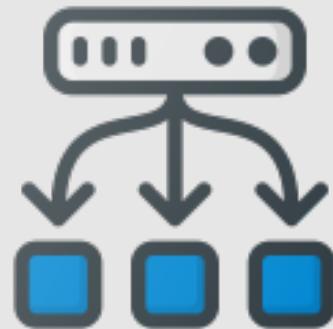
Amazon EC2 Auto Scaling

- EC2 Auto Scaling **launches** and **terminates** instances dynamically
- Scaling is horizontal (scales out)
- Provides **elasticity** and **scalability**
- Responds to EC2 status checks and CloudWatch metrics
- Can scale based on demand (performance) or on a schedule
- Scaling policies define how to respond to changes in demand

Create an Auto Scaling Group

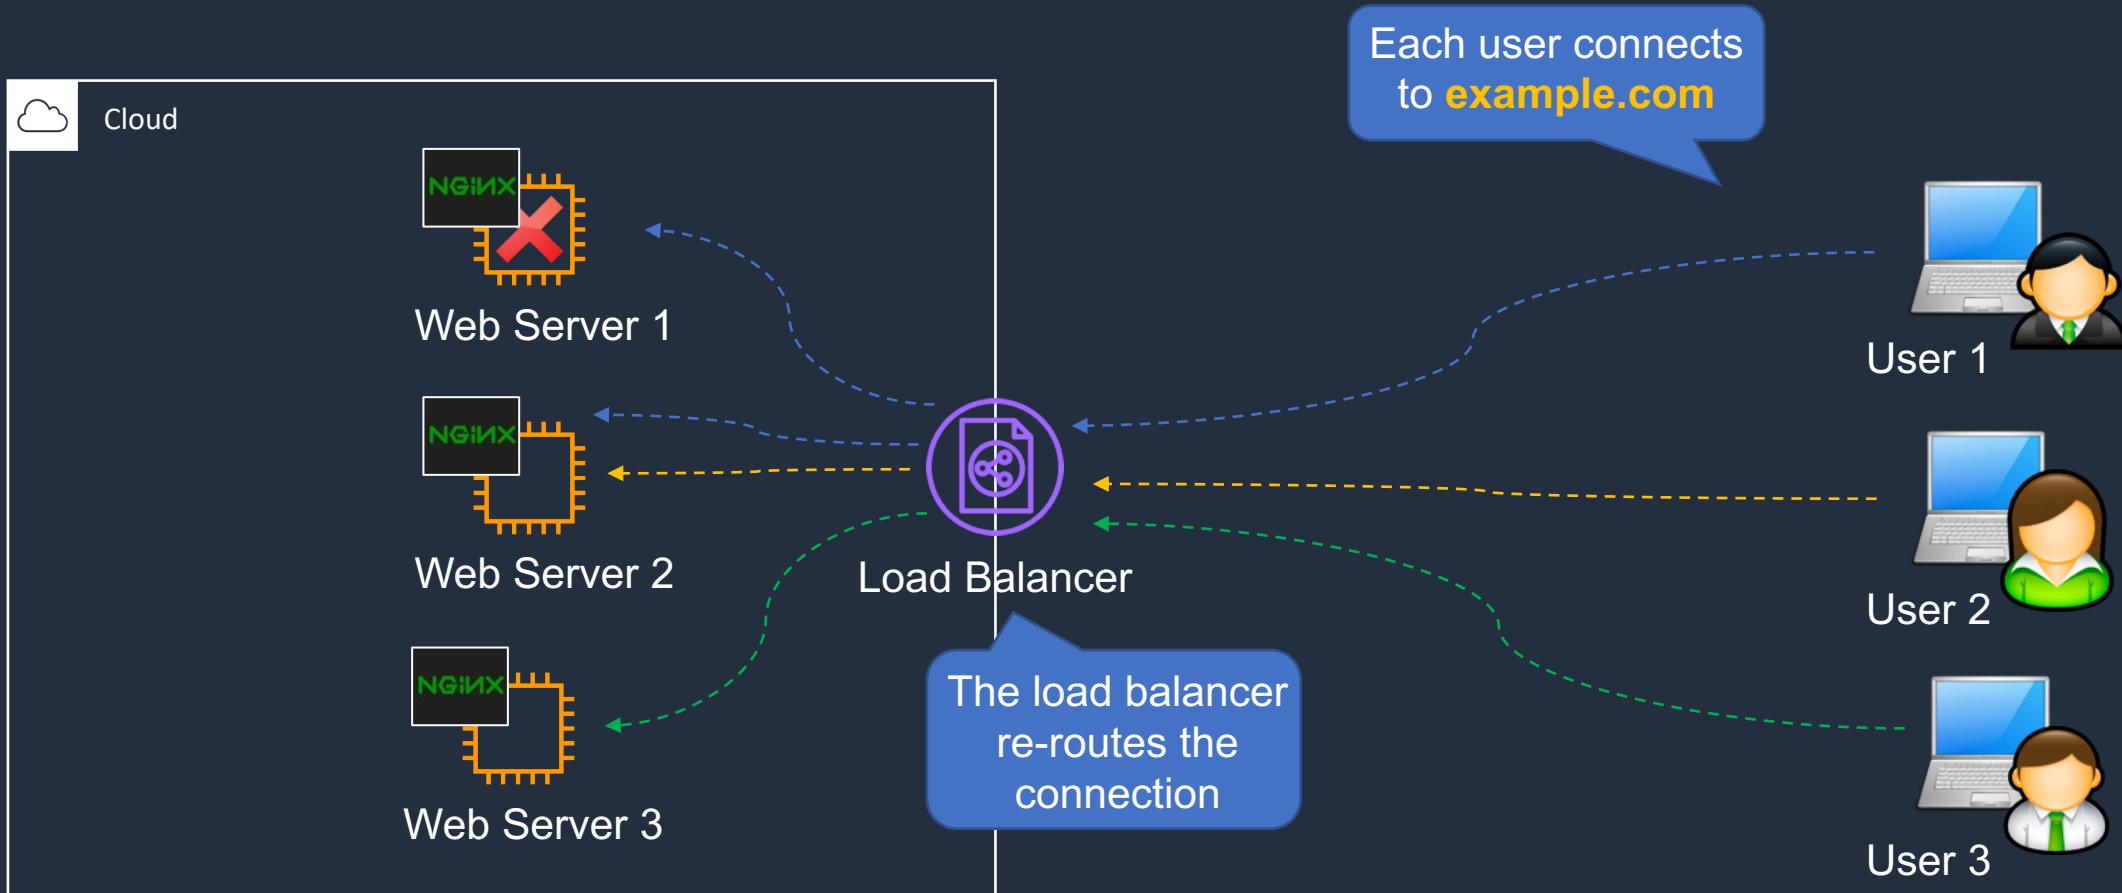


Load Balancing and High Availability





Load Balancing and High Availability





Fault Tolerance

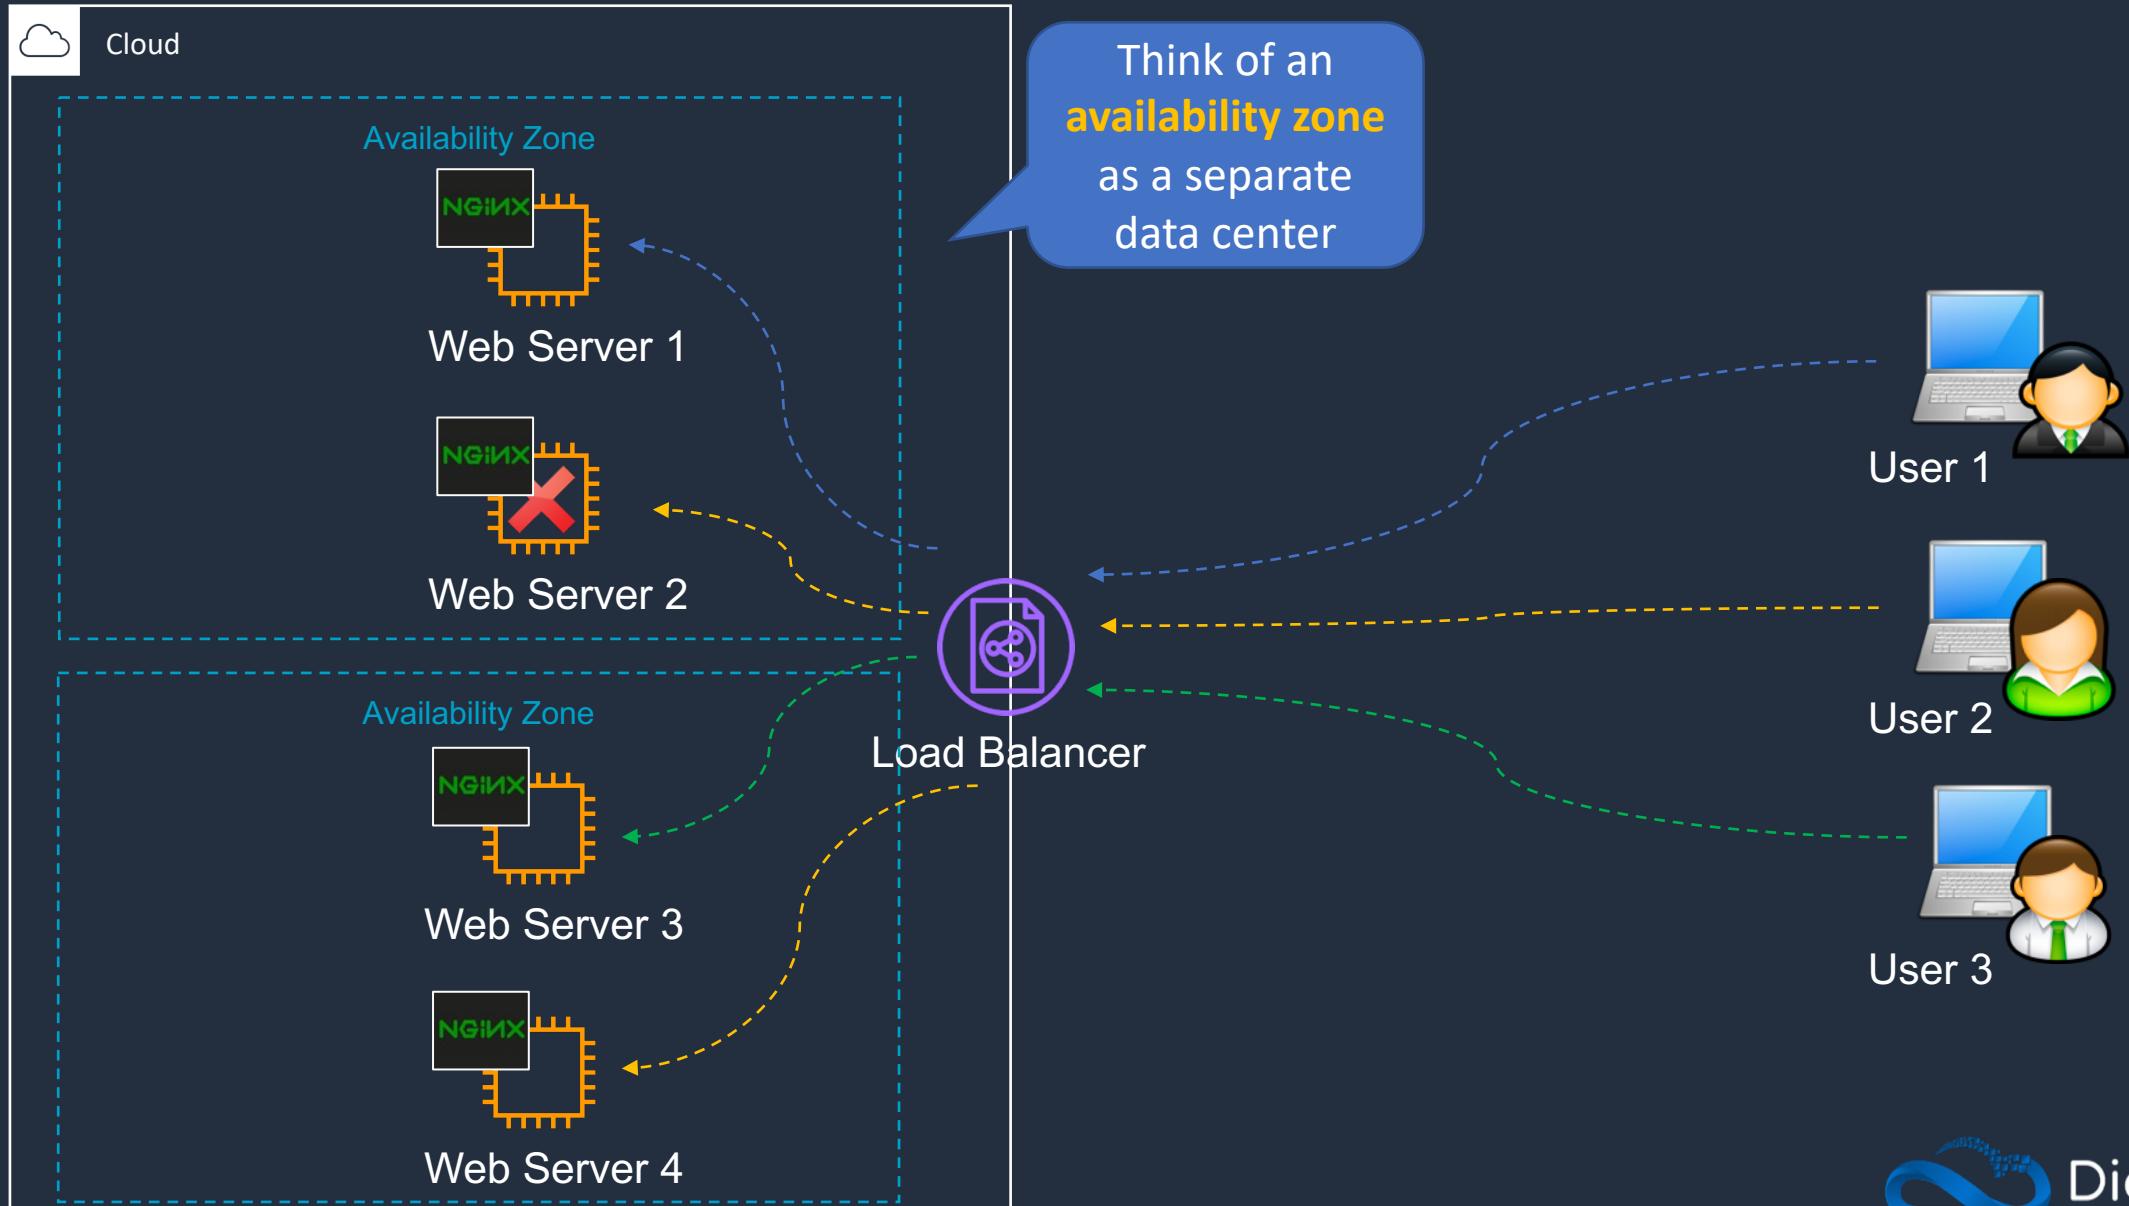
Redundant components allow the system to continue to operate



The system may fail if there is no built-in **redundancy**

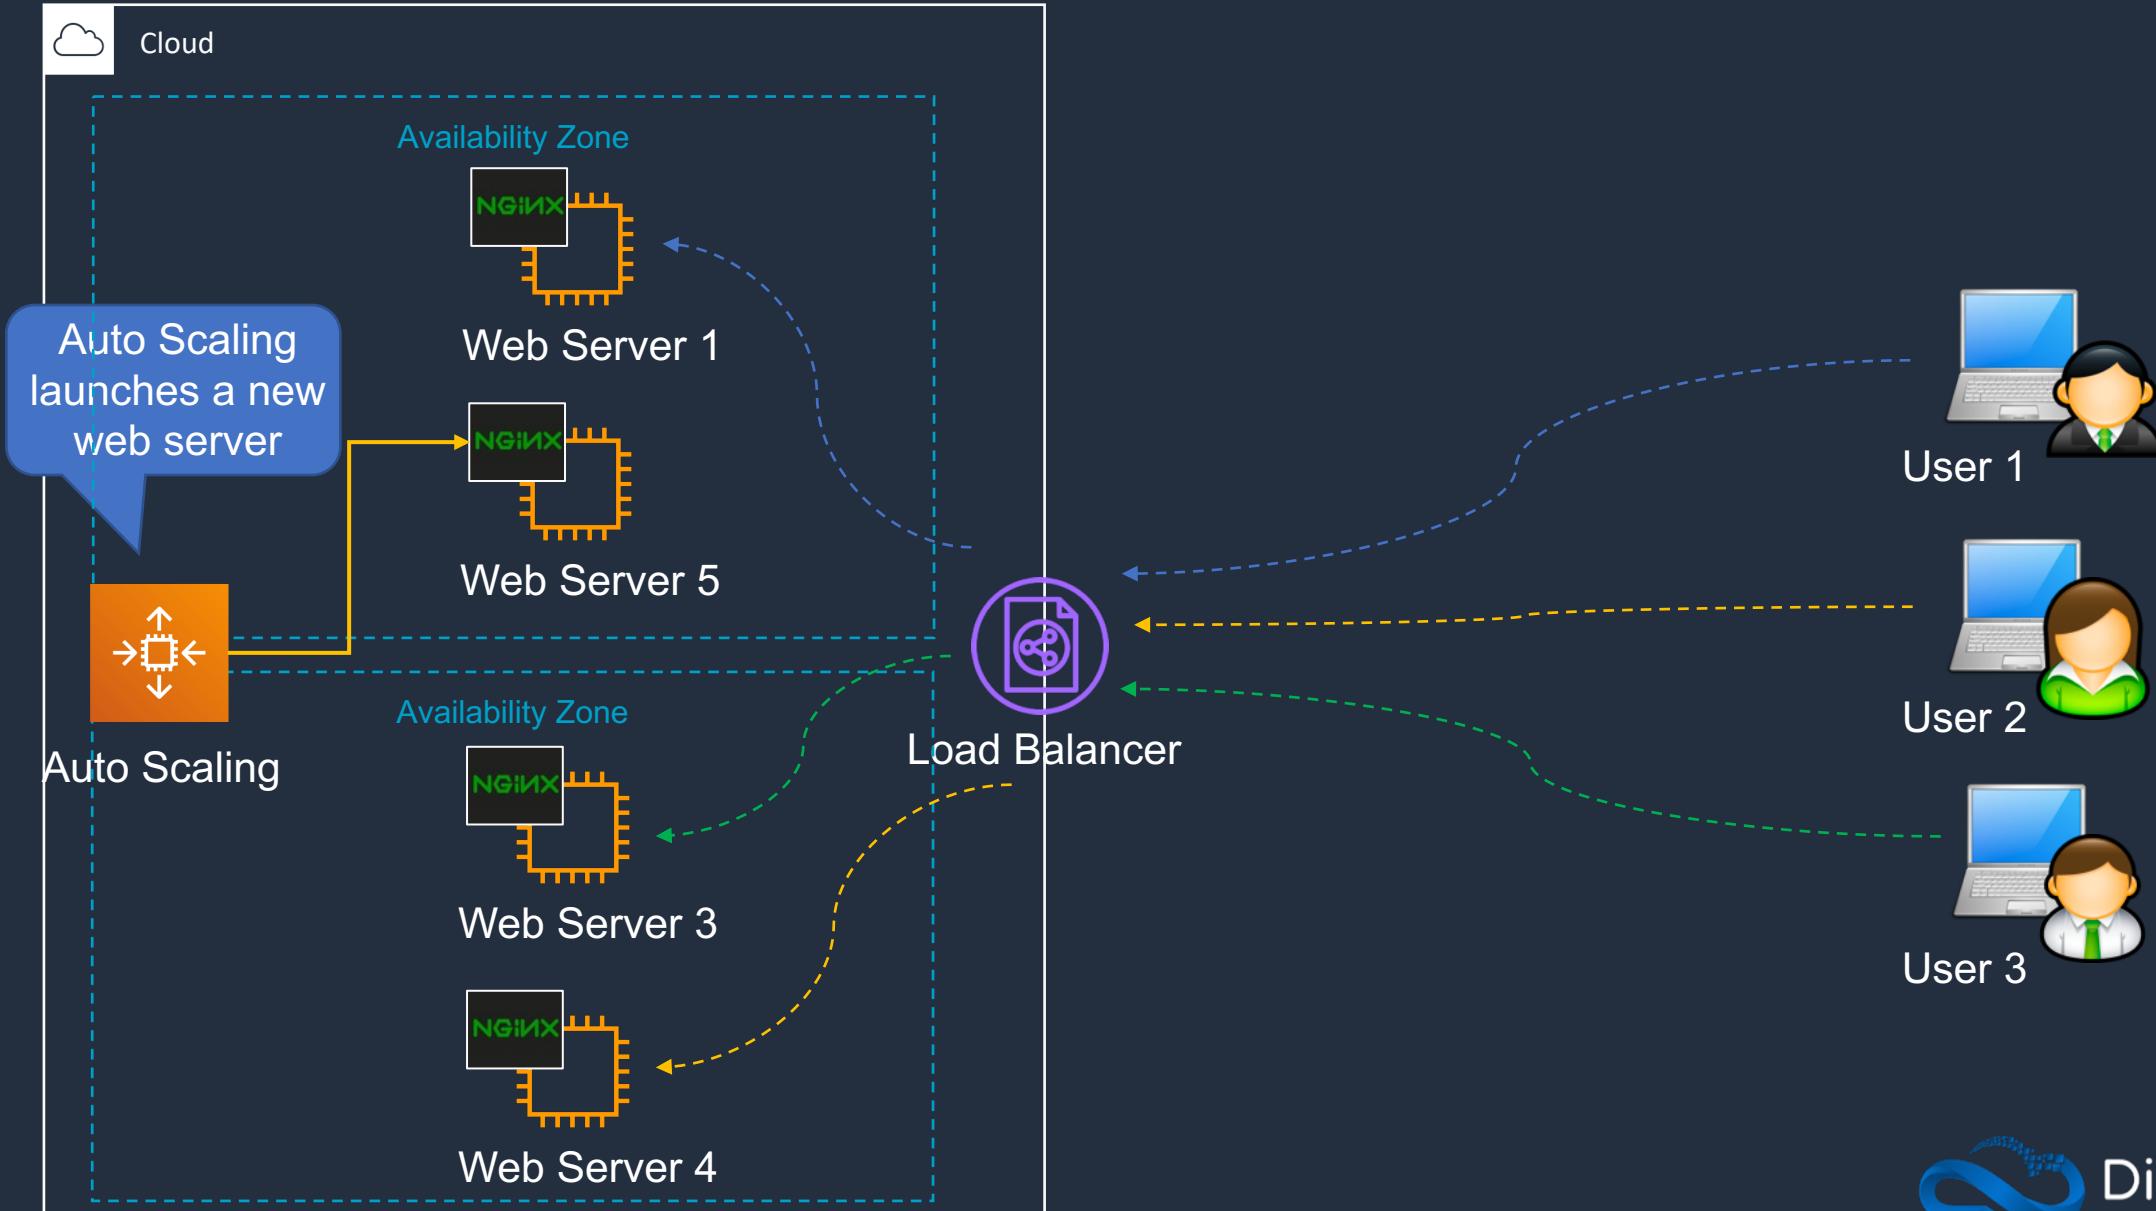


High Availability and Fault Tolerance



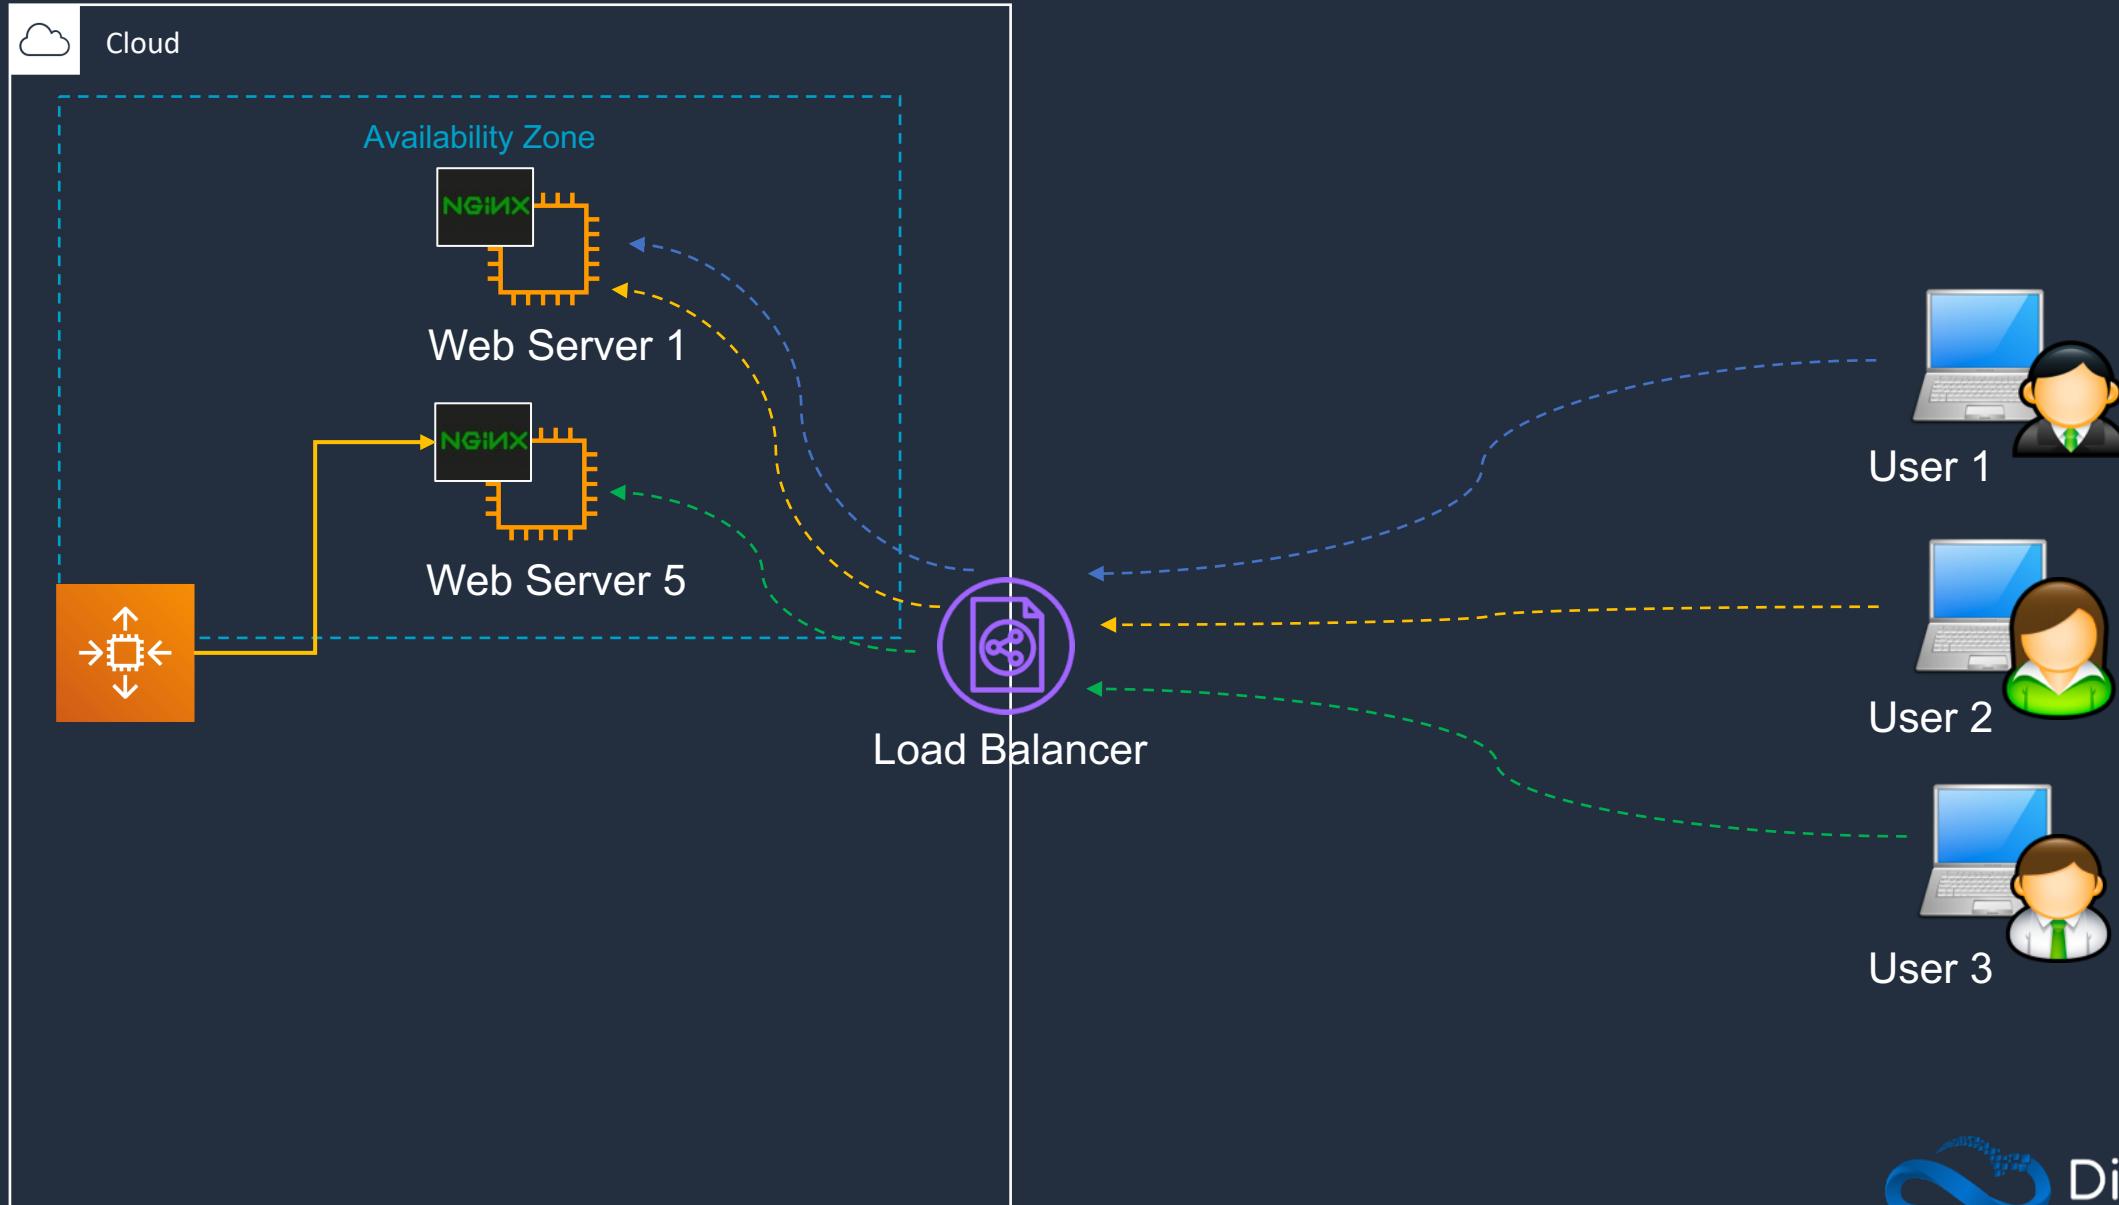


High Availability and Fault Tolerance

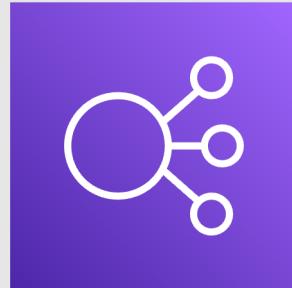




High Availability and Fault Tolerance

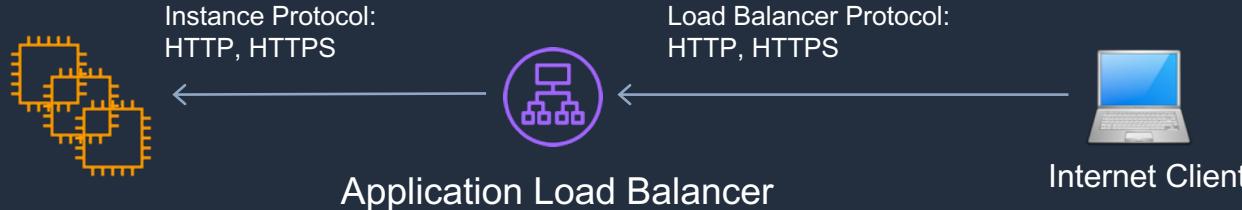


Amazon Elastic Load Balancer (ELB)



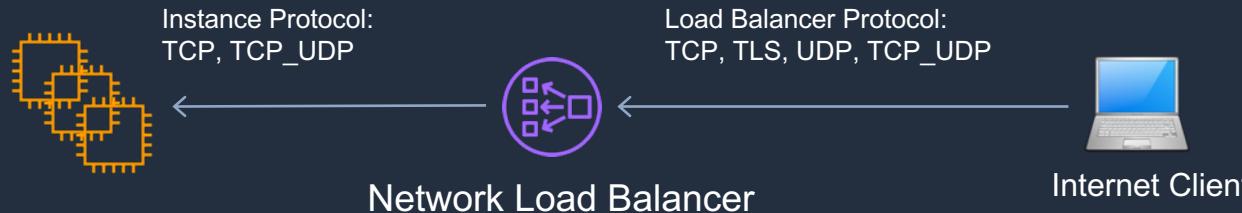


Types of Elastic Load Balancer (ELB)



Application Load Balancer

- Operates at the request level
- Routes based on the content of the request (layer 7)
- Supports advanced routing



Network Load Balancer

- Operates at the connection level
- Routes connections based on IP protocol data (layer 4)
- Offers ultra high performance, low latency and TLS offloading at scale



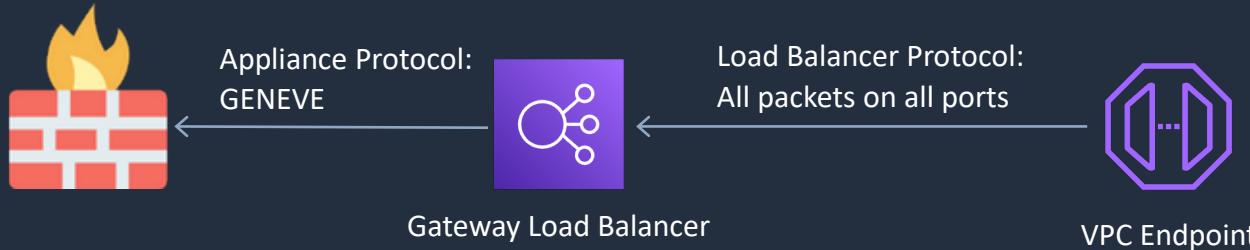
Types of Elastic Load Balancer (ELB)

Old and **shouldn't** be
the exam anymore



Classic Load Balancer

- Old generation; not recommended for new applications
- Performs routing at Layer 4 and Layer 7
- Use for existing applications running in EC2-Classic



Gateway Load Balancer

- Used in front of virtual appliances such as firewalls, IDS/IPS, and deep packet inspection systems

New and **not** yet
on the exam

Attach an Application Load Balancer

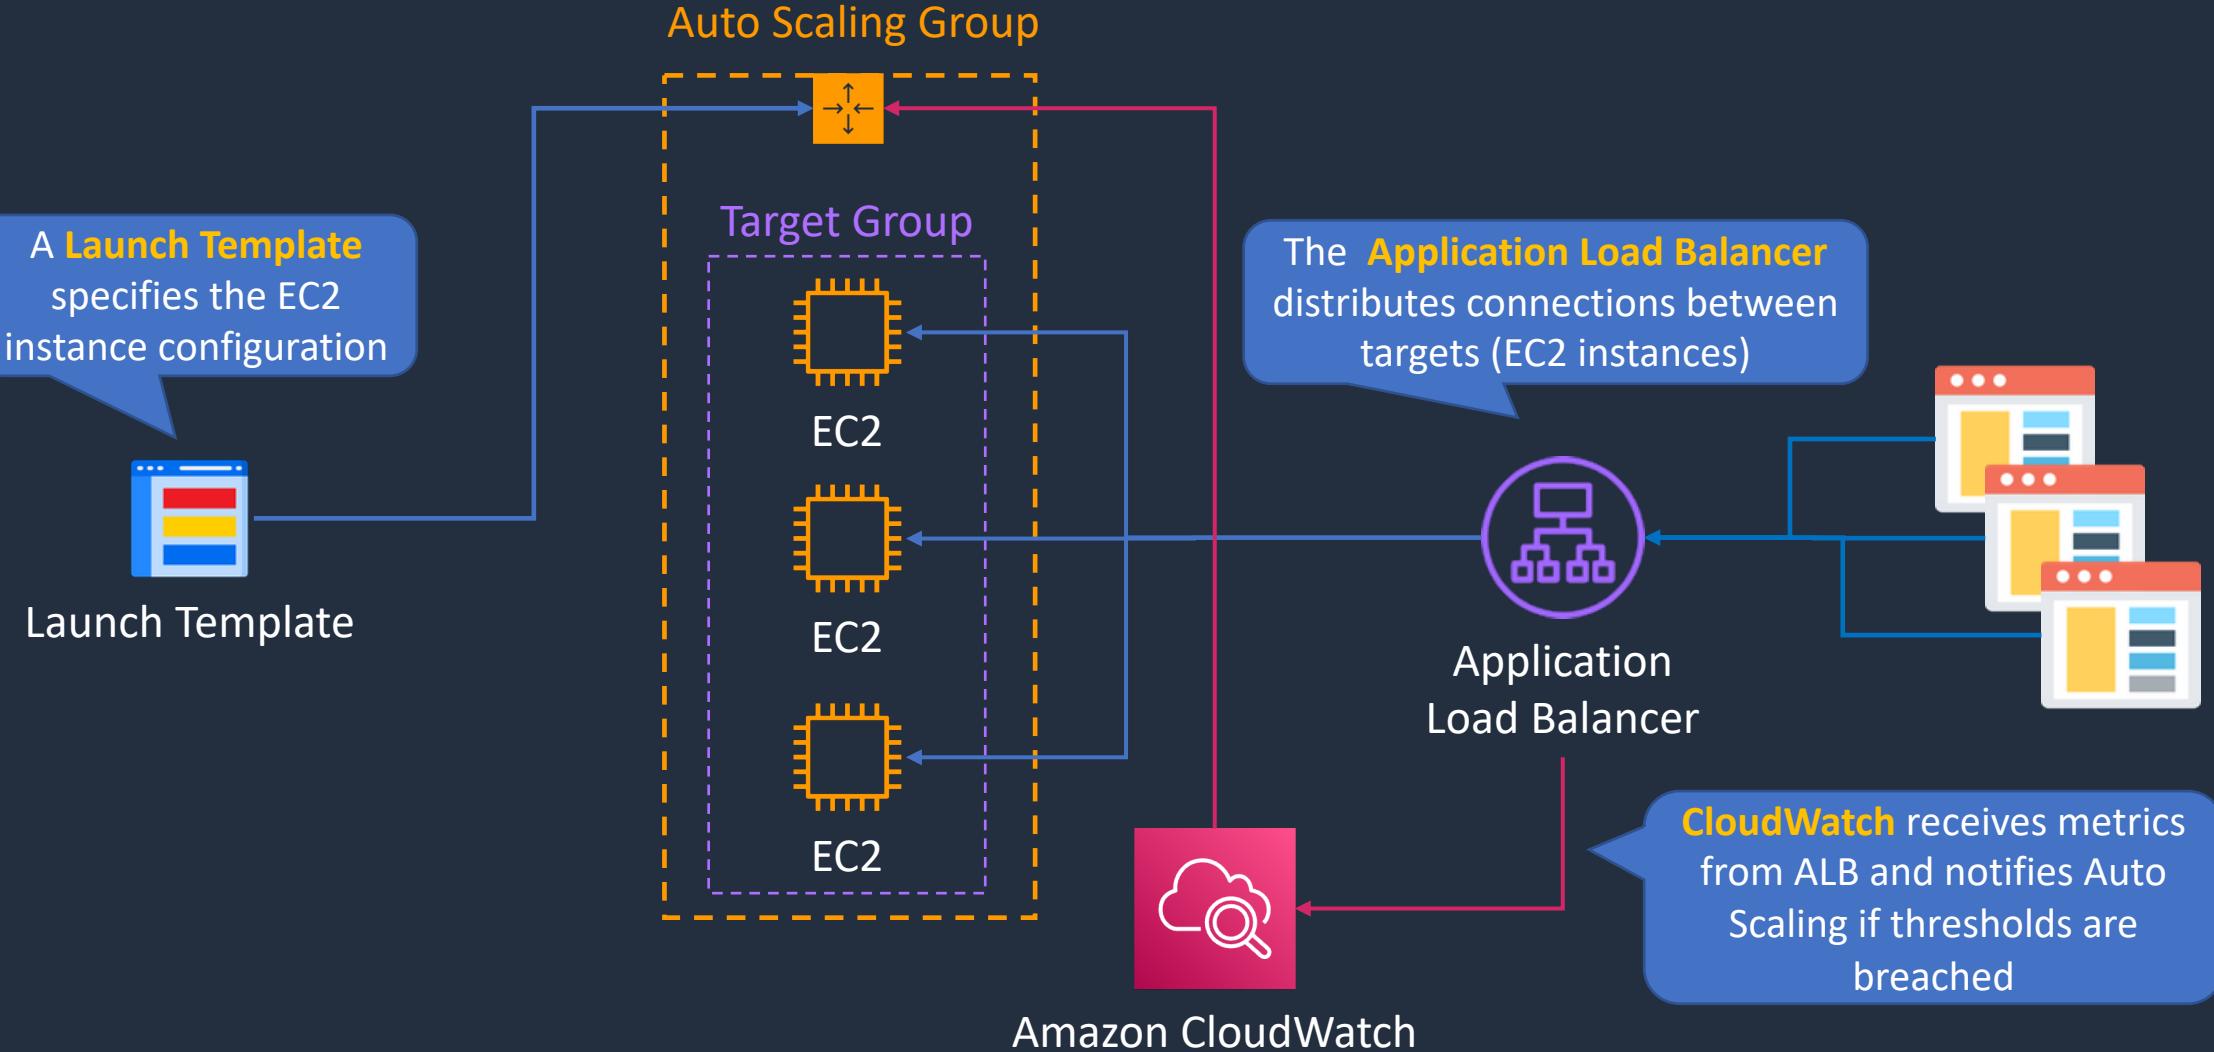


Elastically Scale the Application

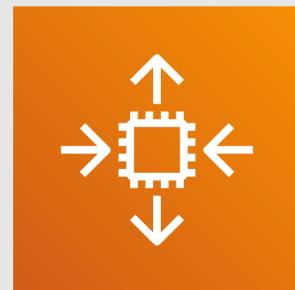




Elastically Scale the Application



Scaling Policies





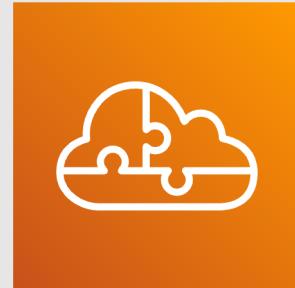
Scaling Policies

- **Target Tracking** – Attempts to keep the group at or close to the metric
- **Simple Scaling** – Adjust group size based on a metric
- **Step Scaling** – Adjust group size based on a metric – adjustments vary based on the size of the alarm breach
- **Scheduled Scaling** – Adjust the group size at a specific time

SECTION 8

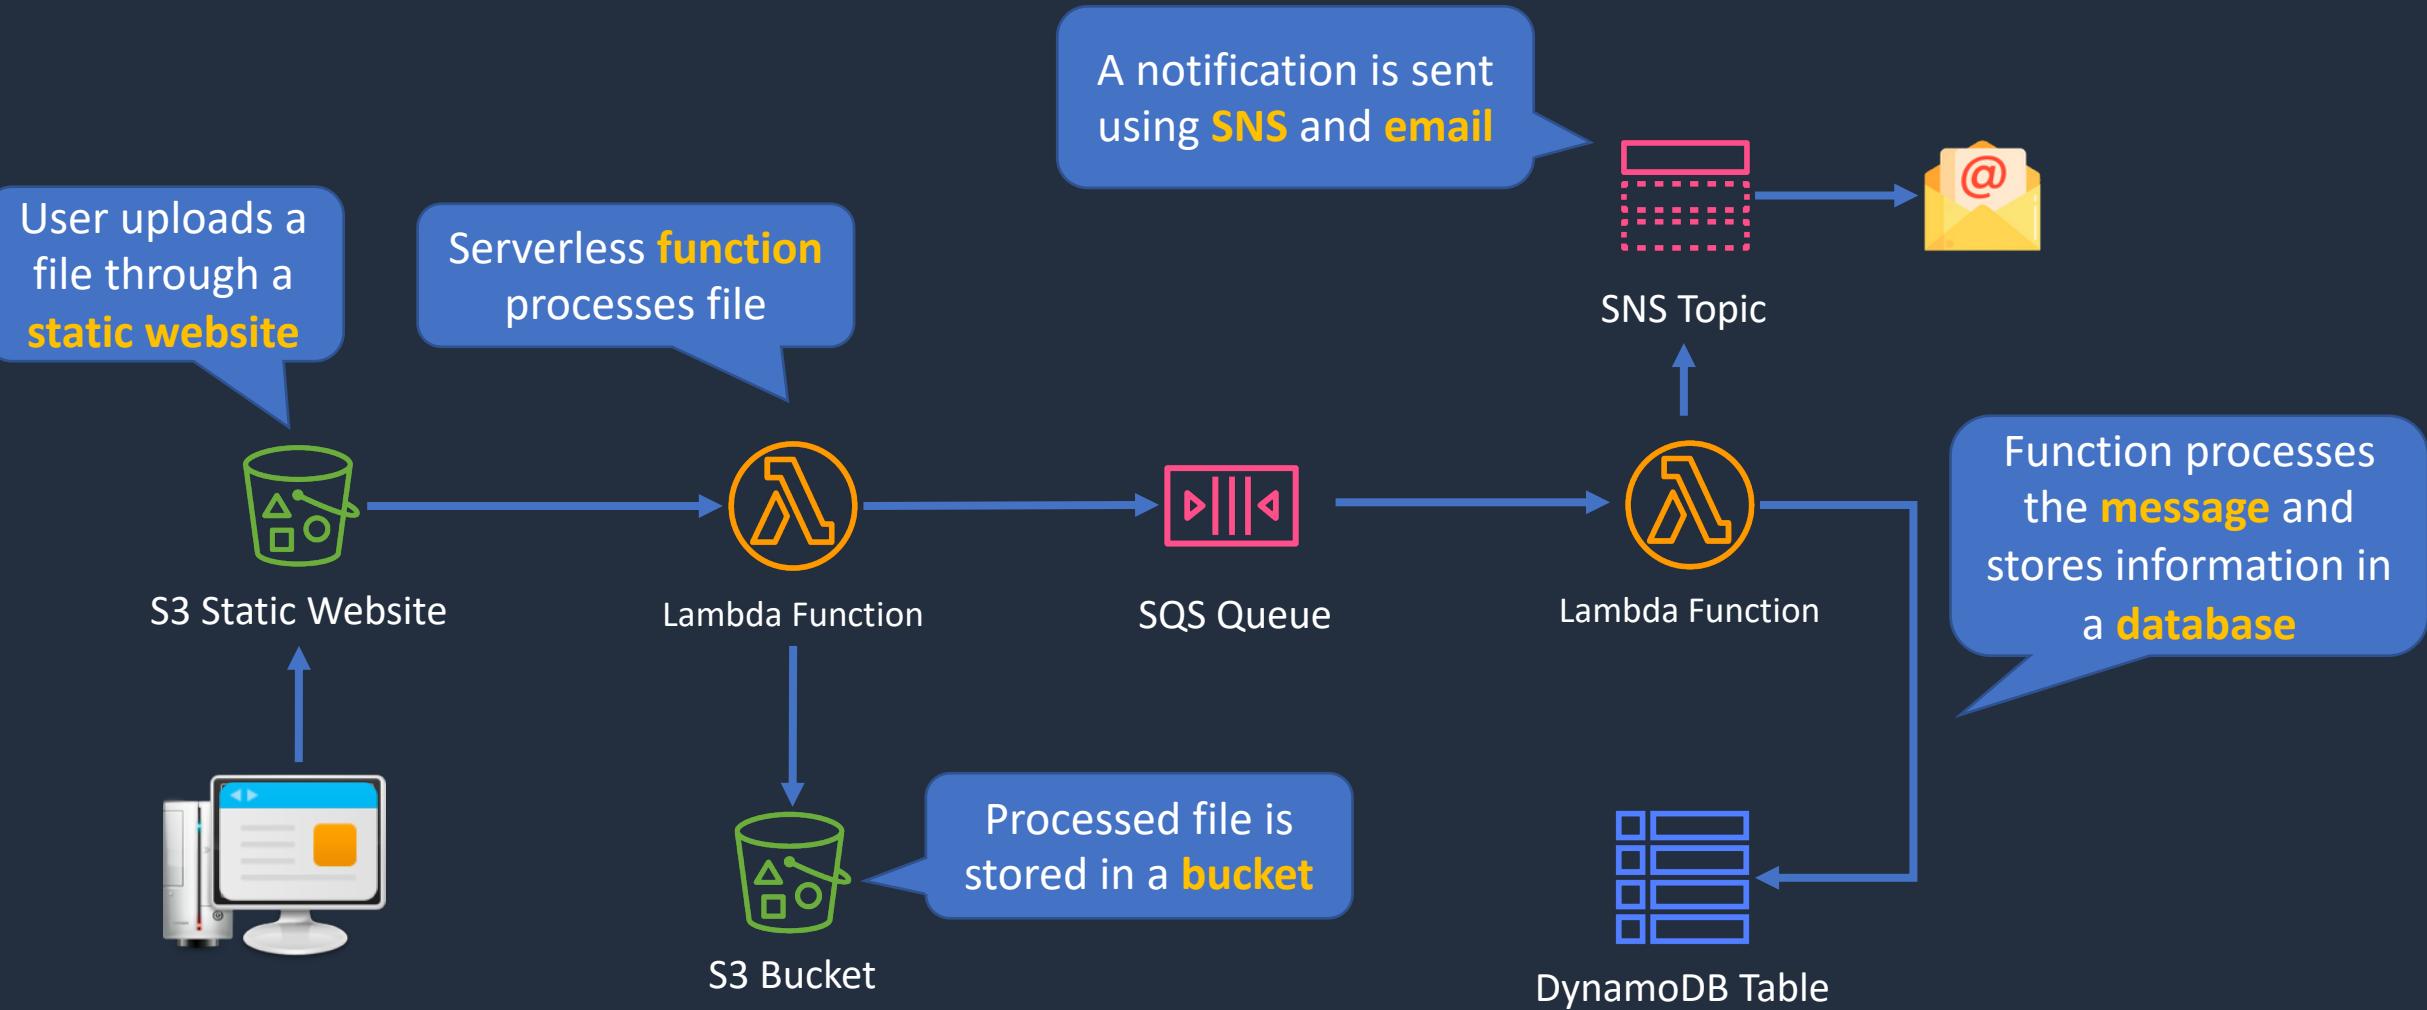
Application Services

Serverless Services





Serverless Services





Serverless Services

- With serverless there are **no instances** to manage
- You don't need to provision hardware
- There is no management of operating systems or software
- Capacity provisioning and patching is handled automatically
- Provides automatic scaling and high availability
- Can be very cheap!



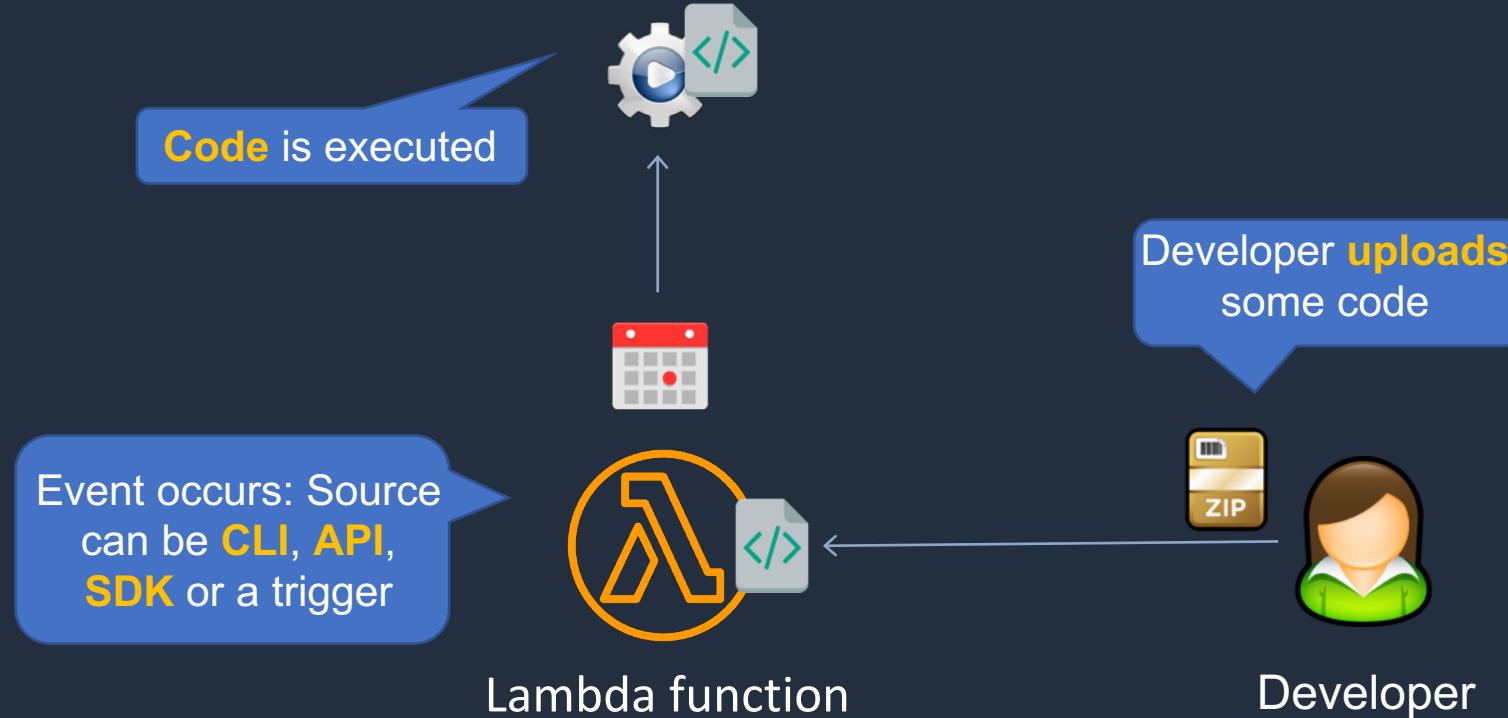
Serverless Services

- Serverless services include:
 - AWS Lambda
 - AWS Fargate
 - Amazon EventBridge
 - AWS Step Functions
 - Amazon SQS
 - Amazon SNS
 - Amazon API Gateway
 - Amazon S3
 - Amazon DynamoDB

AWS Lambda Functions



AWS Lambda Functions





AWS Lambda Functions

- AWS Lambda executes code only when needed and scales automatically
- You pay only for the compute time you consume (you pay nothing when your code is not running)
- Benefits of AWS Lambda:
 - No servers to manage
 - Continuous scaling
 - Millisecond billing
 - Integrates with almost all other AWS services



AWS Lambda Functions

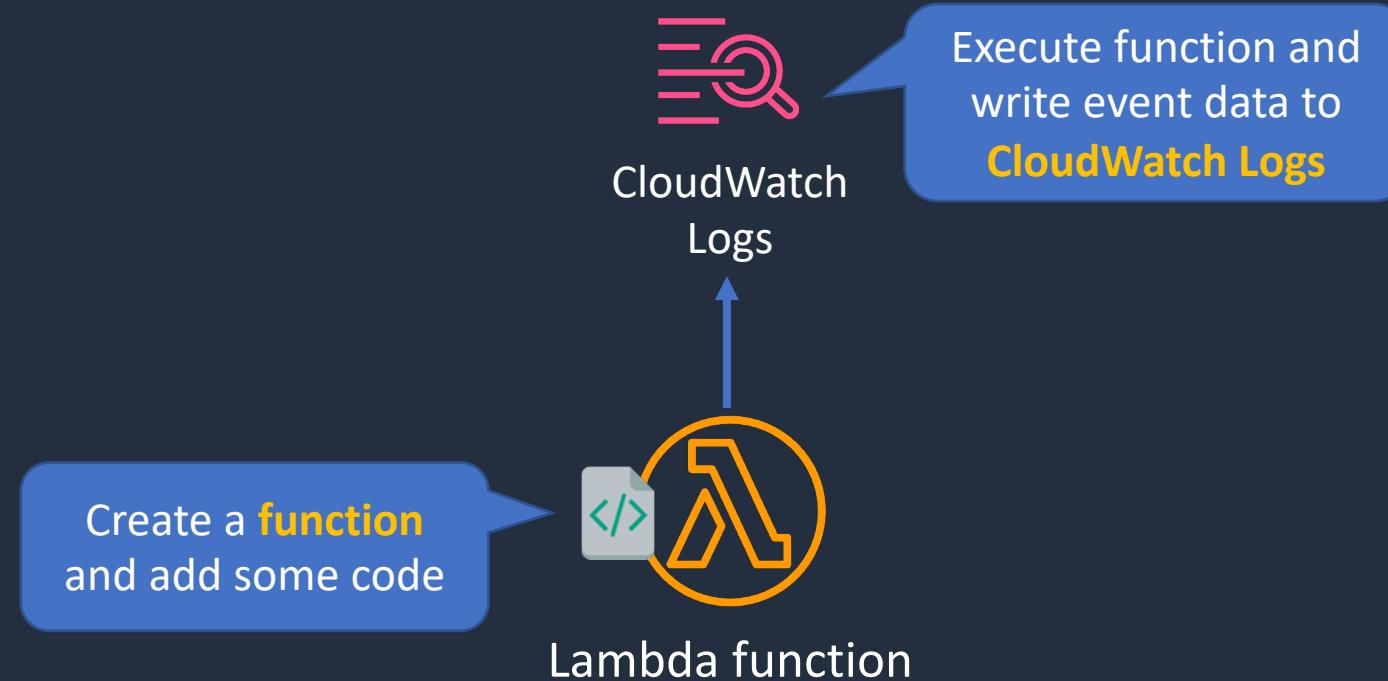
- Primary use cases for AWS Lambda:
 - Data processing
 - Real-time file processing
 - Real-time stream processing
 - Build serverless backends for web, mobile, IOT, and 3rd party API requests

Create a Simple Lambda Function





Create a Simple Lambda Function

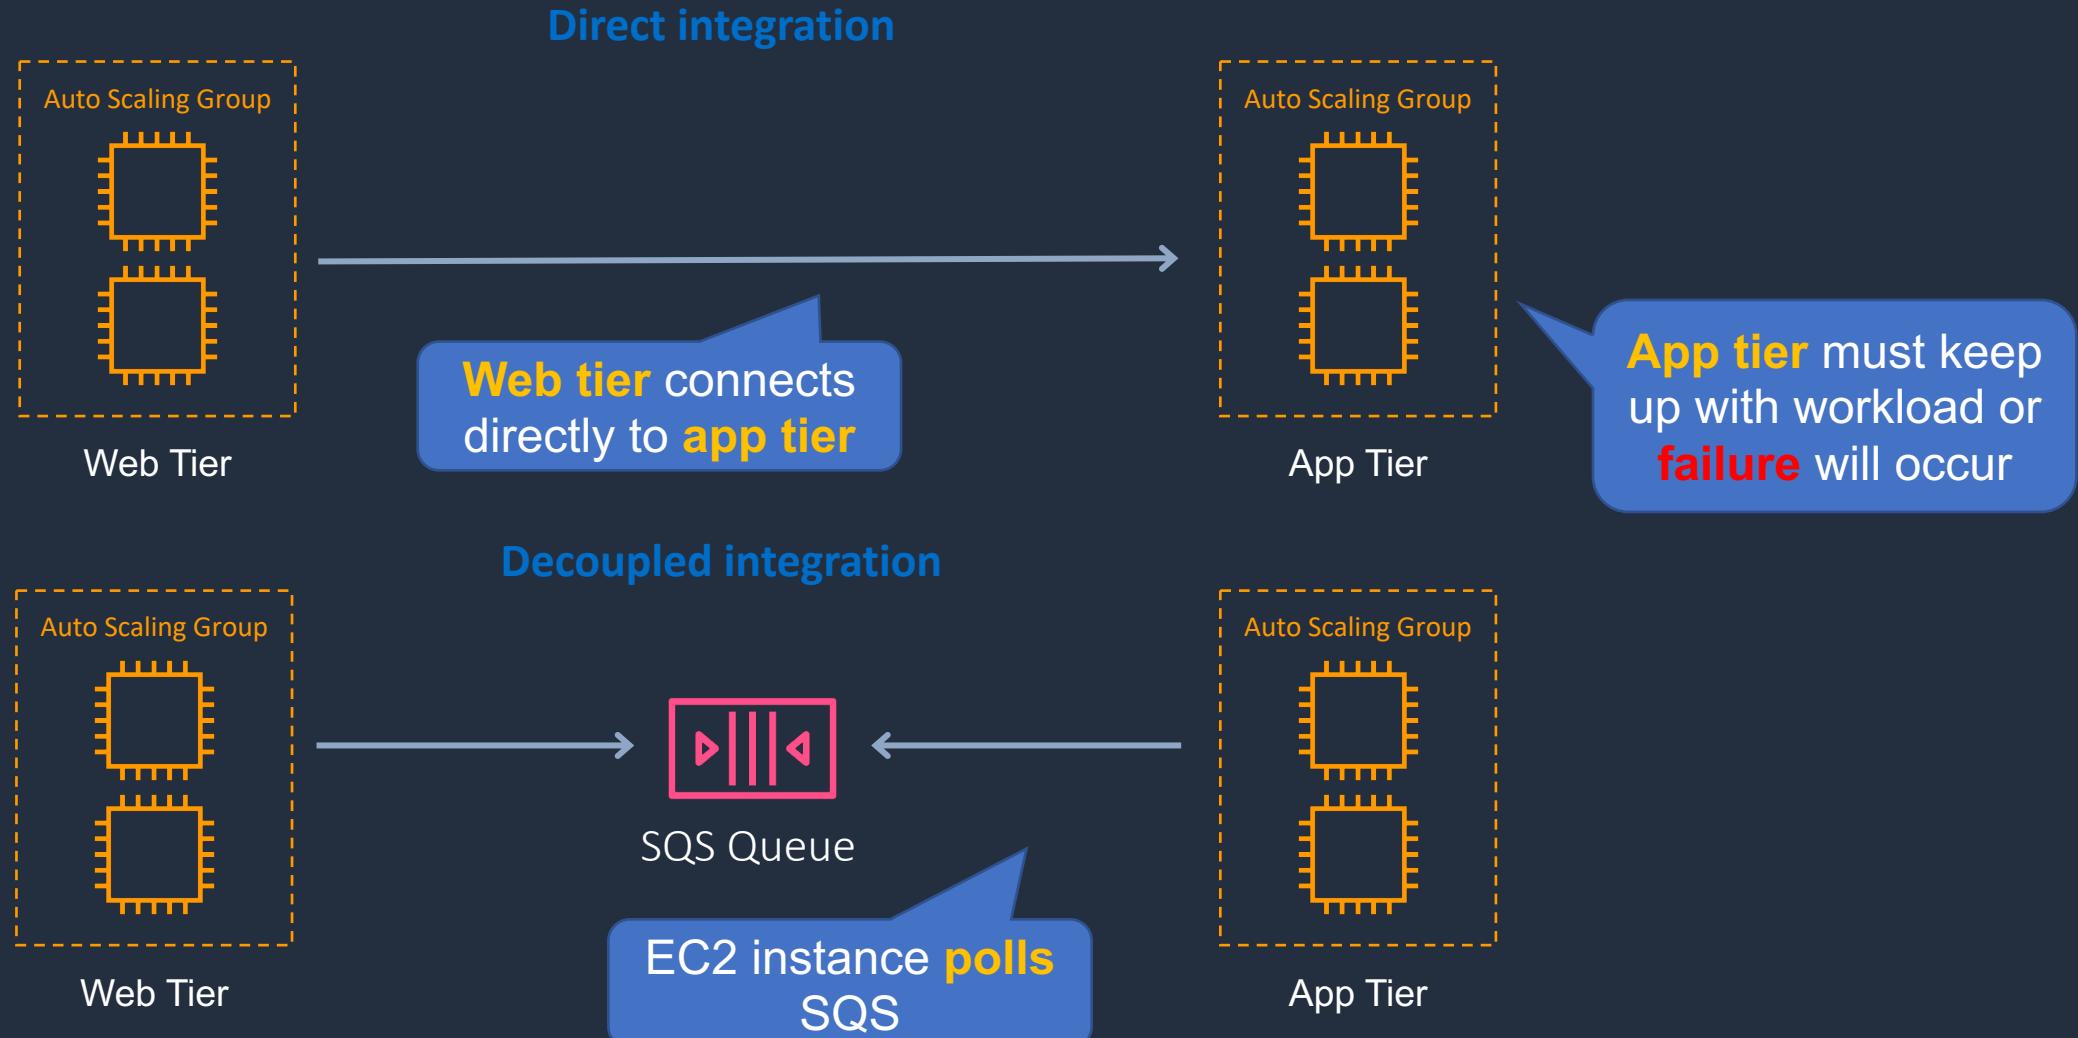


Application Integration Services





Amazon Simple Queue Service (SQS)





Amazon SQS

- SQS offers a reliable, highly-scalable, hosted queue for storing messages in transit between computers
- SQS is used for distributed/decoupled applications
- SQS uses a message-oriented API
- SQS uses pull based (polling) not push based



Amazon MQ

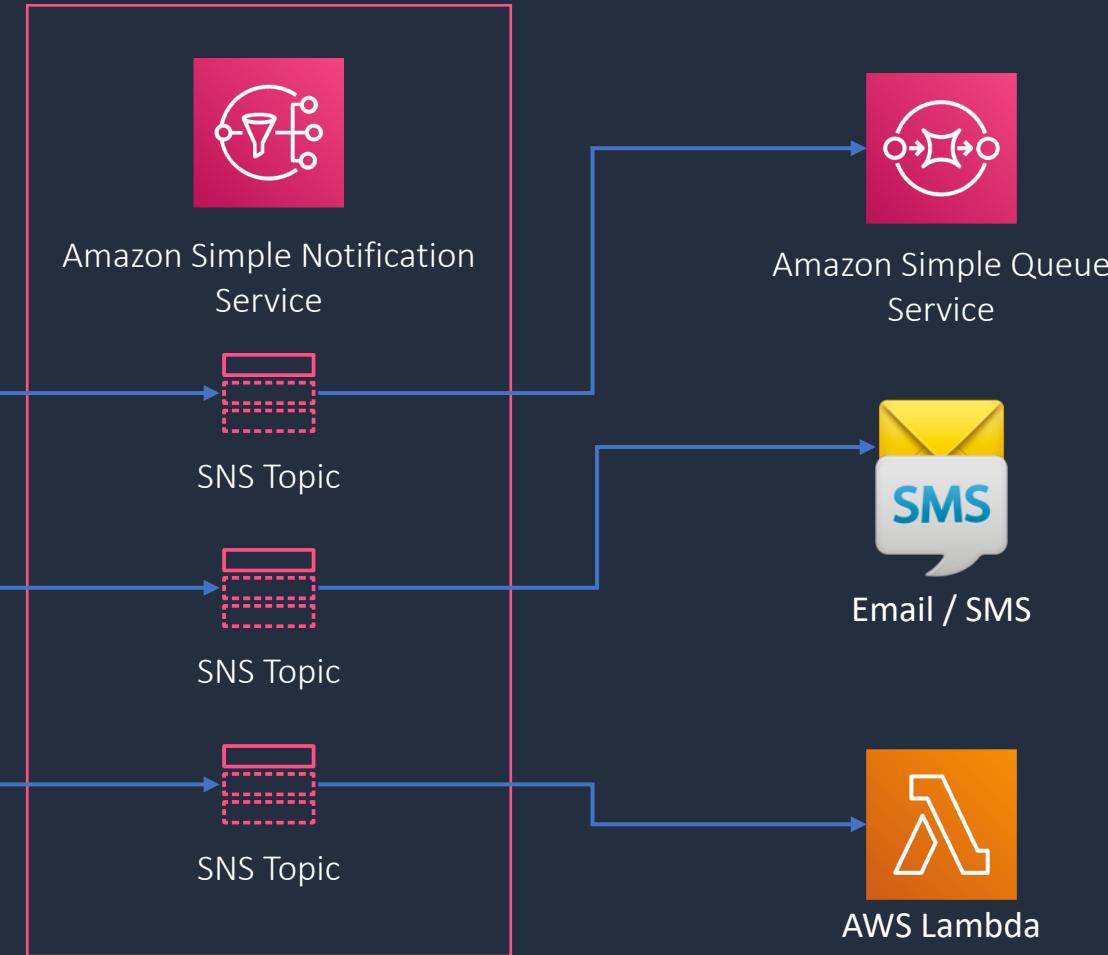
- Message broker service
- Similar to Amazon SQS
- Based on Apache Active MQ and RabbitMQ
- Used when customers require industry standard APIs and protocols
- Useful when migrating existing queue-based applications into the cloud



Amazon Simple Notification Service (SNS)

PUBLISHERS

- Amazon EC2
- Amazon CloudWatch
- Amazon Simple Storage Service



SUBSCRIBERS

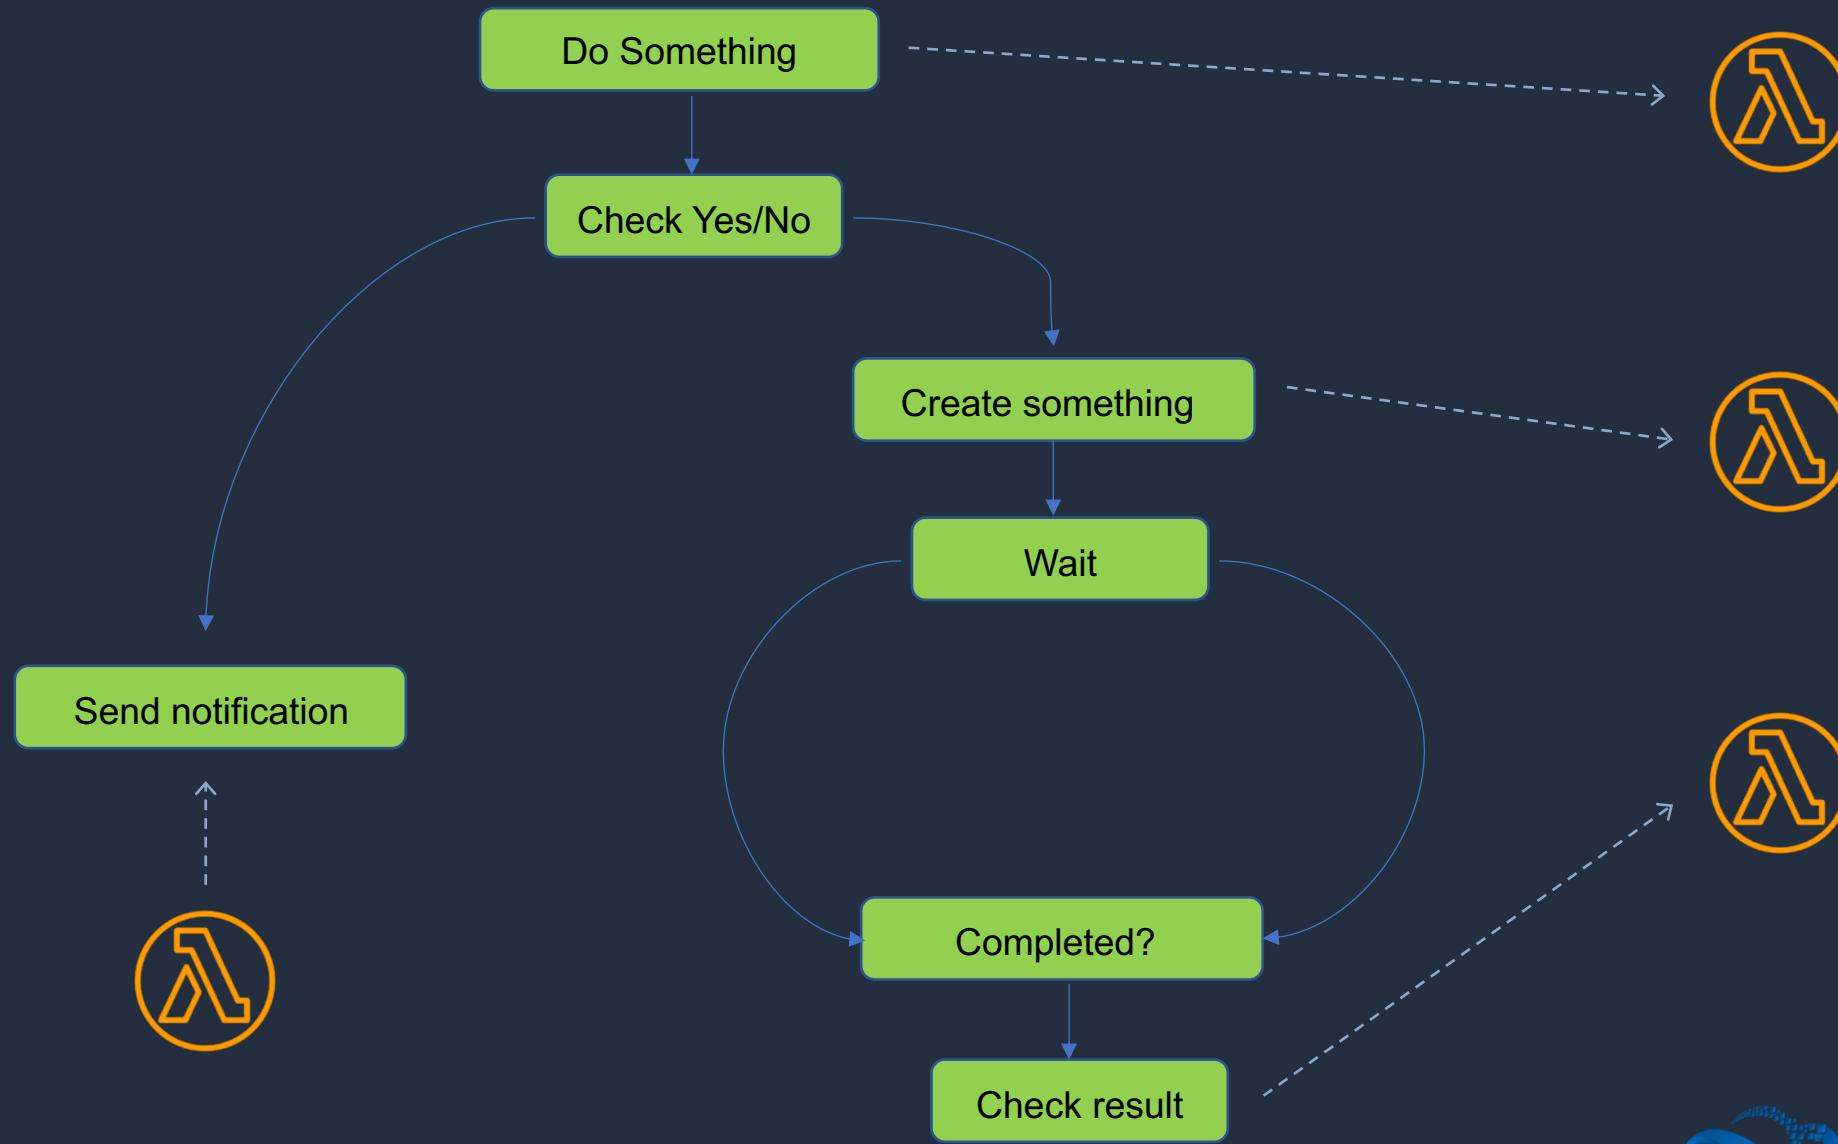


Amazon SNS

- Amazon SNS is used for building and integrating loosely-coupled, distributed applications
- Provides instantaneous, push-based delivery (no polling)
- Uses simple APIs and easy integration with applications
- Offered under an inexpensive, pay-as-you-go model with no up-front costs



AWS Step Functions



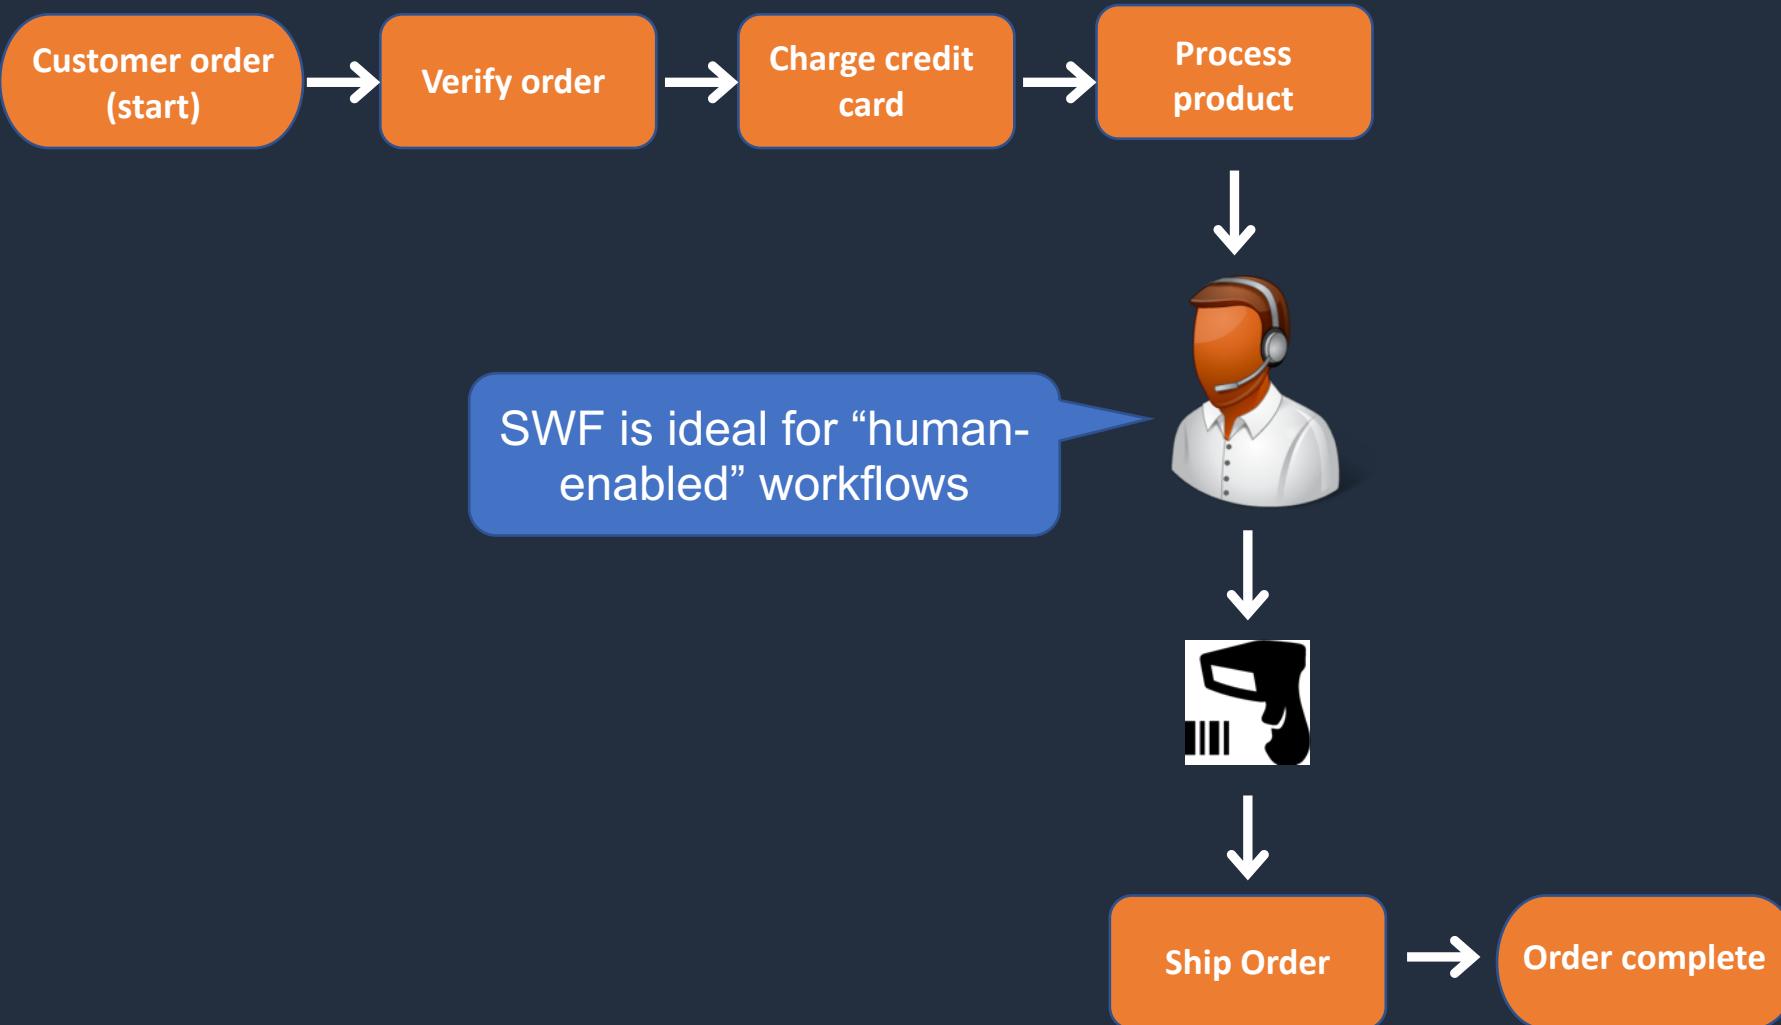


AWS Step Functions

- AWS Step Functions makes it easy to coordinate the components of distributed applications as a series of steps in a visual workflow
- You can quickly build and run state machines to execute the steps of your application in a reliable and scalable fashion



AWS Simple Workflow Service (SWF)





Amazon SWF

- Amazon Simple Workflow Service (SWF) is a web service that makes it easy to coordinate work across distributed application components
- Create distributed asynchronous systems as workflows
- Best suited for human-enabled workflows like an order fulfilment system or for procedural requests
- AWS recommends that for new applications customers consider Step Functions instead of SWF



Application Integration Services Comparison

Service	What it does	Example use cases
Simple Queue Service	Messaging queue; store and forward patterns	Building distributed / decoupled applications
Simple Notification Service	Set up, operate, and send notifications from the cloud	Send email notification when CloudWatch alarm is triggered
Step Functions	Out-of-the-box coordination of AWS service components with visual workflow	Order processing workflow
Simple Workflow Service	Need to support external processes or specialized execution logic	Human-enabled workflows like an order fulfilment system or for procedural requests Note: AWS recommends that for new applications customers consider Step Functions instead of SWF
Amazon MQ	Message broker service for Apache Active MQ and RabbitMQ	Need a message queue that supports industry standard APIs and protocols; migrate queues to AWS

Amazon EventBridge / CloudWatch Events



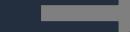
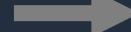


Amazon EventBridge

EventBridge used to be known as **CloudWatch Events**

Event Source

EC2 instance
terminated event

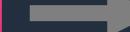
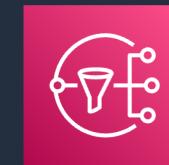


EventBridge
event bus

Event

Rule

Send **SNS** notification



Target



Amazon EventBridge

Event matching pattern
You can use pre-defined pattern provided by a service or create a custom pattern

Pre-defined pattern by service
 Custom pattern

Service provider
AWS services or custom/partner services

AWS

Service name
The name of partner service selected as the event source

EC2

Event type
The type of events as the source of the matching pattern

EC2 Instance State-change Notification

Any state
 Specific state(s)

terminated X

Any instance
 Specific instance Id(s)

i-1234567890abcdef0

Event pattern

Copy Edit

```
1 {
2     "source": ["aws.ec2"],
3     "detail-type": ["EC2 Instance State-change N
4     "detail": {
5         "state": ["terminated"],
6         "instance-id": ["i-1234567890abcdef0"]
7     }
8 }
```

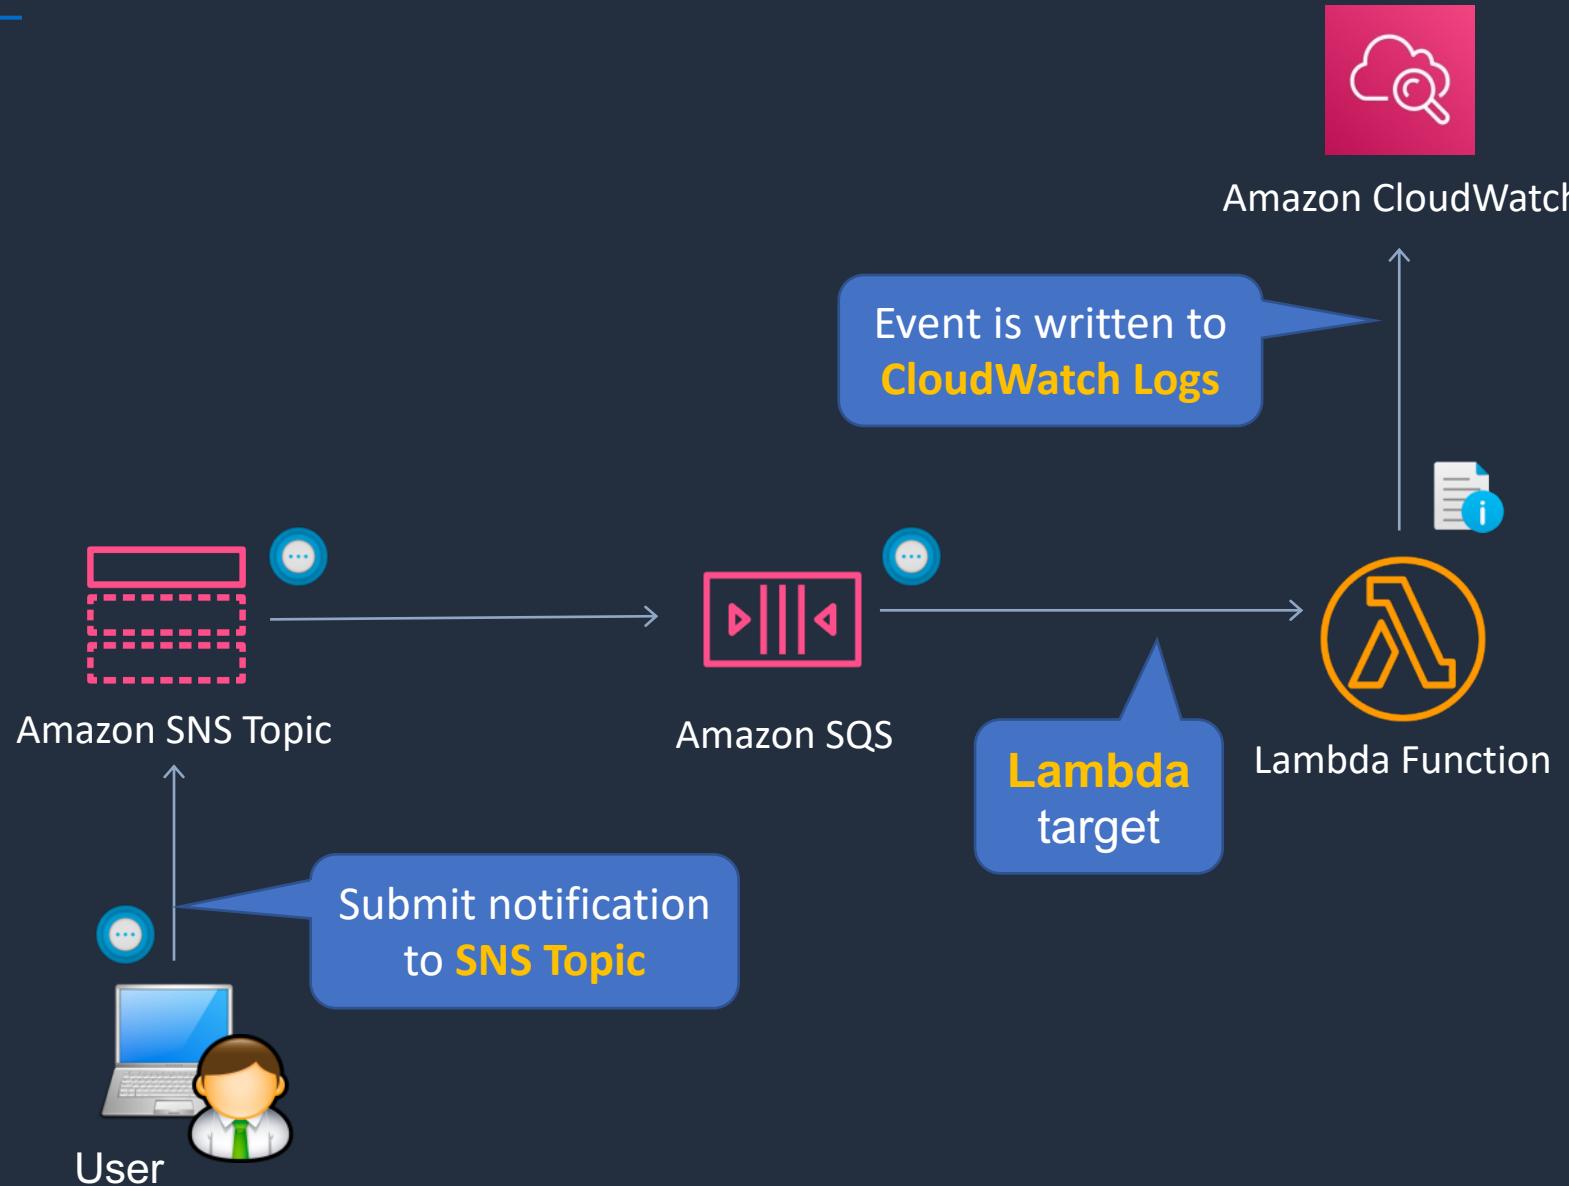
```
{
    "version": "0",
    "id": "6a7e8feb-b491-4cf7-a9f1-bf3703467718",
    "detail-type": "EC2 Instance State-change Notification",
    "source": "aws.ec2",
    "account": "111122223333",
    "time": "2017-12-22T18:43:48Z",
    "region": "us-west-1",
    "resources": [
        "arn:aws:ec2:us-west-1:123456789012:instance/i-1234567890abcdef0"
    ],
    "detail": {
        "instance-id": "i-1234567890abcdef0",
        "state": "terminated"
    }
}
```

Create an Event-Driven Application

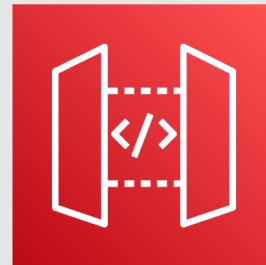




Simple Event-Driven Application

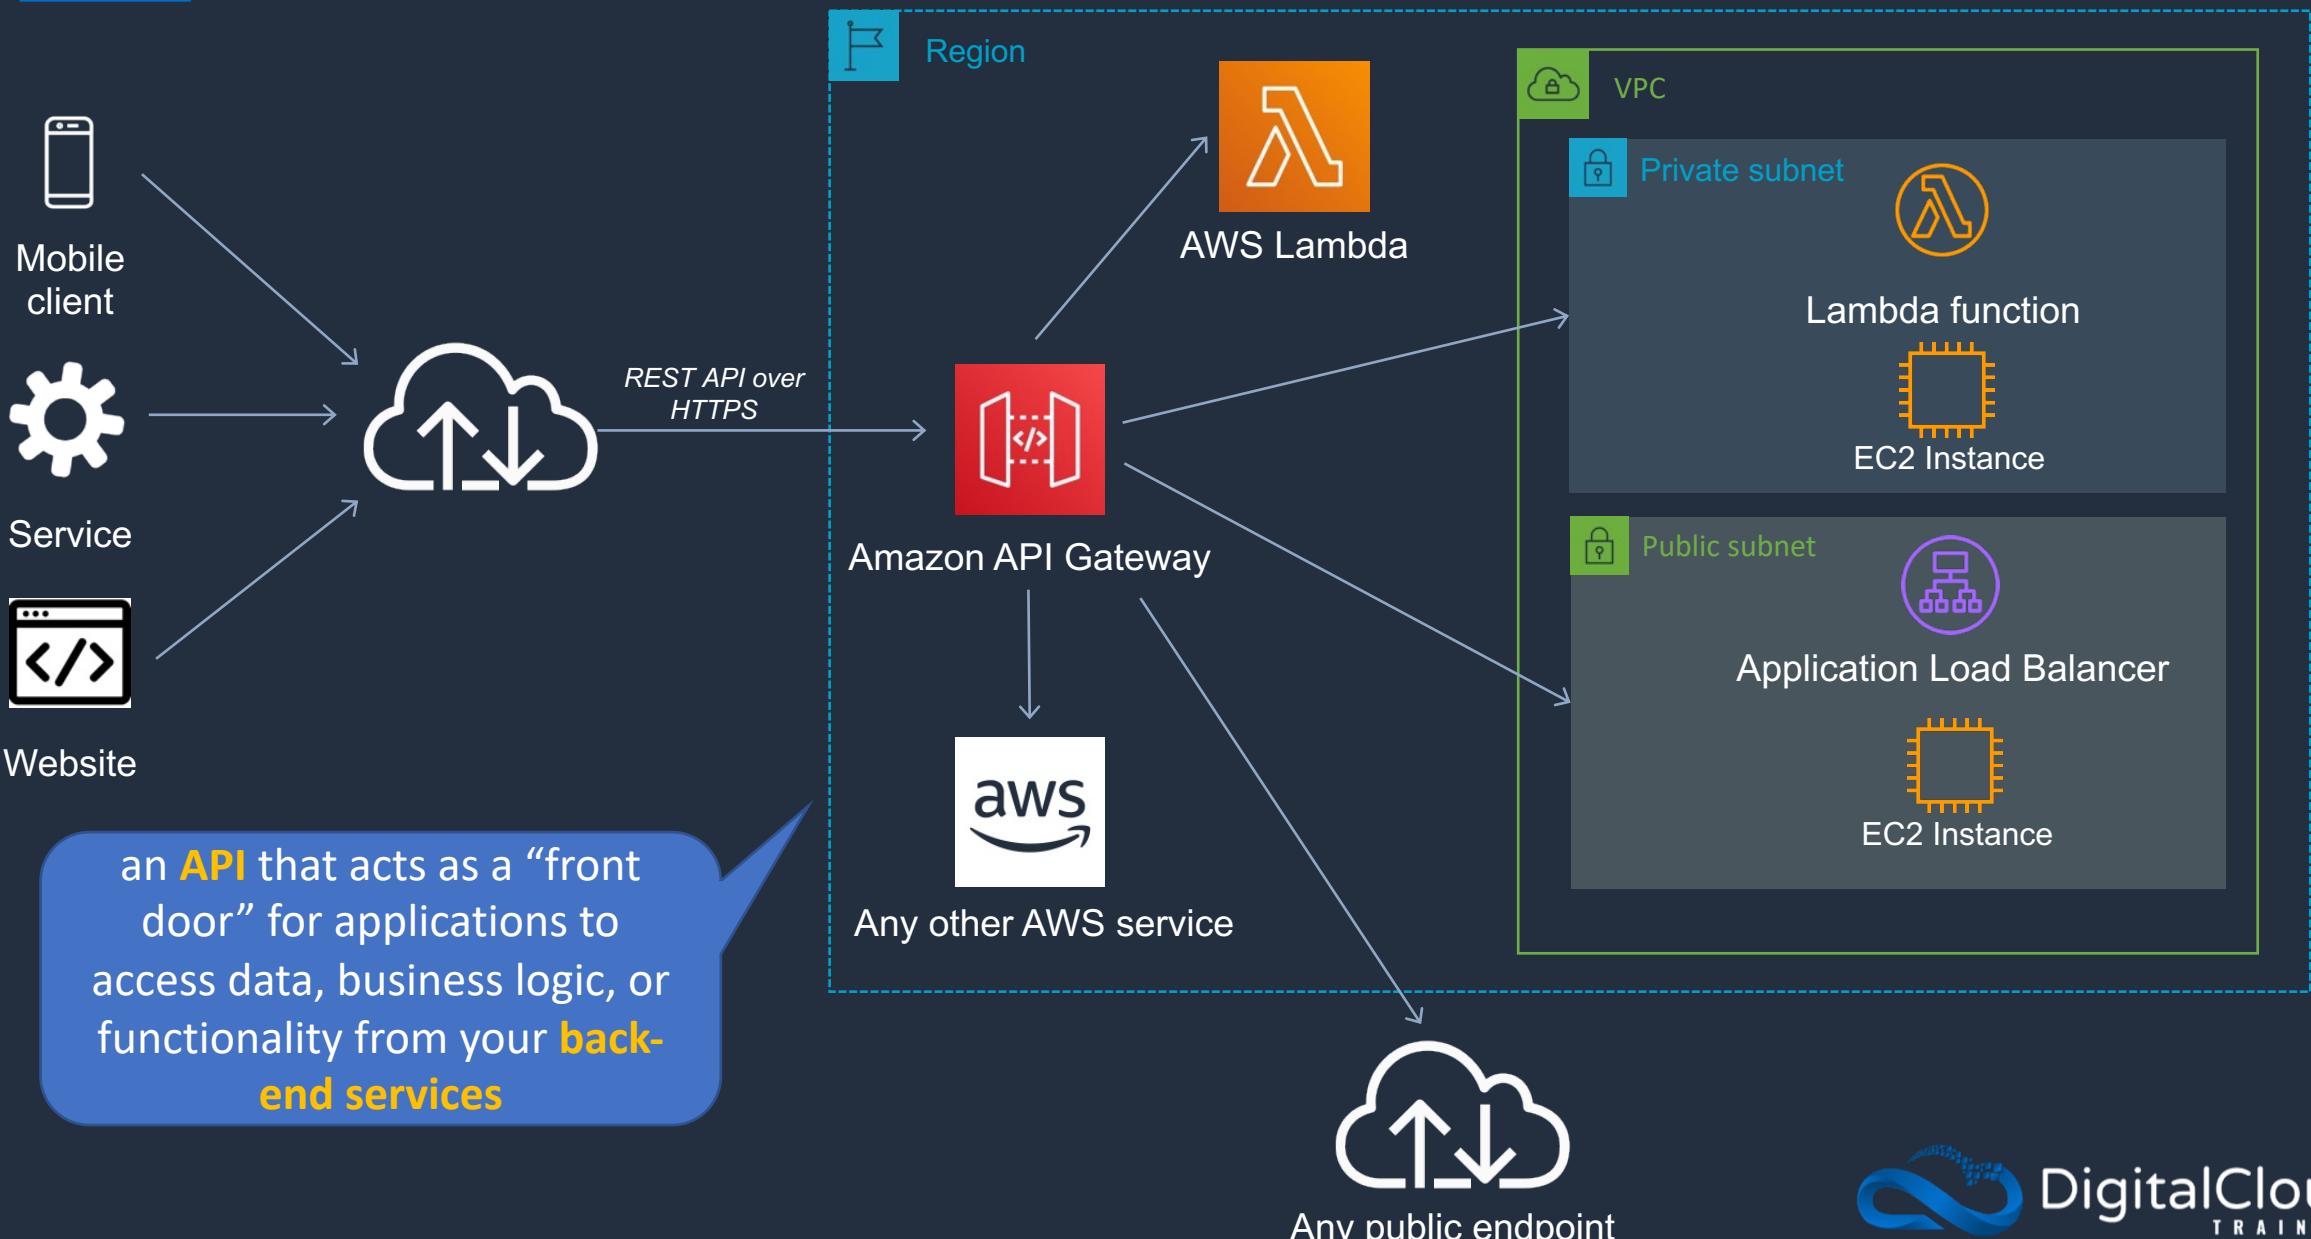


Amazon API Gateway





Amazon API Gateway



SECTION 9

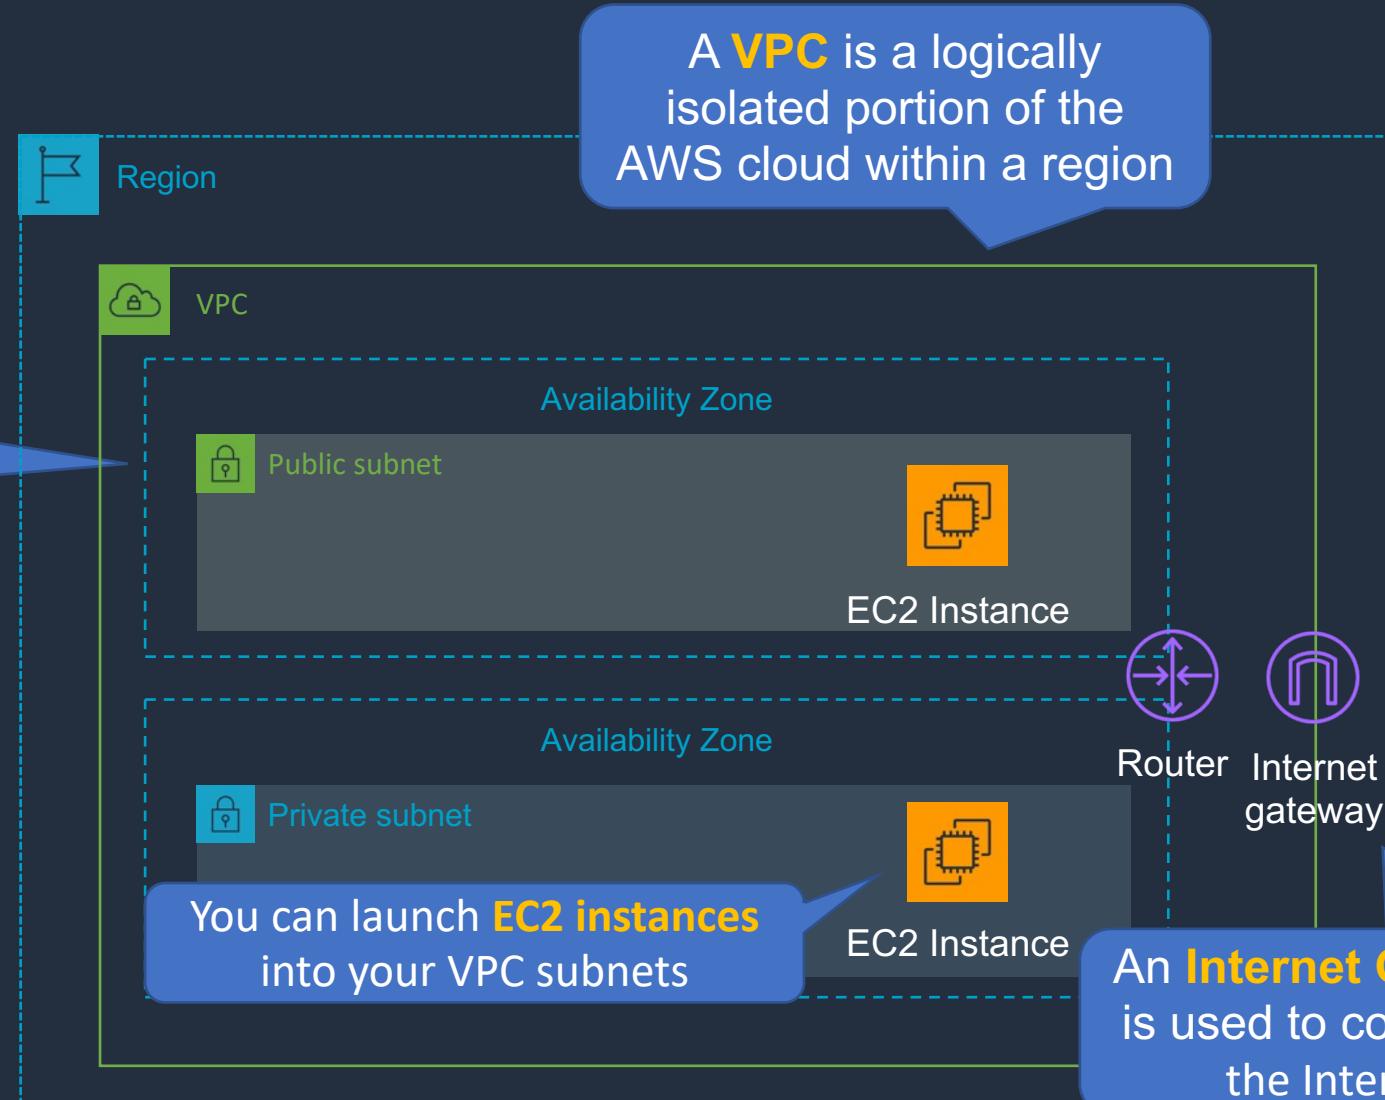
Amazon VPC, Networking, and Hybrid

Amazon Virtual Private Cloud (VPC)





Amazon VPC





Amazon VPC



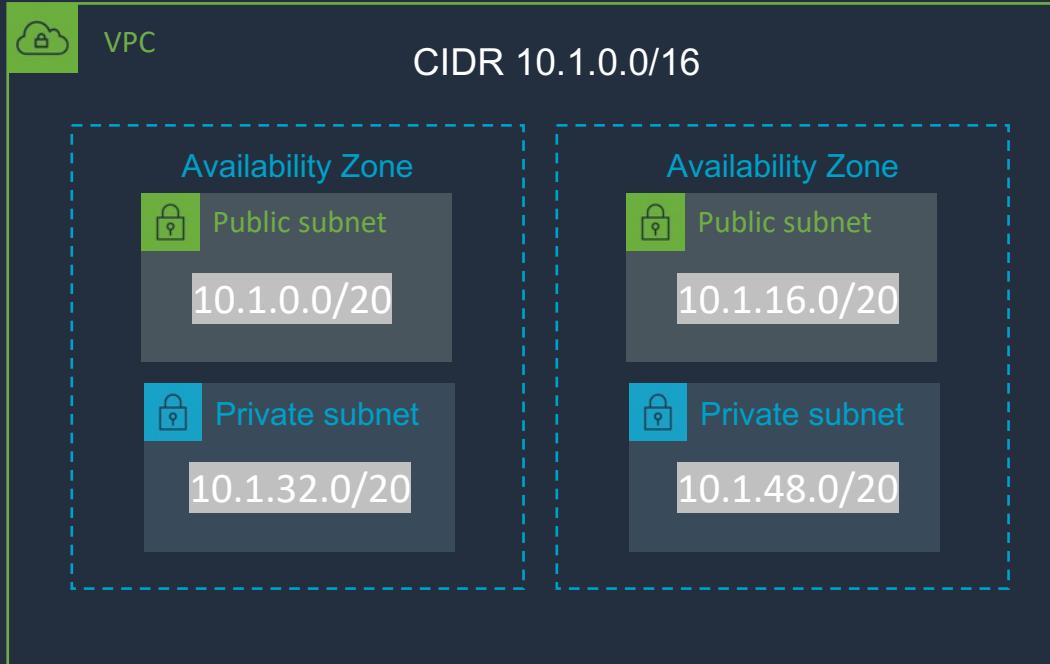
Region

Each **VPC** has a different block of IP addresses

CIDR stands for Classless Interdomain Routing



Each subnet has a block of **IP addresses** from the CIDR block



You can create **multiple VPCs** within each region



Amazon VPC

VPC Component	What it is
Virtual Private Cloud (VPC)	A logically isolated virtual network in the AWS cloud
Subnet	A segment of a VPC's IP address range where you can place groups of isolated resources
Internet Gateway/Egress-only Internet Gateway	The Amazon VPC side of a connection to the public Internet for IPv4/IPv6
Router	Routers interconnect subnets and direct traffic between Internet gateways, virtual private gateways, NAT gateways, and subnets
Peering Connection	Direct connection between two VPCs
VPC Endpoints	Private connection to public AWS services
NAT Instance	Enables Internet access for EC2 instances in private subnets managed by you
NAT Gateway	Enables Internet access for EC2 instances in private subnets (managed by AWS)
Virtual Private Gateway	The Amazon VPC side of a Virtual Private Network (VPN) connection
Customer Gateway	Customer side of a VPN connection
AWS Direct Connect	High speed, high bandwidth, private network connection from customer to aws
Security Group	Instance-level firewall
Network ACL	Subnet-level firewall



Amazon VPC

- A virtual private cloud (VPC) is a virtual network dedicated to your AWS account
- Analogous to having your own DC inside AWS
- It is logically isolated from other virtual networks in the AWS Cloud
- Provides complete control over the virtual networking environment including selection of IP ranges, creation of subnets, and configuration of route tables and gateways
- You can launch your AWS resources, such as Amazon EC2 instances, into your VPC



Amazon VPC

- When you create a VPC, you must specify a range of IPv4 addresses for the VPC in the form of a Classless Inter-Domain Routing (CIDR) block; for example, 10.0.0.0/16
- A VPC spans all the Availability Zones in the region
- You have full control over who has access to the AWS resources inside your VPC
- By default you can create up to 5 VPCs per region
- A default VPC is created in each region with a subnet in each AZ

Create a Custom VPC

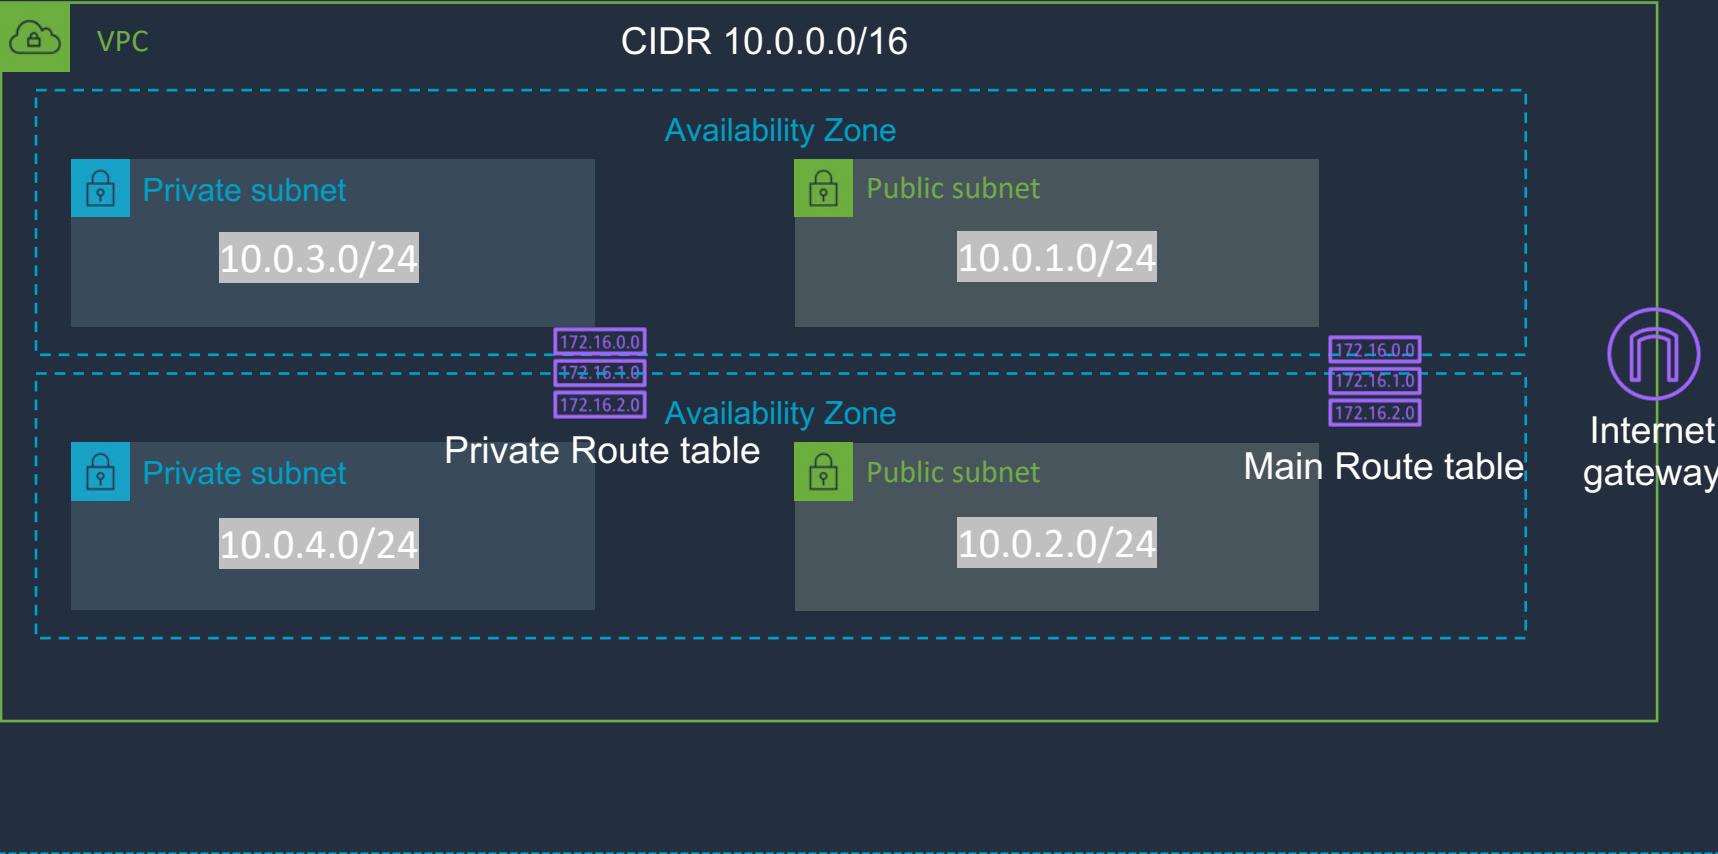




Create a Custom VPC



Region



Main Route Table

Destination	Target
10.0.0.0/16	Local
0.0.0.0/0	igw-id

Private Route Table

Destination	Target
10.0.0.0/16	Local

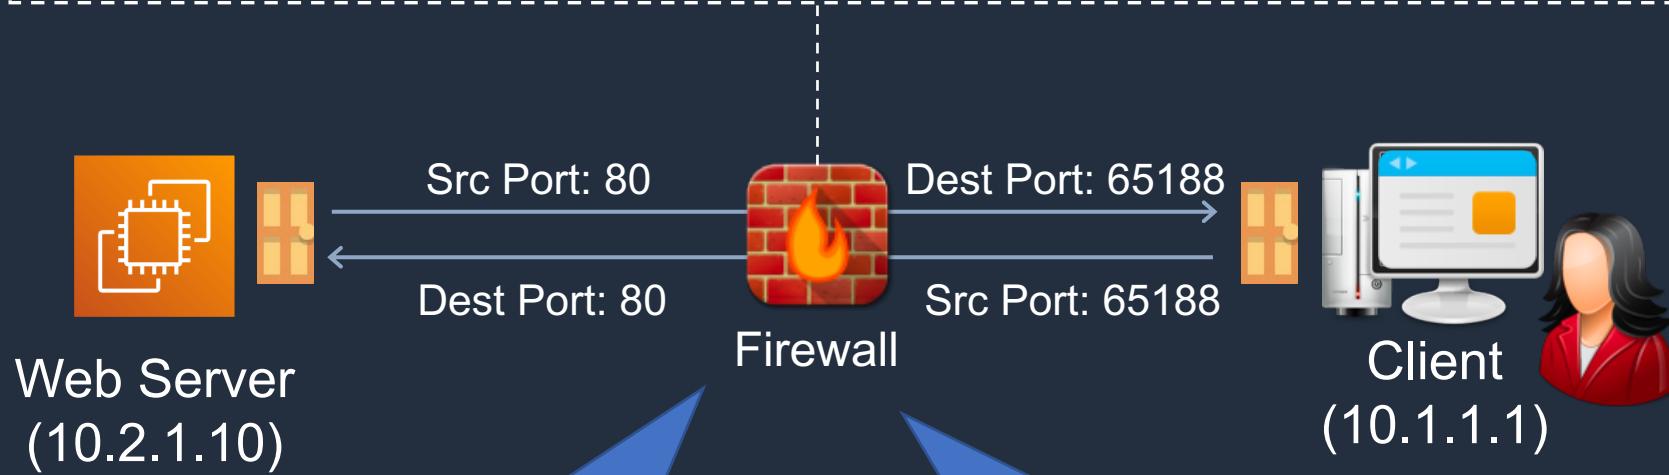
Security Groups and Network ACLs





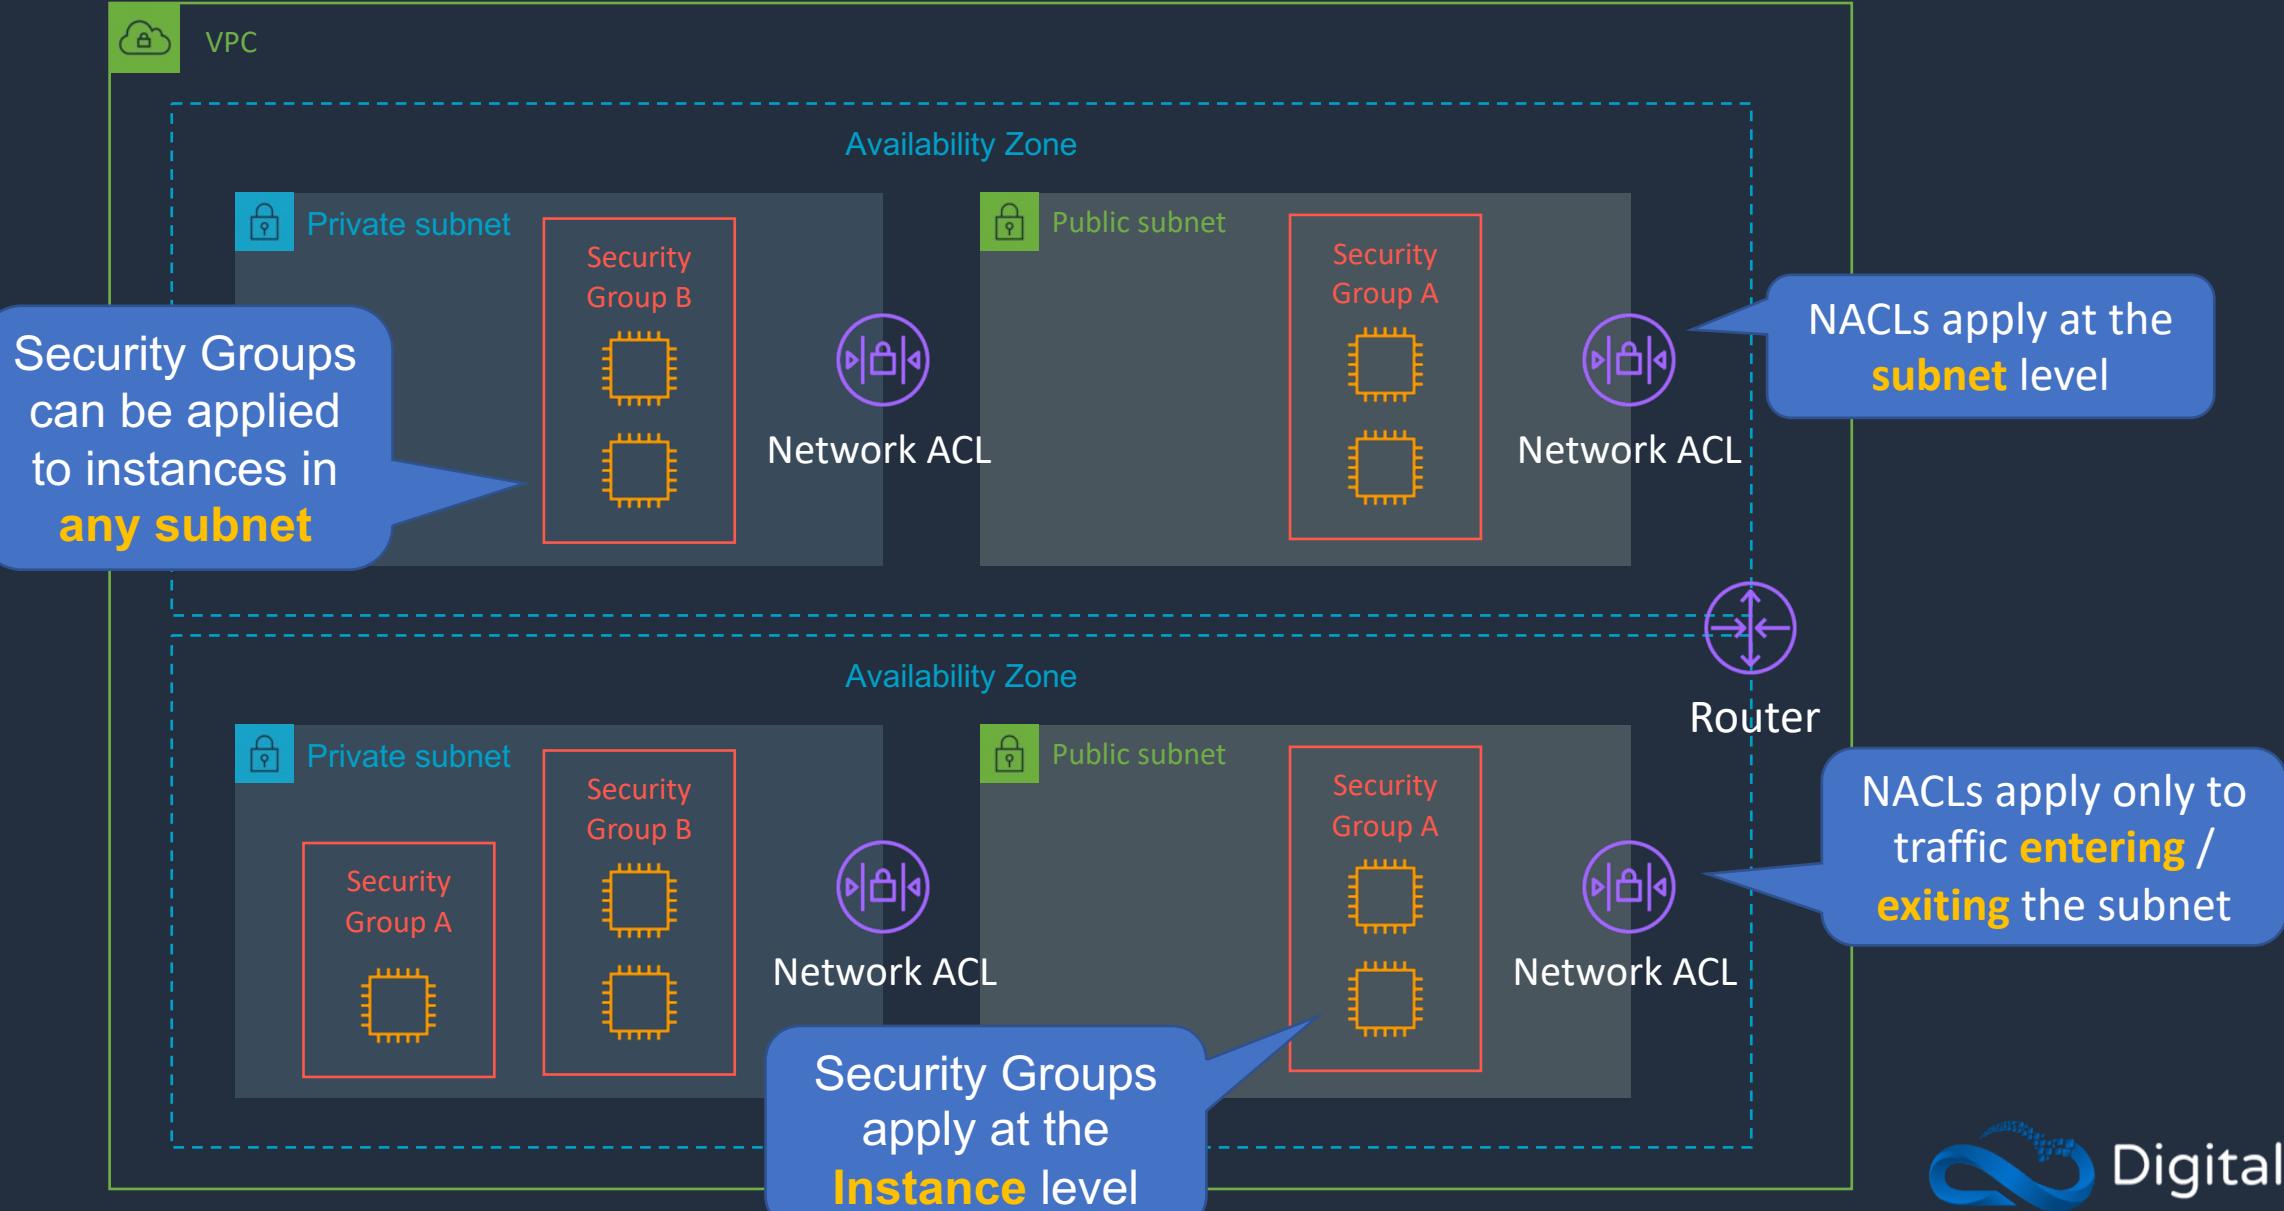
Stateful vs Stateless Firewalls

PROTOCOL	SOURCE IP	DESTINATION IP	SOURCE PORT	DESTINATION PORT
HTTP	10.1.1.1	10.2.1.10	65188	80
HTTP	10.2.1.10	10.1.1.1	80	65188





Security Groups and Network ACLs





Security Group Rules

Security groups support
allow rules only

Inbound rules

Type	Protocol	Port range	Source
SSH	TCP	22	0.0.0.0/0
RDP	TCP	3389	0.0.0.0/0
RDP	TCP	3389	::/0
HTTPS	TCP	443	0.0.0.0/0
HTTPS	TCP	443	::/0
All ICMP - IPv4	ICMP	All	0.0.0.0/0

Separate rules
are defined for
outbound traffic

A source can be an **IP
address or security
group ID**



Network ACLs

Inbound Rules

Rule #	Type	Protocol	Port Range	Source	Allow / Deny
100	ALL Traffic	ALL	ALL	0.0.0.0/0	ALLOW
101	ALL Traffic	ALL	ALL	::/0	ALLOW
*	ALL Traffic	ALL	ALL	0.0.0.0/0	DENY
*	ALL Traffic	ALL	ALL	::/0	DENY

Outbound Rules

Rule #	Type	Protocol	Port Range	Destination	
100	ALL Traffic	ALL	ALL	0.0.0.0/0	ALLOW
101	ALL Traffic	ALL	ALL	::/0	ALLOW
*	ALL Traffic	ALL	ALL	0.0.0.0/0	DENY
*	ALL Traffic	ALL	ALL	::/0	DENY

NACLs have an explicit deny

Rules are processed in order

Configure Security Groups and NACLs



Public, Private and Elastic IP Addresses



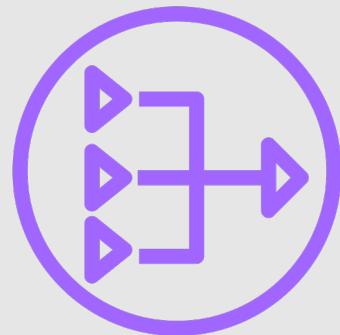
→ Public, Private and Elastic IP addresses

Name	Description
Public IP address	<p>Lost when the instance is stopped</p> <p>Used in Public Subnets</p> <p>No charge</p> <p>Associated with a private IP address on the instance</p> <p>Cannot be moved between instances</p>
Private IP address	<p>Retained when the instance is stopped</p> <p>Used in Public and Private Subnets</p>
Elastic IP address	<p>Static Public IP address</p> <p>You are charged if not used</p> <p>Associated with a private IP address on the instance</p> <p>Can be moved between instances and Elastic Network Adapters</p>

Working with IP Addresses

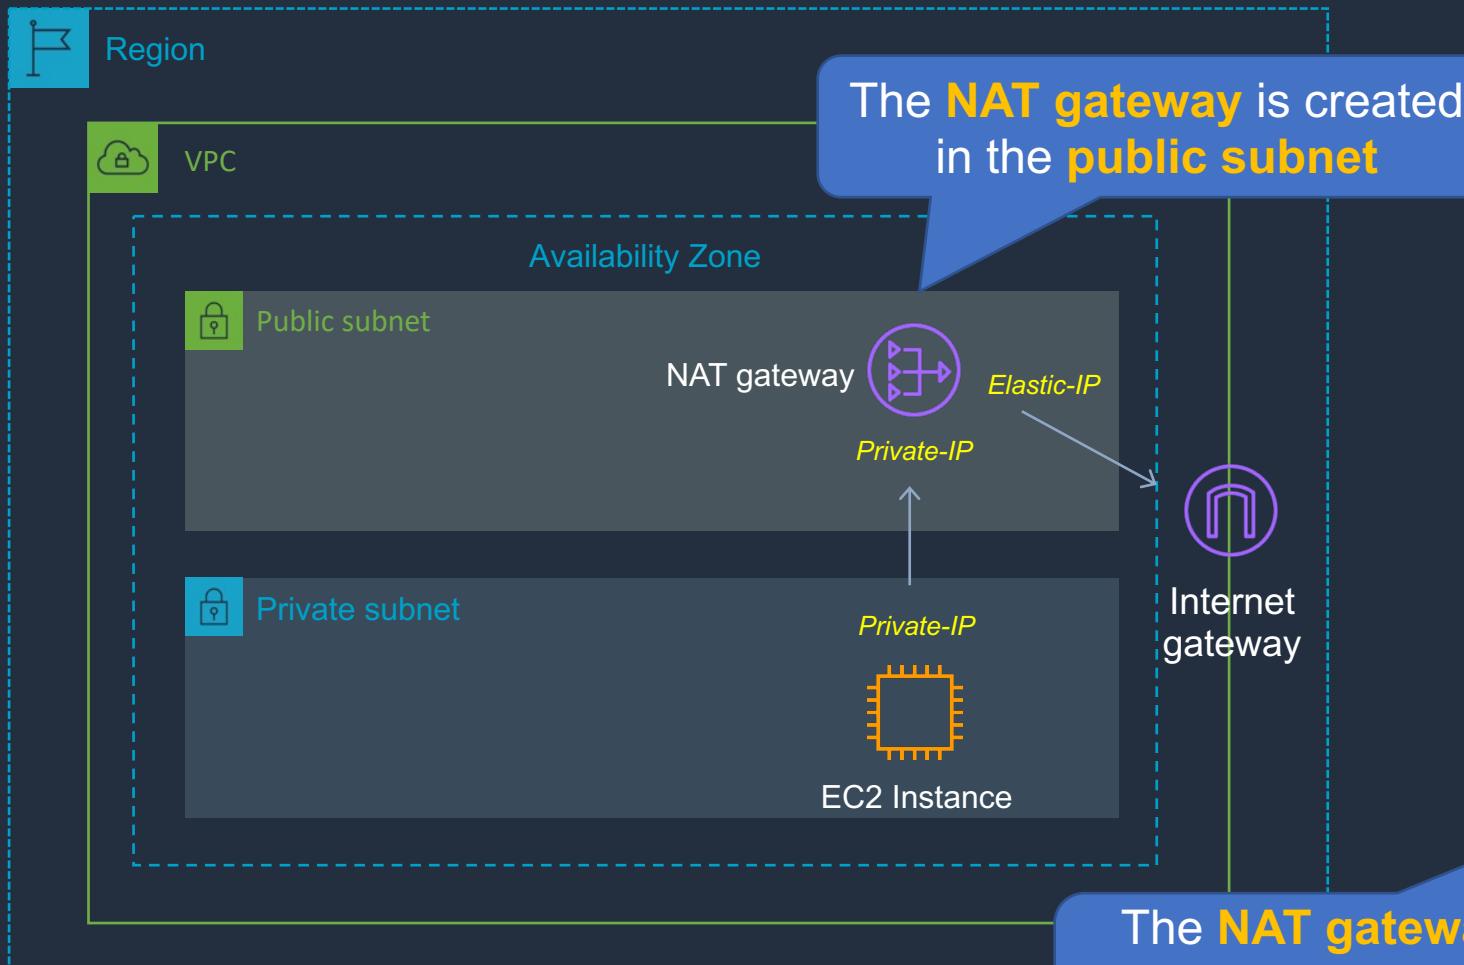


NAT Gateways and NAT Instances



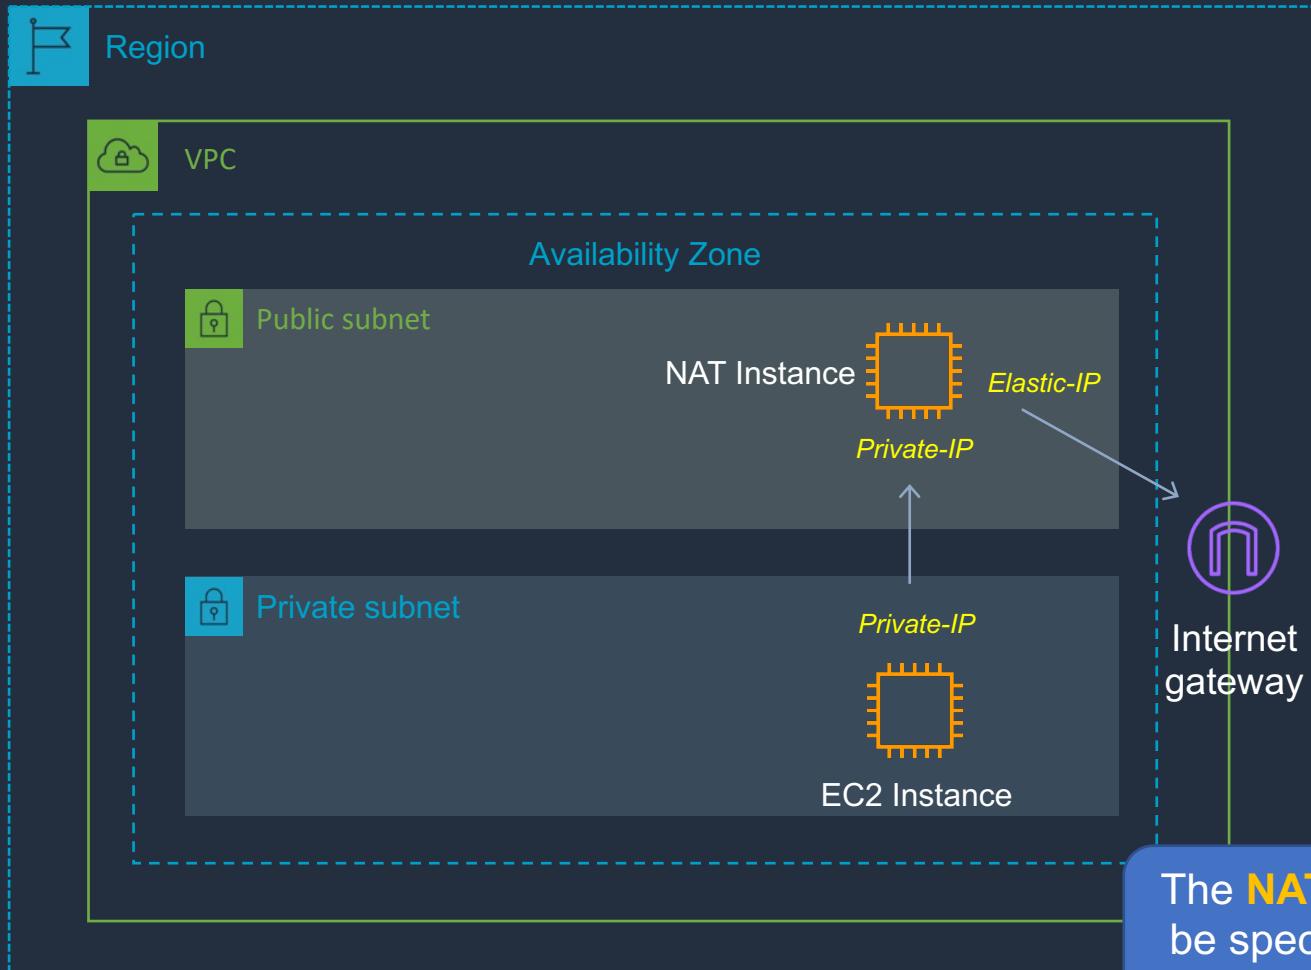


NAT Gateways





NAT Instances



Main Route Table

Destination	Target
10.0.0.0/16	Local
0.0.0.0/0	igw-id

Private Route Table

Destination	Target
10.0.0.0/16	Local
0.0.0.0/0	nat-instance-id

The **NAT instance ID** must
be specified in the **private
subnet RT**



NAT Instance vs NAT Gateway

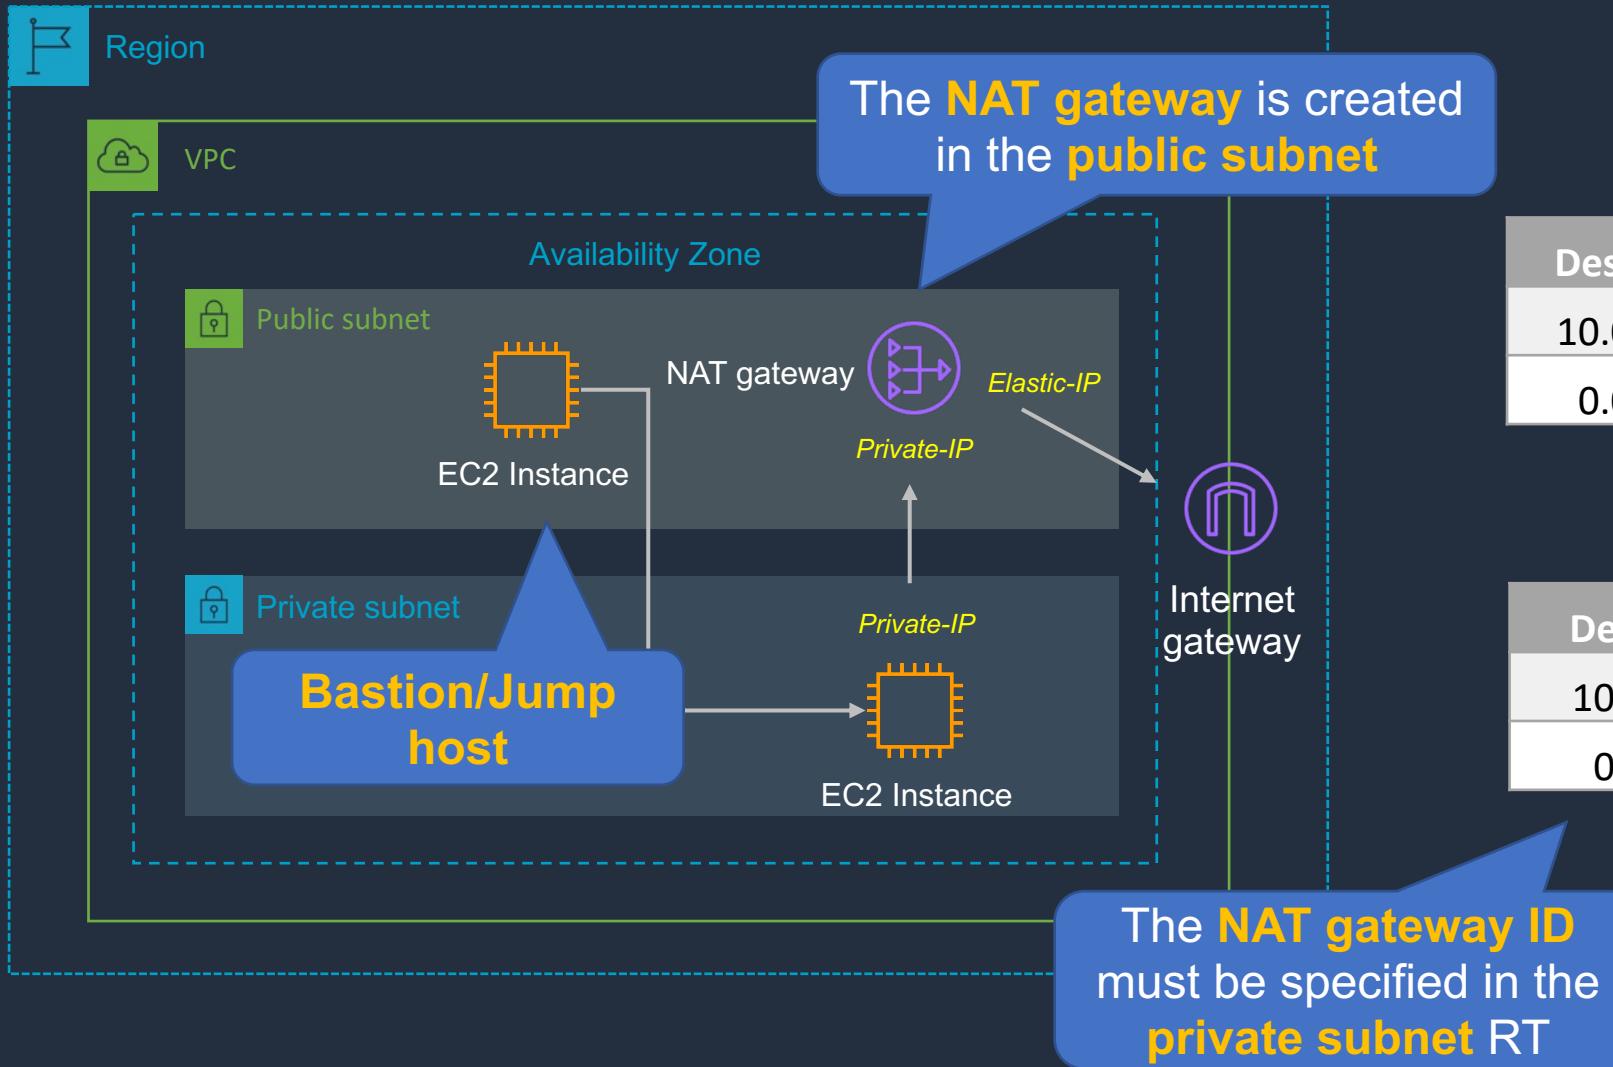
NAT Instance	NAT Gateway
Managed by you (e.g. software updates)	Managed by AWS
Scale up (instance type) manually and use enhanced networking	Elastic scalability up to 45 Gbps
No high availability – scripted/auto-scaled HA possible using multiple NATs in multiple subnets	Provides automatic high availability within an AZ and can be placed in multiple AZs

Deploy a NAT Gateway





NAT Gateways

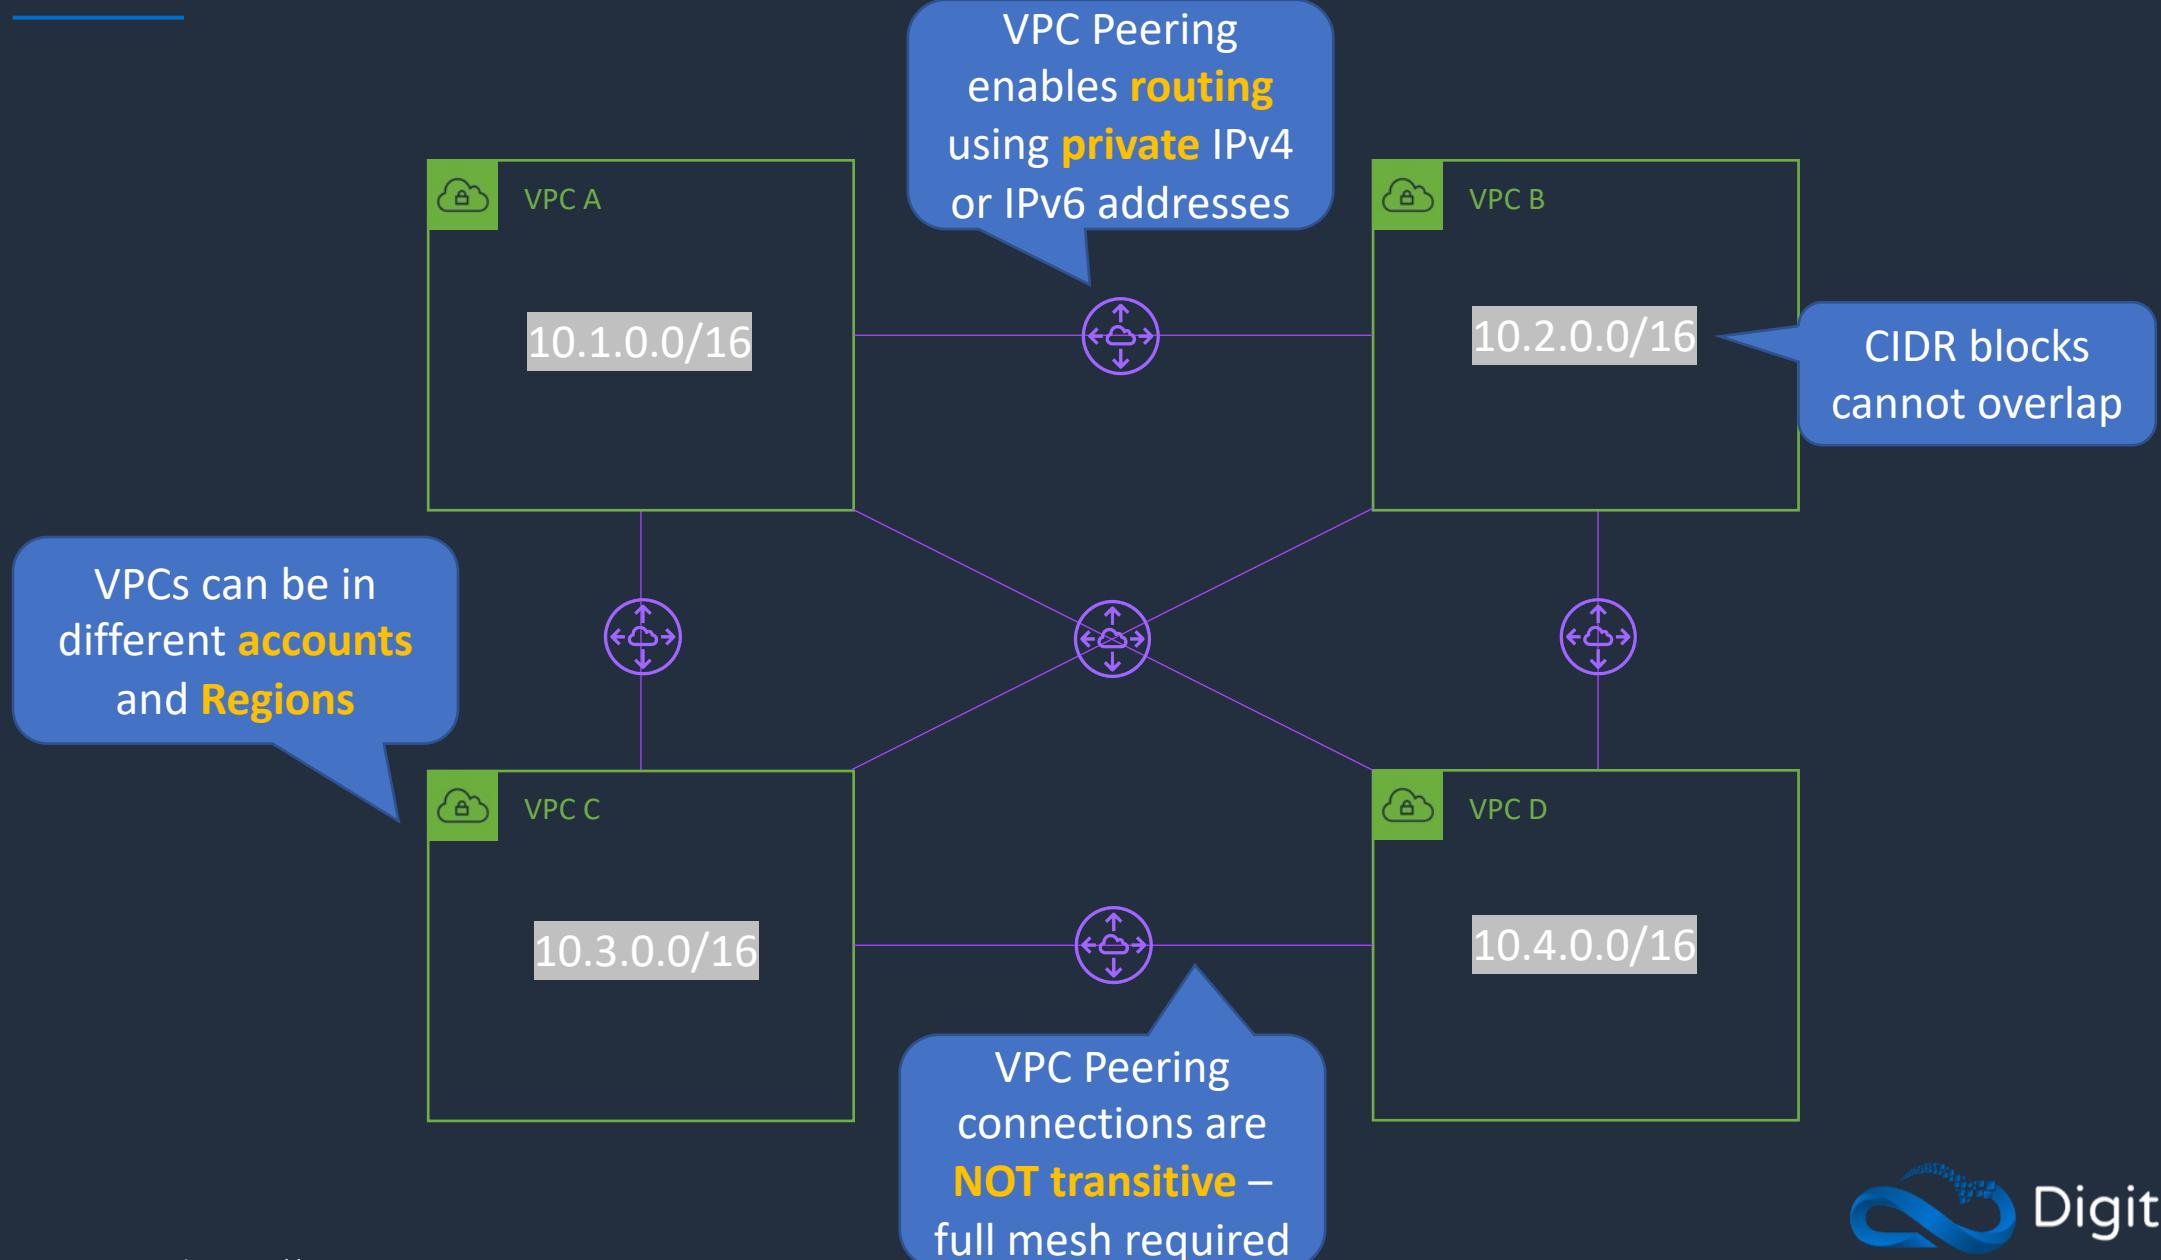


Amazon VPC Peering

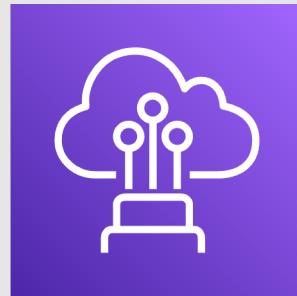




VPC Peering

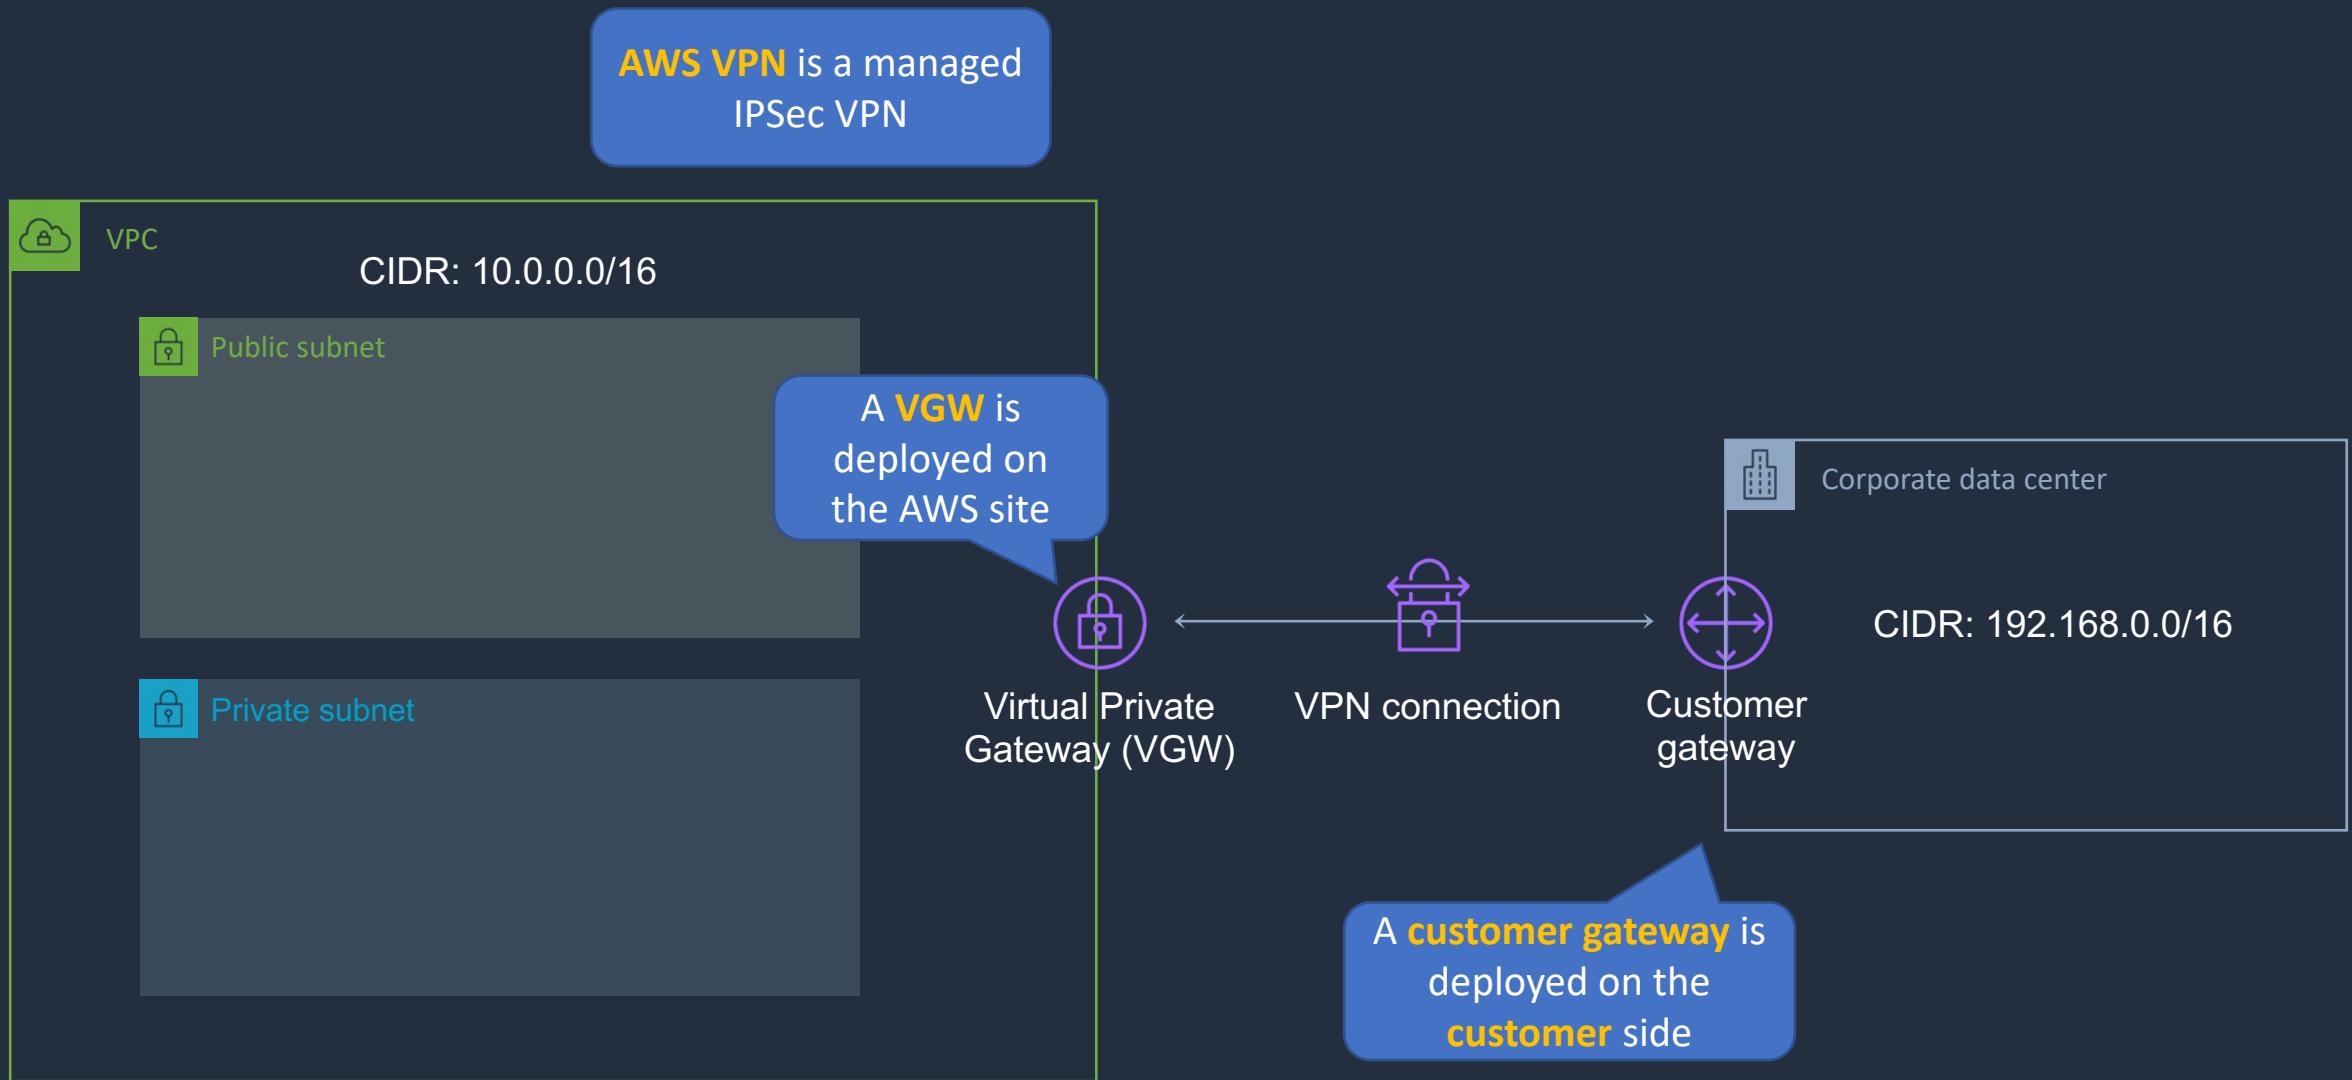


Amazon VPN and AWS Direct Connect



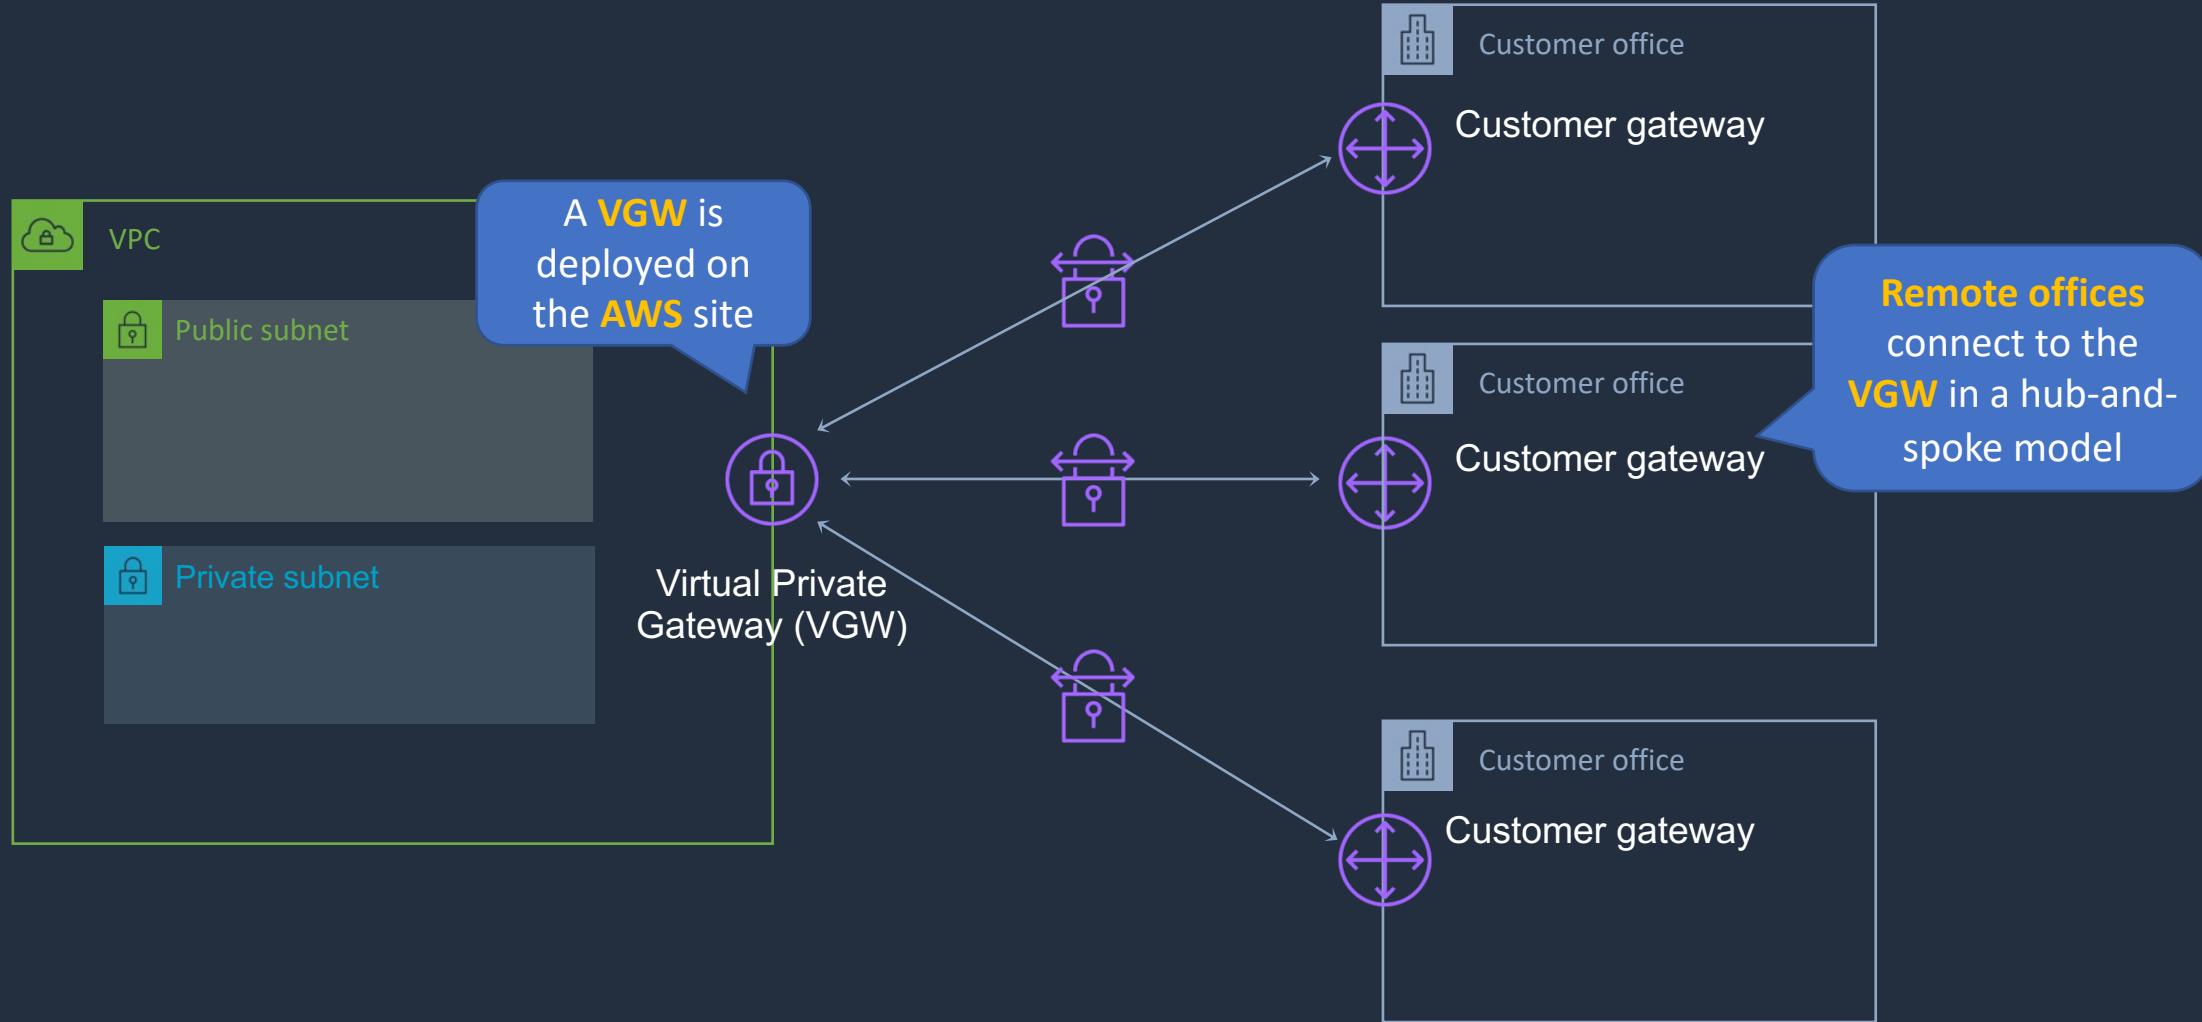


AWS Site-to-Site VPN





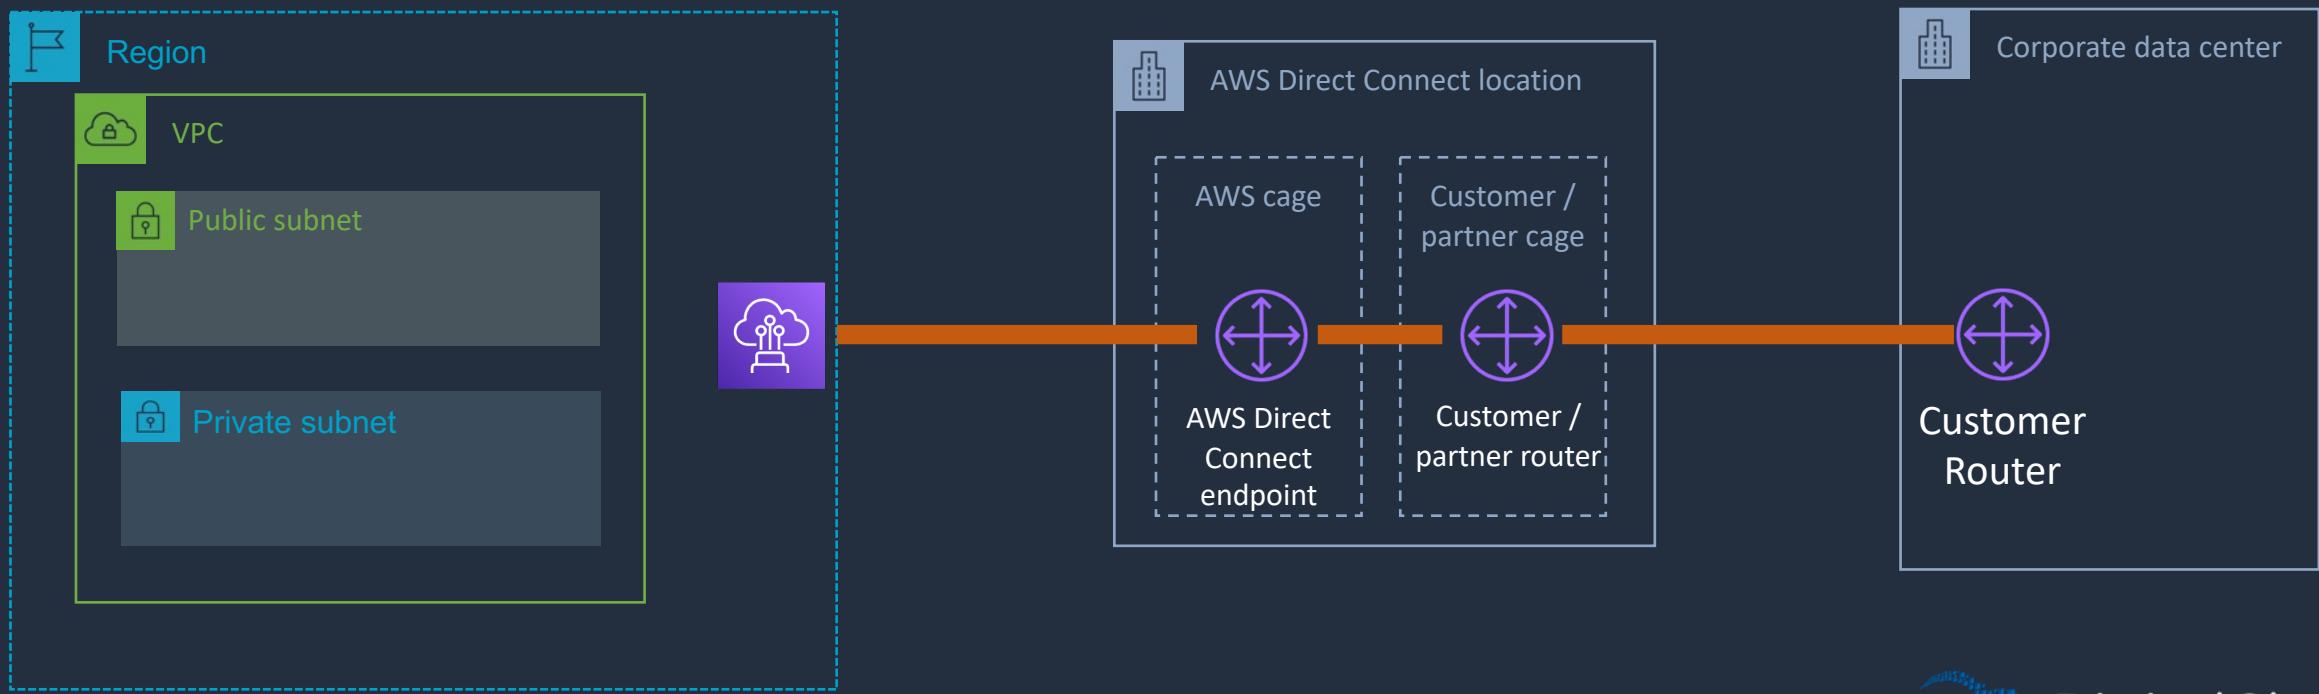
AWS VPN CloudHub



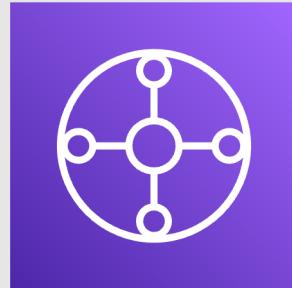


AWS Direct Connect

- **Private** connectivity between AWS and your data center / office
- Consistent network experience – increased **speed/latency** & **bandwidth/throughput**
- Lower costs for organizations that transfer **large** volumes of data



AWS Transit Gateway





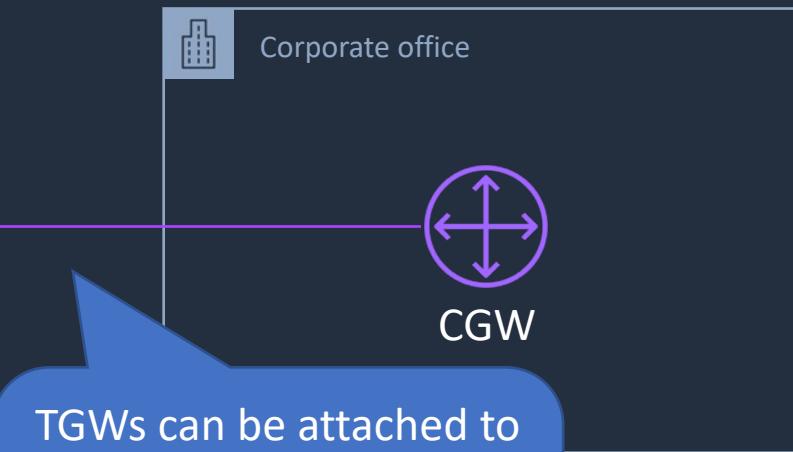
AWS Transit Gateway



Transit Gateway is a network transit hub that interconnects **VPCs** and **on-premises** networks

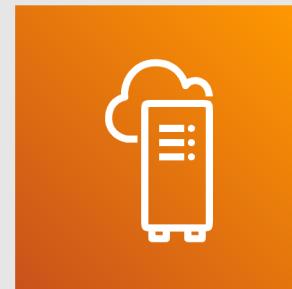


VPCs are attached to Transit Gateway



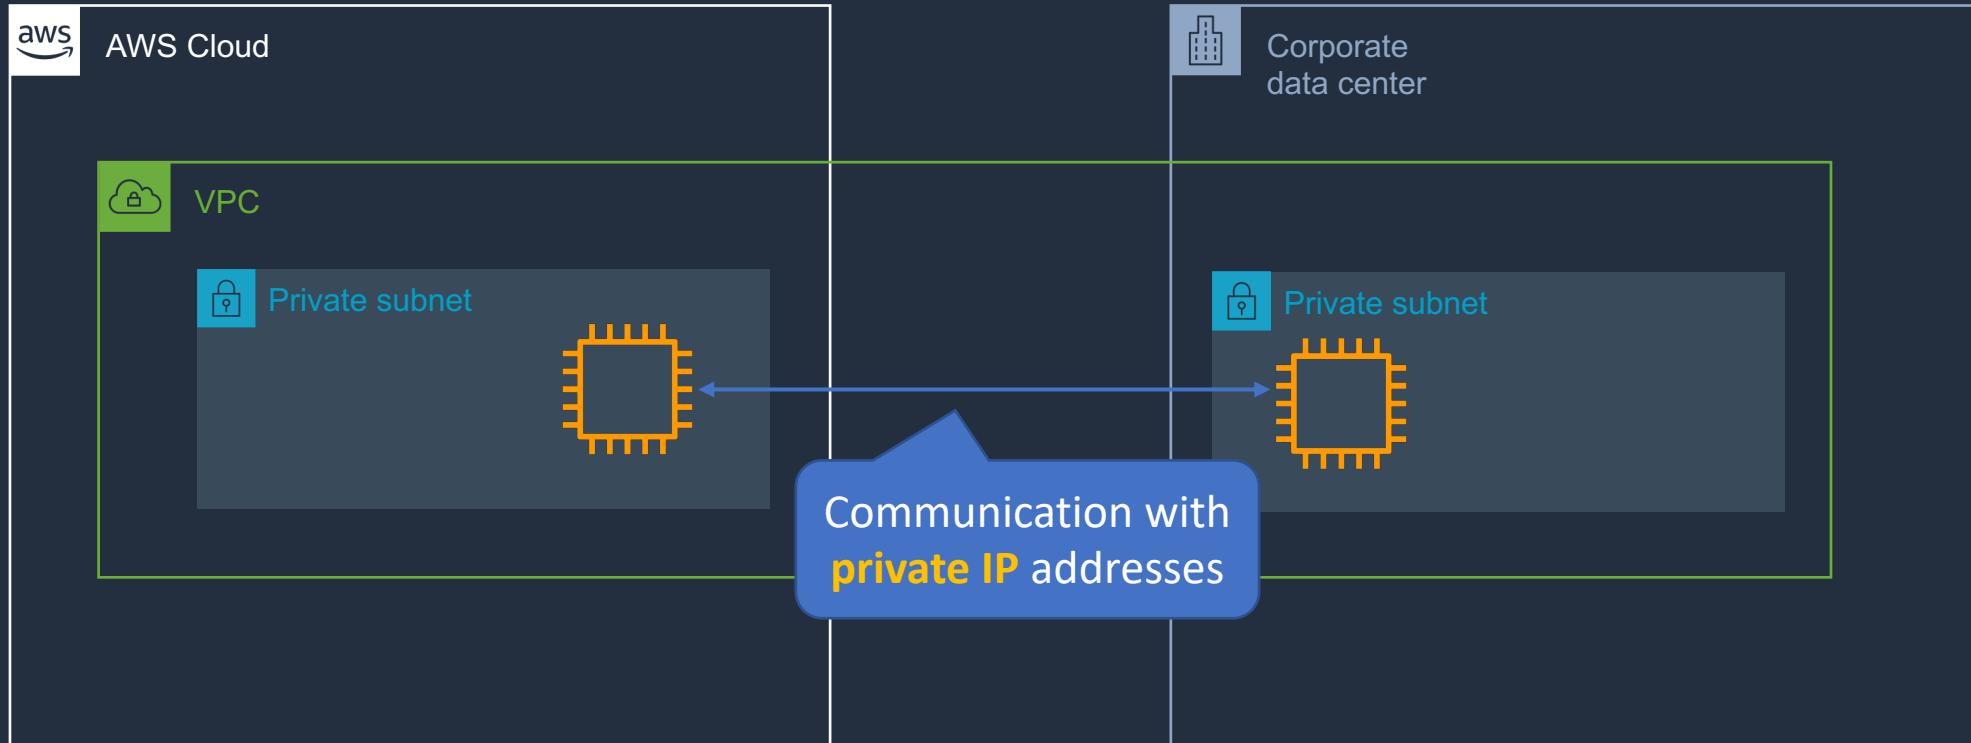
TGWs can be attached to **VPNs, Direct Connect Gateways, 3rd party appliances and TGWs** in other Regions/accounts

AWS Outposts





AWS Outposts





AWS Outposts

Services you can run on AWS Outposts include:

- Amazon EC2
- Amazon EBS
- Amazon S3
- Amazon VPC
- Amazon ECS/EKS
- Amazon RDS
- Amazon EMR

SECTION 10

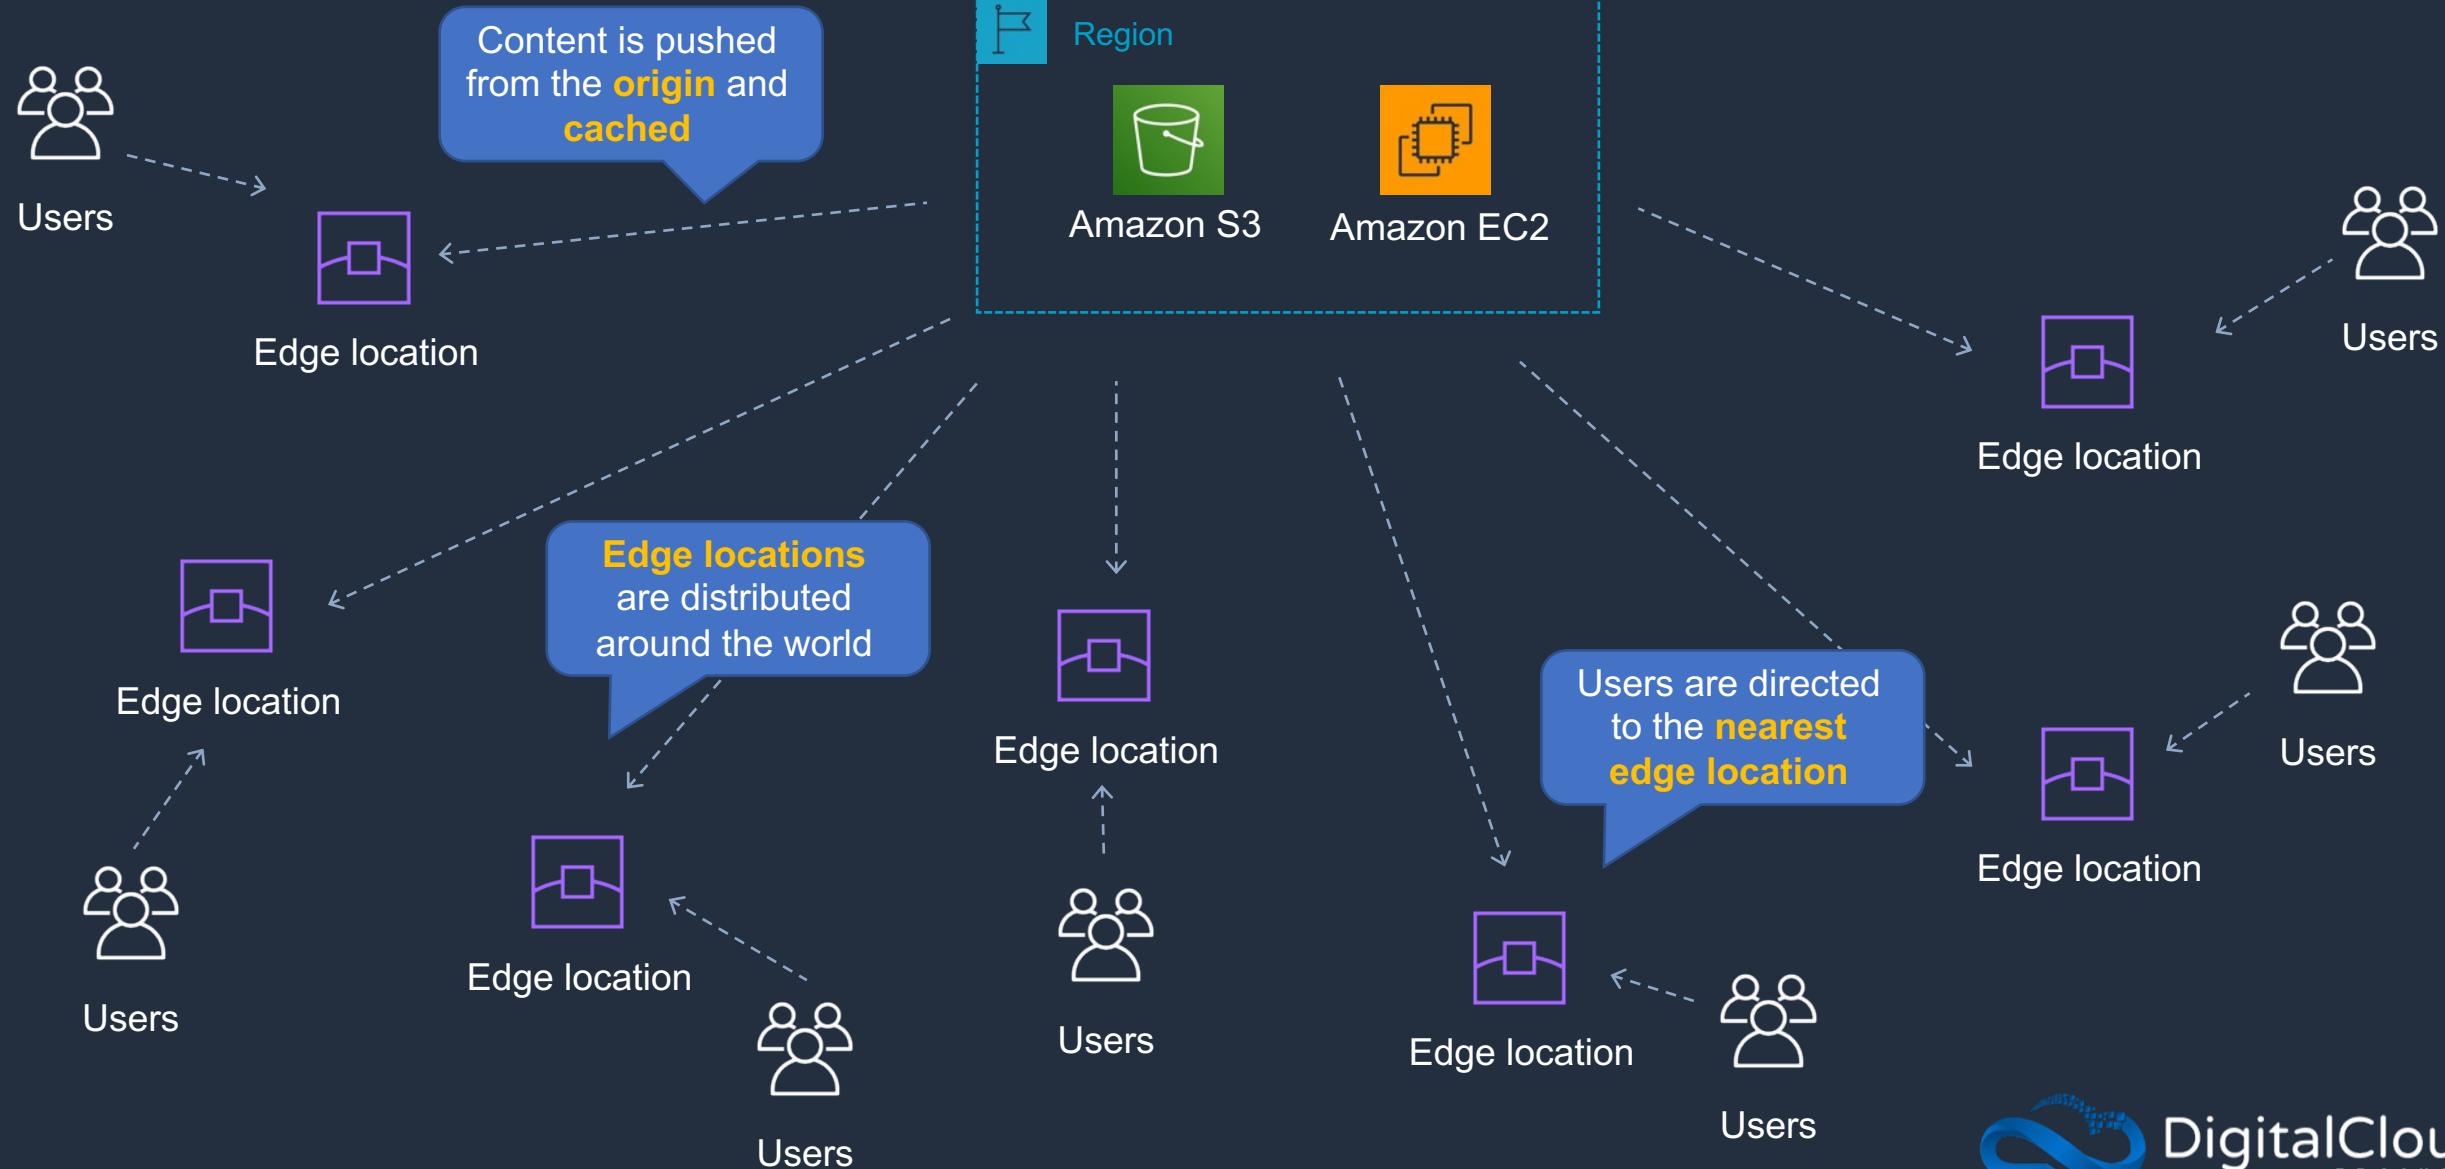
Deployment and Automation

Caching and Amazon CloudFront





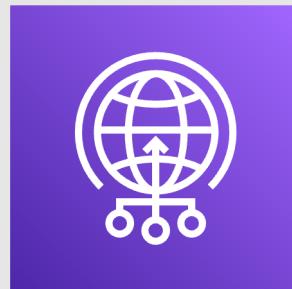
Amazon CloudFront



S3 Static Website with CloudFront



AWS Global Accelerator



AWS Global Accelerator



Users in US

Connect via Edge Location

Requests are routed to the **optimal endpoint**



us-east-1



Edge location

User traffic ingresses using the closest **Edge Location**



Addresses:
51.45.2.12
53.58.31.89

Static **anycast** IP addresses

AWS Global Network

Traffic traverses the **AWS global network**



ap-southeast-2



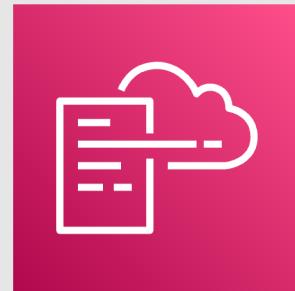
Users are **redirected** to another **endpoint**



AWS Global Accelerator vs CloudFront

- Both use the AWS global network and edge locations
- CloudFront improves performance for cacheable content and dynamic content
- GA improves performance for a wide range of applications over TCP and UDP
- GA proxies connections to applications in one or more AWS Regions
- GA provides failover between AWS Regions

AWS CloudFormation





AWS CloudFormation

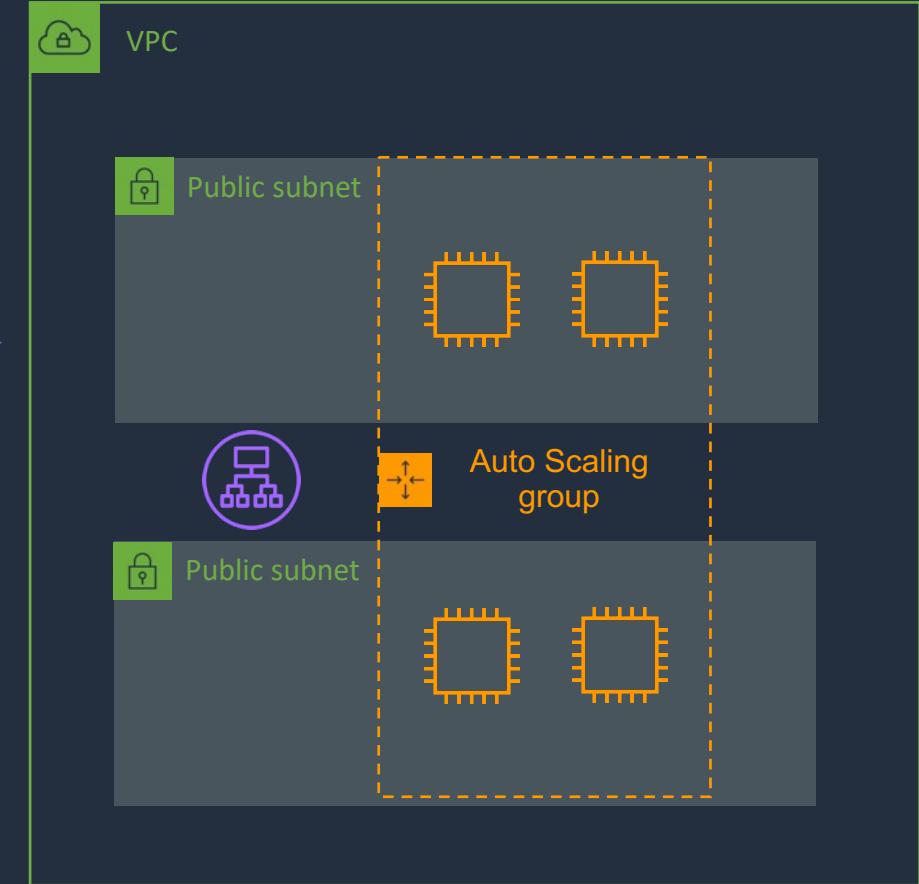
Infrastructure patterns are defined in a **template** file using **code**



CloudFormation **builds** your infrastructure according to the **template**

AWS CloudFormation

```
1 "AWSTemplateFormatVersion": "2010-09-09",
2
3 "Description": "AWS CloudFormation Sample Template WordPress_Multi_AZ: WordPress is web
4
5 "Parameters": {
6   "VpcId": {
7     "Type": "AWS::EC2::VPC::Id",
8     "Description": "VpcId of your existing Virtual Private Cloud (VPC)",
9     "ConstraintDescription": "must be the VPC Id of an existing Virtual Private Cloud."
10 },
11
12 "Subnets": {
13   "Type": "List<AWS::EC2::Subnet::Id>",
14   "Description": "The list of SubnetIds in your Virtual Private Cloud (VPC)",
15   "ConstraintDescription": "must be a list of at least two existing subnets associated
16 },
```





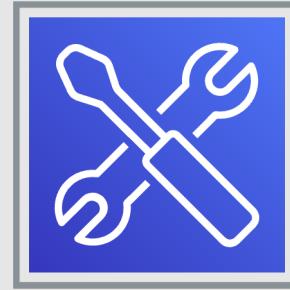
AWS CloudFormation

- Infrastructure is provisioned consistently, with fewer mistakes (human error)
- Less time and effort than configuring resources manually
- Free to use (you're only charged for the resources provisioned)
- A template is a YAML or JSON template used to describe the end-state of the infrastructure you are either provisioning or changing
- CloudFormation creates a Stack based on the template
- Can easily rollback and delete the entire stack as well

Create CloudFormation Stack



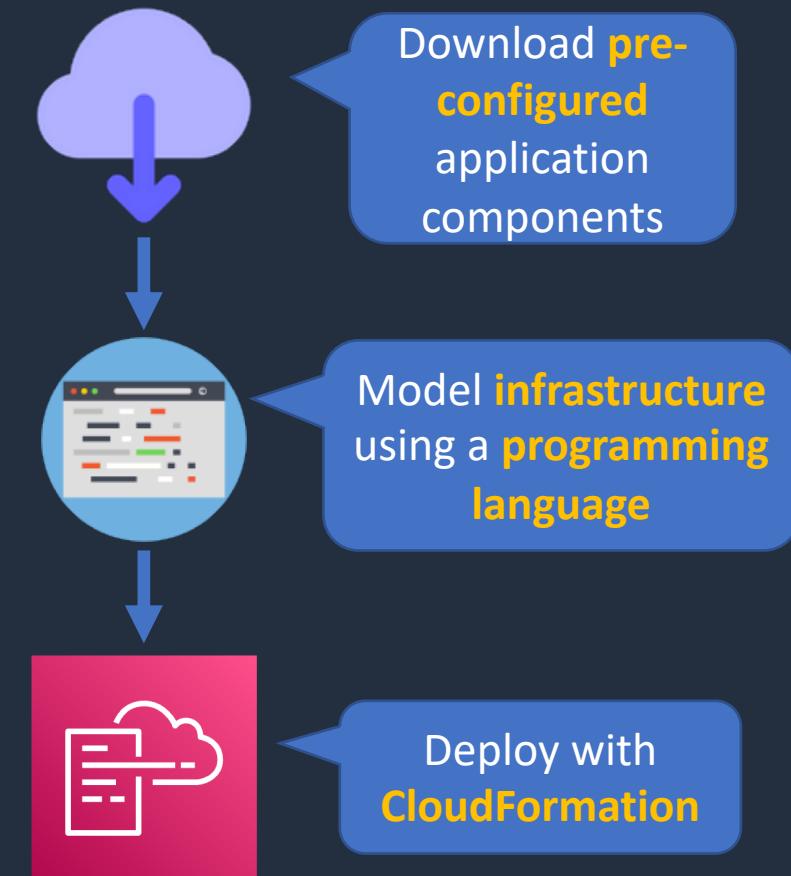
AWS Cloud Development Kit (CDK)





AWS Cloud Development Kit

- Open-source software development framework to define your cloud application resources using **familiar programming languages**
- Preconfigures cloud resources with proven defaults using **constructs**
- Provisions your resources using **AWS CloudFormation**
- Enables you to model application infrastructure using TypeScript, Python, Java, and .NET
- Use existing IDE, testing tools, and workflow patterns

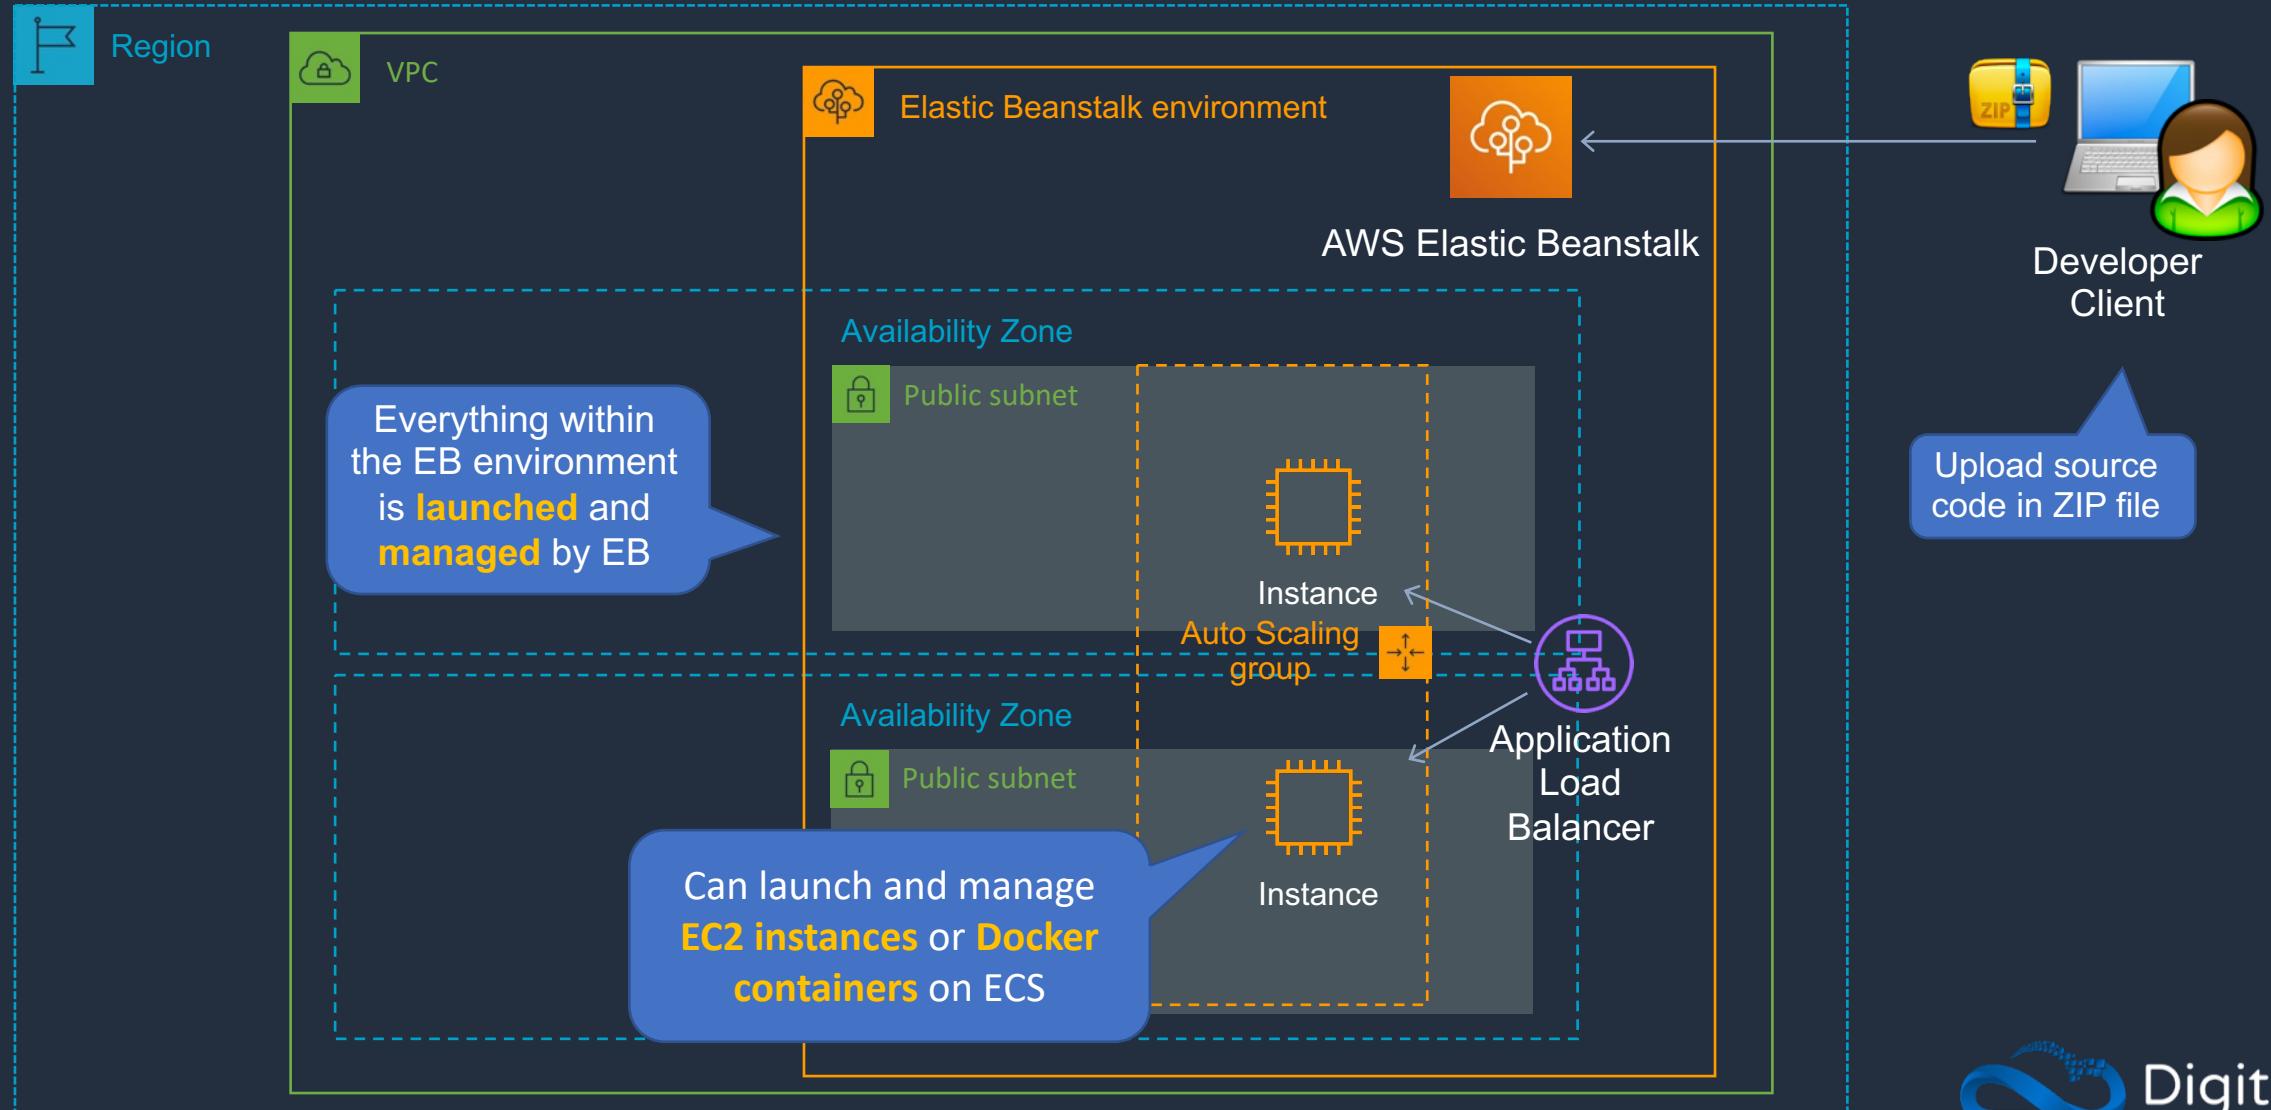


AWS Elastic Beanstalk





AWS Elastic Beanstalk





AWS Elastic Beanstalk

- Supports Java, .NET, PHP, Node.js, Python, Ruby, Go, and Docker web applications
- Integrates with VPC
- Integrates with IAM
- Can provision most database instances
- Allows full control of the underlying resources
- Code is deployed using a WAR file or Git repository



AWS Elastic Beanstalk

CloudFormation	Elastic Beanstalk
“Template-driven provisioning”	“Web apps made easy”
Deploys infrastructure using code	Deploys applications on EC2 (PaaS)
Can be used to deploy almost any AWS service	Deploys web applications based on Java, .NET, PHP, Node.js, Python, Ruby, Go, and Docker
Uses JSON or YAML template files	Uses ZIP or WAR files (or Git)
Similar to Terraform	Similar to Google App Engine

Deploy a Web Application

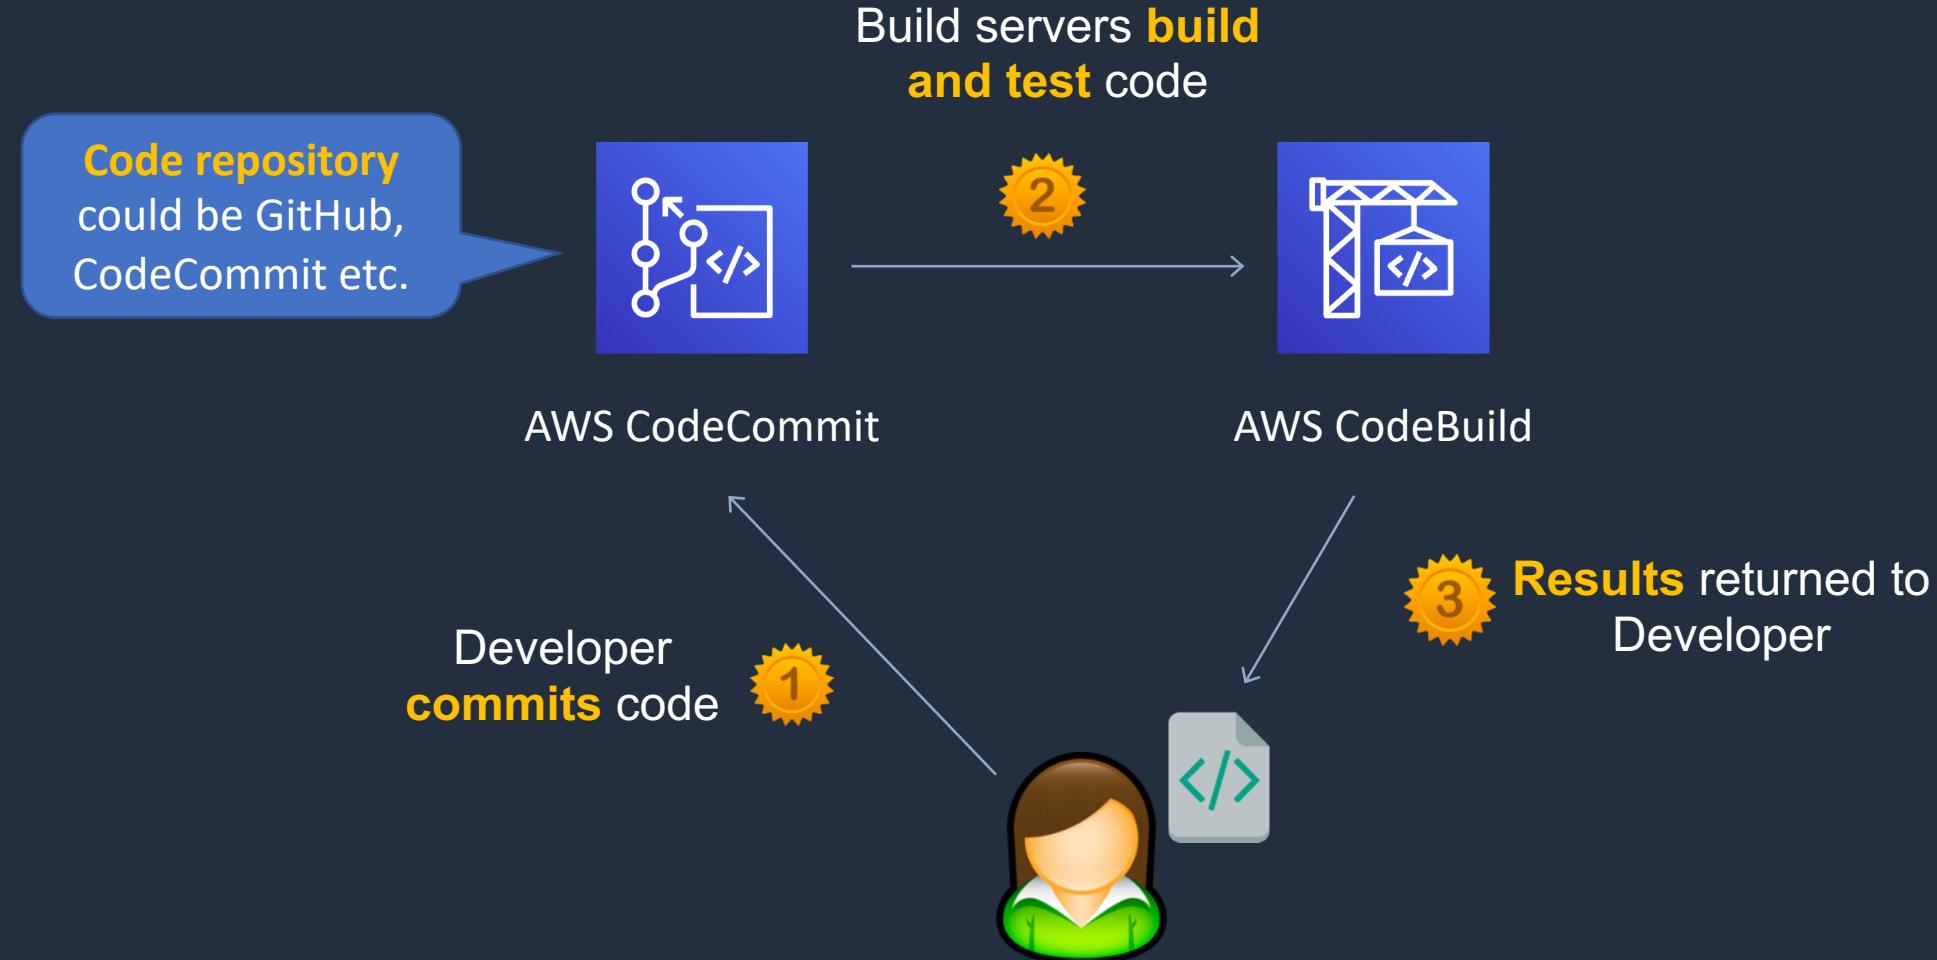


AWS Developer Tools (Code*)



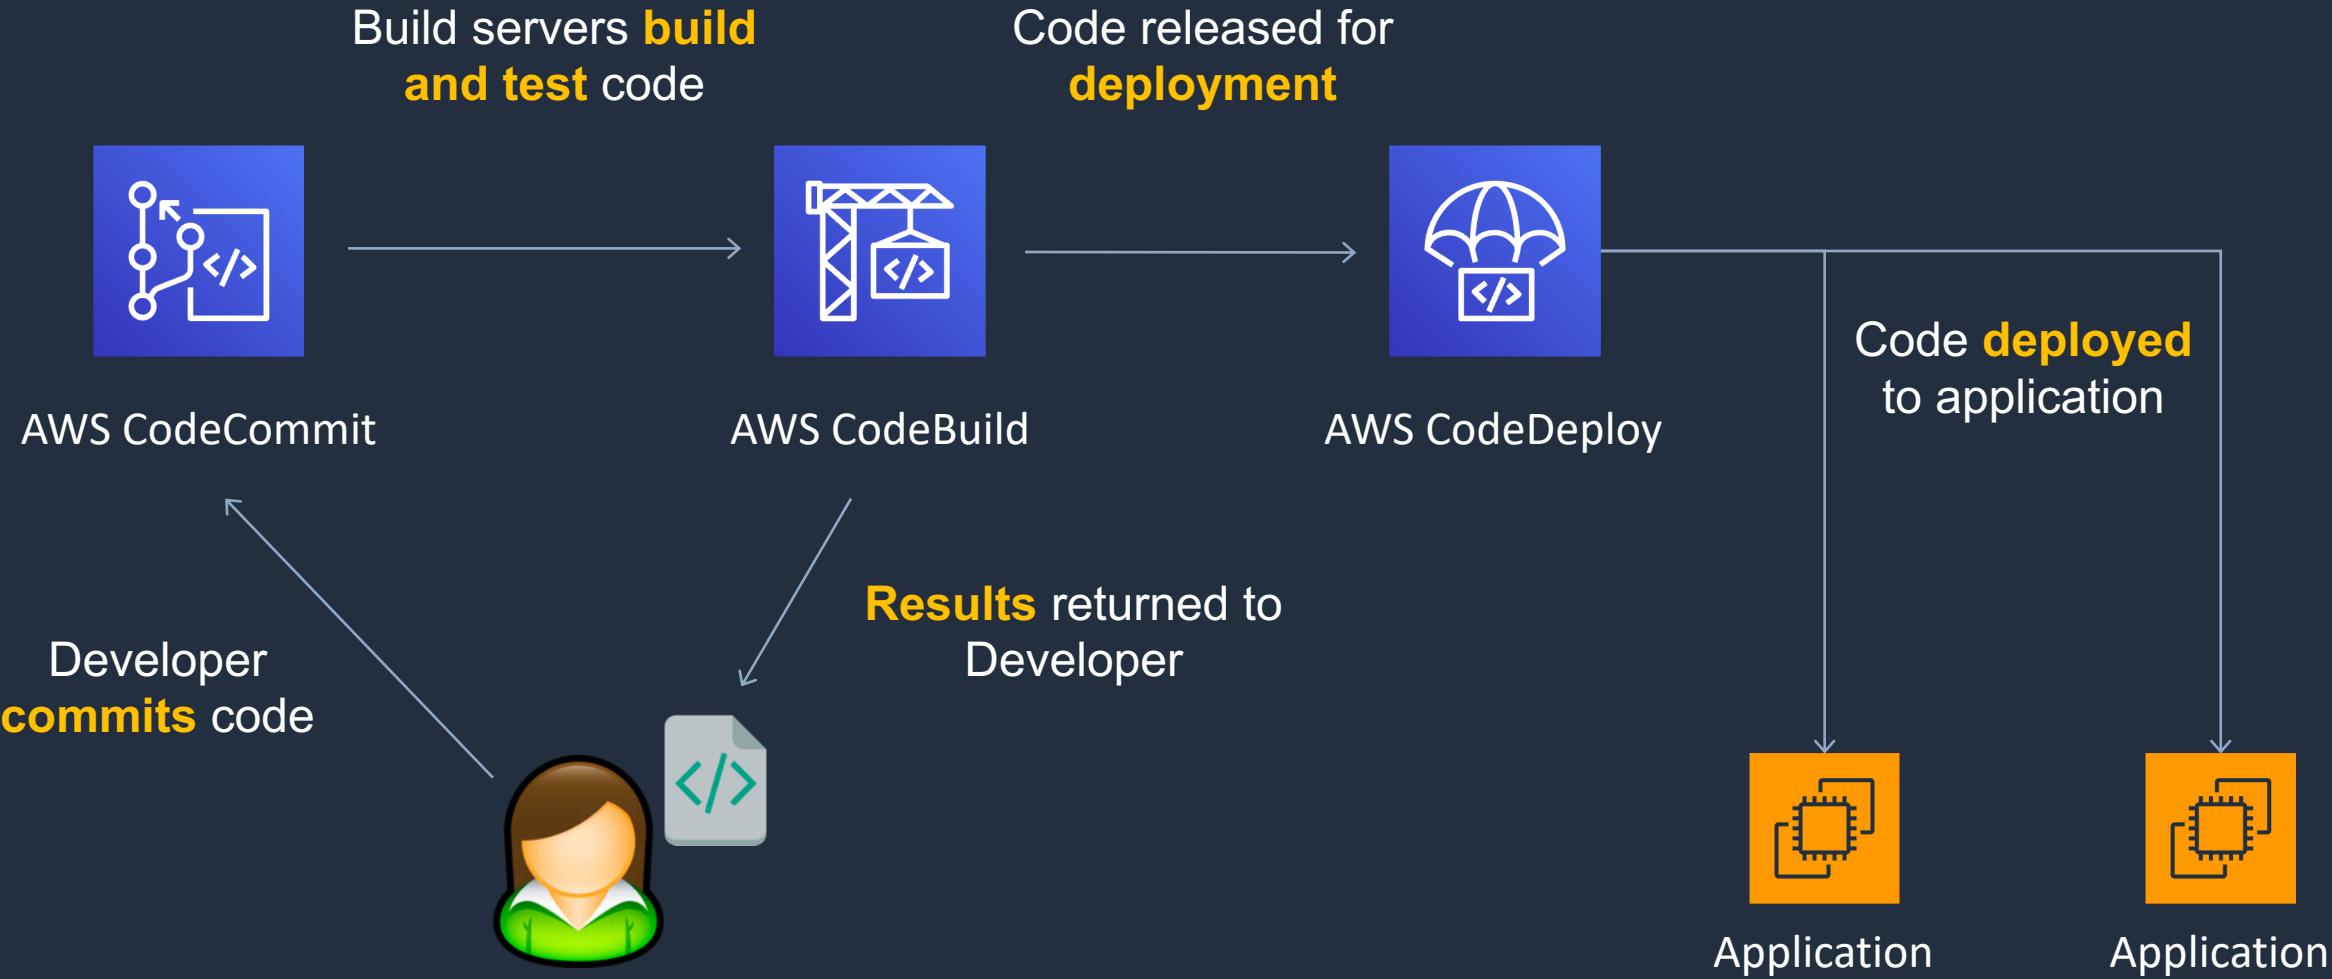


Continuous Integration





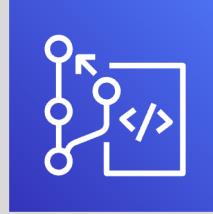
Continuous Integration and Continuous Delivery





Continuous Integration and Continuous Delivery

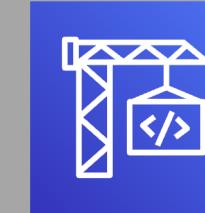
CODE



AWS CodeCommit

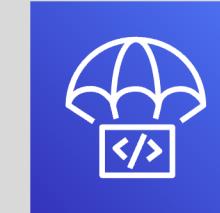
BUILD & TEST

AWS CodePipeline



AWS CodeBuild

DEPLOY

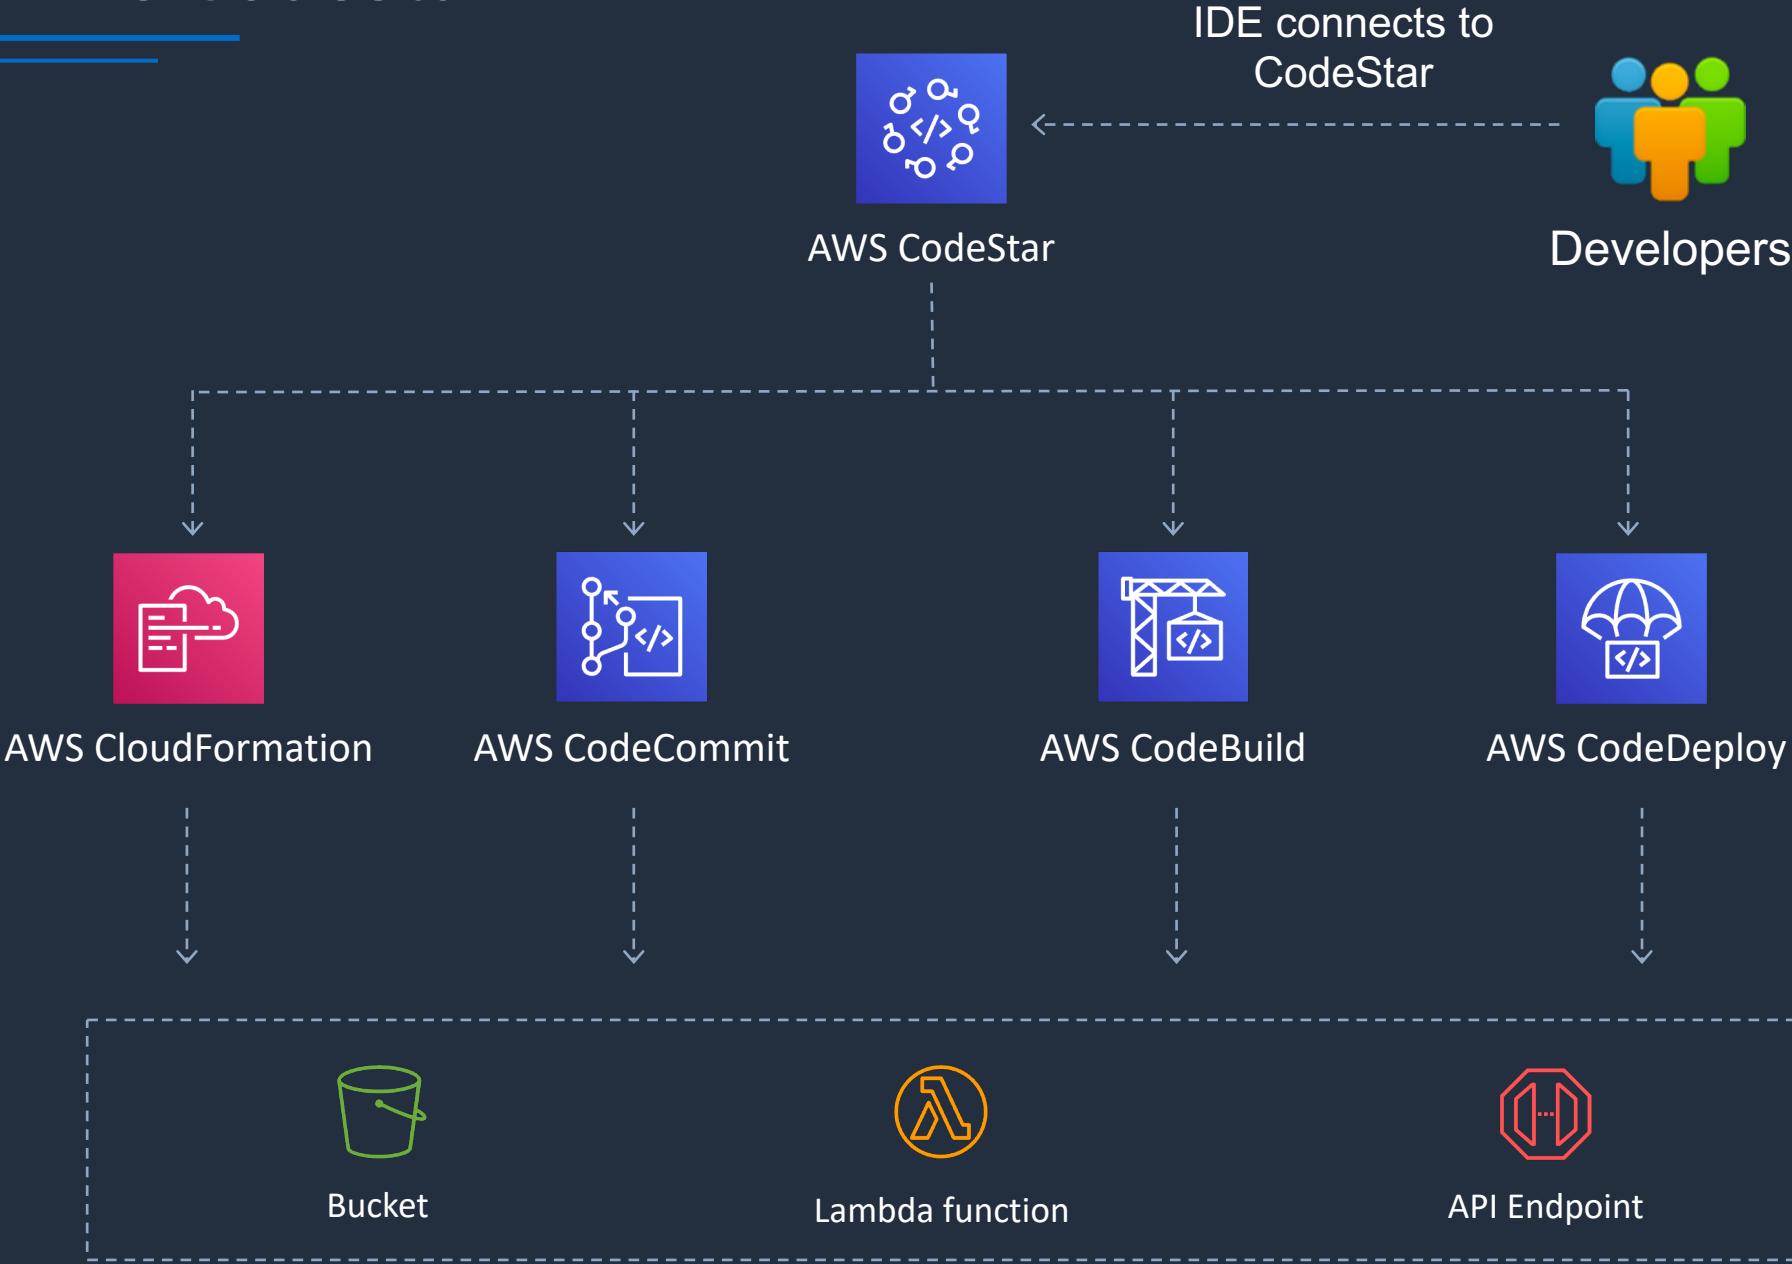


AWS CodeDeploy





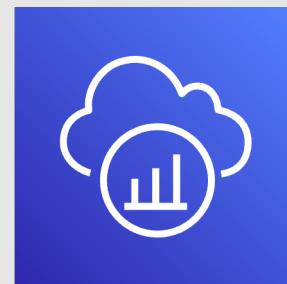
AWS CodeStar



AWS CodeStar



AWS X-Ray





AWS X-Ray

Records **latency** writing to a DynamoDB table



Client



Records **latency** from client to application



AWS X-Ray

- AWS X-Ray helps developers analyze and debug production, distributed applications, such as those built using a microservices architecture
- AWS X-Ray supports applications running on:
 - Amazon EC2
 - Amazon ECS
 - AWS Lambda
 - AWS Elastic Beanstalk
- Need to integrate the X-Ray SDK with your application and install the X-Ray agent

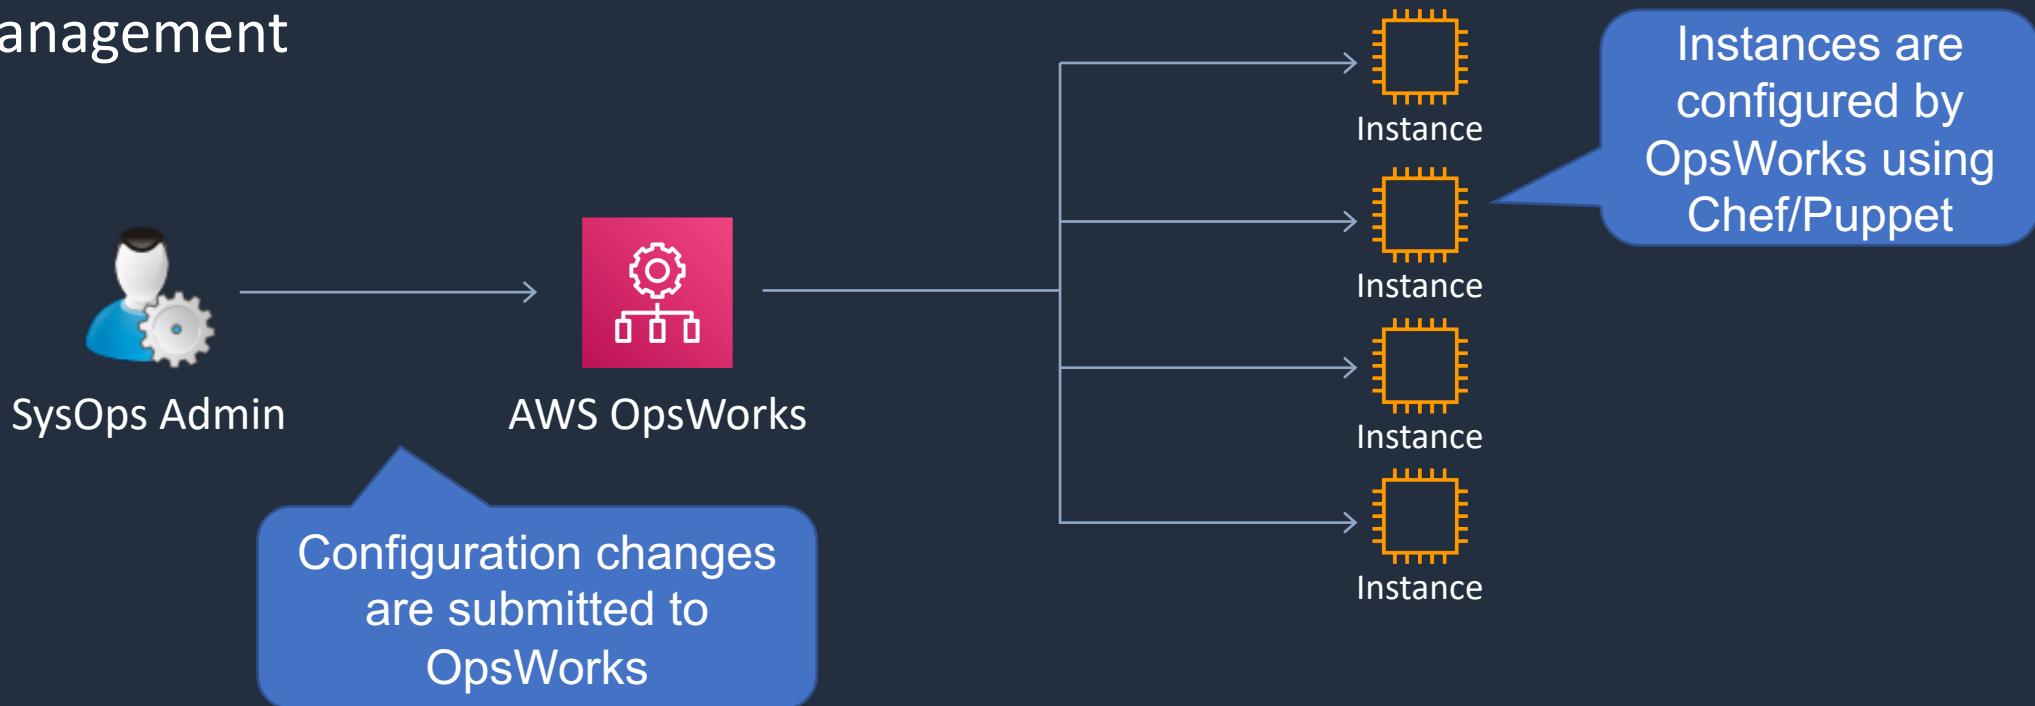
AWS OpsWorks





AWS OpsWorks

- AWS OpsWorks is a configuration management service that provides managed instances of **Chef** and **Puppet**
- Updates include patching, updating, backup, configuration and compliance management



SECTION 11

Databases and Analytics

Types of Database





Relational vs Non-Relational

Key differences are how data are **managed** and how data are **stored**

Relational	Non-Relational
Organized by tables, rows and columns	Varied data storage models
Rigid schema (SQL)	Flexible schema (NoSQL) – data stored in key-value pairs, columns, documents or graphs
Rules enforced within database	Rules can be defined in application code (outside database)
Typically scaled vertically	Scales horizontally
Supports complex queries and joins	Unstructured, simple language that supports any kind of schema
Amazon RDS, Oracle, MySQL, IBM DB2, PostgreSQL	Amazon DynamoDB, MongoDB, Redis, Neo4j



Relational Database

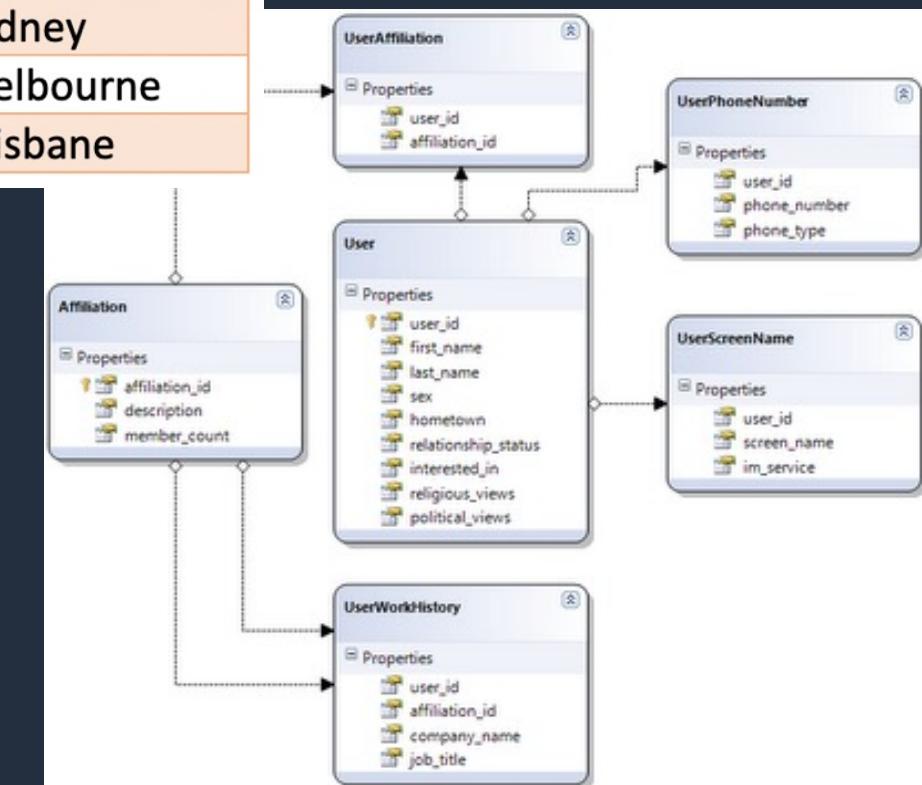
EmployeeID	FirstName	LastName	JobRole	Location
00001	Paul	Peterson	Senior Developer	Sydney
00002	Kaleigh	Annette	Assistant Manager	Brisbane
00003	Carl	Wood	Sales Support	Sydney
00004	Vinni	Jones	Customer Services	Melbourne
00005	Stefanie	Howard	IT Architect	Brisbane

Relational Database

Structured Query Language (SQL) query:

```
SELECT FirstName  
FROM employees  
WHERE Location = Sydney
```

Relational Database – Multiple Tables



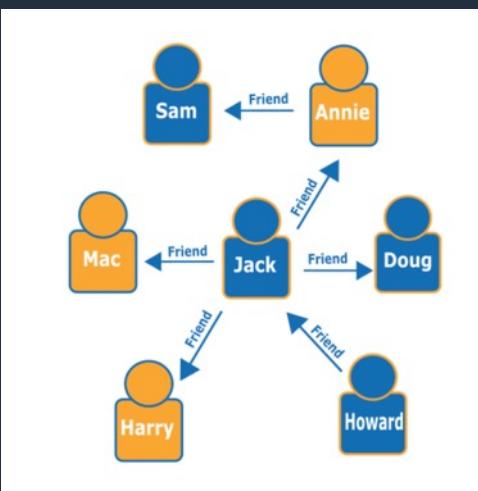


Types of Non-Relational DB (NoSQL)

Key-value – e.g. Amazon DynamoDB



Graph – e.g. Amazon Neptune



Document – e.g. MongoDB

```
JSON
1 [           ...
2 {             ...
3   "year" : 2013,
4   "title" : "Turn It Down, Or Else!",
5   "info" : {
6     "directors" : [ "Alice Smith", "Bob Jones" ],
7     "release_date" : "2013-01-18T00:00:00Z",
8     "rating" : 6.2,
9     "genres" : [ "Comedy", "Drama" ],
10    "image_url" : "http://ia.media-imdb.com/images/N/09ERWAU7FS797AJ7LU8HN09AMUP908RLlo5JF90EWR7LJKQ7@._V1_SX400_.jpg",
11    "plot" : "A rock band plays their music at high volumes, annoying the neighbors.",
12    "actors" : [ "David Matthewman", "Jonathan G. Neff" ]
13  }
14 },
15 {
16   "year": 2015,
17   "title": "The Big New Movie",
18   "info": {
19     "plot": "Nothing happens at all.",
20     "rating": 0
21   }
22 }
23 ]
```



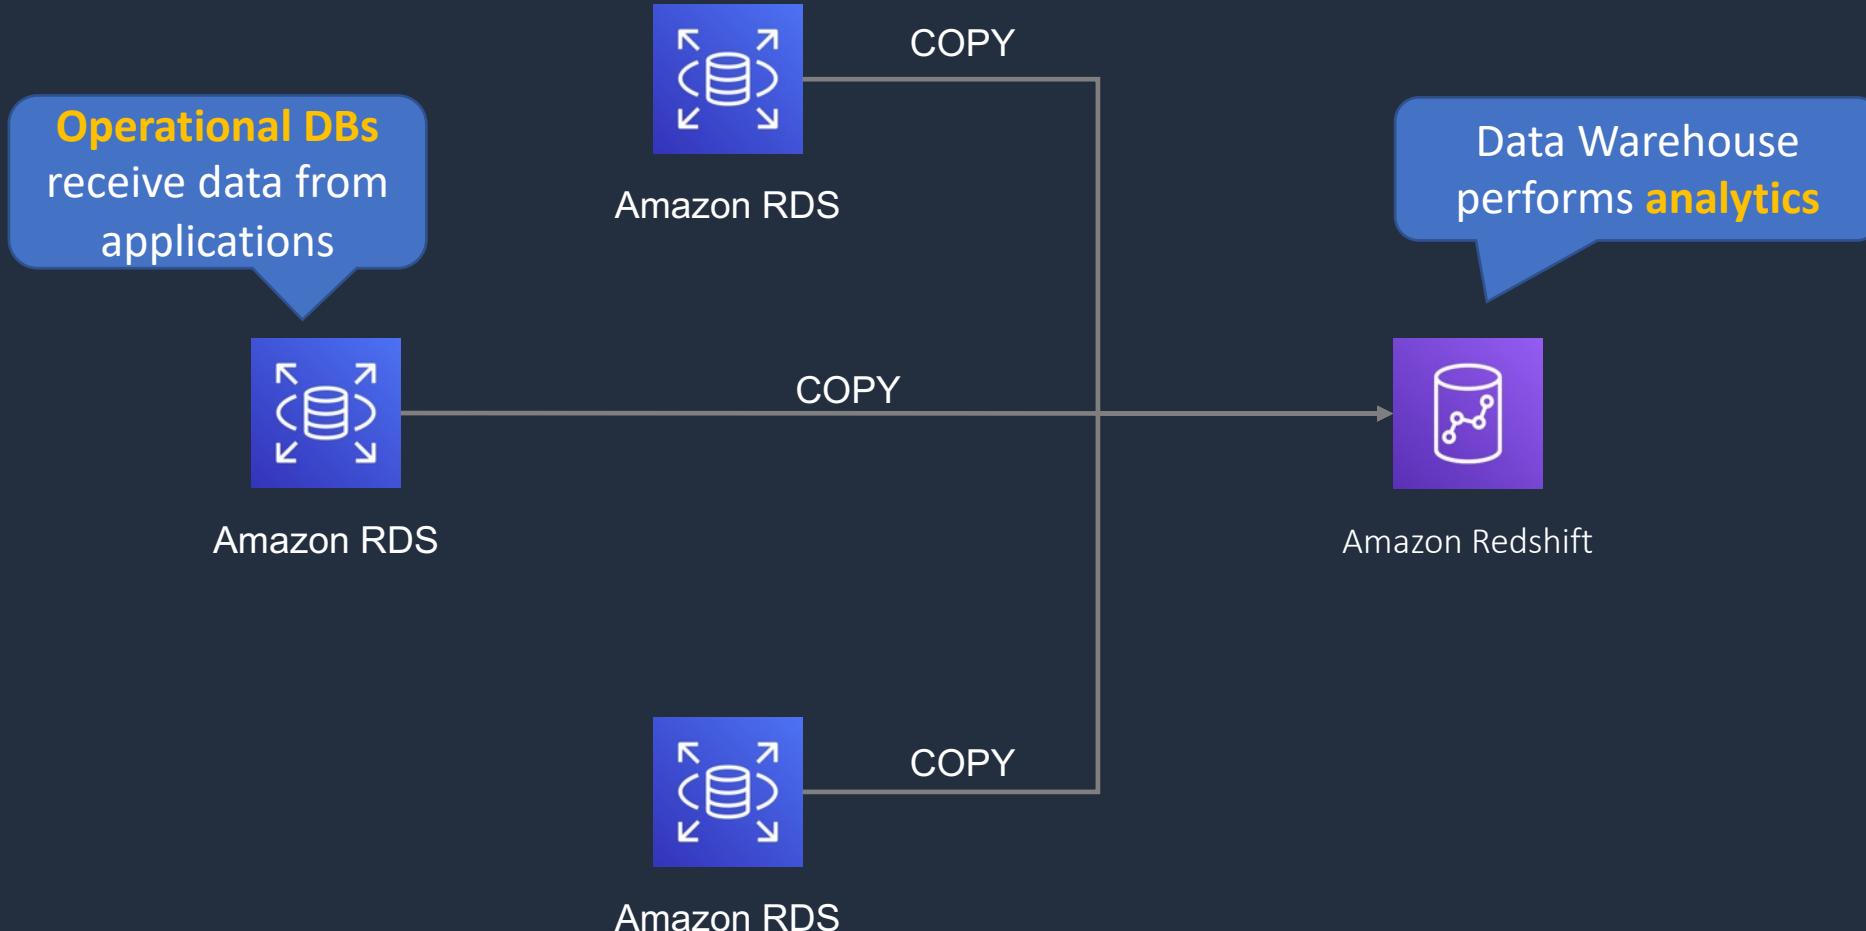
Operational vs Analytical

Key differences are **use cases** and how the database is **optimized**

Operational / transactional	Analytical
Online Transaction Processing (OLTP)	Online Analytics Processing (OLAP) – the source data comes from OLTP DBs
Production DBs that process transactions. E.g. adding customer records, checking stock availability (INSERT, UPDATE, DELETE)	Data warehouse. Typically, separated from the customer facing DBs. Data is extracted for decision making
Short transactions and simple queries	Long transactions and complex queries
Relational examples: Amazon RDS, Oracle, IBM DB2, MySQL	Relational examples: Amazon RedShift, Teradata, HP Vertica
Non-relational examples: MongoDB, Cassandra, Neo4j, HBase	Non-relational examples: Amazon EMR, MapReduce



Operational vs Analytical





AWS Databases

Data Store	Use Case
Database on EC2	<ul style="list-style-type: none">• Need full control over instance and database• Third-party database engine (not available in RDS)
Amazon RDS	<ul style="list-style-type: none">• Need traditional relational database• e.g. Oracle, PostgreSQL, Microsoft SQL, MariaDB, MySQL• Data is well-formed and structured
Amazon DynamoDB	<ul style="list-style-type: none">• NoSQL database• In-memory performance• High I/O needs• Dynamic scaling
Amazon RedShift	<ul style="list-style-type: none">• Data warehouse for large volumes of aggregated data
Amazon ElastiCache	<ul style="list-style-type: none">• Fast temporary storage for small amounts of data• In-memory database
Amazon EMR	<ul style="list-style-type: none">• Analytics workloads using the Hadoop framework

Amazon Relational Database Service (RDS)





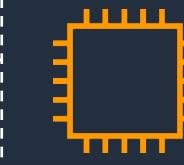
Amazon RDS

RDS is a **managed**, relational database



Amazon RDS

RDS runs on **EC2 instances**, so you must choose an **instance type**



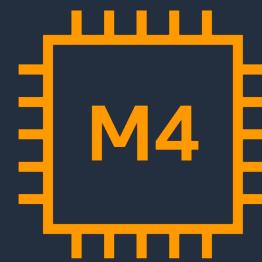
EC2

RDS supports the following database engines:

- Amazon Aurora
- MySQL
- MariaDB
- Oracle
- Microsoft SQL Server
- PostgreSQL

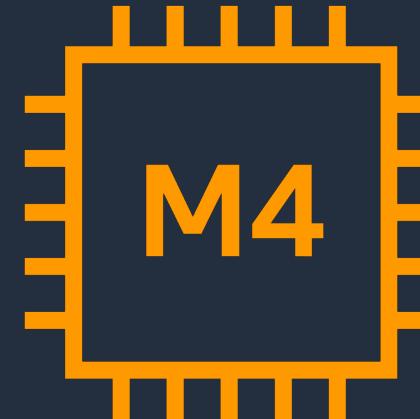


Amazon RDS Scaling Up (vertically)



M4 instance

db.m4.large 2
vCPUs, 8 GiB
RAM

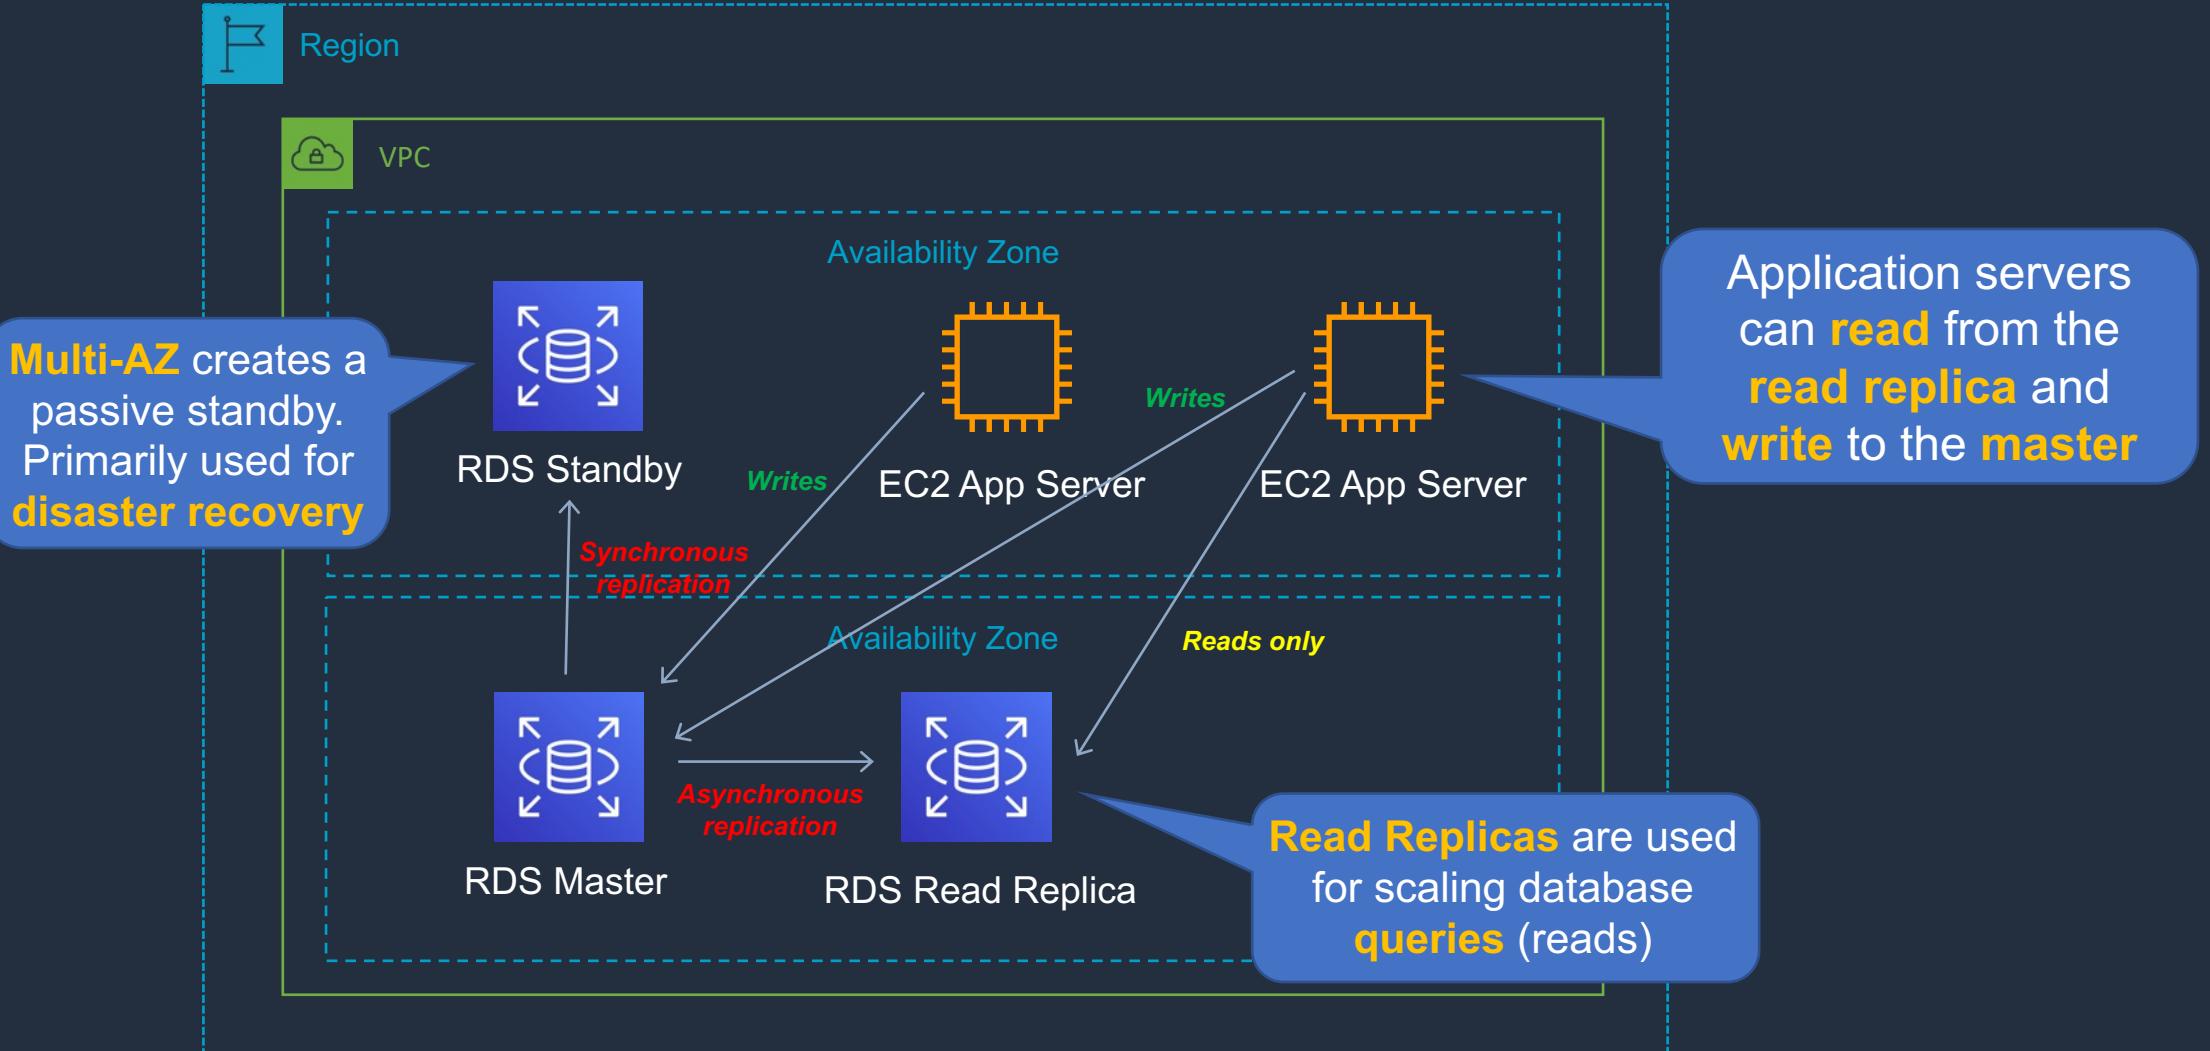


M4 Instance

db.m4.2xlarge
4 vCPUs, 32
GiB RAM



Disaster Recovery (DR) and Scaling Out (Horizontally)





Amazon RDS

- RDS uses EC2 instances, so you must choose an instance family/type
- Relational databases are known as Structured Query Language (SQL) databases
- RDS is an Online Transaction Processing (OLTP) type of database
- Easy to setup, highly available, fault tolerant, and scalable
- Common use cases include online stores and banking systems
- You can encrypt your Amazon RDS instances and snapshots at rest by enabling the encryption option for your Amazon RDS DB instance (during creation)
- Encryption uses AWS Key Management Service (KMS)



Amazon RDS

- Amazon RDS supports the following database engines:
 - SQL Server
 - Oracle
 - MySQL Server
 - PostgreSQL
 - Aurora
 - MariaDB
- Scales up by increasing instance size (compute and storage)
- Read replicas option for read heavy workloads (scales out for reads/queries only)
- Disaster recovery with Multi-AZ option

Create Amazon RDS Database



Amazon Aurora





Amazon Aurora

- Amazon Aurora is an AWS database offering in the RDS family
- Amazon Aurora is a MySQL and PostgreSQL-compatible relational database built for the cloud
- Amazon Aurora is up to five times faster than standard MySQL databases and three times faster than standard PostgreSQL databases
- Amazon Aurora features a distributed, fault-tolerant, self-healing storage system that auto-scales up to 64TB per database instance

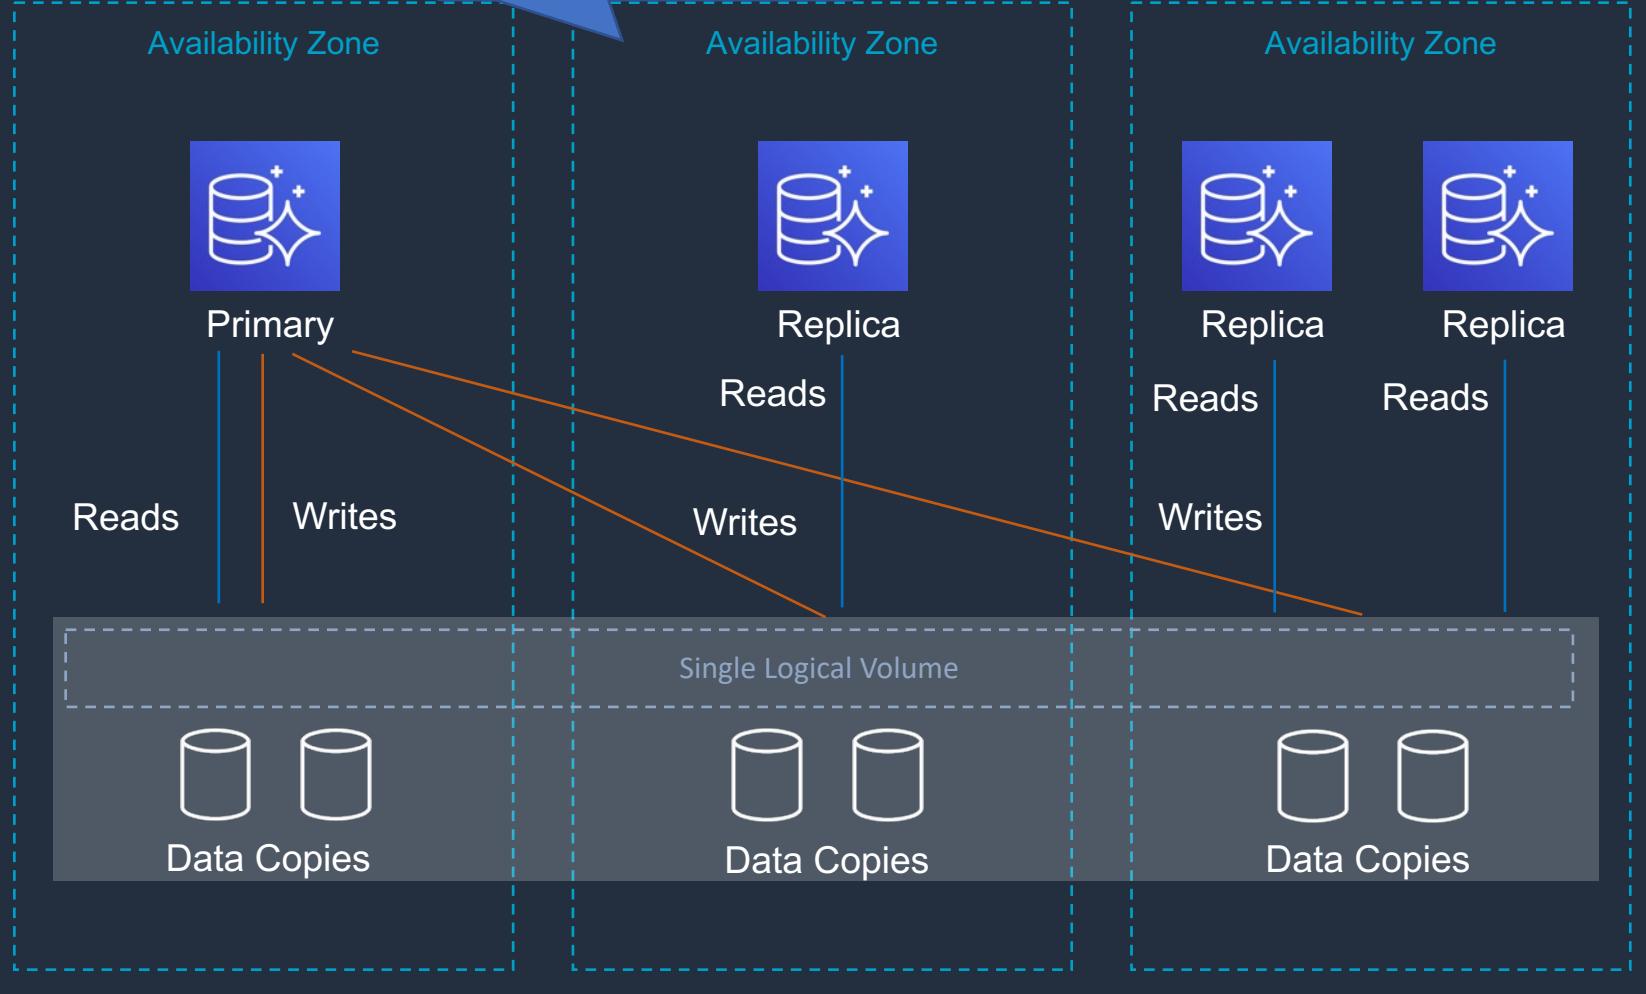


Amazon Aurora



Region

Aurora Replicas are
within a region



Aurora Fault Tolerance

- Fault tolerance across 3 AZs
- Single logical volume
- Aurora Replicas scale-out read requests
- Can **promote** Aurora Replica to be a new primary or create new primary
- Can use **Auto Scaling** to add replicas



Amazon Aurora Key Features

Aurora Feature	Benefit
High performance and scalability	Offers high performance, self-healing storage that scales up to 64TB, point-in-time recovery and continuous backup to S3
DB compatibility	Compatible with existing MySQL and PostgreSQL open source databases
Aurora Replicas	In-region read scaling and failover target – up to 15 (can use Auto Scaling)
MySQL Read Replicas	Cross-region cluster with read scaling and failover target – up to 5 (each can have up to 15 Aurora Replicas)
Global Database	Cross-region cluster with read scaling (fast replication / low latency reads). Can remove secondary and promote
Multi-Master	Scales out writes within a region. In preview currently and will not appear on the exam
Serverless	On-demand, autoscaling configuration for Amazon Aurora - does not support read replicas or public IPs (can only access through VPC or Direct Connect - not VPN)

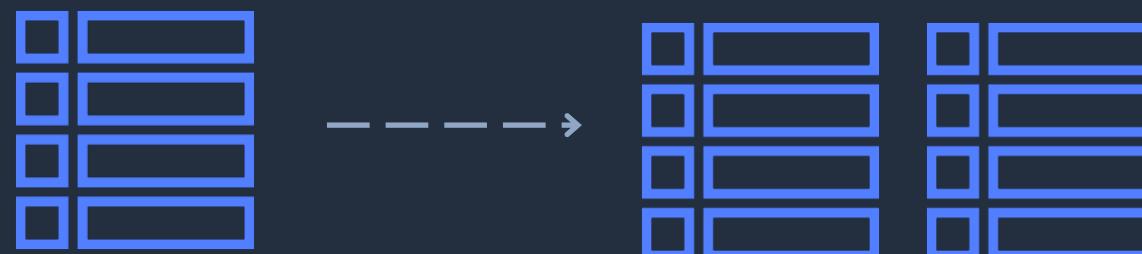
Amazon DynamoDB





Amazon DynamoDB

- Fully managed NoSQL database service
- Key/value store and document store
- It is a non-relational, key-value type of database
- Fully serverless service
- Push button scaling



DynamoDB Table



Amazon DynamoDB

- DynamoDB is made up of:

- Tables
- Items
- Attributes

userid	orderid	book	price	date
user001	1000092	ISBN100..	9.99	2020.04..
user002	1000102	ISBN100..	24.99	2020.03..
user003	1000168	ISBN2X0..	12.50	2020.04..



Amazon DynamoDB Key Features

DynamoDB Feature	Benefit
Serverless	Fully managed, fault tolerant, service
Highly available	99.99% availability SLA – 99.999% for Global Tables!
NoSQL type of database with Name / Value structure	Flexible schema, good for when data is not well structured or unpredictable
Horizontal scaling	Seamless scalability to any scale with push button scaling or Auto Scaling
DynamoDB Accelerator (DAX)	Fully managed in-memory cache for DynamoDB that increases performance (microsecond latency)
Backup	Point-in-time recovery down to the second in last 35 days; On-demand backup and restore
Global Tables	Fully managed multi-region, multi-master solution

Create Amazon DynamoDB Table

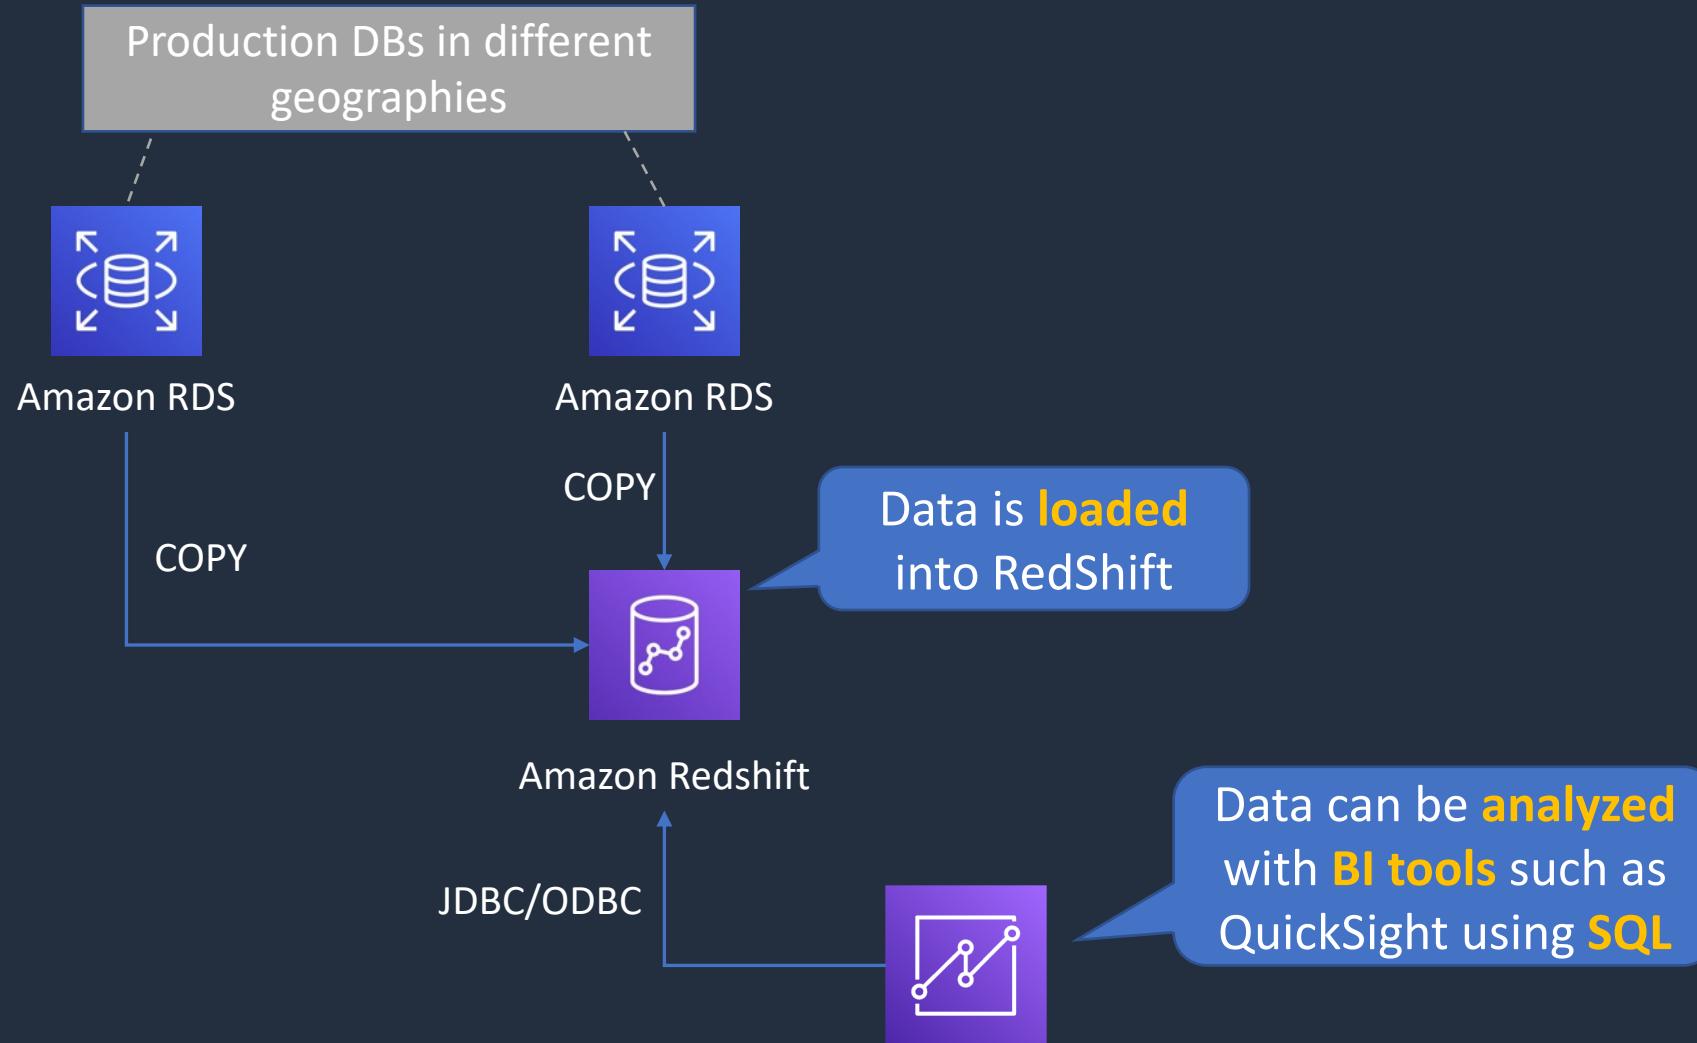


Amazon RedShift





Amazon Redshift

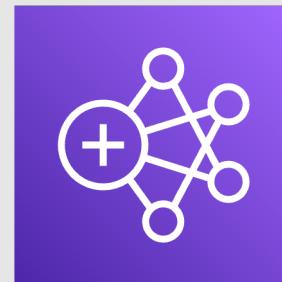




Amazon RedShift

- Amazon Redshift is a fast, fully managed data warehouse that makes it simple and cost-effective to analyze all your data using standard SQL and existing Business Intelligence (BI) tools
- RedShift is a SQL based data warehouse used for analytics applications
- RedShift is a relational database that is used for Online Analytics Processing (OLAP) use cases
- RedShift uses Amazon EC2 instances, so you must choose an instance family/type
- RedShift always keeps three copies of your data
- RedShift provides continuous/incremental backups

Amazon Elastic Map Reduce (EMR)



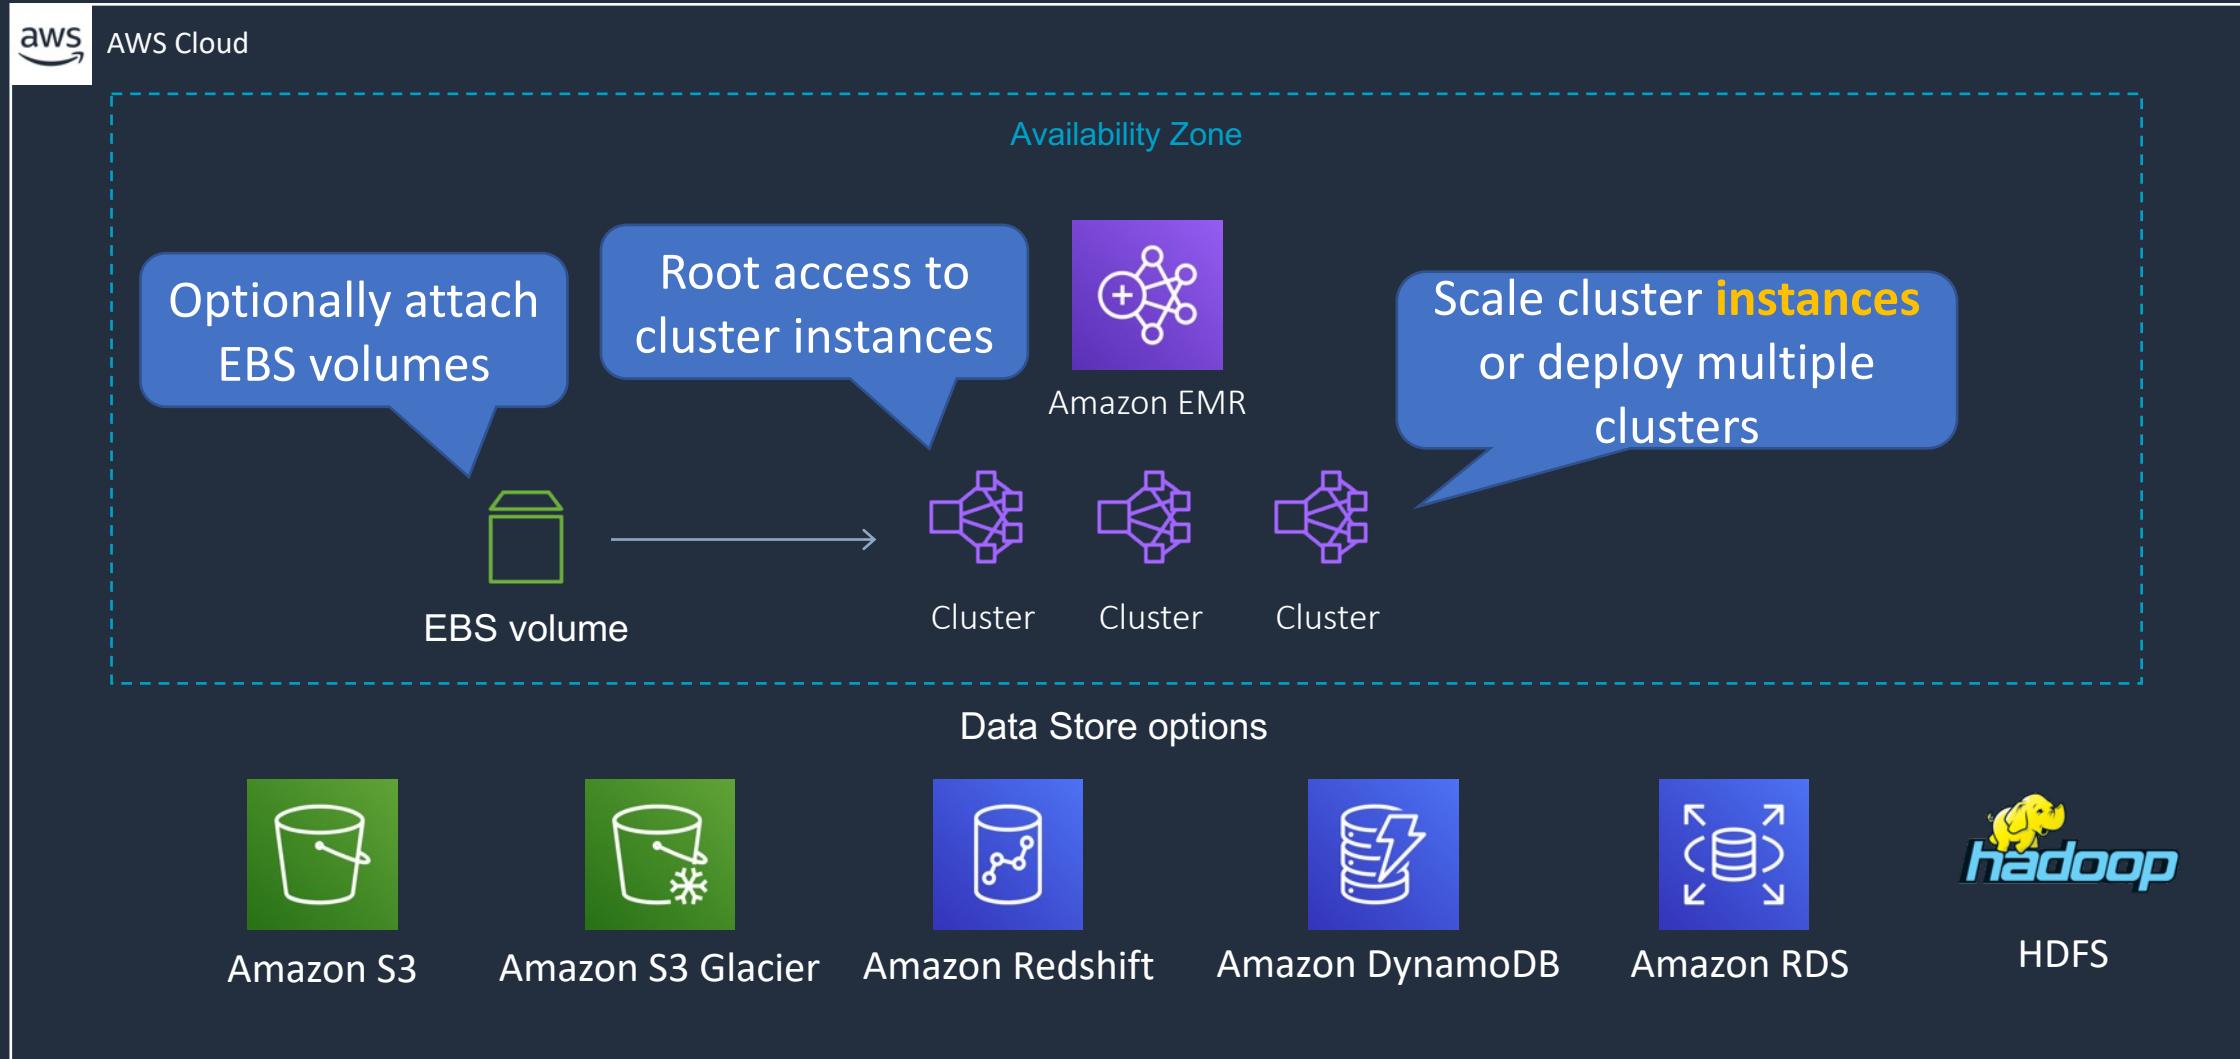


Amazon EMR

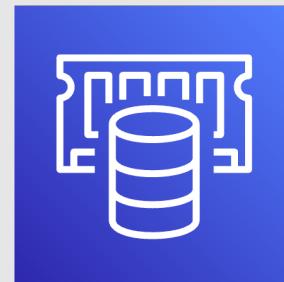
- Managed cluster platform that simplifies running big data frameworks including **Apache Hadoop** and **Apache Spark**
- Used for processing data for analytics and business intelligence
- Can also be used for transforming and moving large amounts of data
- Performs extract, transform, and load (ETL) functions



Amazon EMR



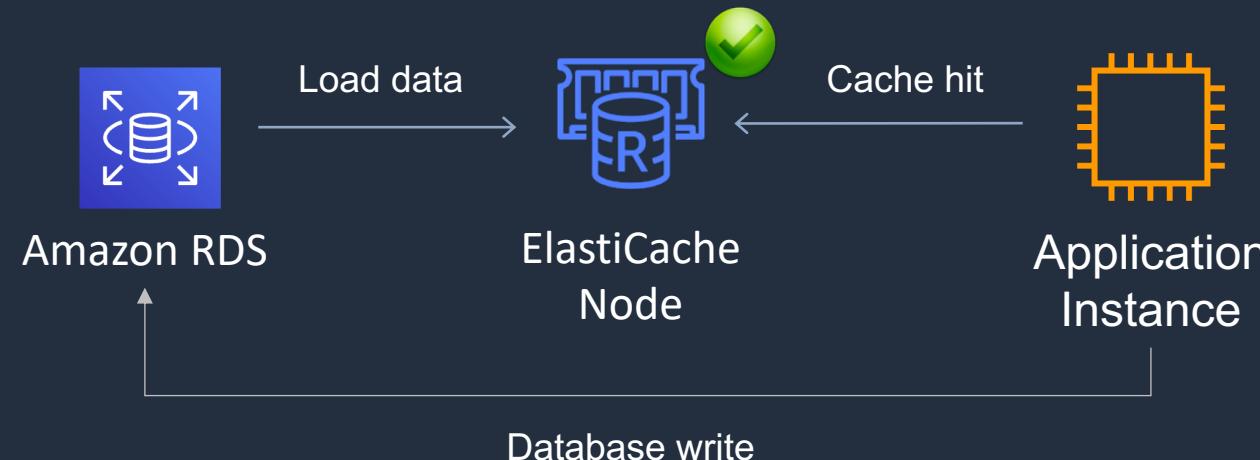
Amazon ElastiCache





Amazon ElastiCache

- Fully managed implementations **Redis** and **Memcached**
- ElastiCache is a **key/value** store
- In-memory database offering high performance and low latency
- Can be put in front of databases such as RDS and DynamoDB



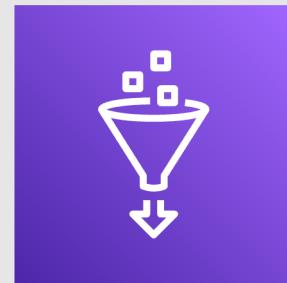
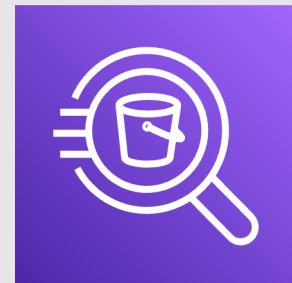


Amazon ElastiCache

- ElastiCache nodes run on Amazon EC2 instances, so you must choose an instance family/type

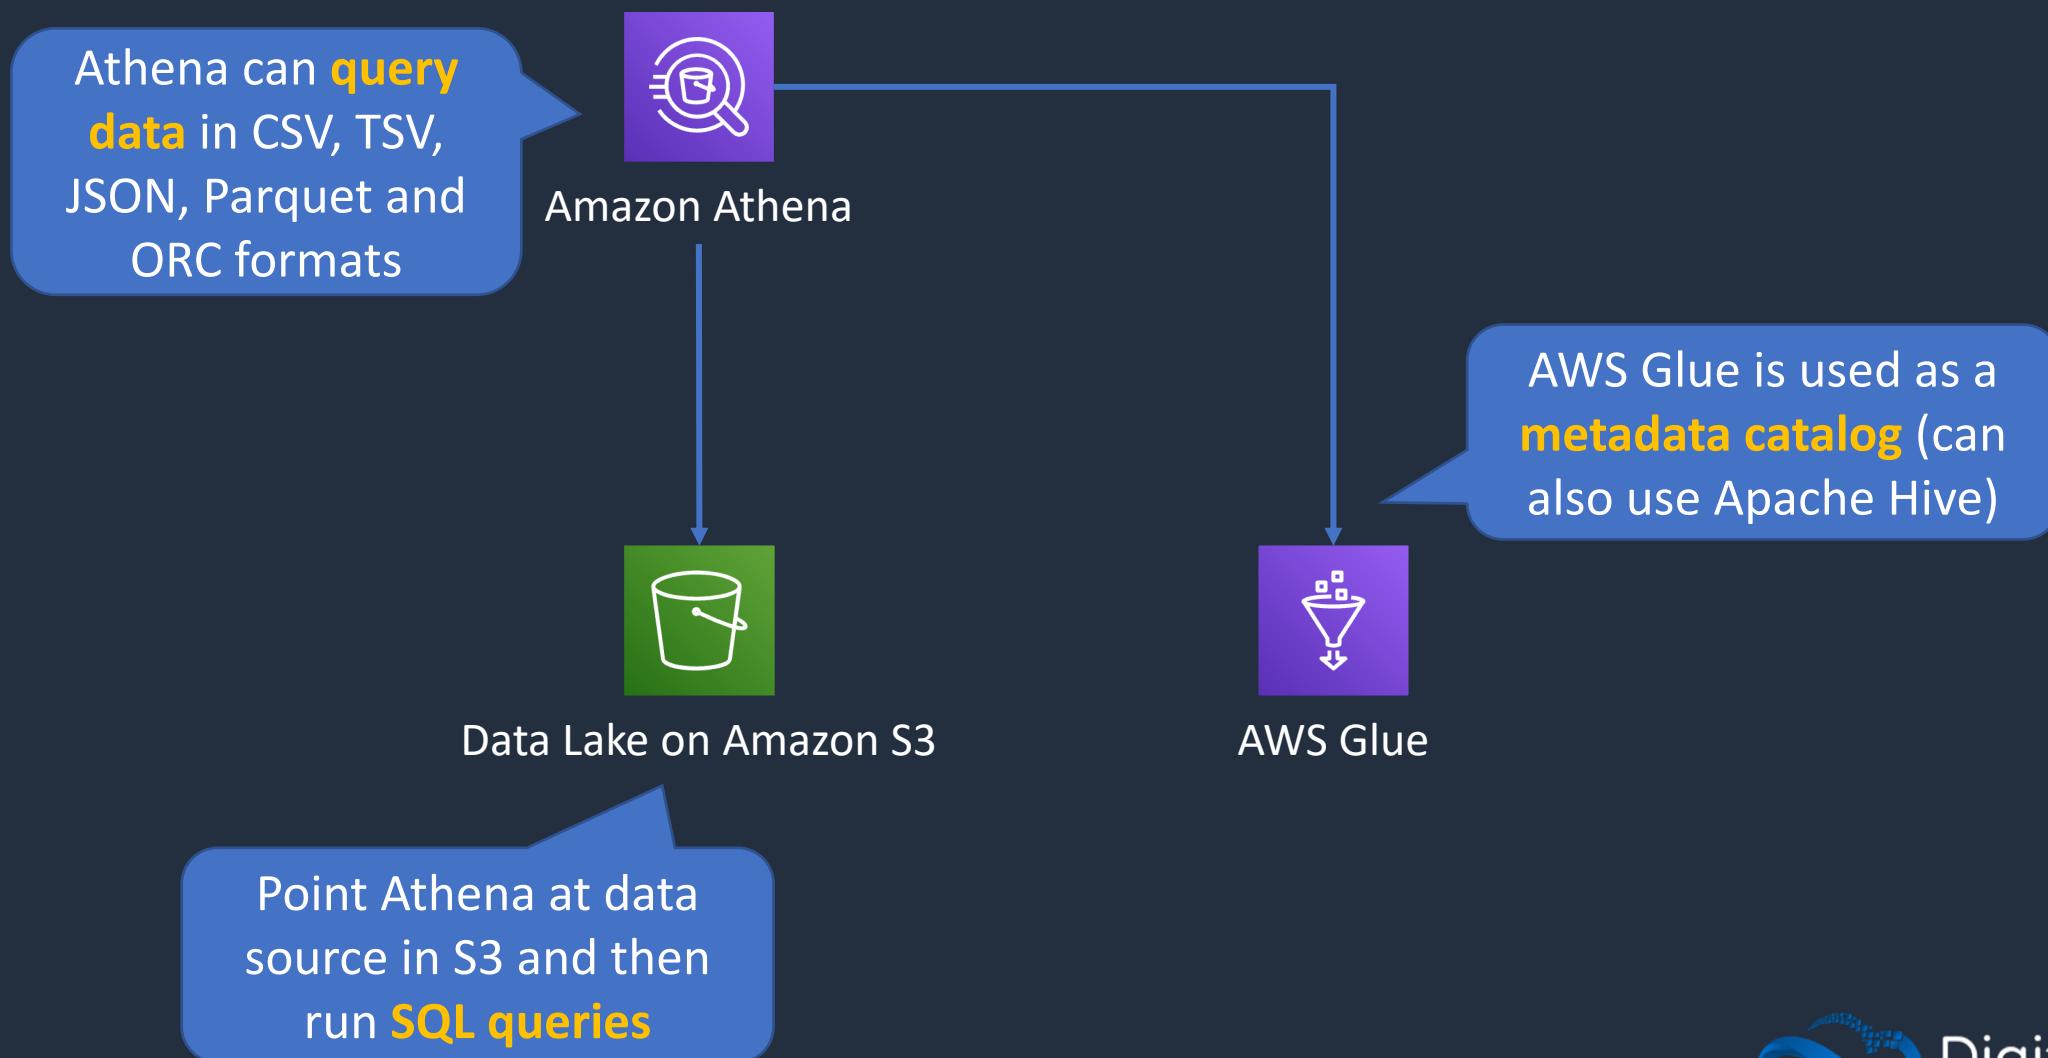
Use Case	Benefit
Web session store	In cases with load-balanced web servers, store web session information in Redis so if a server is lost, the session info is not lost, and another web server can pick it up
Database caching	Use Memcached in front of AWS RDS to cache popular queries to offload work from RDS and return results faster to users
Leaderboards	Use Redis to provide a live leaderboard for millions of users of your mobile app
Streaming data dashboards	Provide a landing spot for streaming sensor data on the factory floor, providing live real-time dashboard displays

Amazon Athena and AWS Glue





Amazon Athena and AWS Glue





Amazon Athena

- Athena queries data in S3 using SQL
- Can be connected to other data sources with Lambda
- Data can be in CSV, TSV, JSON, Parquet and ORC formats
- Uses a managed Data Catalog (AWS Glue) to store information and schemas about the databases and tables



AWS Glue

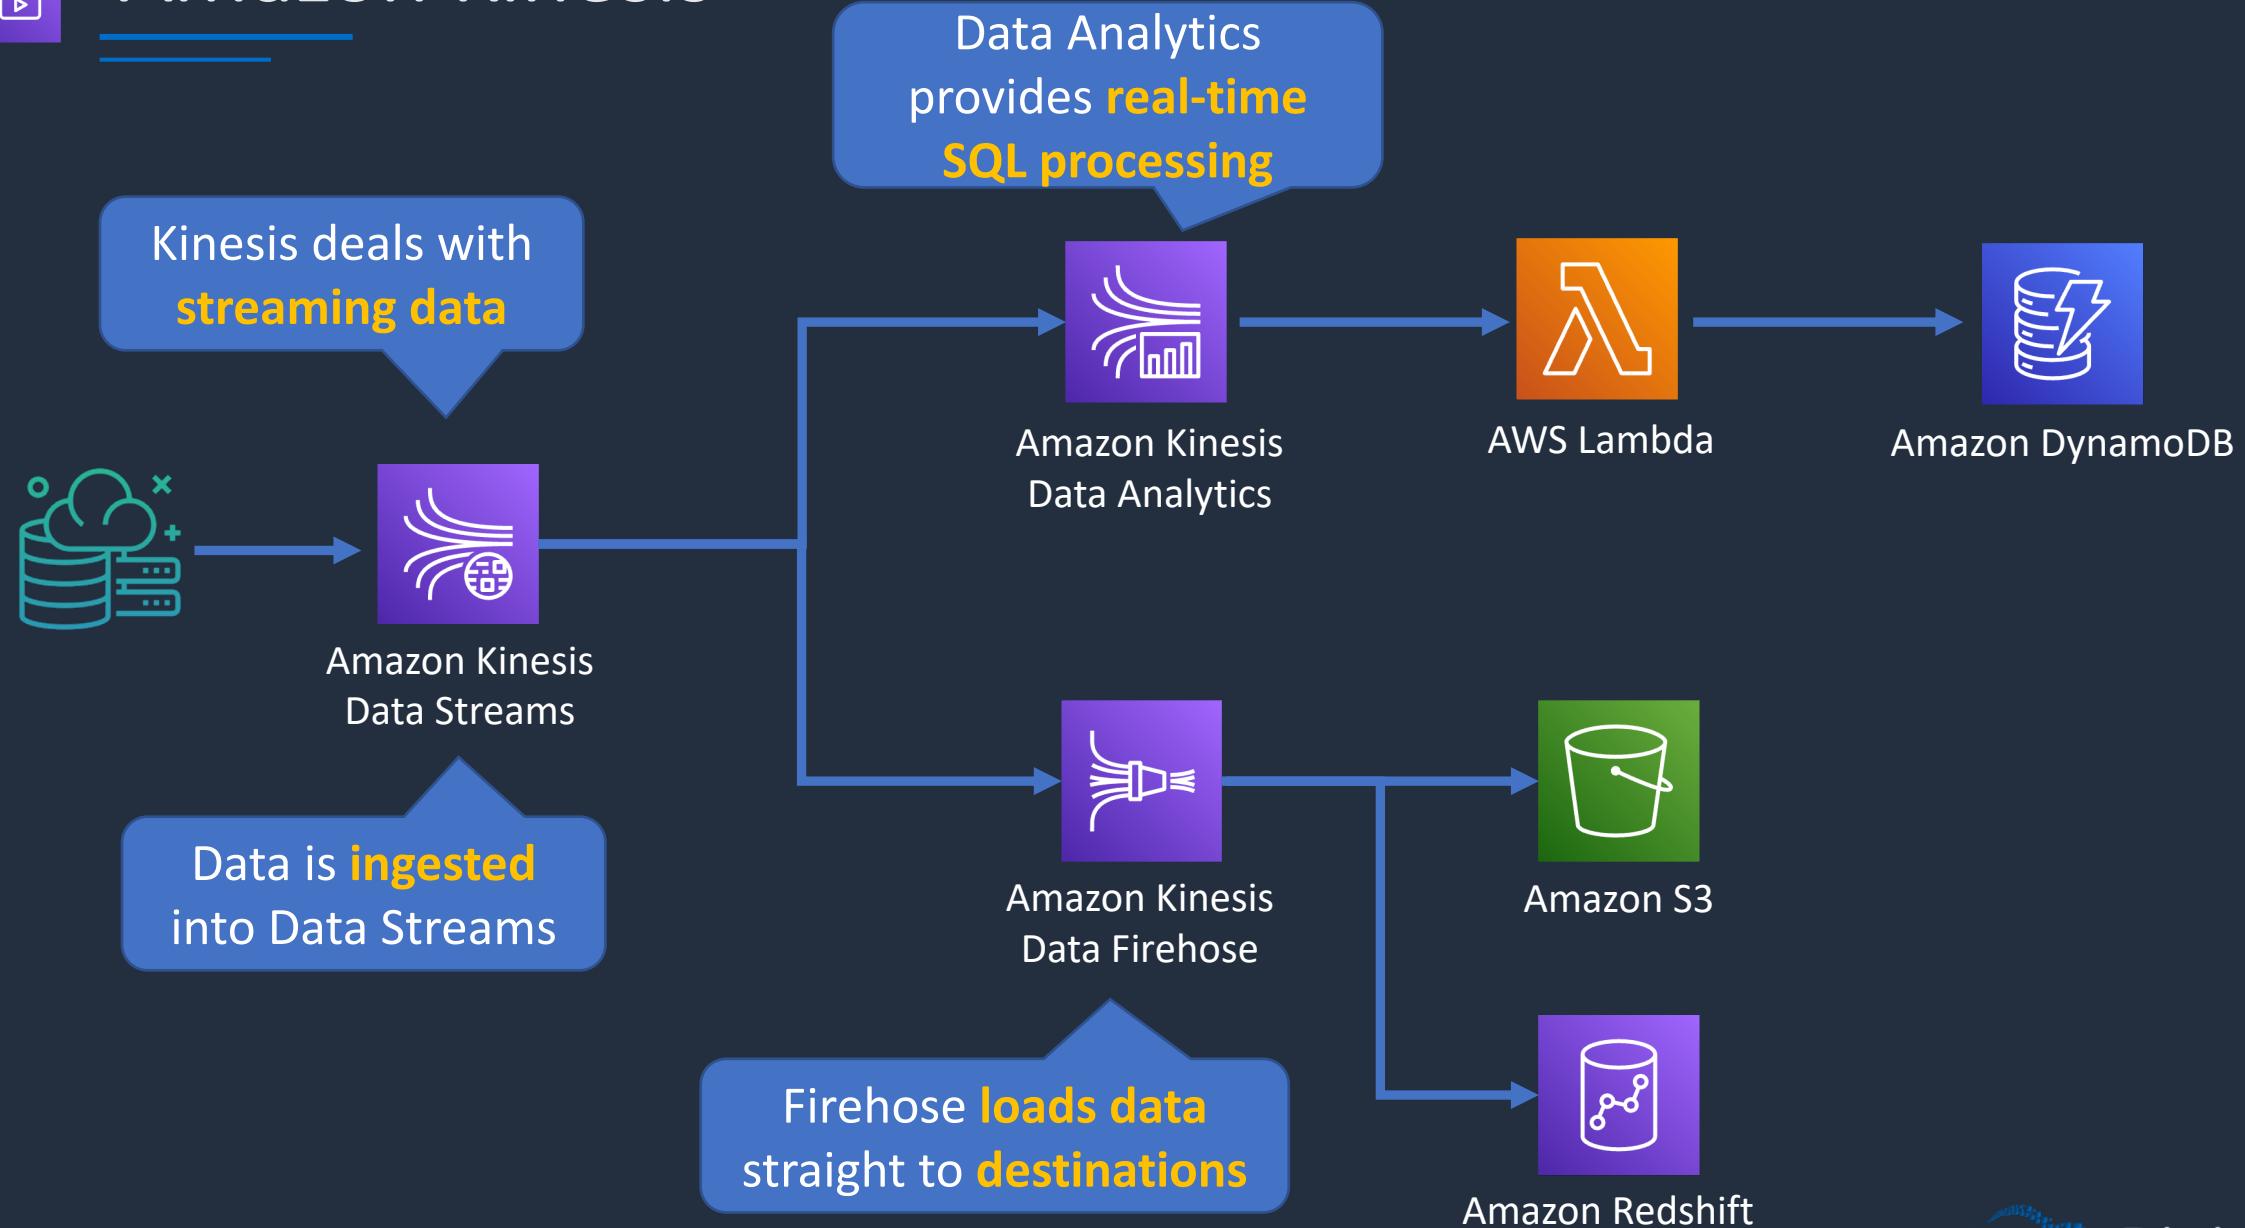
- Fully managed extract, transform and load (ETL) service
- Used for preparing data for analytics
- AWS Glue runs the ETL jobs on a fully managed, scale-out Apache Spark environment
- Works with data lakes (e.g. data on S3), data warehouses (including RedShift), and data stores (including RDS or EC2 databases)

Amazon Kinesis





Amazon Kinesis





Amazon Kinesis

Examples of streaming data use cases include:

- Purchases from online stores
- Stock prices
- Game data (statistics and results as the gamer plays)
- Social network data
- Geospatial data (think uber.com)
- IoT sensor data



Amazon Kinesis

Kinesis Data Streams

- Producers send data which is stored in shards for up to 7 days
- Consumers process the data and save to another service

Amazon Kinesis Data Firehose

- No shards, completely automated and elastically scalable
- Saves data directly to another service such as S3, Splunk, RedShift, or Elasticsearch

Amazon Kinesis Data Analytics

- Provides real-time SQL processing for streaming data

Other Databases and Analytics Services





Other Databases and Analytics Services

AWS Data Pipeline

- Processes and moves data between different AWS compute and storage services
- Save results to services including S3, RDS, DynamoDB, and EMR

Amazon QuickSight

- Business intelligence (BI) service
- Create and publish interactive BI dashboards for Machine Learning-powered insights

Amazon Neptune

- Fully managed graph database service



Other Databases and Analytics Services

Amazon DocumentDB

- Fully managed document database service (non-relational)
- Supports MongoDB workloads
- Queries and indexes JSON data

Amazon QLDB

- Fully managed ledger database for immutable change history
- Provides cryptographically verifiable transaction logging

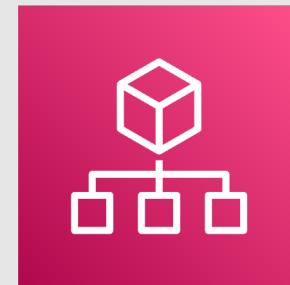
Amazon Managed Blockchain

- Fully managed service for joining public and private networks using Hyperledger Fabric and Ethereum

SECTION 12

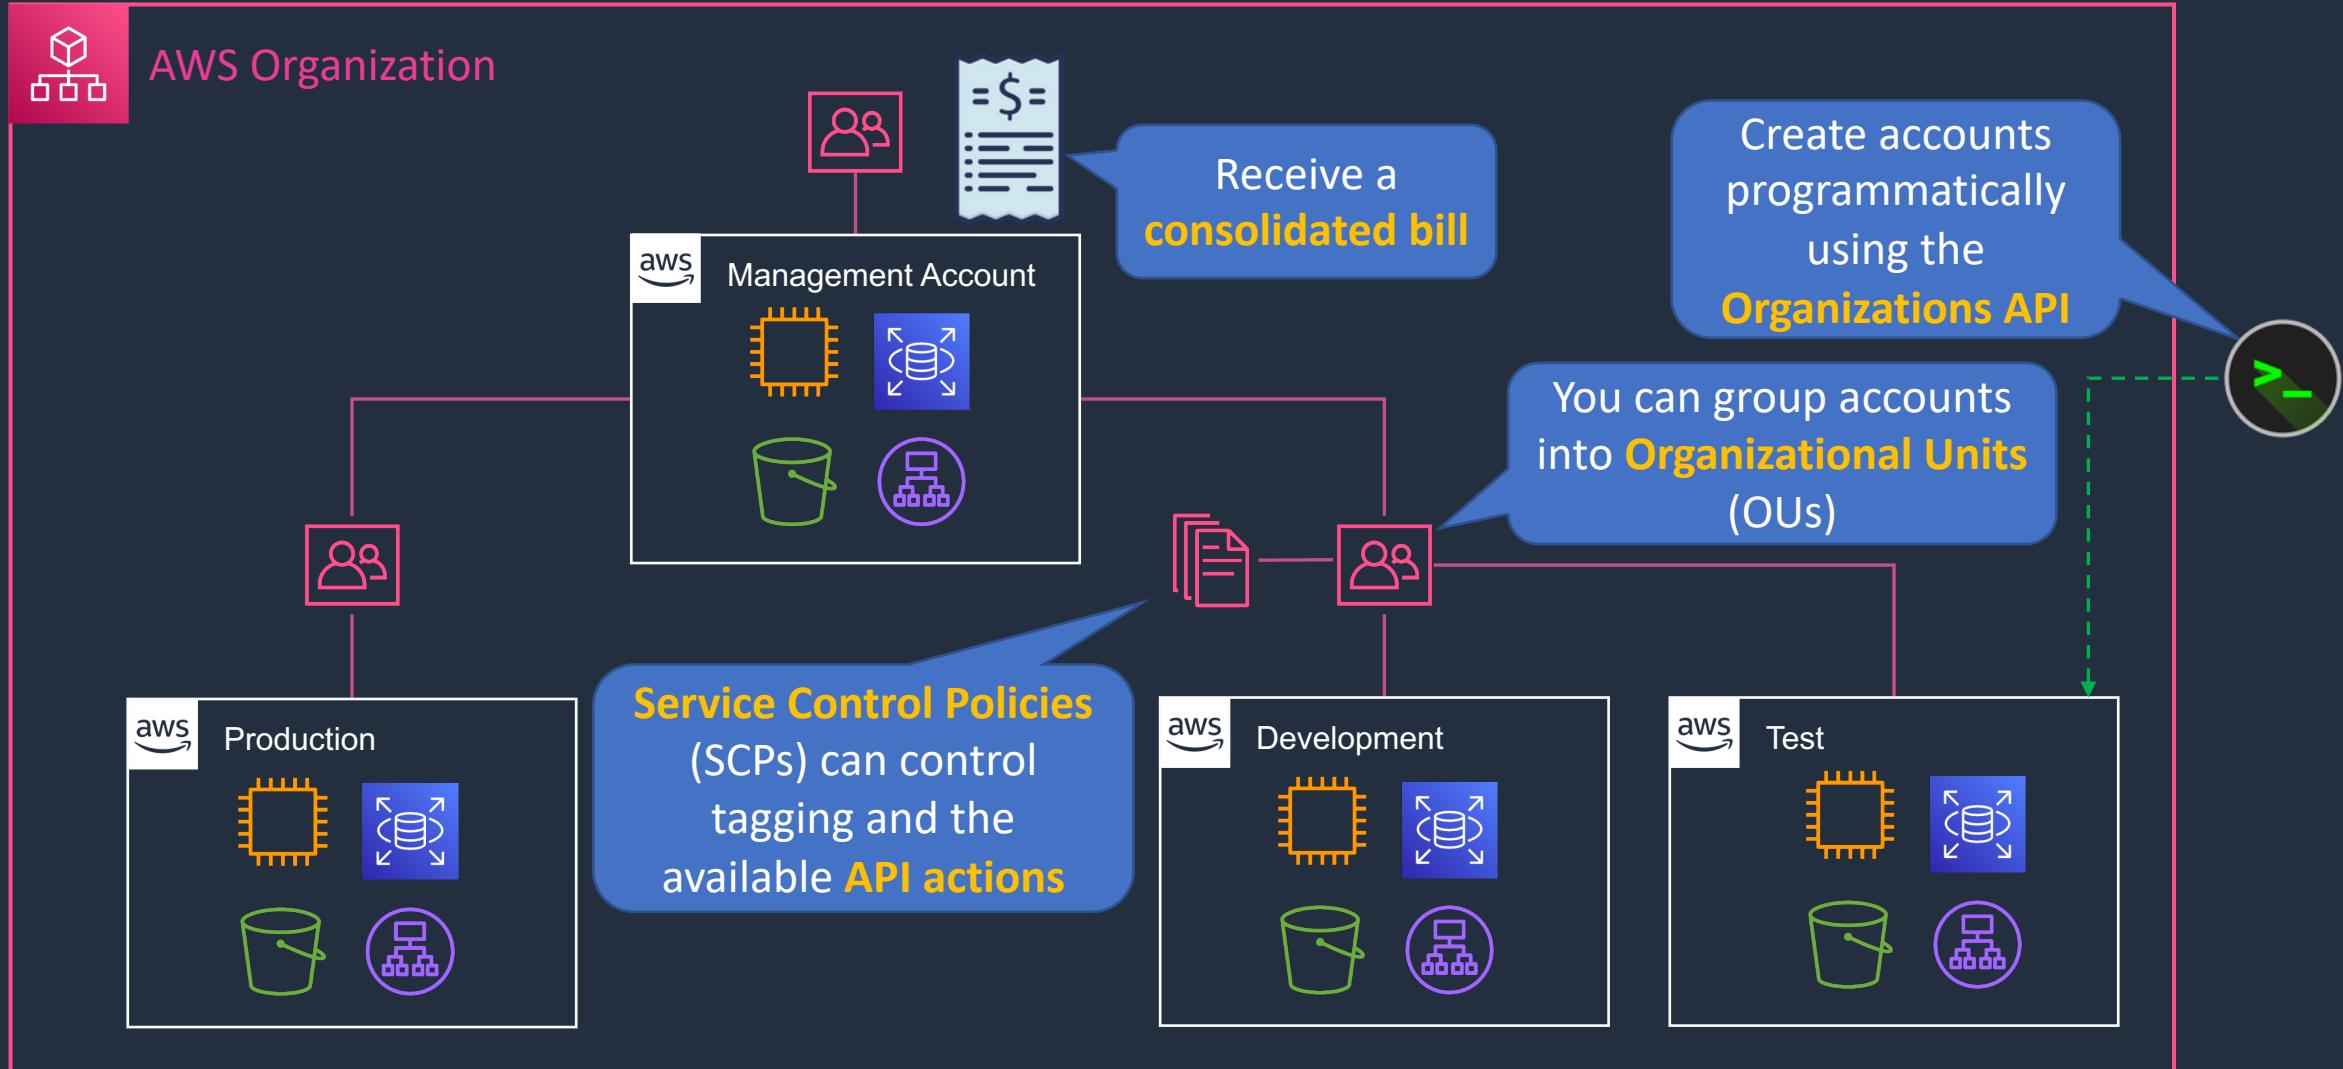
Management and Governance

AWS Organizations





AWS Organizations





AWS Organizations

- AWS organizations allows you to consolidate multiple AWS accounts into an organization that you create and centrally manage
- Available in two feature sets:
 - **Consolidated Billing**
 - **All features**
- Includes root accounts and organizational units
- Policies are applied to root accounts or OUs
- Consolidated billing includes:
 - **Paying Account** – independent and cannot access resources of other accounts
 - **Linked Accounts** – all linked accounts are independent

AWS Organizations



AWS Control Tower





AWS Control Tower

- Simplifies the process of creating multi-account environments
- Sets up governance, compliance, and security guardrails for you
- Integrates with other services and features to setup the environment for you including:
 - AWS Organizations, SCPs, OUs, AWS Config, AWS CloudTrail, Amazon S3, Amazon SNS, AWS CloudFormation, AWS Service Catalog, AWS Single Sign-On (SSO)



AWS Control Tower

Examples of guardrails AWS Control Tower can configure for you include:

- Disallowing public write access to Amazon Simple Storage Service (Amazon S3) buckets
- Disallowing access as a root user without multi-factor authentication
- Enabling encryption for Amazon EBS volumes attached to Amazon EC2 instances

AWS Systems Manager





AWS Systems Manager

- Manages many AWS resources including Amazon EC2, Amazon S3, Amazon RDS etc.
- Systems Manager Components:
 - **Automation**
 - **Run Command**
 - **Inventory**
 - **Patch Manager**
 - **Session Manager**
 - **Parameter Store**



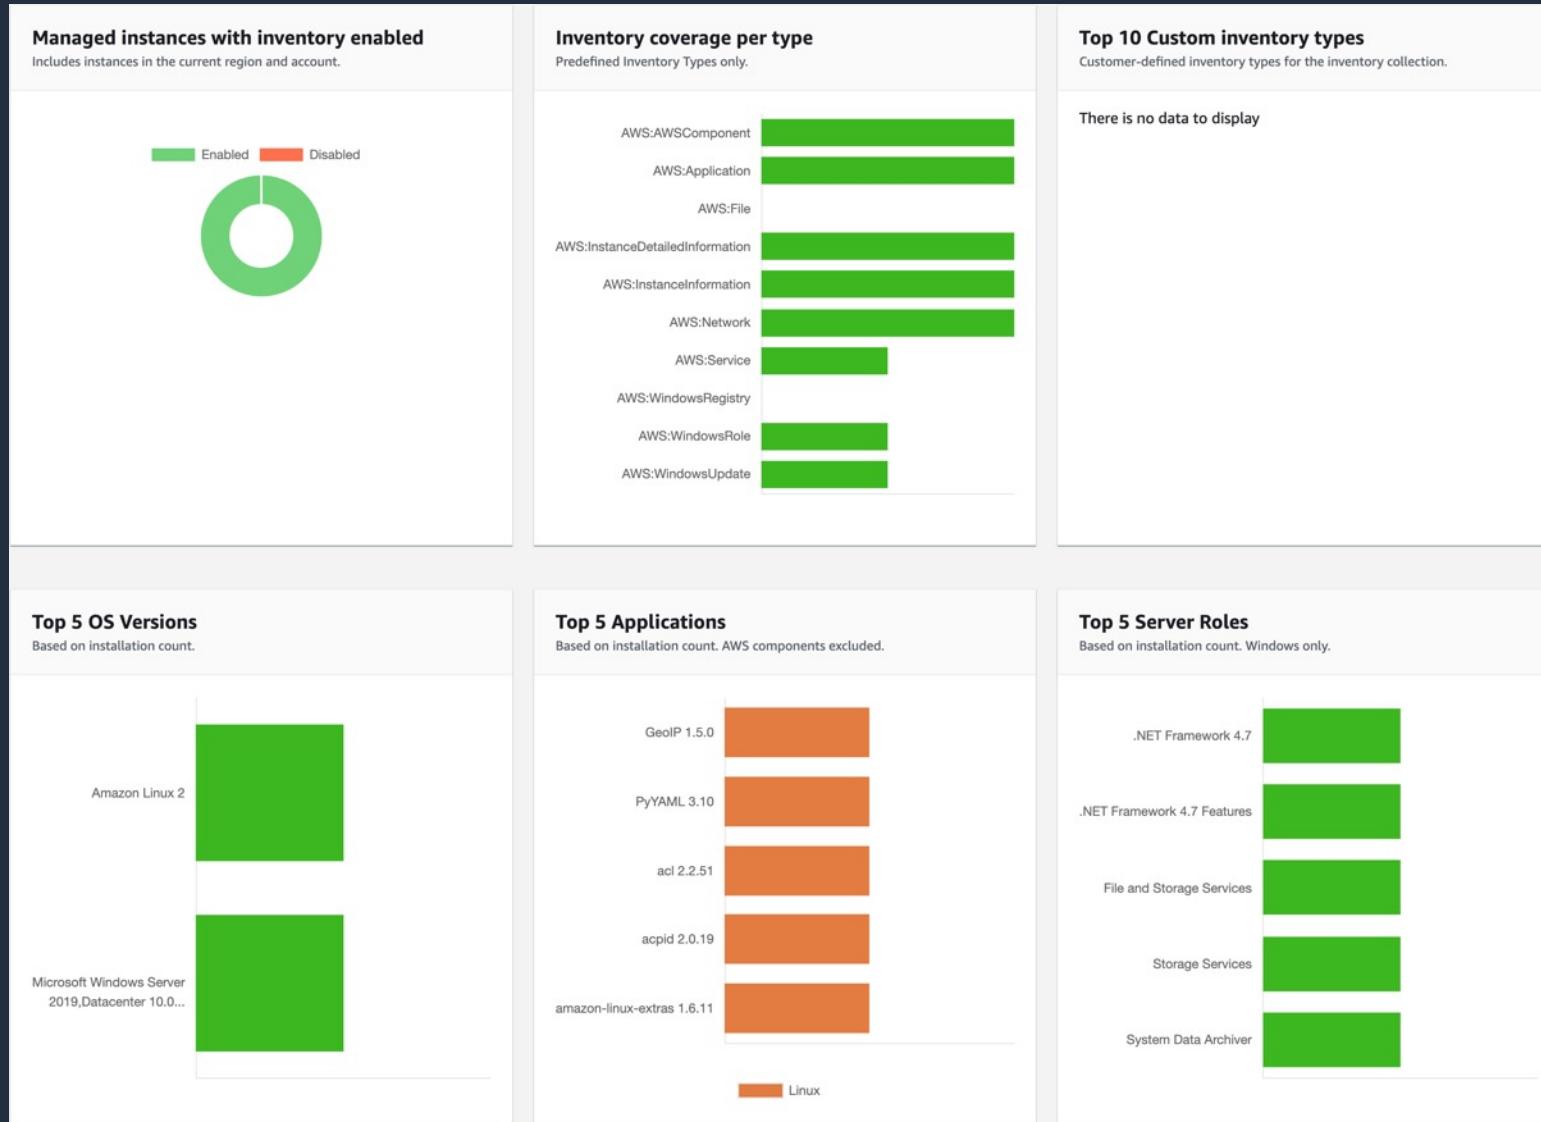
AWS Systems Manager



This command checks
for missing **updates**



AWS Systems Manager





AWS Systems Manager

Patch Manager

- Deploy operating system and software patches automatically across large groups of Amazon EC2 or on-premises instances

Compliance

- Scan managed instances for patch compliance and configuration inconsistencies



AWS Systems Manager

Session Manager

- Secure remote management of your instances at scale without logging into your servers
- Replaces the need for bastion hosts, SSH, or remote PowerShell

Parameter Store

- Parameter Store provides secure, hierarchical storage for configuration data management and secrets management

Manage EC2 Instances



AWS Service Catalog



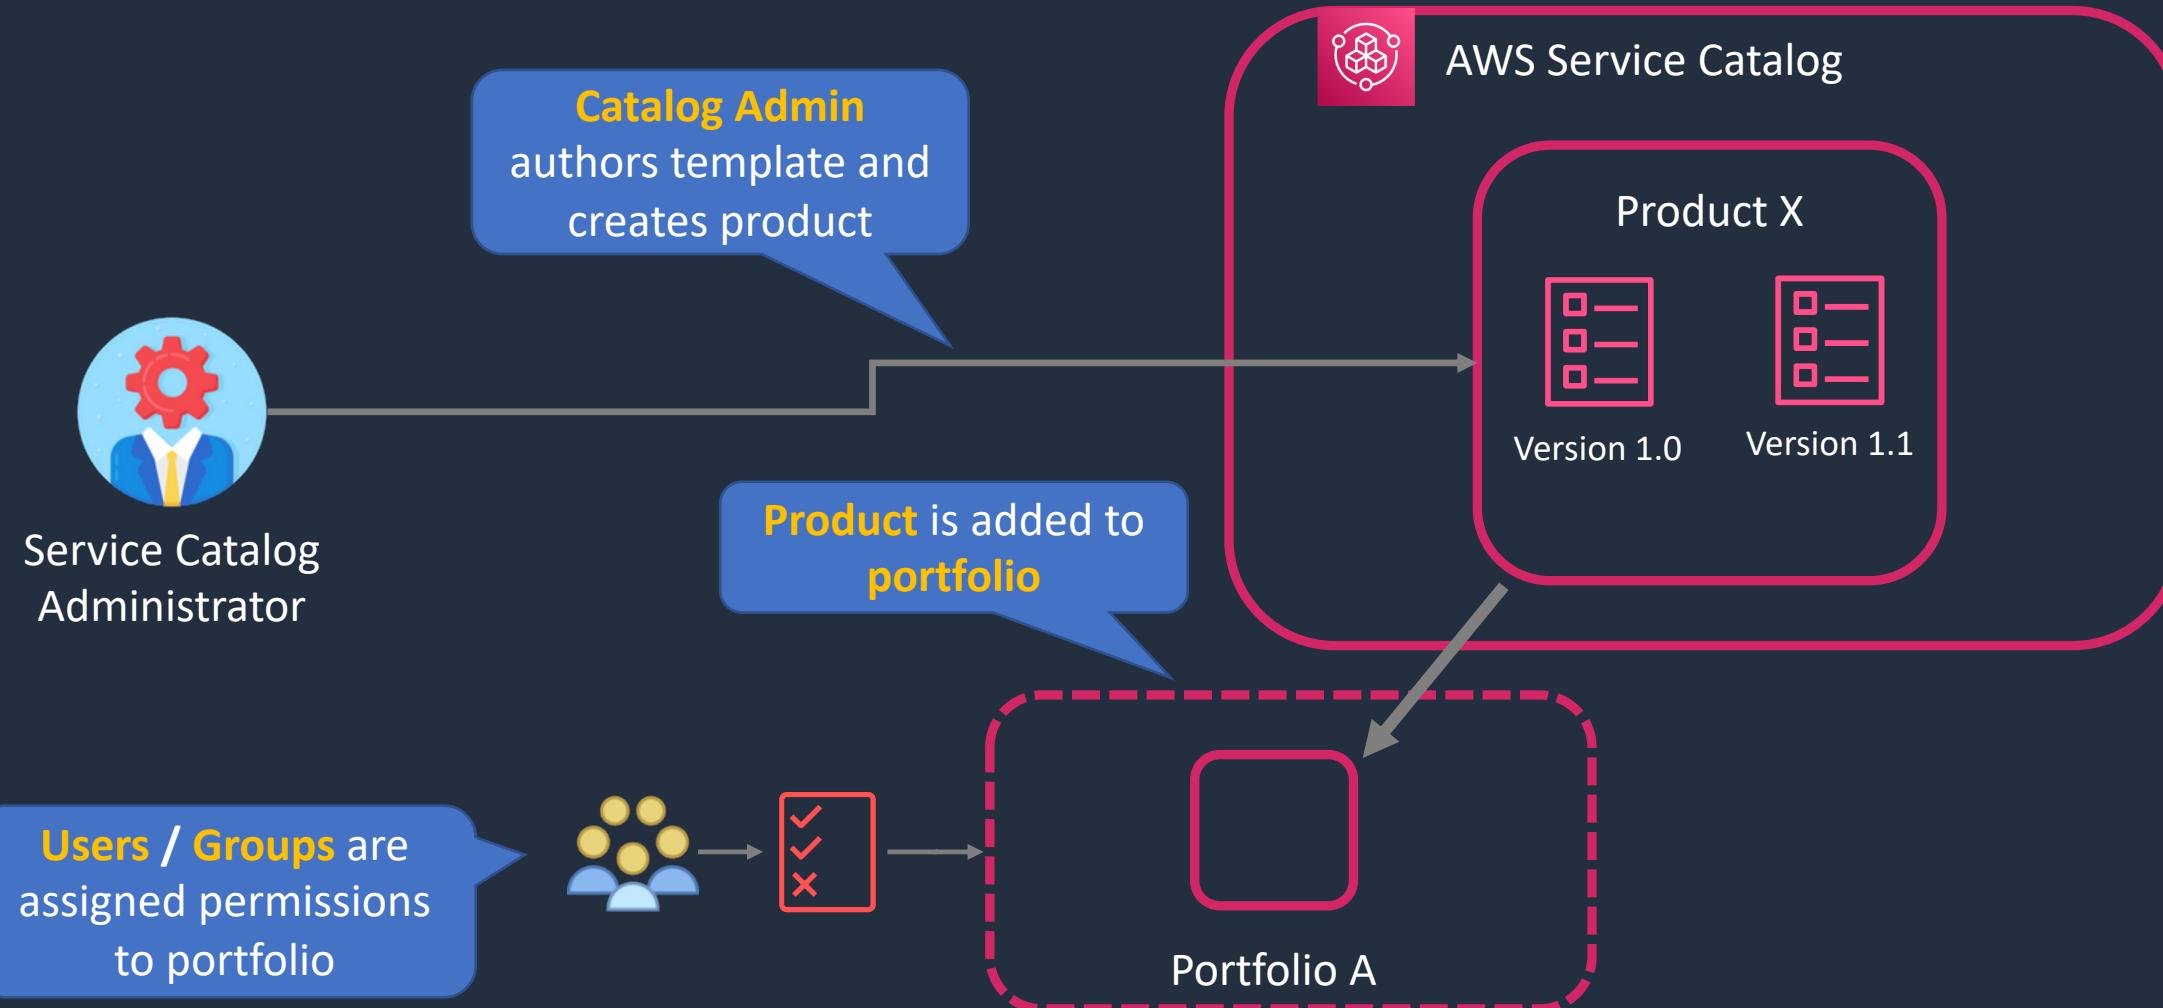


AWS Service Catalog

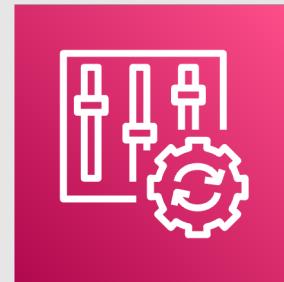
- AWS Service Catalog allows organizations to create and manage **catalogs of IT services** that are approved for use on AWS
- AWS Service Catalog allows you to **centrally manage** commonly deployed IT services
- IT services can include virtual machine images, servers, software, and databases and multi-tier application architectures
- Enables users to quickly deploy only the approved IT services they need



AWS Service Catalog



AWS Config

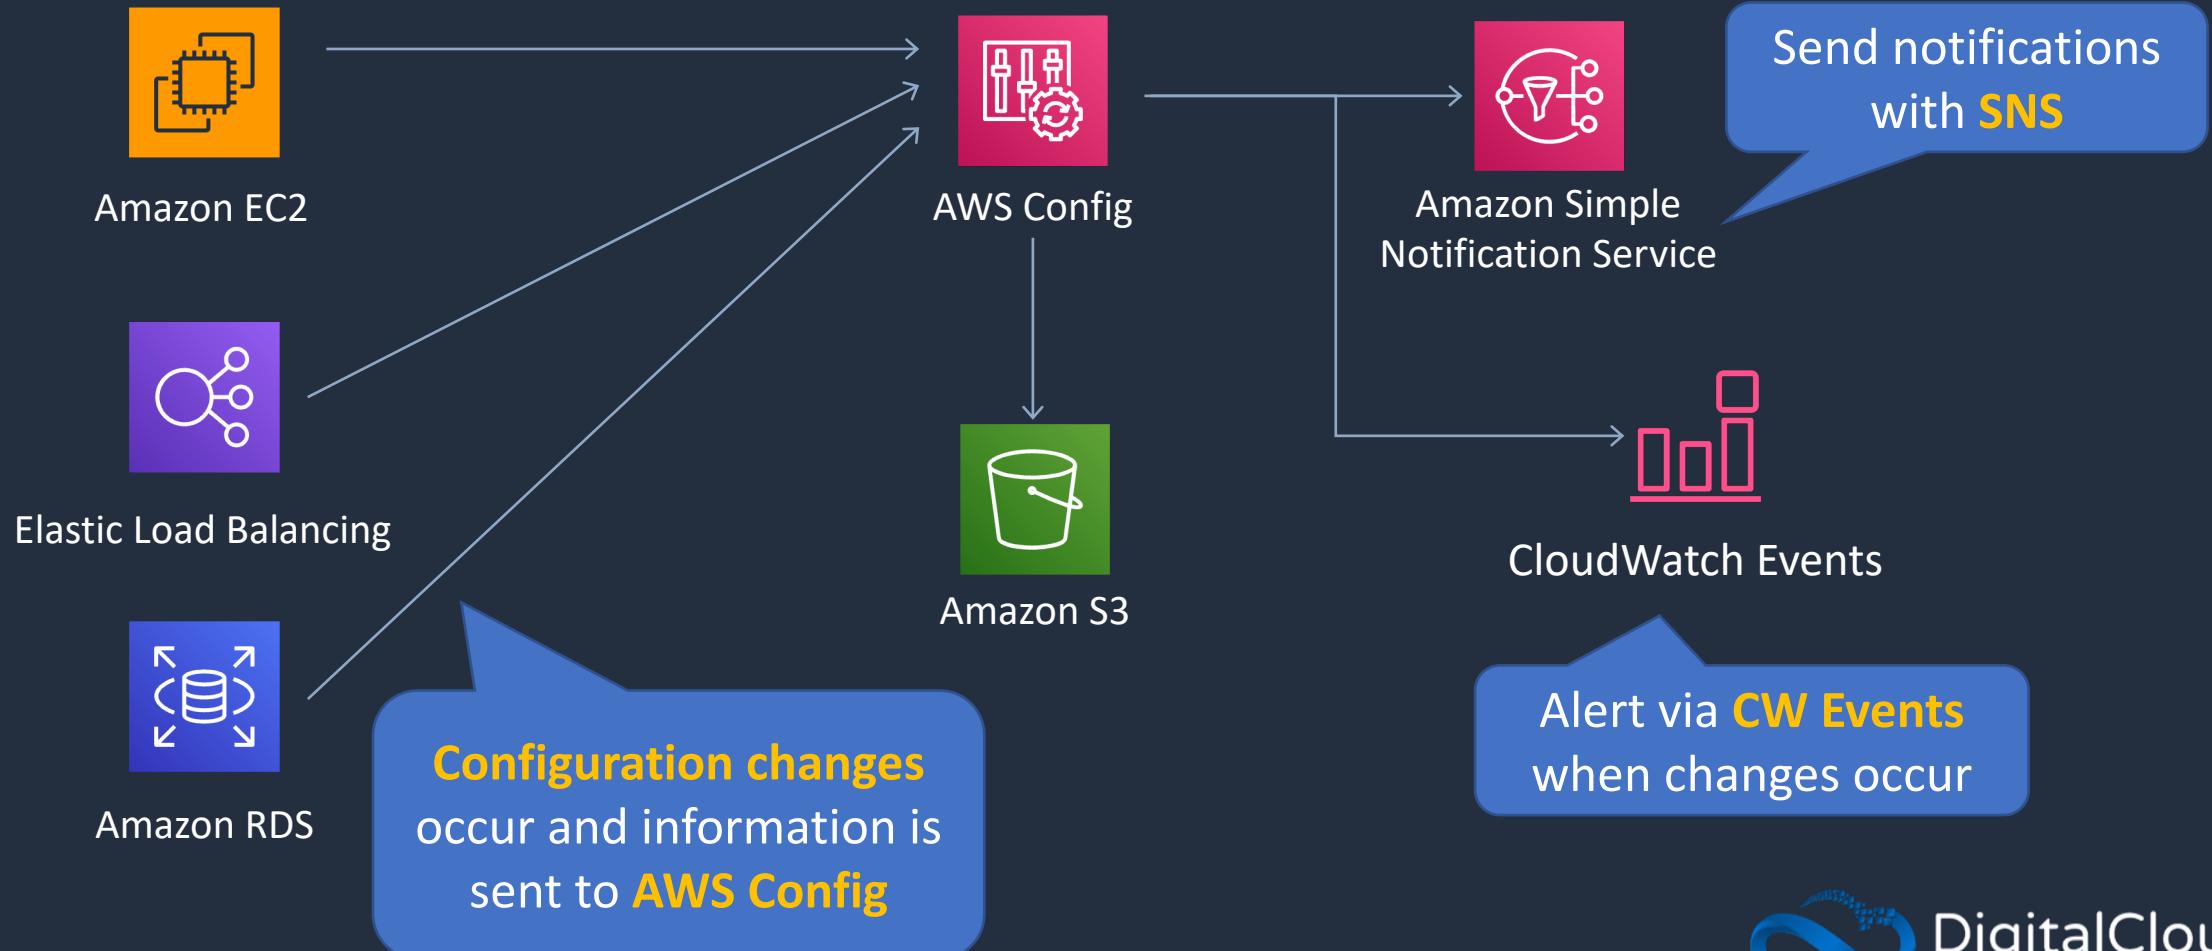




AWS Config

AWS Config evaluates the **configuration** against desired configurations

Example Services:





AWS Config

Example Rule	Description
s3-bucket-server-side-encryption-enabled	Checks that your Amazon S3 bucket either has S3 default encryption enabled or that the S3 bucket policy explicitly denies put-object requests without server side encryption
restricted-ssh	Checks whether security groups that are in use disallow unrestricted incoming SSH traffic
rds-instance-public-access-check	Checks whether the Amazon Relational Database Service (RDS) instances are not publicly accessible
cloudtrail-enabled	Checks whether AWS CloudTrail is enabled in your AWS account

Configuration Compliance



AWS Trusted Advisor





AWS Trusted Advisor

- Trusted Advisor is an online resource that helps to reduce cost, increase performance and improve security by optimizing your AWS environment
- Trusted Advisor provides real time guidance to help you provision your resources following best practices
- Advisor will advise you on **Cost Optimization**, **Performance**, **Security**, and **Fault Tolerance**

AWS Health API and Dashboards





AWS Personal Health Dashboard

- AWS Personal Health Dashboard provides alerts and remediation guidance when AWS is experiencing events that may **impact you**
- Personal Health Dashboard gives you a **personalized** view into the performance and availability of the AWS services underlying your AWS resources
- Also provides proactive notification to help you plan for scheduled activities



AWS Service Health Dashboard

Not personalized information so may not be relevant to you

No proactive notification of scheduled activities

Current Status - Jun 7, 2020 PDT

Amazon Web Services publishes our most up-to-the-minute information on service availability in the table below. Check back here any time to get current status information, or subscribe to an RSS feed to be notified of interruptions to each individual service. If you are experiencing a real-time, operational issue with one of our services that is not described below, please inform us by clicking on the "Contact Us" link to submit a service issue report. All dates and times are Pacific Time (PST/PDT).

North America	South America	Europe	Africa	Asia Pacific	Middle East	Contact Us
Recent Events	Details				RSS	
No recent events.						
Remaining Services		Details			RSS	
Alexa for Business (N. Virginia)		Service is operating normally				
Amazon API Gateway (Montreal)		Service is operating normally				
Amazon API Gateway (N. California)		Service is operating normally				
Amazon API Gateway (N. Virginia)		Service is operating normally				
Amazon API Gateway (Ohio)		Service is operating normally				
Amazon API Gateway (Oregon)		Service is operating normally				
Amazon AppStream 2.0 (N. Virginia)		Service is operating normally				
Amazon AppStream 2.0 (Oregon)		Service is operating normally				
Amazon Athena (Montreal)		Service is operating normally				
Amazon Athena (N. Virginia)		Service is operating normally				
Amazon Athena (Ohio)		Service is operating normally				
Amazon Athena (Oregon)		Service is operating normally				

Shows current status information on service availability

SECTION 13

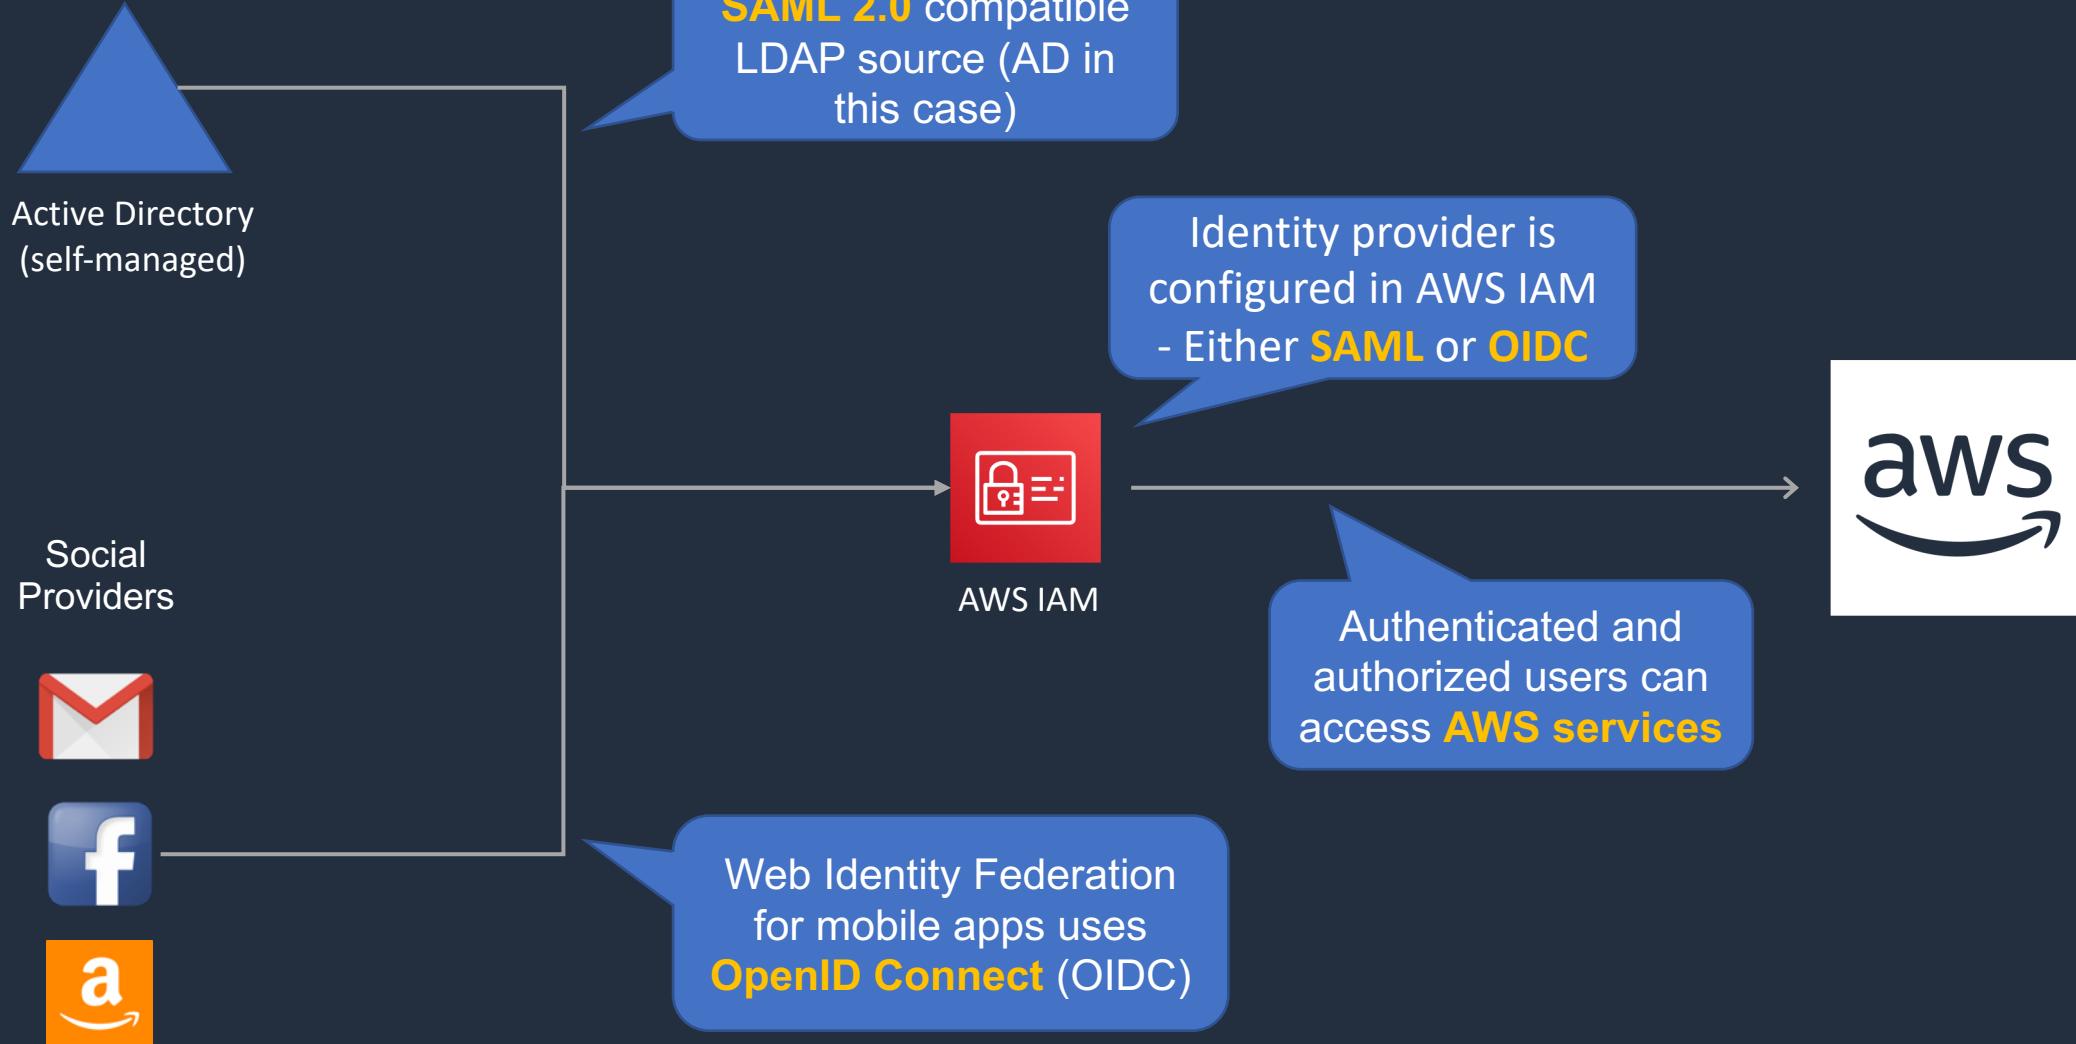
AWS Cloud Security and Identity

Identity Providers and Federation



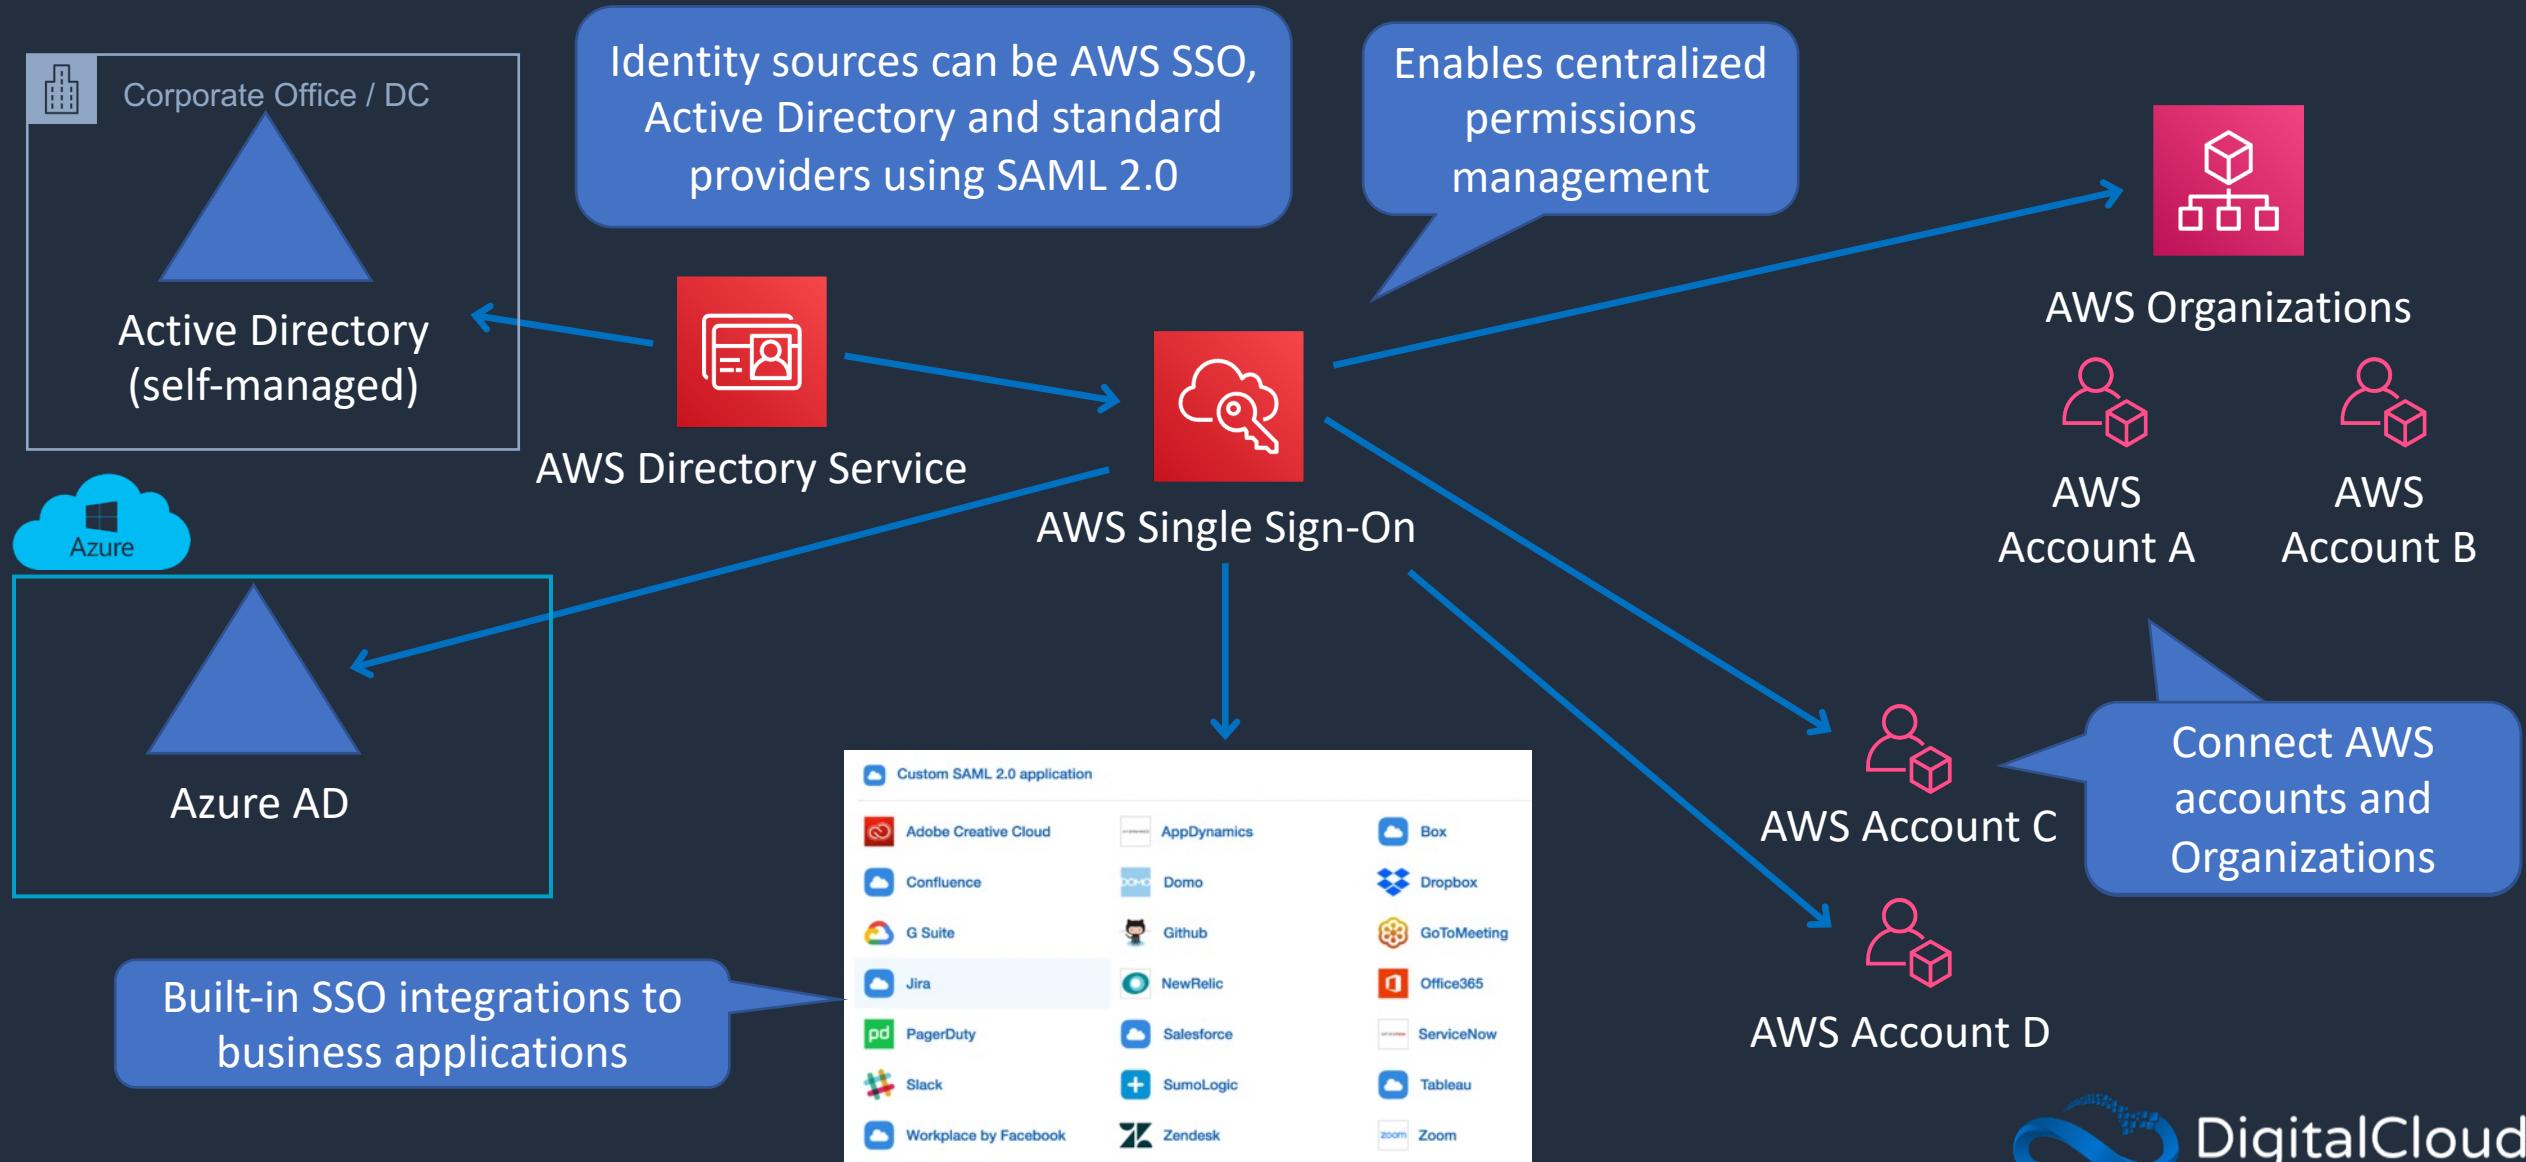


Federation with IAM



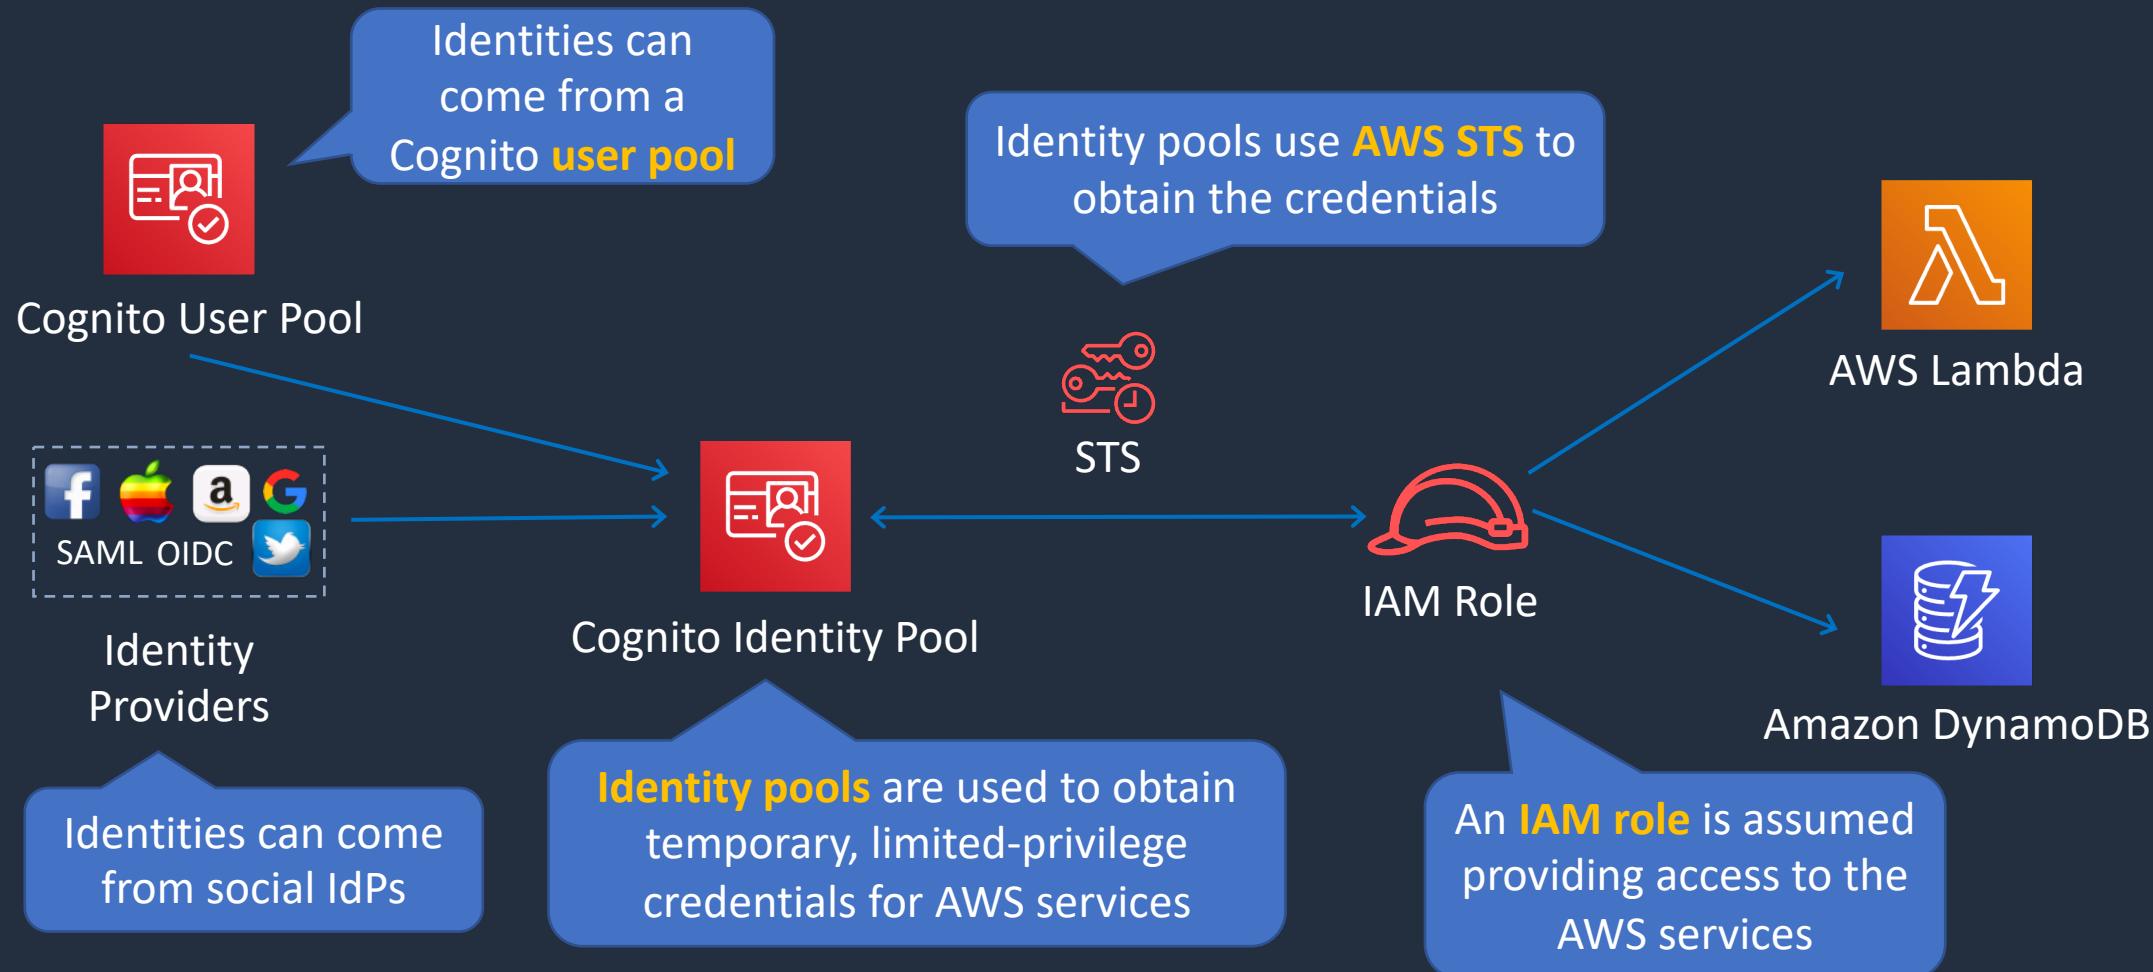


AWS Single Sign-On (SSO)

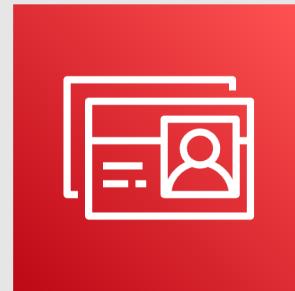




AWS Cognito (Mobile and Web Apps)



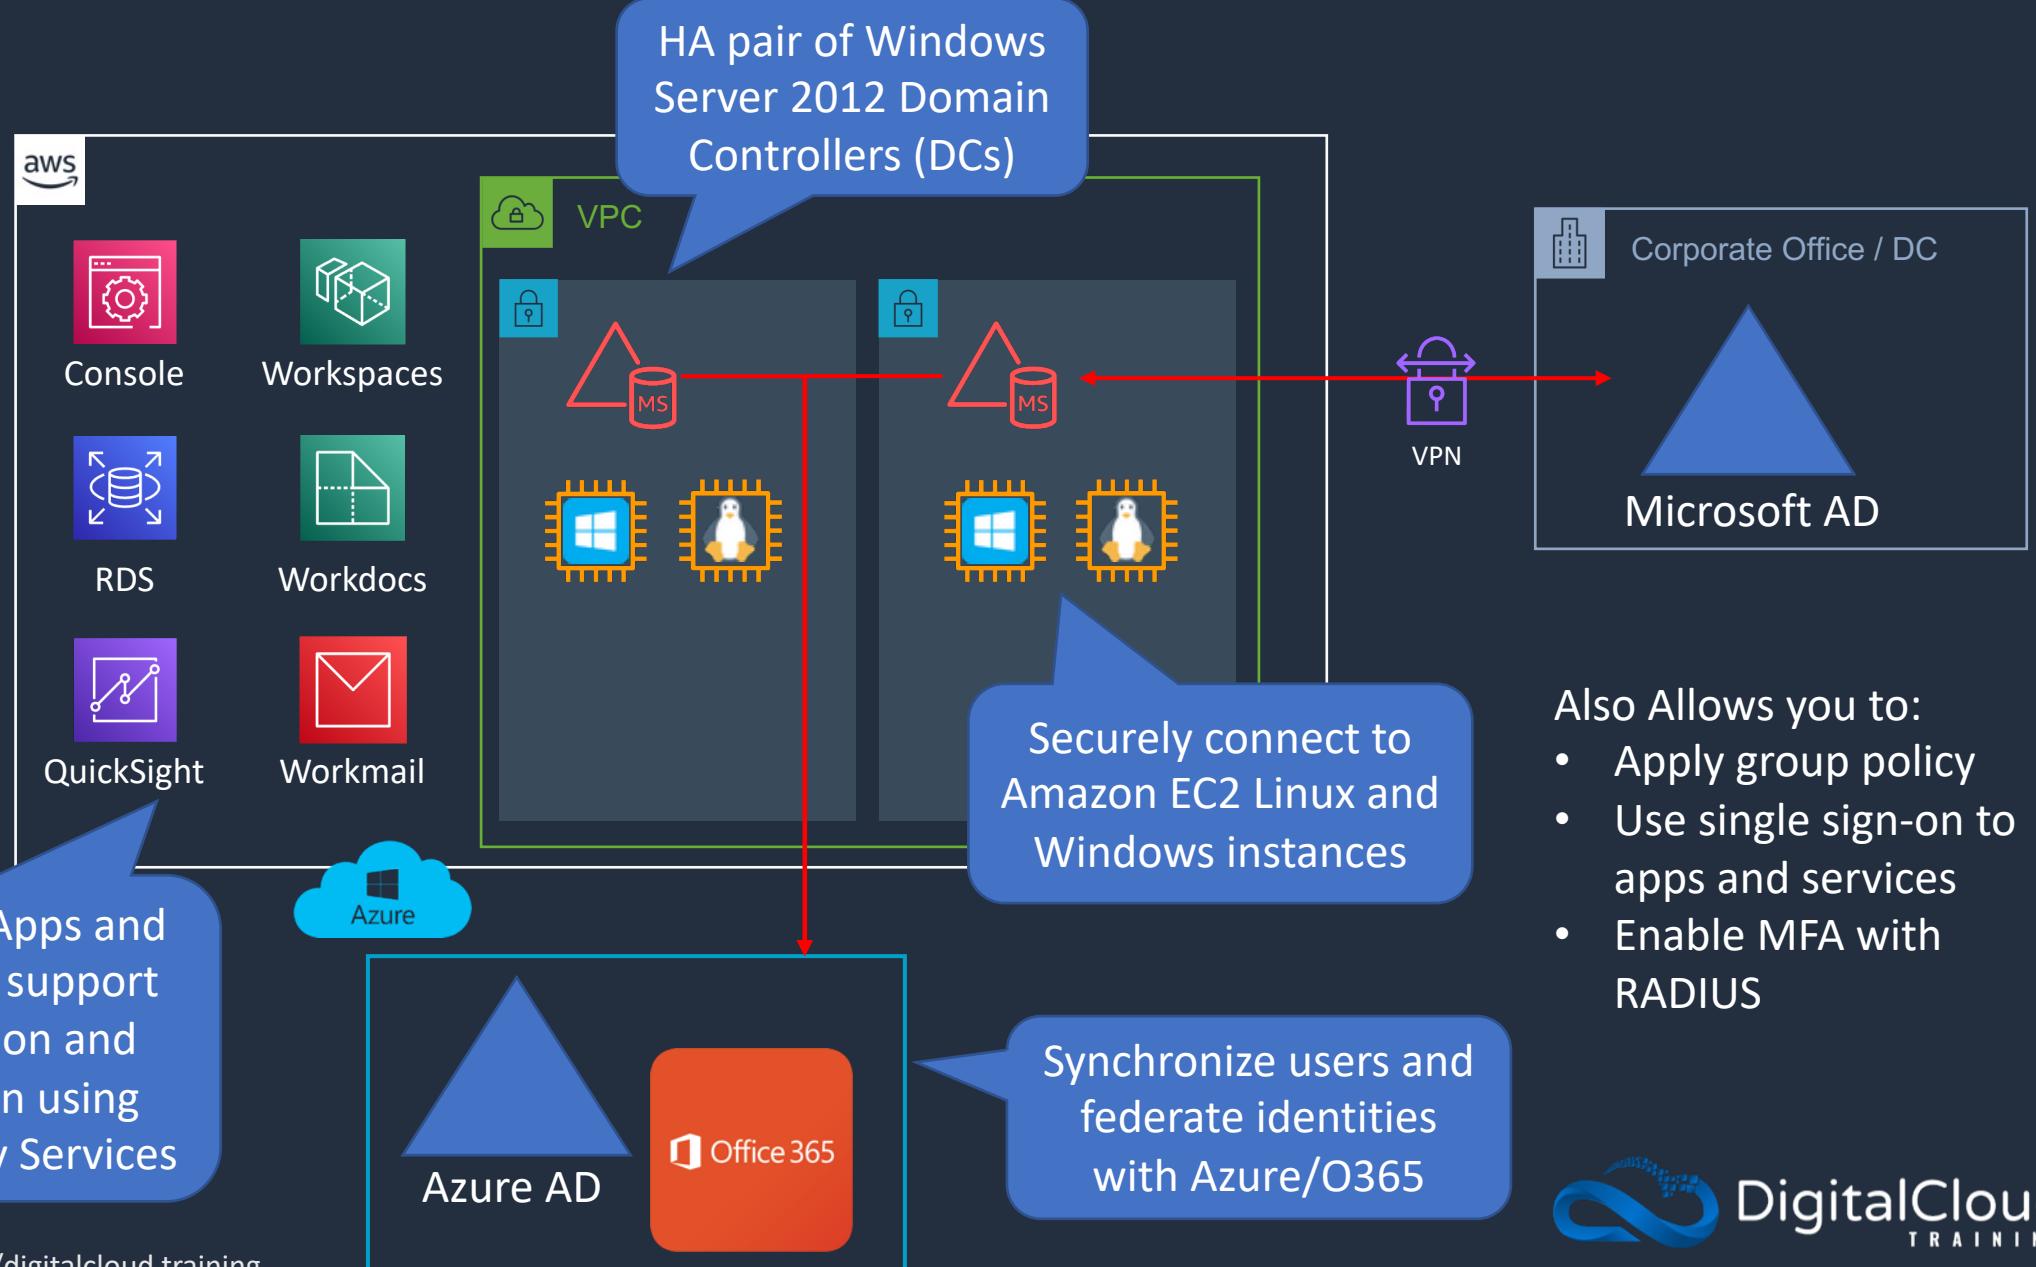
AWS Directory Service





AWS Managed Microsoft Active Directory

Managed implementation of Microsoft Active Directory running on Windows Server 2012 R2



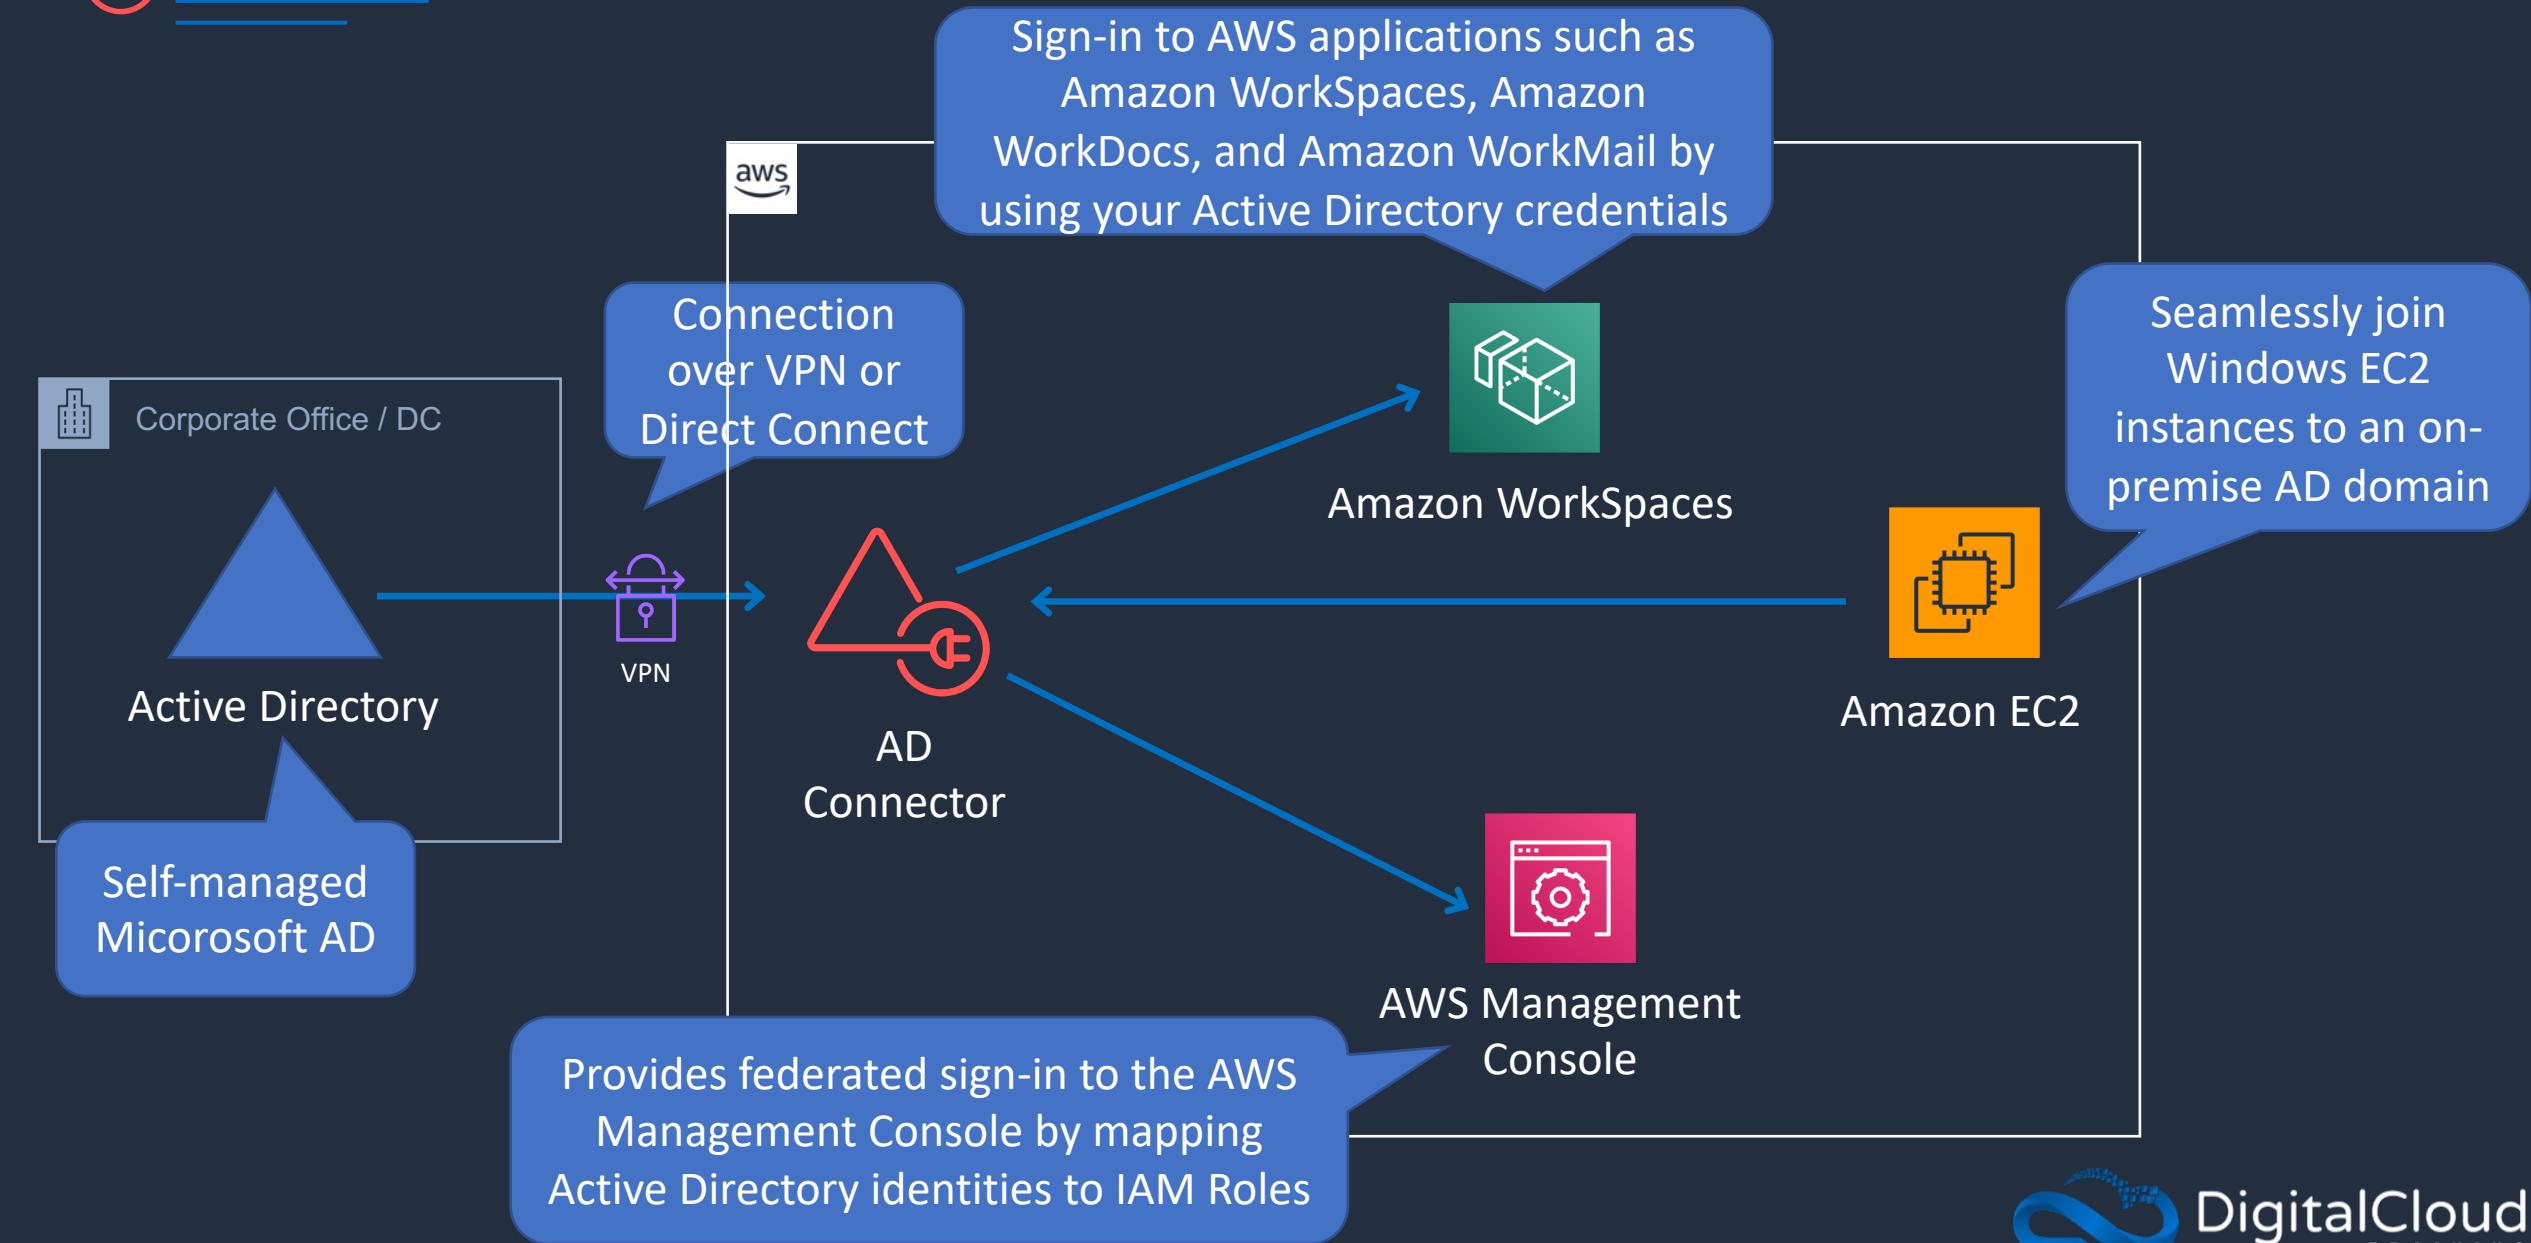


AWS Managed Microsoft Active Directory

- Fully managed AWS services on AWS infrastructure
- Best choice if you have more than 5000 users and/or need a trust relationship set up
- You can setup trust relationships to extend authentication from on-premises Active Directories into the AWS cloud
- On-premise users and groups can access resources in either domain using SSO
- Can be used as a standalone AD in the AWS cloud



AD Connector





AD Connector

- AD Connector is a directory gateway for redirecting directory requests to your on-premise Active Directory.
- AD Connector eliminates the need for directory synchronization and the cost and complexity of hosting a federation infrastructure
- Connects your existing on-premise AD to AWS
- Best choice when you want to use an existing Active Directory with AWS services.



AWS Directory Services Options

Directory Service	Service Description	Use Case
AWS Directory Service for Microsoft Active Directory	AWS-managed full Microsoft AD running on Windows Server 2012 R2	Enterprises that want hosted Microsoft Active Directory
AD Connector	Allows on-premises users to log into AWS services with their existing AD credentials	Single sign-on for on-premises employees
Simple AD	Low scale, low cost, AD implementation based on Samba	Simple user directory, or you need LDAP compatibility

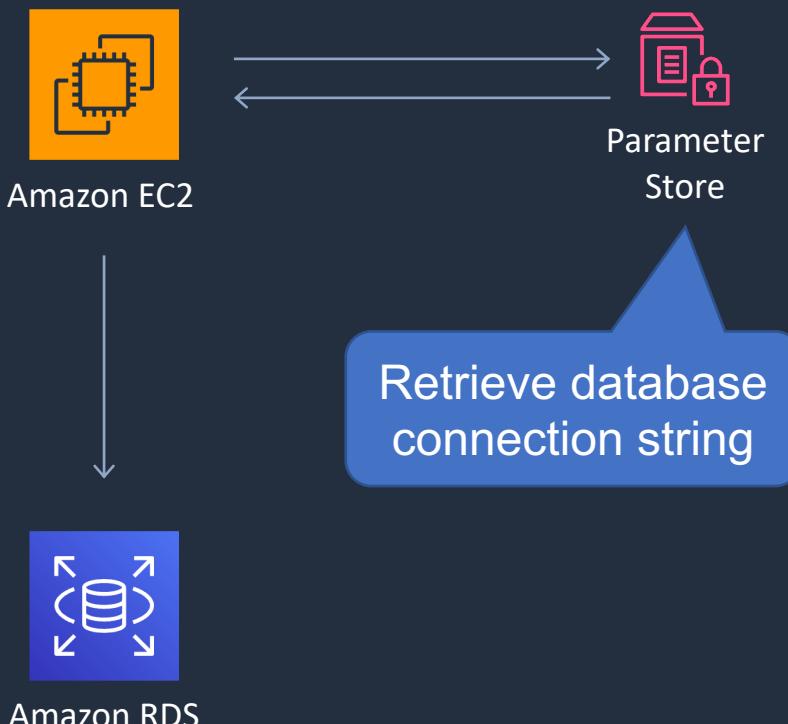
Protecting Secrets





Systems Manager Parameter Store

- Provides secure, hierarchical storage for configuration data management and secrets management
- It is highly scalable, available, and durable
- You can store data such as passwords, database strings, and license codes as parameter values
- You can store values as plaintext (unencrypted data) or ciphertext (encrypted data)
- You can then reference values by using the unique name that you specified when you created the parameter

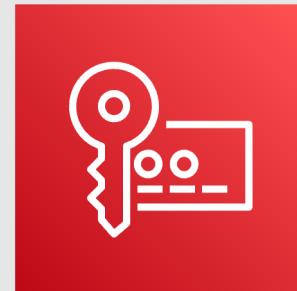




AWS Secrets Manager

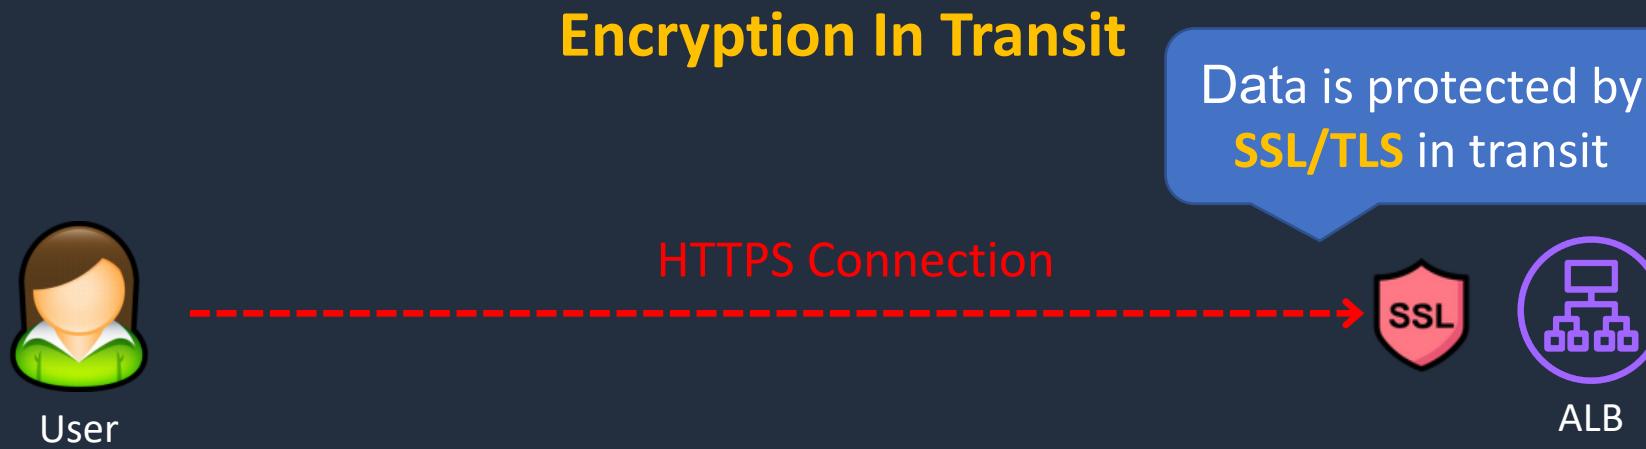
- Similar to Parameter Store
- Allows native and automatic rotation of keys
- Fine-grained permissions
- Central auditing for secret rotation

Encryption



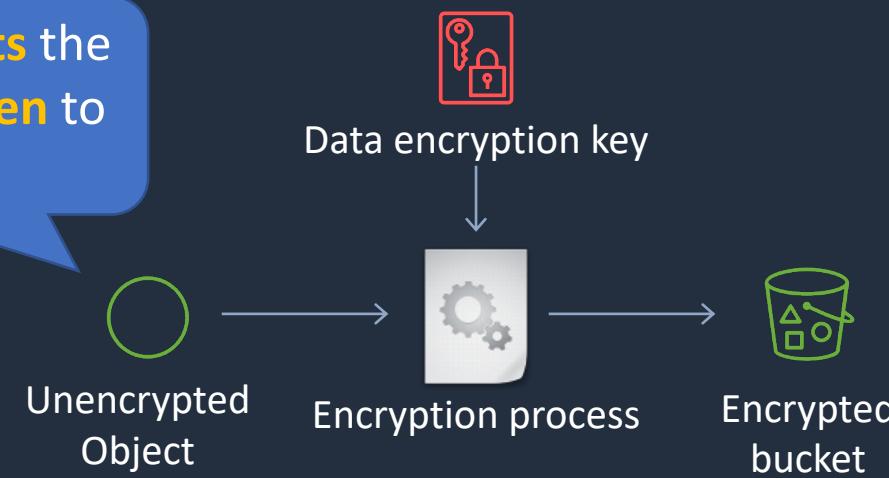


Encryption In Transit vs At Rest



Encryption At Rest

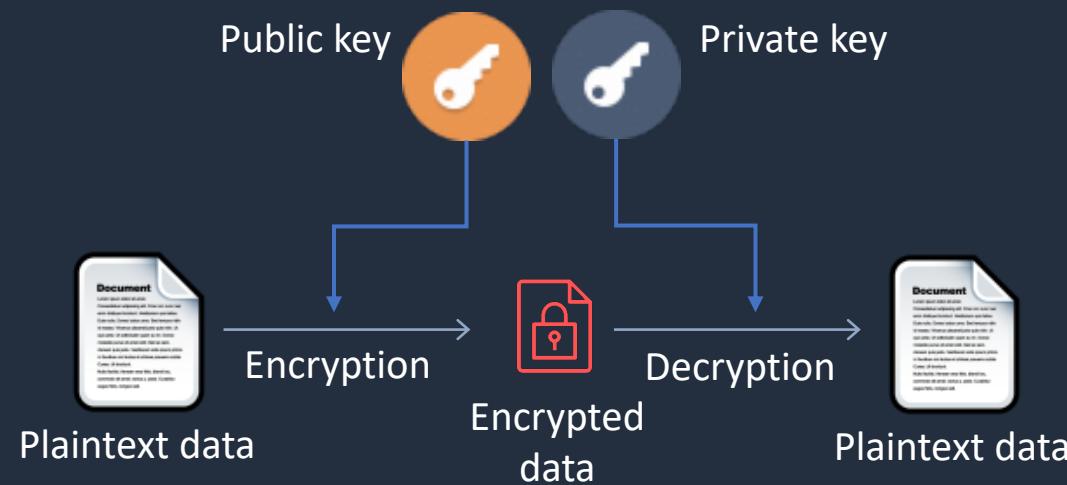
Amazon S3 encrypts the object as it is written to the bucket it





Asymmetric Encryption

- Asymmetric encryption is also known as public key cryptography
- Messages encrypted with the public key can only be decrypted with the private key
- Messages encrypted with the private key can be decrypted with the public key
- Examples include SSL/TLS and SSH





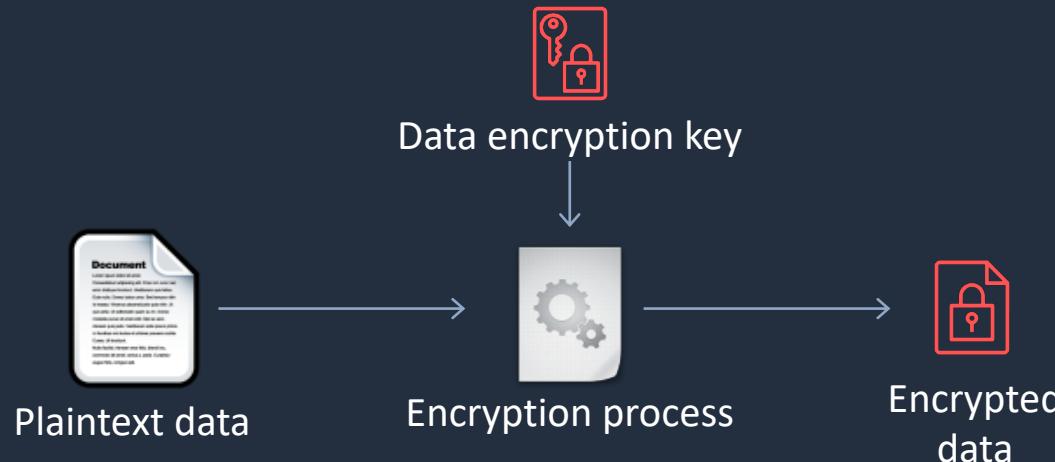
AWS Certificate Manager (ACM)

- Create, store and renew SSL/TLS X.509 certificates
- Single domains, multiple domain names and wildcards
- Integrates with several AWS services including:
 - **Elastic Load Balancing**
 - **Amazon CloudFront**
 - **AWS Elastic Beanstalk**
 - **AWS Nitro Enclaves**
 - **AWS CloudFormation**

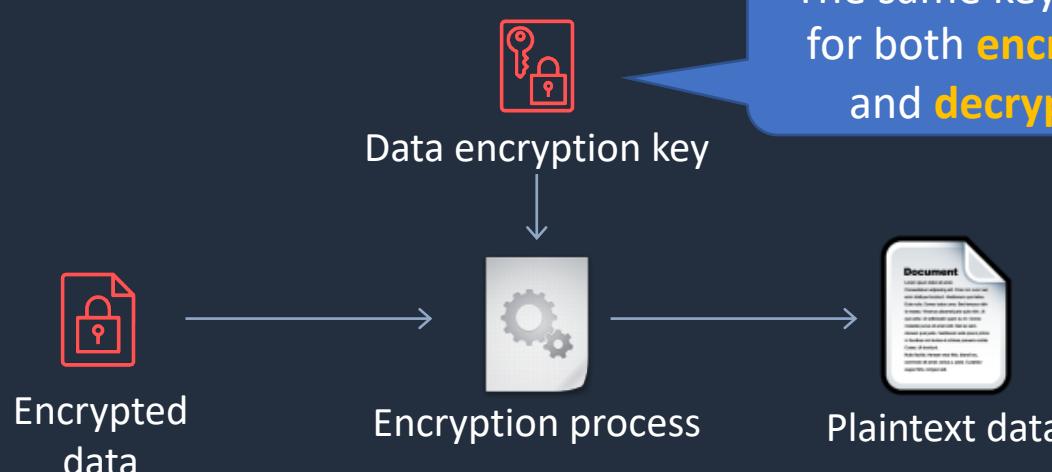


Symmetric Encryption

Encryption



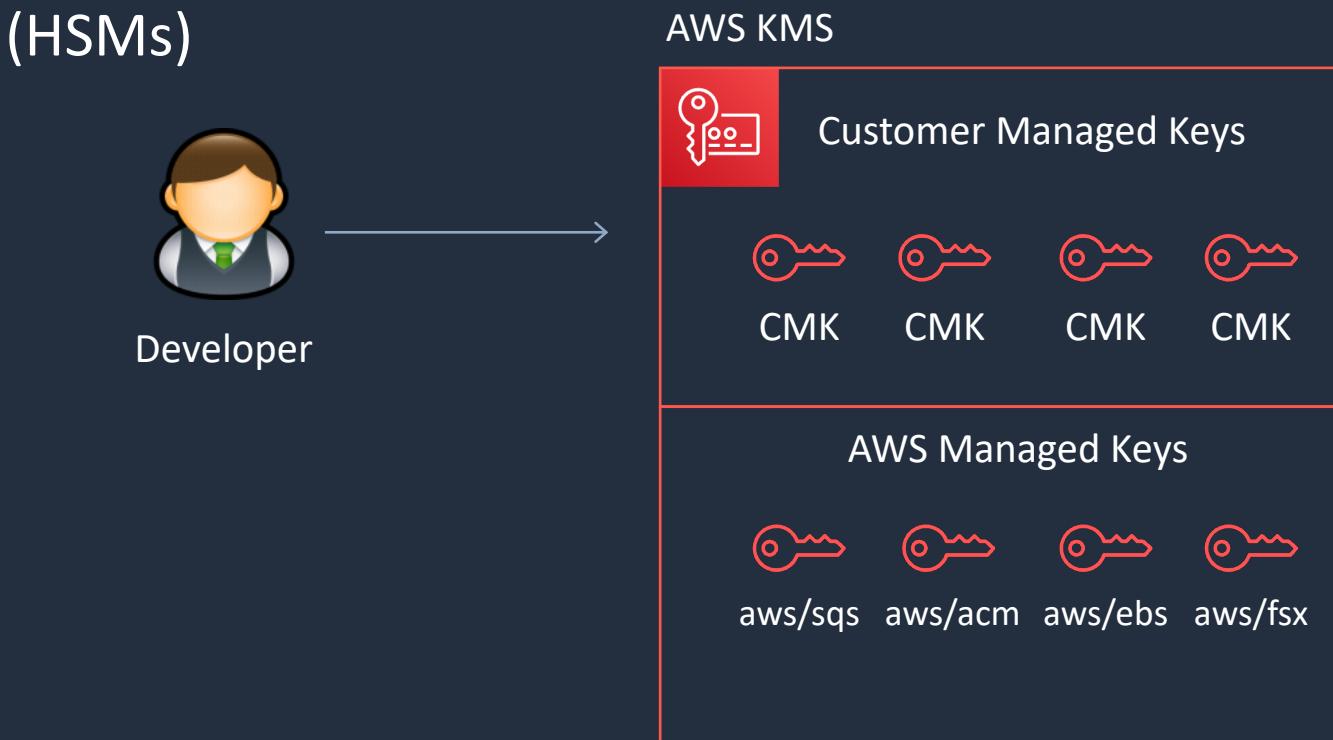
Decryption





AWS Key Management Service (KMS)

- Create and manage **symmetric** and **asymmetric** encryption keys
- The **customer master keys** (CMKs) are protected by hardware security modules (HSMs)





AWS CloudHSM

- AWS CloudHSM is a cloud-based hardware security module (HSM)
- Generate and use your own encryption keys on the AWS Cloud
- Manage your own encryption keys using FIPS 140-2 Level 3 validated HSMs
- CloudHSM runs in your VPC

	CloudHSM	AWS KMS
Tenancy	Single-tenant HSM	Multi-tenant AWS service
Availability	Customer-managed durability and available	Highly available and durable key storage and management
Root of Trust	Customer managed root of trust	AWS managed root of trust
FIPS 140-2	Level 3	Level 2 / Level 3 in some areas
3 rd Party Support	Broad 3 rd Party Support	Broad AWS service support

Encryption on AWS



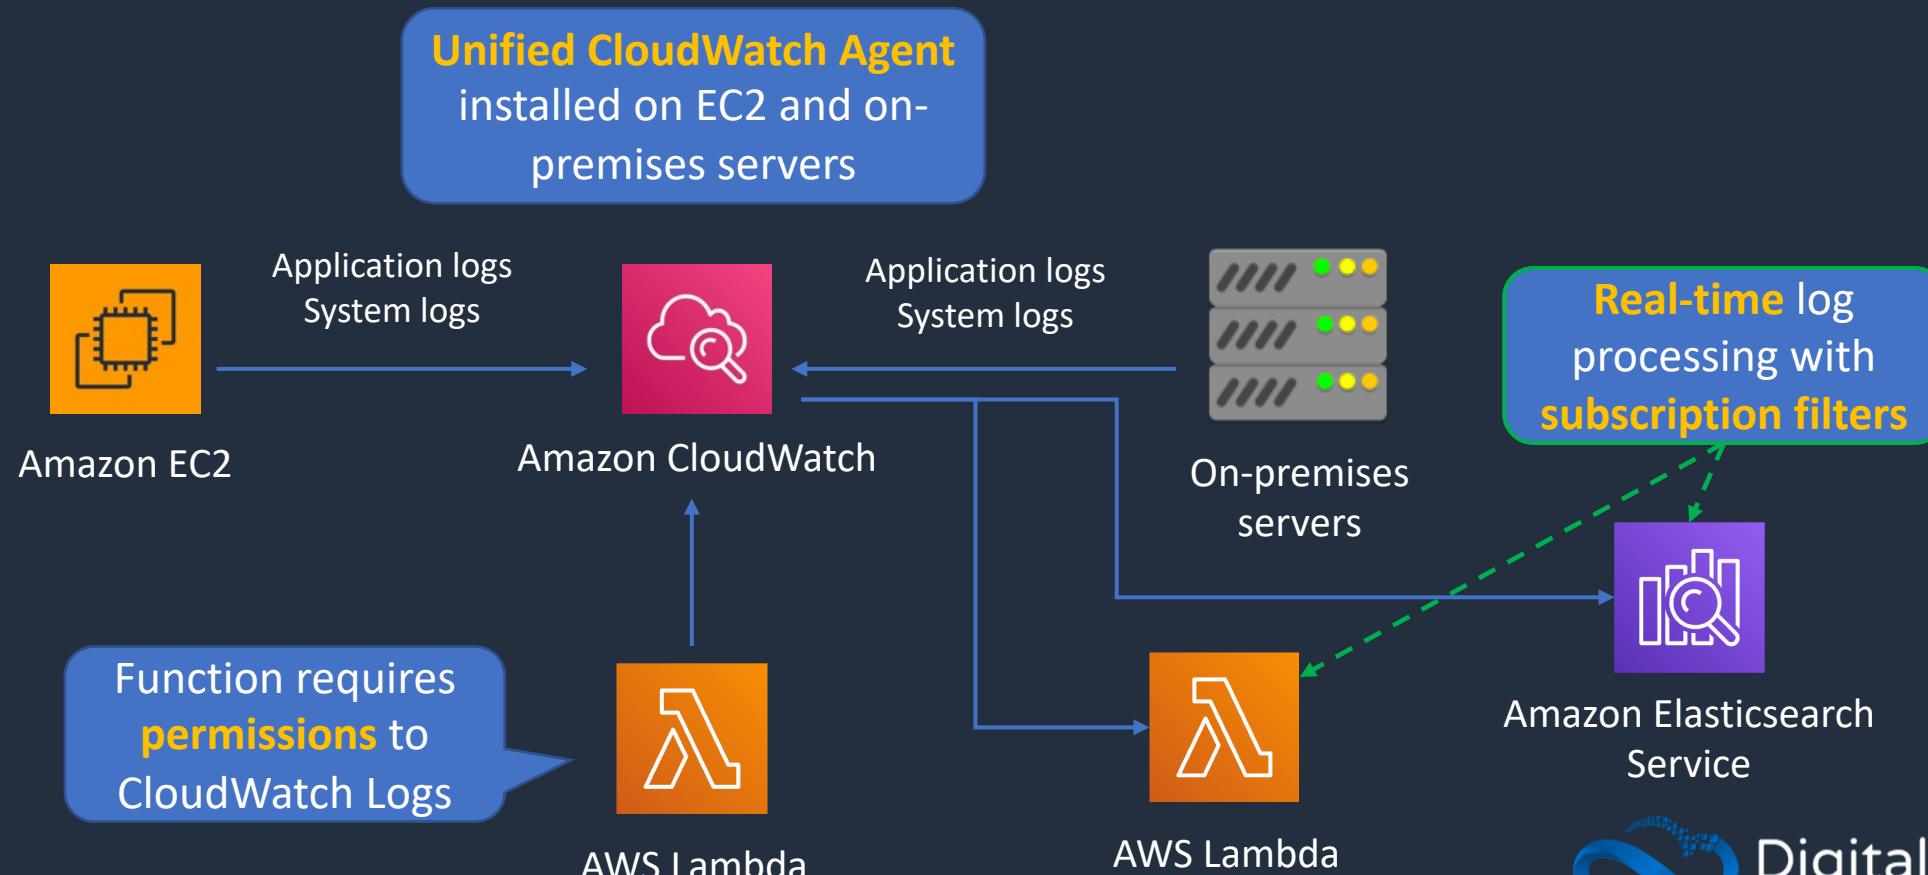
Logging and Auditing





Amazon CloudWatch Logs

- Gather application and system logs in CloudWatch
- Defined expiration policies and KMS encryption





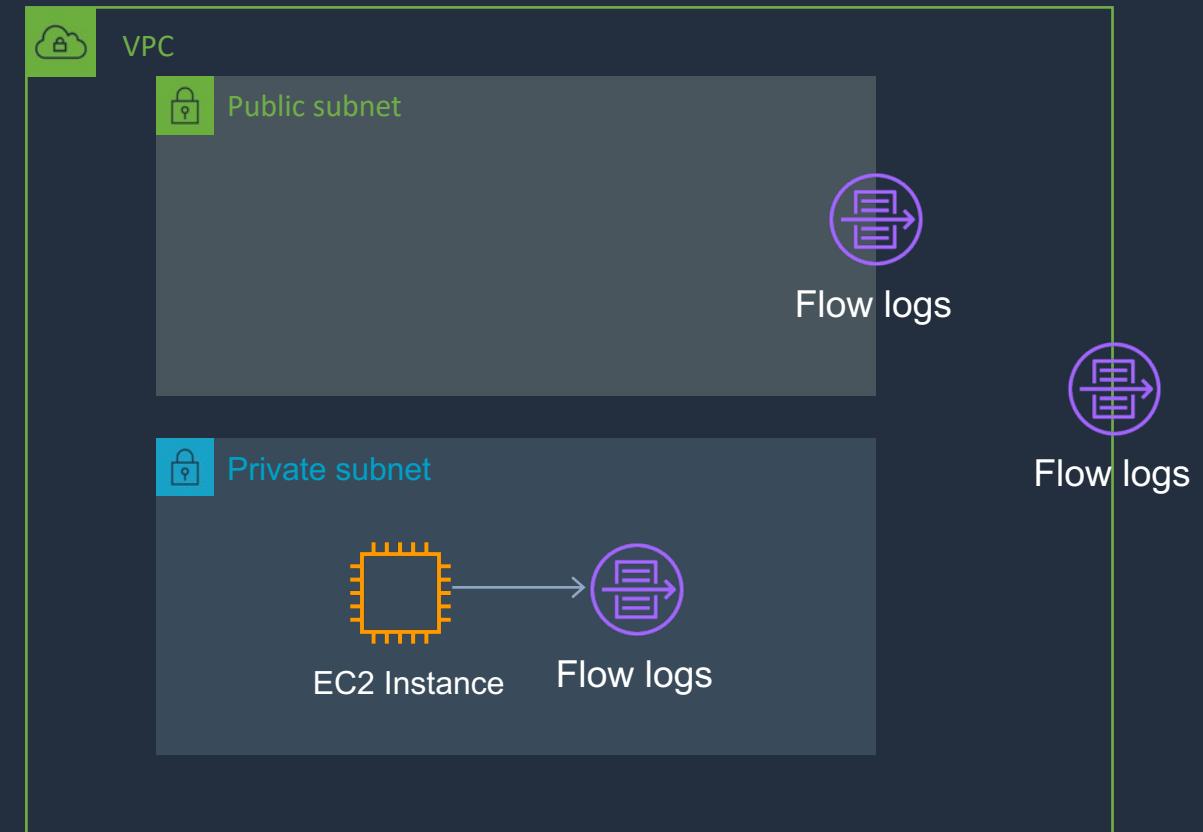
AWS CloudTrail

- CloudTrail logs **API activity** for auditing
- By default, management events are logged and retained for 90 days
- A **CloudTrail Trail** logs any events to S3 for indefinite retention
- Trail can be within Region or all Regions
- CloudWatch Events can be triggered based on API calls in CloudTrail
- Events can be streamed to CloudWatch Logs



VPC Flow Logs

- Flow Logs capture information about the IP traffic going to and from network interfaces in a VPC
- Flow log data is stored using Amazon CloudWatch Logs
- Flow logs can be created at the following levels:
 - VPC
 - Subnet
 - Network interface





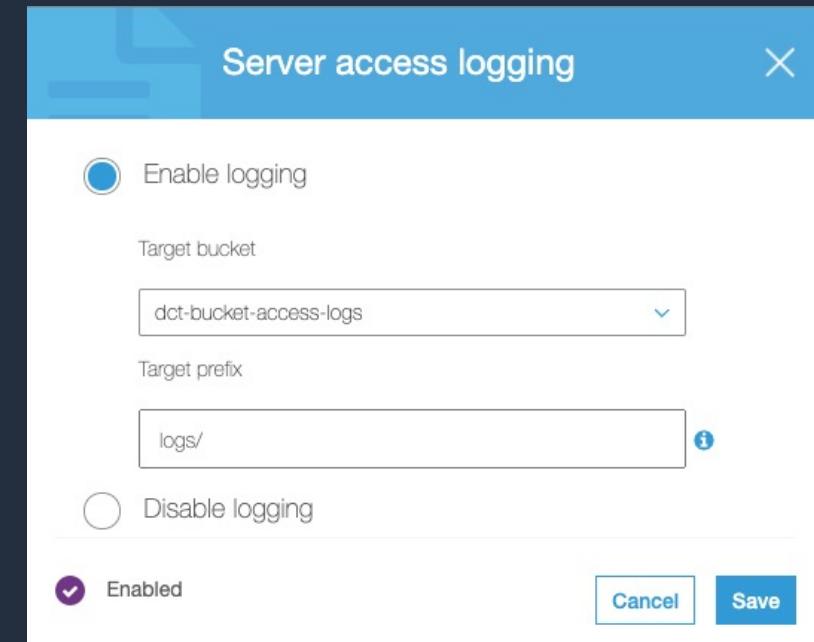
Access Logs

Elastic Load Balancing Access Logs

- capture detailed information about requests sent to the load balancer
- Use to analyze traffic patterns and troubleshoot issues
- Can identify requester, IP, request type etc.
- Can be optionally stored and retained in S3.

S3 Access Logs

- Provides detailed records for the requests that are made to a bucket
- Details include the requester, bucket name, request time, request action, response status, and error code (if applicable)
- Disabled by default



AWS CloudTrail



Detect and Respond





Amazon Detective

- Analyze, investigate, and quickly identify the root cause of potential security issues or suspicious activities
- Automatically collects data from AWS resources
- Uses machine learning, statistical analysis, and graph theory
- Creates a unified, interactive view of resources, users and interactions between them
- Data sources include VPC Flow Logs, CloudTrail, and GuardDuty



AWS GuardDuty

- Intelligent threat detection service
- Detects account compromise, instance compromise, malicious reconnaissance, and bucket compromise
- Continuous monitoring for events across:
 - **AWS CloudTrail Management Events**
 - **AWS CloudTrail S3 Data Events**
 - **Amazon VPC Flow Logs**
 - **DNS Logs**

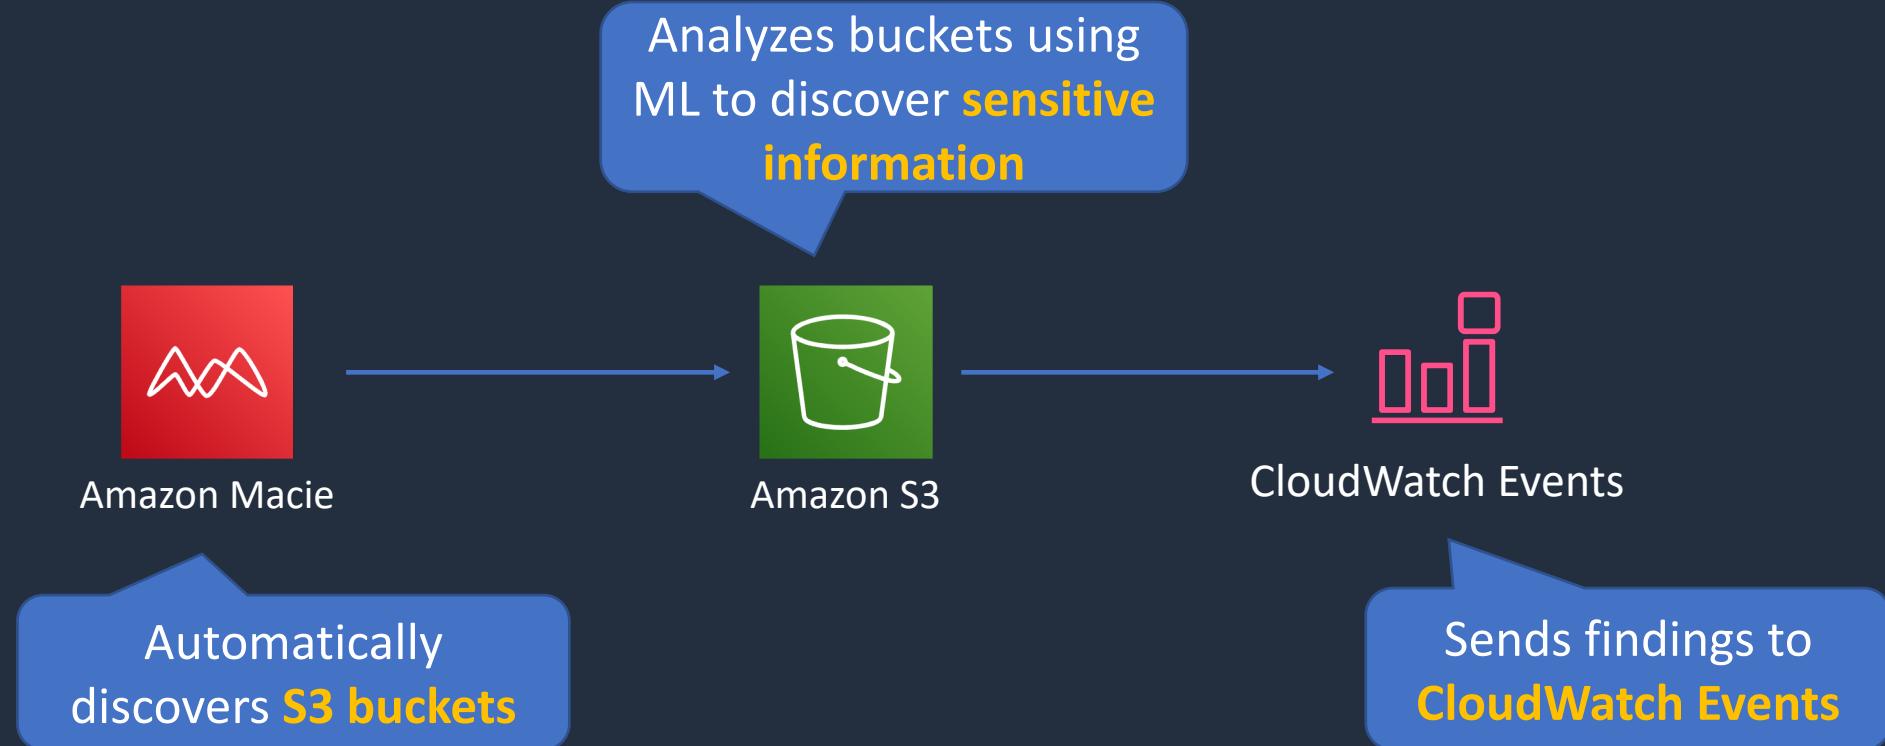


Amazon Macie

- Macie is a fully managed data security and data privacy service
- Uses machine learning and pattern matching to discover, monitor, and help you protect your sensitive data on Amazon S3
- Macie enables security compliance and preventive security
- Can Identify a variety of data types, including PII, Protected Health Information (PHI), regulatory documents, API keys, and secret keys



Amazon Macie



Firewalls and DDoS Protection



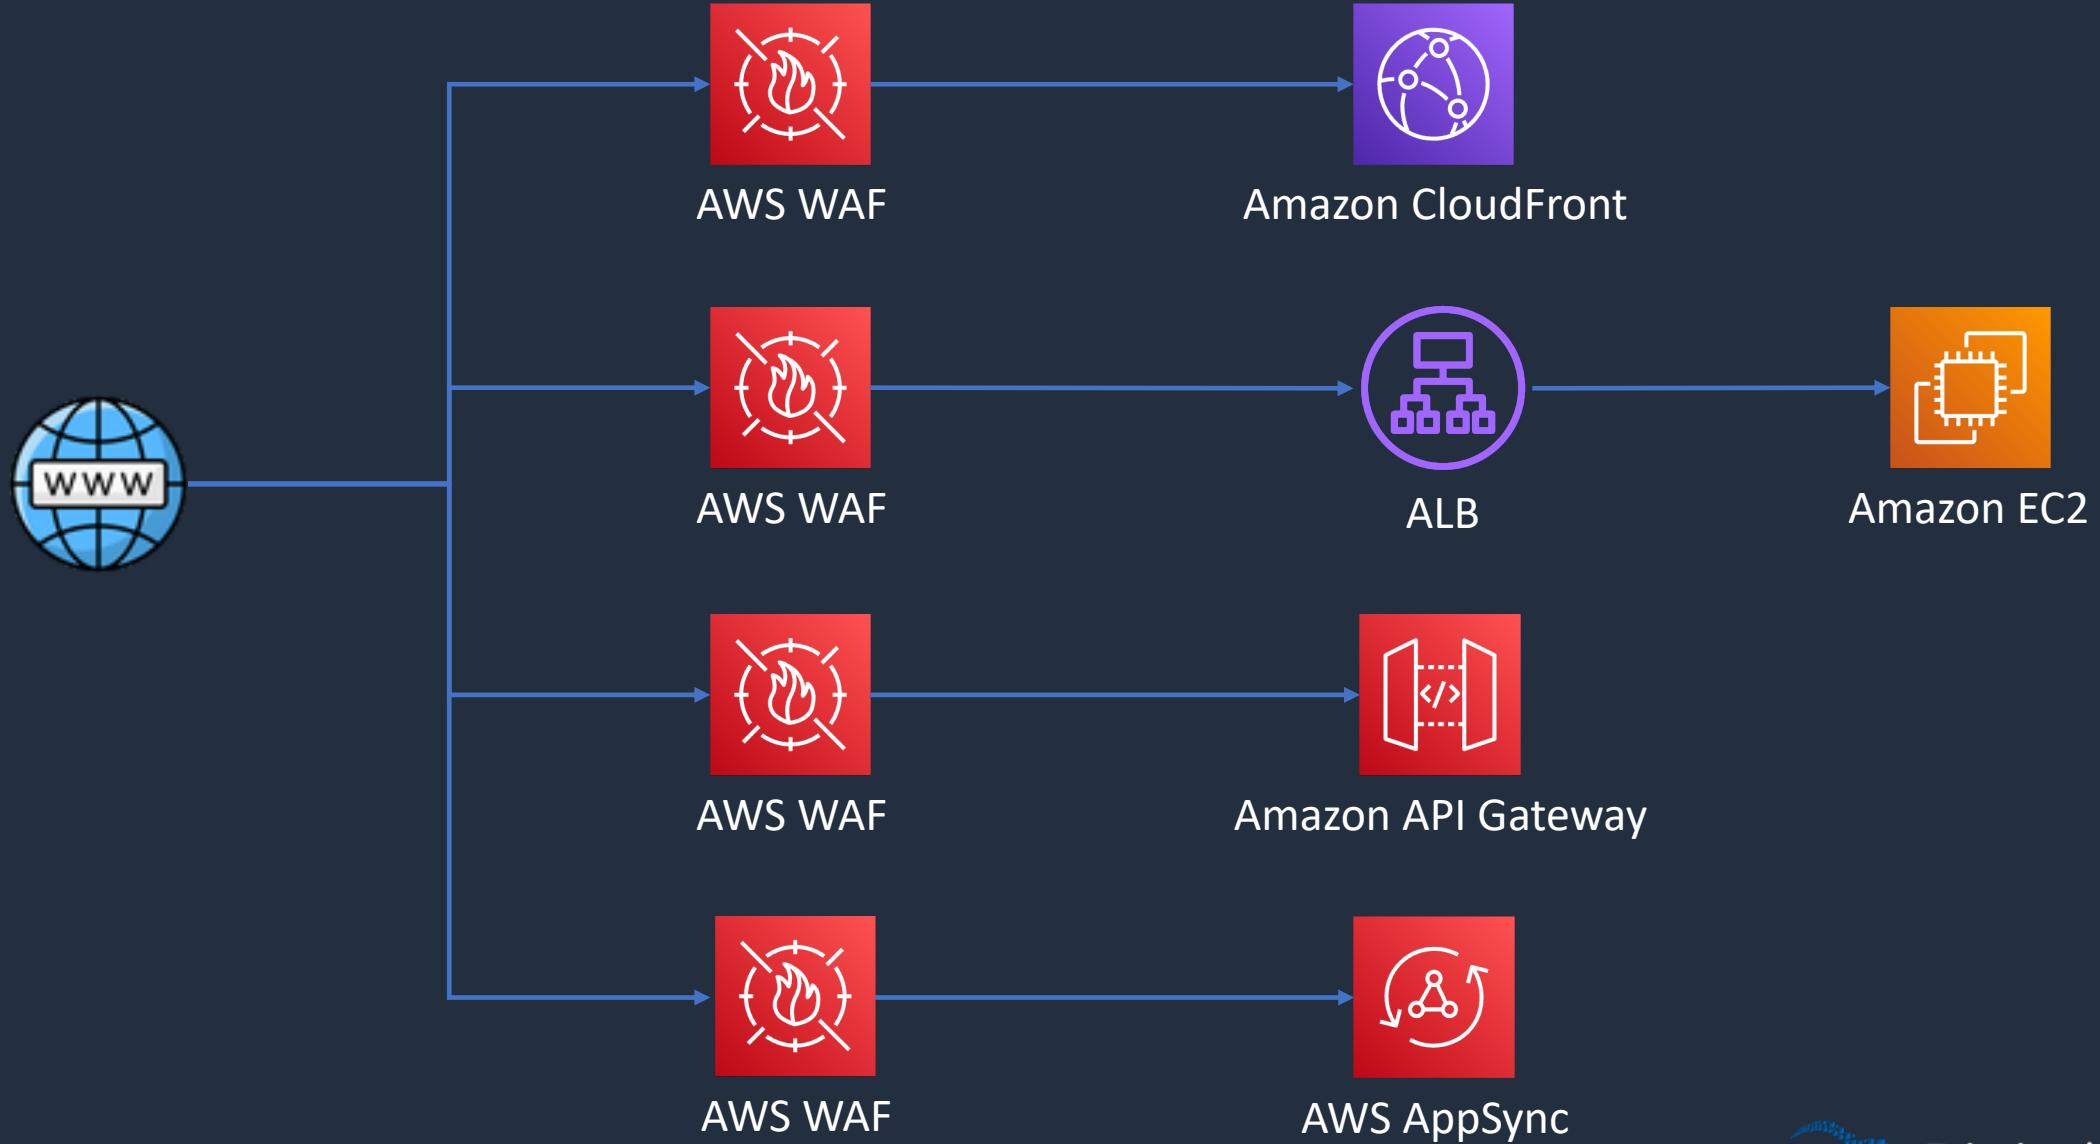


AWS Web Application Firewall (WAF)

- AWS WAF is a **web application firewall**
- WAF lets you create rules to filter web traffic based on conditions that include IP addresses, HTTP headers and body, or custom URIs
- WAF makes it easy to create rules that block common web exploits like **SQL injection** and **cross site scripting**
- The rules are known as Web ACLs



AWS Web Application Firewall (WAF)





AWS Shield

- AWS Shield is a managed **Distributed Denial of Service** (DDoS) protection service
- Safeguards web application running on AWS with always-on detection and automatic inline mitigations
- Helps to minimize application downtime and latency
- Two tiers –
 - **Standard** – no cost
 - **Advanced** - \$3k USD per month and 1 year commitment
- Integrated with Amazon CloudFront (standard included by default)

Compliance Services





AWS Artifact

- AWS Artifact provides on-demand access to AWS' security and compliance reports and select online agreements
- Reports available in AWS Artifact include:
 - Service Organization Control (SOC) reports
 - Payment Card Industry (PCI) reports
- Provides certifications from accreditation bodies across geographies and compliance verticals that validate the implementation and operating effectiveness of **AWS security controls**
- Agreements available in AWS Artifact include the Business Associate Addendum (BAA) and the Nondisclosure Agreement (NDA)

Security Management and Support





AWS Security Hub

- Provides a comprehensive view of security alerts and security posture **across AWS accounts**
- Aggregates, organizes, and prioritizes security alerts, or findings, from multiple AWS services
- Continuously monitors your environment using automated security checks
- Configure security standards to validate against
 - AWS Foundational Security Best Practices v1.0.0
 - CIS AWS Foundations Benchmark v1.2.0
 - PCI DSS v3.2.1



AWS Security Bulletins

- Security and privacy events affecting AWS services are published (also has an RSS feed)

Content Type	
<input checked="" type="checkbox"/> Important	Sudo Security Issue (CVE-2021-3156) AWS-2021-001, 01/27/2021
<input type="checkbox"/> Informational	Xen Security Advisory (XSA-286) AWS-2020-005, 10/23/2020
Year	
<input type="checkbox"/> 2021	Container Networking Security Issue (CVE-2020-8558) AWS-2020-002v2, 07/09/2020
<input type="checkbox"/> 2020	Minimum Version of TLS 1.2 Required for FIPS Endpoints by March 31, 2021 AWS-2020-001, 03/31/2020
<input type="checkbox"/> 2019	Kubernetes Security Issue (CVE-2019-11249) AWS-2019-007, 08/15/2019
<input type="checkbox"/> 2018	Kubernetes Security Issue (CVE-2019-11246) AWS-2019-006, 07/02/2019
<input type="checkbox"/> 2017	Linux Kernel TCP SACK Denial of Service Issues AWS-2019-005, 06/17/2019
<input type="checkbox"/> 2016	
<input type="checkbox"/> 2015	
<input type="checkbox"/> 2014	



AWS Trust & Safety Team

- Contact the **AWS Trust & Safety** team if AWS resources are being used for:
 - Spam
 - Port scanning
 - Denial-of-service attacks
 - Intrusion attempts
 - Hosting of objectionable or copyrighted content
 - Distributing malware
- Email address is: abuse@amazonaws.com

Penetration testing





Penetration Testing

- Penetration testing is the practice of testing one's own application's security for vulnerabilities by simulating an attack
- AWS allows penetration testing without prior approval for 8 AWS services



Penetration Testing

Permitted services

- Amazon EC2 instances, NAT Gateways, and Elastic Load Balancers
- Amazon RDS
- Amazon CloudFront
- Amazon Aurora
- Amazon API Gateways
- AWS Lambda and Lambda Edge functions
- Amazon Lightsail resources
- Amazon Elastic Beanstalk environments

Prohibited Activities

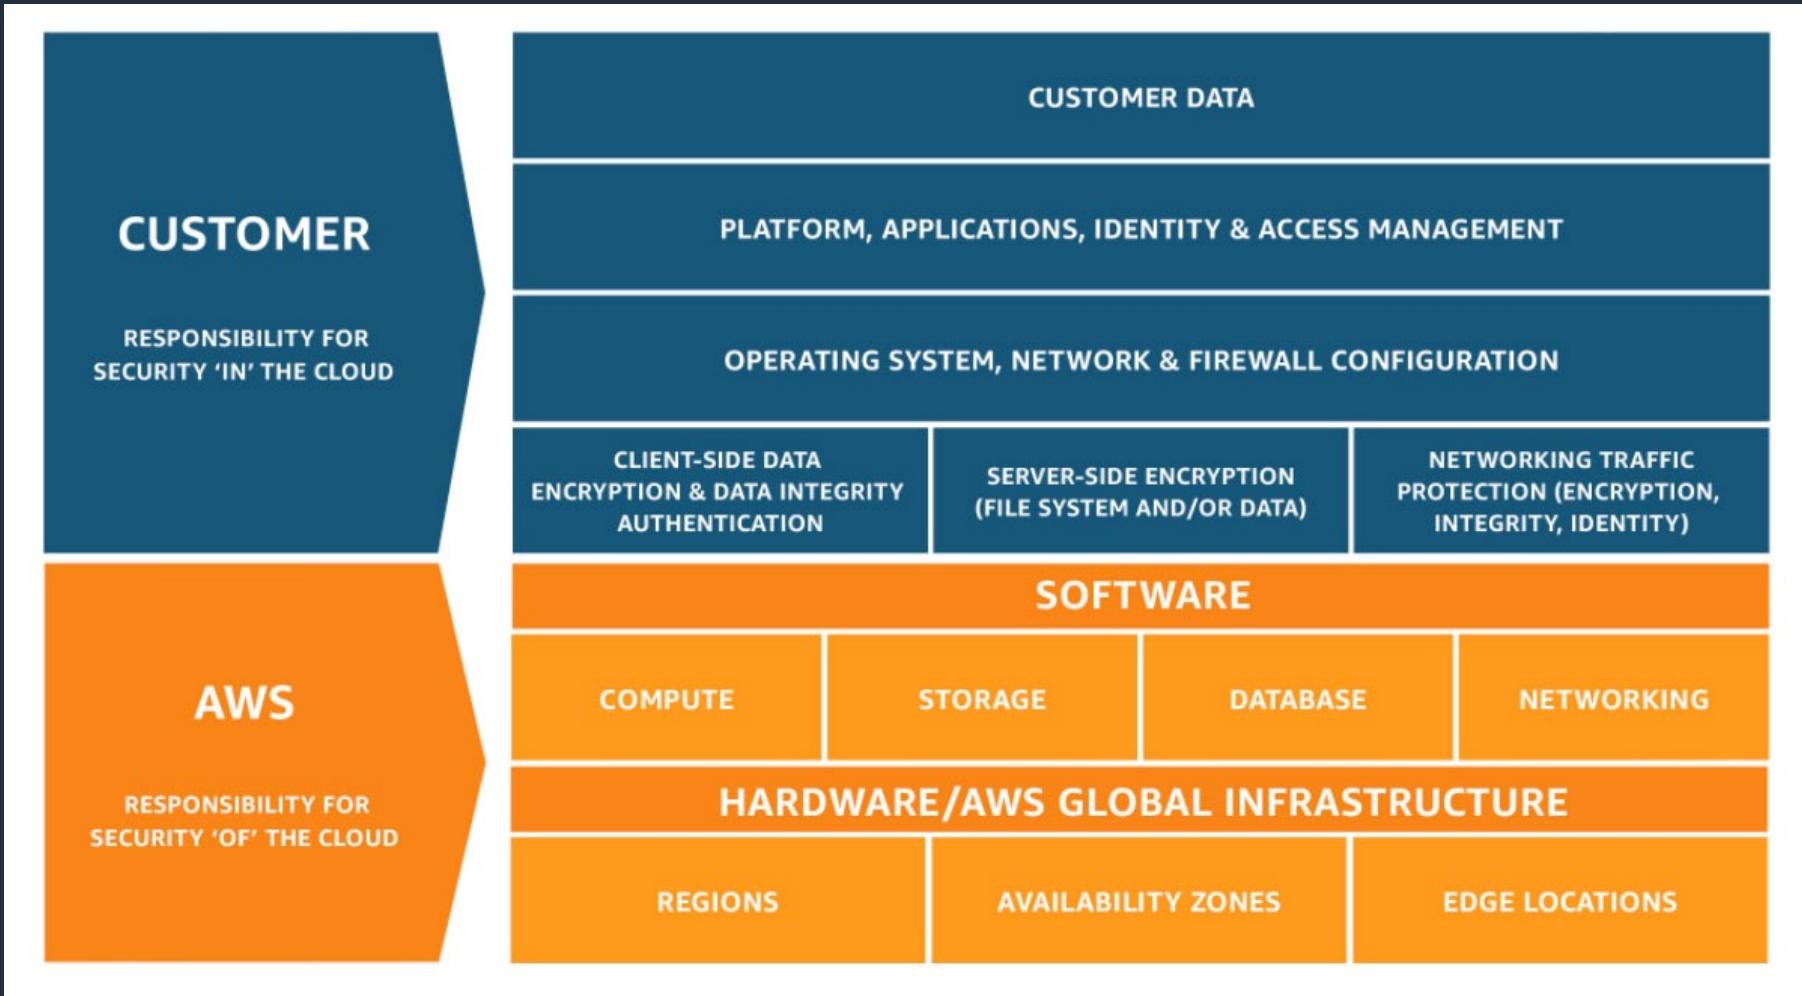
- DNS zone walking via Amazon Route 53 Hosted Zones
- Denial of Service (DoS), Distributed Denial of Service (DDoS), Simulated DoS, Simulated DDoS Port flooding
- Protocol flooding
- Request flooding (login request flooding, API request flooding)

Shared Responsibility Model Review





The AWS Shared Responsibility Model





The AWS Shared Responsibility Model

CUSTOMER RESPONSIBILITY



Bucket with objects



Role



Multi-Factor Authentication



Security Group



Patch management



Staff training



Data encryption



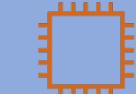
IAM User



Network ACL



SSL encryption



EC2 Instance



Auto Scaling



Elastic load balancer

AWS RESPONSIBILITY



Data center



Data center security



Network router



Network switch



Server



Storage



Database Server



Disk drive

SECTION 14

Architecting for the Cloud

AWS Well-Architected





AWS Well-Architected

- AWS Well-Architected helps cloud architects build secure, high-performing, resilient, and efficient infrastructure for their applications and workloads
- Based on 5 pillars:
 - **Operational Excellence**
 - **Security**
 - **Reliability**
 - **Performance Efficiency**
 - **Cost Optimization**



AWS Well-Architected

Consists of:

- AWS Well-Architected Guidance
- AWS Well-Architected Tool
- AWS Well-Architected Lenses
- AWS Architecture Center

AWS Well-Architected Framework





AWS Well-Architected Framework

- Helps you understand the pros and cons of decisions you make while building systems on AWS
- Based on 5 pillars:

Operational Excellence Pillar

- Support development and run workloads effectively
- Gain insight into workload operations
- Continuously improve processes and procedures to deliver business value



AWS Well-Architected

- Best practices for operational excellence:
 - Perform operations as code
 - Make frequent, small, reversible changes
 - Refine operations procedures frequently
 - Anticipate failure
 - Learn from all operational failures



Security Pillar

- Protect data, systems, and assets to take advantage of cloud technologies to improve your security
- Best practices for security:
 - Implement a strong identity foundation
 - Enable traceability
 - Apply security at all layers
 - Automate security best practices
 - Protect data in transit and at rest
 - Keep people away from data
 - Prepare for security events



Reliability Pillar

- Ensuring a workload can perform its intended function correctly and consistently when it's expected to
- This includes the ability to operate and test the workload through its total lifecycle
- Best practices for reliability:
 - Automatically recover from failure
 - Test recovery procedures
 - Scale horizontally to increase aggregate workload availability
 - Stop guessing capacity
 - Manage change in automation



Performance Efficiency Pillar

- The ability to use computing resources efficiently to meet system requirements, and to maintain that efficiency as demand changes and technologies evolve
- Best practices for performance efficiency:
 - Democratize advanced technologies
 - Go global in minutes
 - Use serverless architectures
 - Experiment more often
 - Consider mechanical sympathy



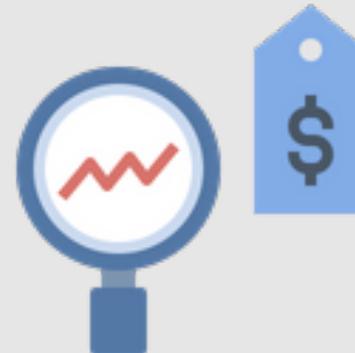
Cost Optimization Pillar

- The ability to run systems to deliver business value at the lowest price point
- Best practices for cost optimization:
 - Implement Cloud Financial Management
 - Adopt a consumption model
 - Measure overall efficiency
 - Stop spending money on undifferentiated heavy lifting
 - Analyze and attribute expenditure

SECTION 15

Accounts, Billing and Support

AWS Pricing Fundamentals





AWS Pricing Fundamentals

Compute



Amount of resources such as CPU and RAM and **duration**

Storage



Quantity of **data stored**

Outbound Data Transfer



Quantity of data that is **transferred out** from all services



AWS Pricing Fundamentals

Pay-as-you-go

- Easily adapt to changing business needs
- Improved responsiveness to change
- Adapt based on needs, not forecasts
- Reduce risk over overpositioning of missing capacity



AWS Pricing Fundamentals

Save when you reserve

- Invest in reserved capacity (e.g. RDS and EC2)
- Save up to 75% compared to on-demand (pay-as-you-go)
- The more you pay upfront the greater the discount



AWS Pricing Fundamentals

Pay less by using more

- Pay less using volume-based discounts
- Tiered pricing means the more you use the lower the unit pricing

Amazon EC2 Pricing Options





Amazon EC2 Pricing Options

On-Demand

Standard rate - no discount; no commitments; dev/test, short-term, or unpredictable workloads

Spot Instances

Bid for unused capacity; up to 90% discount; can be terminated at any time; workloads with flexible start and end times

Dedicated Hosts

Physical server dedicated for your use; Socket/core visibility, host affinity; pay per host; workloads with server-bound software licenses

Reserved

1 or 3-year commitment; up to 75% discount; steady-state, predictable workloads and reserved capacity

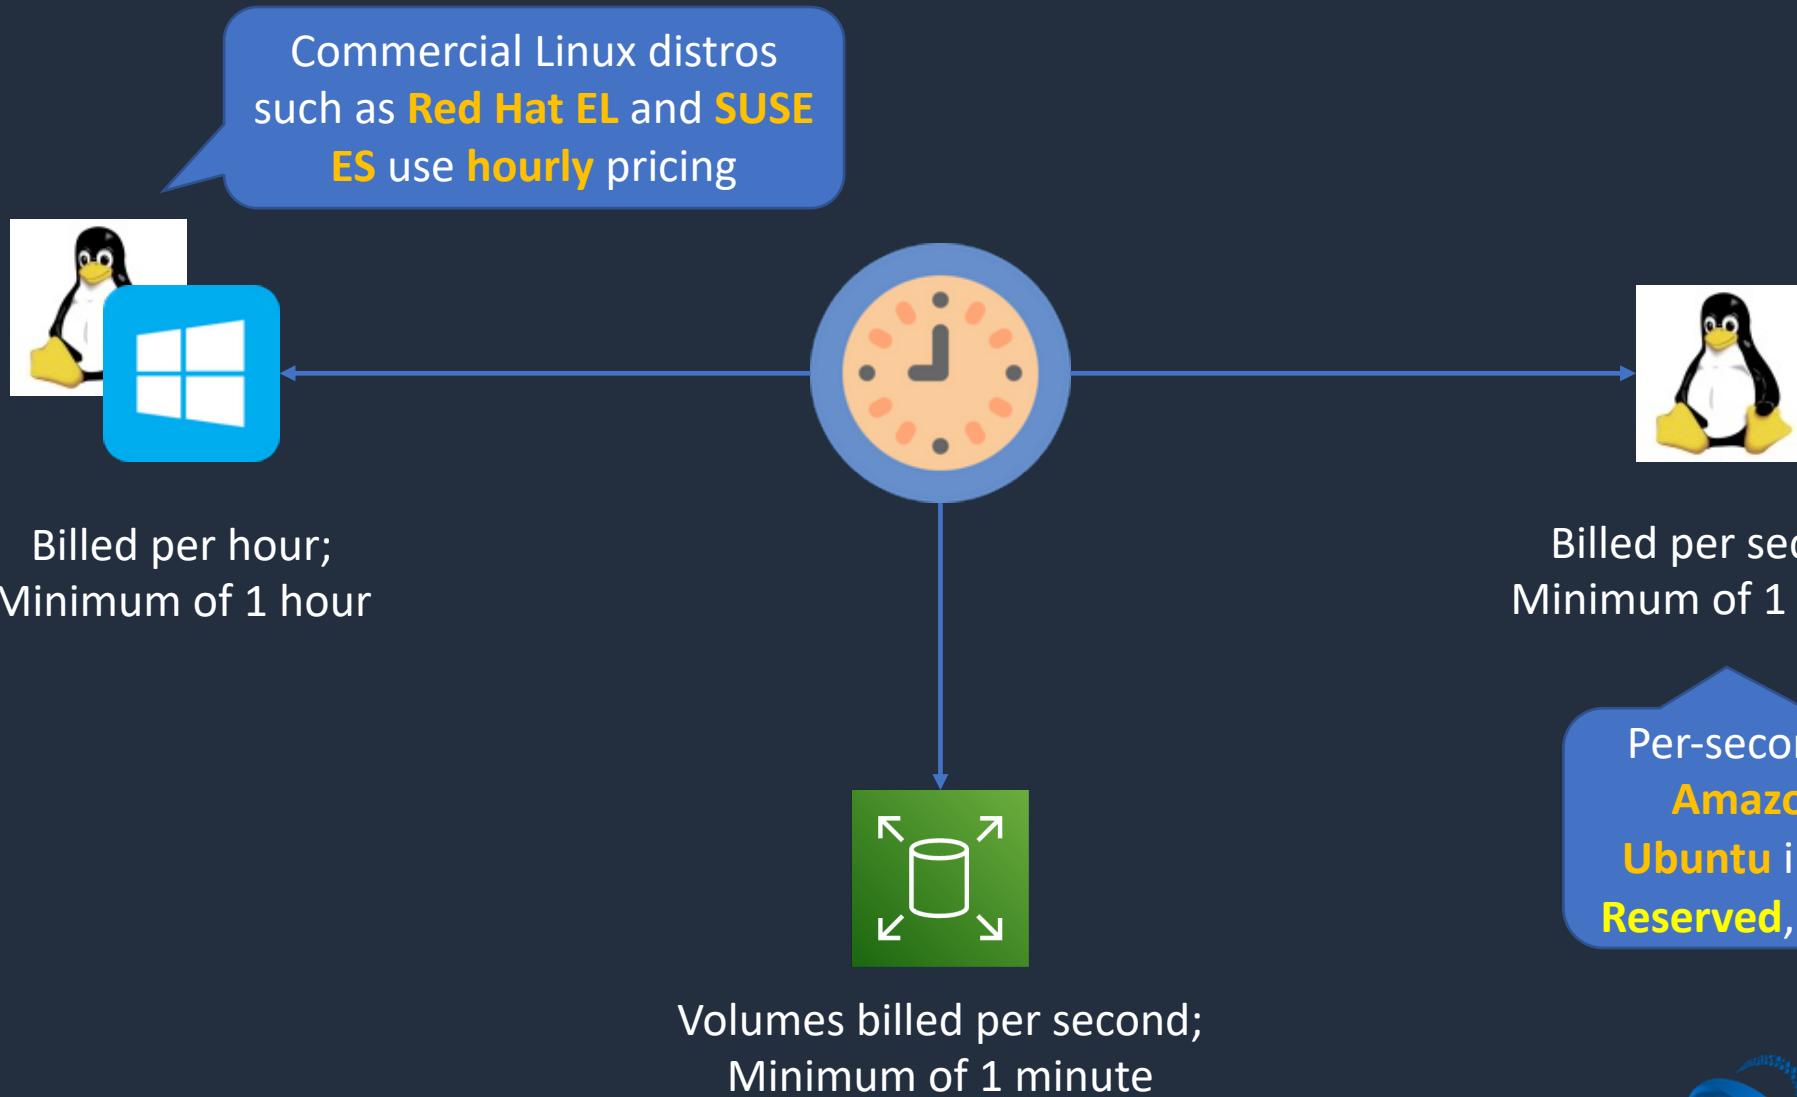
Dedicated Instances

Physical isolation at the host hardware level from instances belonging to other customers; pay per instance

Savings Plans

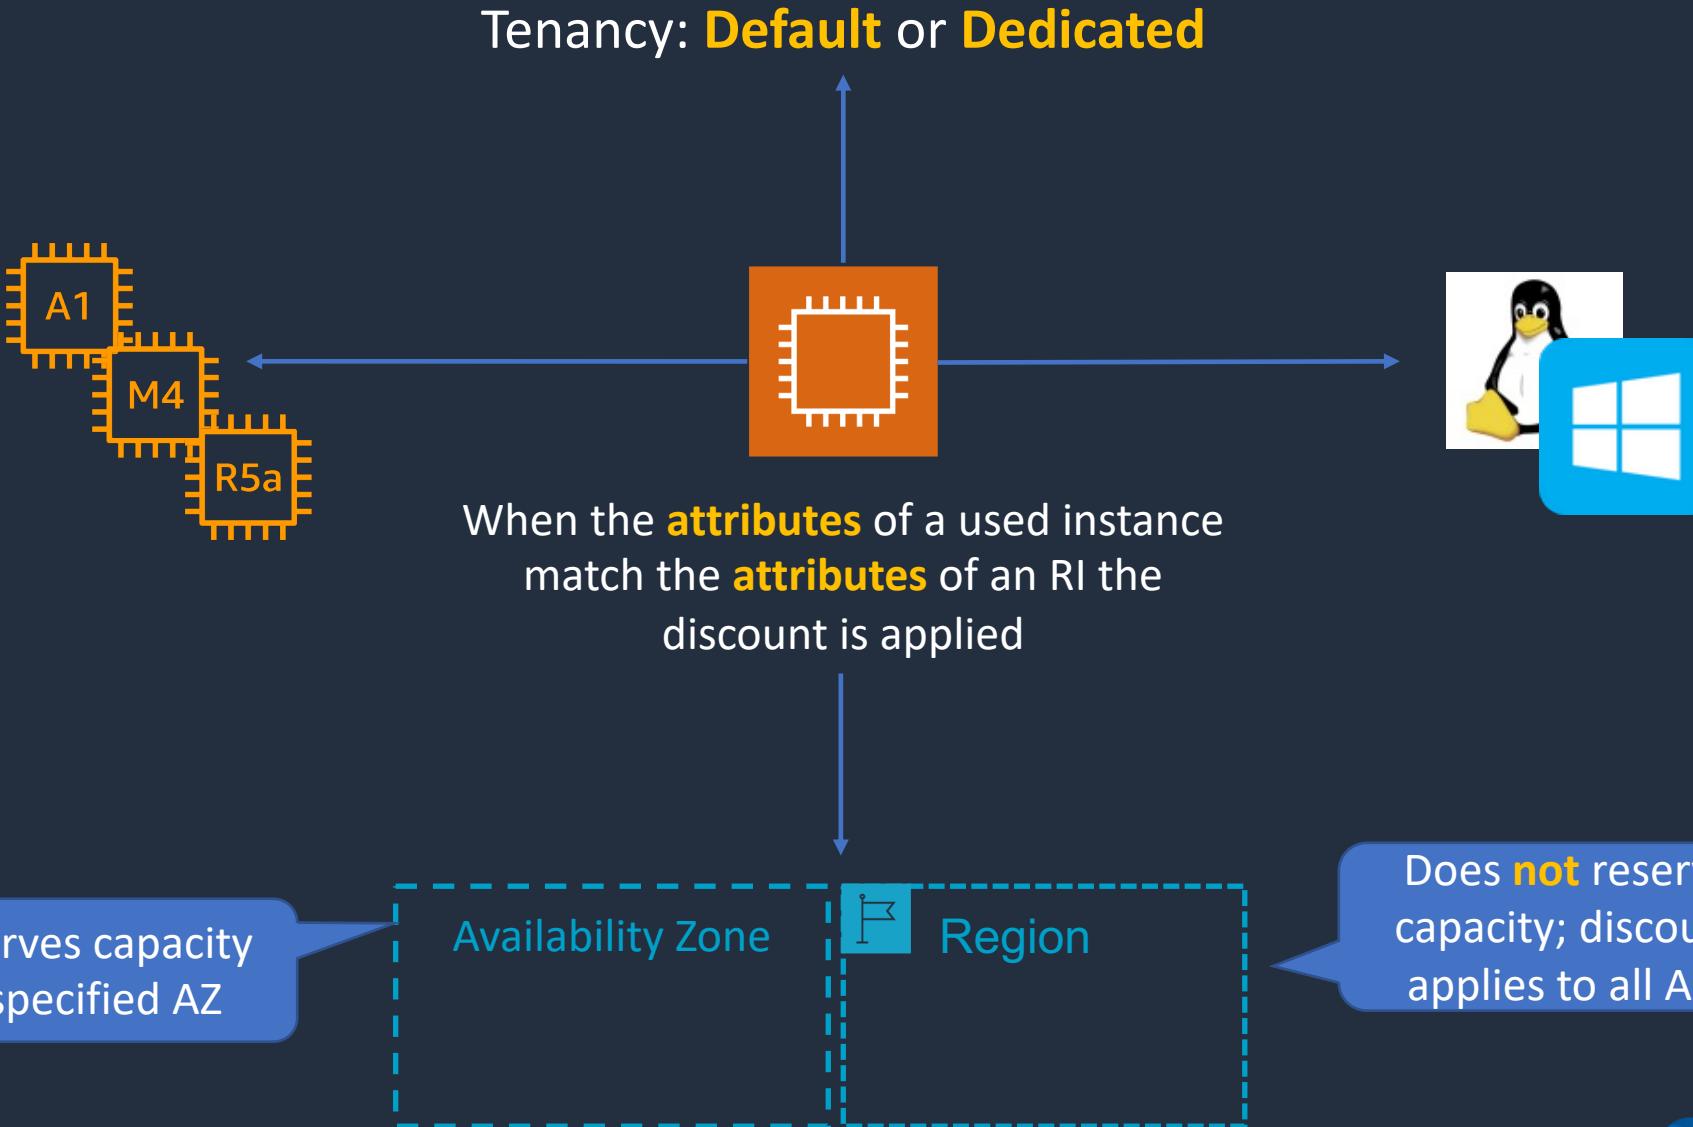
Commitment to a consistent amount of usage (EC2 + Fargate + Lambda); Pay by \$/hour; 1 or 3-year commitment

\$ Amazon EC2 Billing





Amazon EC2 Reserved Instances (RIs)





Amazon EC2 Reserved Instances (RIs)



Scheduled RI

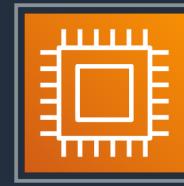
- Match capacity reservation to **recurring schedule**
- Minimum **1200 hours** per year
- Example: Reporting app that runs 6 hours a day 4 days a week = 1248 hours per year



Important
We do not have any capacity for purchasing Scheduled Reserved Instances or any plans to make it available in the future. To reserve capacity, use [On-Demand Capacity Reservations](#) instead. For discounted rates, use [Savings Plans](#).

This message started showing recently but exam may **not** reflect this yet

\$ AWS Savings Plans



Compute Savings Plan



1 or 3-year; hourly commitment to usage of **Fargate**, **Lambda**, and **EC2**; Any Region, family, size, tenancy, and OS



EC2 Savings Plan



1 or 3-year; hourly commitment to usage of **EC2** within a **selected Region** and **Instance Family**; Any size, tenancy and OS



Amazon EC2 Spot Instances



Bid for unused capacity at up to 90% discount



2-minute warning if AWS need to reclaim capacity – available via **instance metadata** and **CloudWatch Events**



Spot Instance: One or more EC2 instances



Spot Fleet: launches and maintains the number of Spot / On-Demand instances to meet specified target capacity



EC2 Fleet: launches and maintains specified number of Spot / On-Demand / Reserved instances in a **single API call**

Can define separate OD/Spot **capacity targets, bids, instance types, and AZs**



Spot Block



Requirement:
Uninterrupted for
1-6 hours

Pricing is **30% - 45%** less
than On-Demand



Solution: **Spot Block**

```
$ aws ec2 request-spot-instances \
--block-duration-minutes 360 \
--instance-count 5 \
--spot-price "0.25" ...
```



Dedicated Instances and Dedicated Hosts

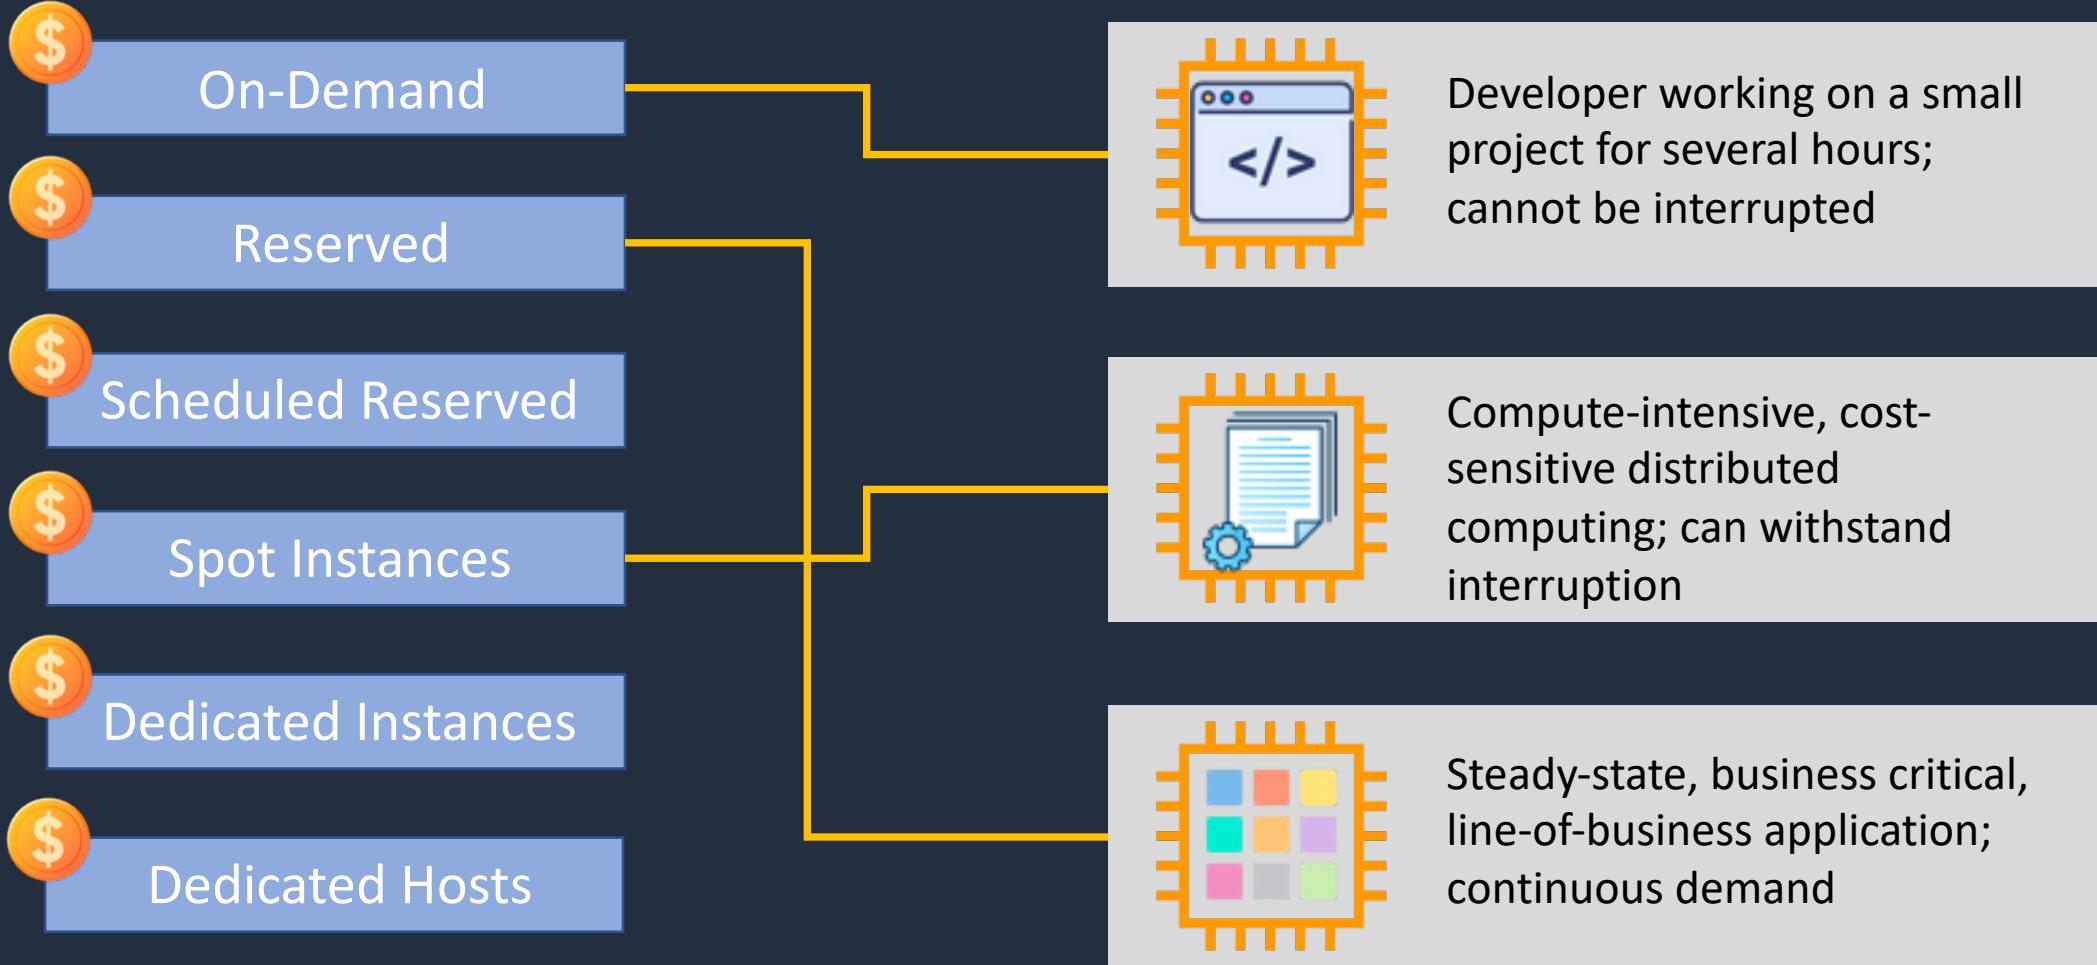
Characteristic	Dedicated Instances	Dedicated Hosts
Enables the use of dedicated physical servers	X	X
Per instance billing (subject to a \$2 per region fee)	X	
Per host billing		X
Visibility of sockets, cores, host ID		X
Affinity between a host and instance		X
Targeted instance placement		X
Automatic instance placement	X	X
Add capacity using an allocation request		X

Amazon EC2 Pricing Use Cases



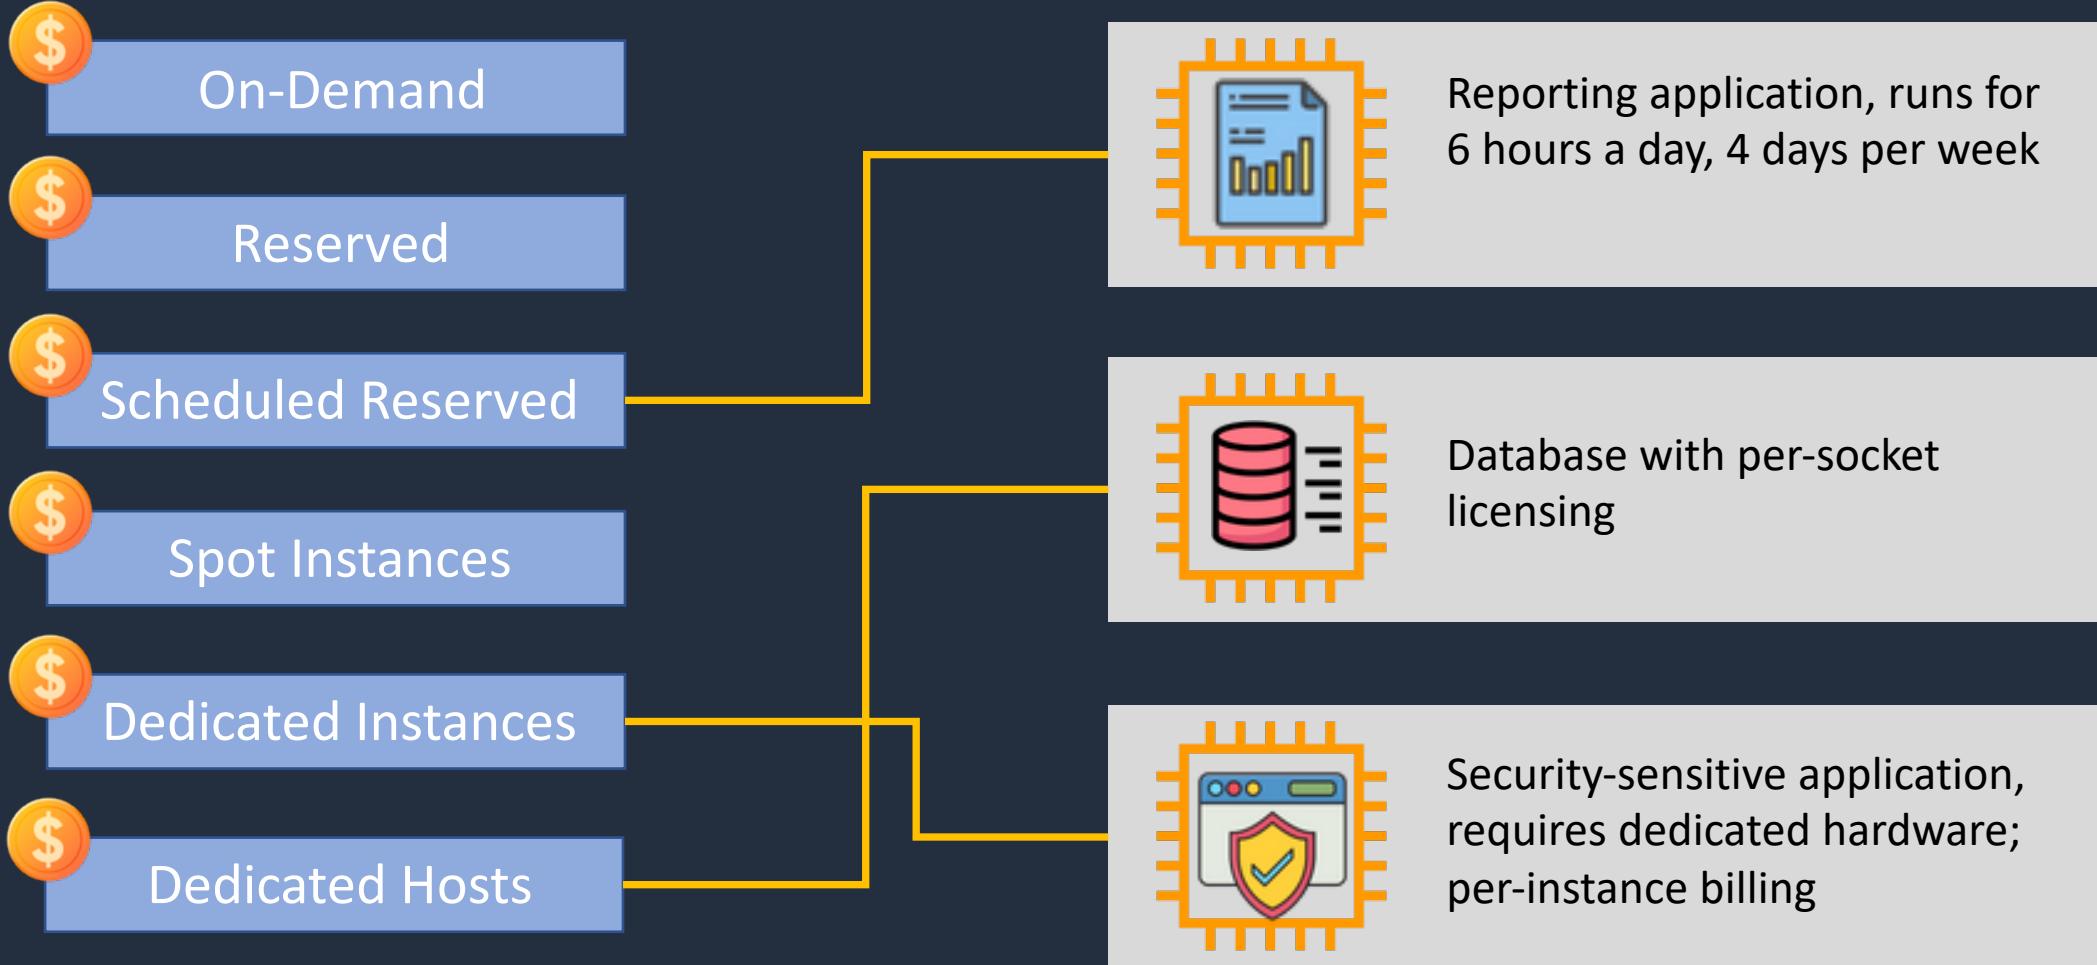


Amazon EC2 Pricing Use Cases





Amazon EC2 Pricing Use Cases



Pricing for other AWS Services





Amazon S3 Pricing

- **Storage class** – e.g. Standard or IA
- **Storage quantity** – data volume stored in your buckets on a per GB basis
- **Number of requests** – the number and type of requests, e.g. GET, PUT, POST, LIST, COPY
- **Lifecycle transitions requests** – moving data between storage classes
- **Data transfer** – data transferred out of an S3 region is charged



Amazon S3 Glacier Pricing

- Three options for access to archives, listed in the table below:

	Expedited	Standard	Bulk
Data access time	1-5 minutes	3-5 hours	5-12 hours
Data retrievals	\$0.03 per GB	\$0.01 per GB	\$0.0025 per GB
Retrieval requests	On-Demand: \$0.01 per request Provisioned: \$100 per Provisioned Capacity Unit	\$0.050 per 1,000 requests	\$0.025 per 1,000 requests



Amazon EBS Pricing

- **Volumes** – volume storage for all EBS volumes type is charged by the amount of GB provisioned per month
- **Snapshots** – based on the amount of space consumed by snapshots in S3. Copying snapshots is charged on the amount of data copied across regions
- **Data transfer** – inbound data transfer is free, outbound data transfer charges are tiered



Amazon RDS Pricing

- **Clock hours of server uptime** – amount of time the DB instance is running
- **Database characteristics** – e.g. database engine, size and memory class
- **Database purchase type** – e.g. On-Demand, Reserved.
- **Number of database instances**
- **Provisioned storage** – backup is included up to 100% of the size of the DB
- **Additional storage** – the amount of storage in addition to the provisioned storage is charged per GB per month



Amazon RDS Pricing

- **Requests** – the number of input and output requests to the DB
- **Deployment type** – single AZ or multi-AZ
- **Reserved Instances** – RDS RIs can be purchased with No Upfront, Partial Upfront, or All Upfront terms. Available for Aurora, MySQL, MariaDB, Oracle and SQL Server



Amazon DynamoDB Pricing

- Charged for reading, writing, and storing data
- **On-demand capacity mode**
 - Charged for reads and writes
 - No need to specify how much capacity is required
 - Good for unpredictable workloads
- **Provisioned capacity mode**
 - Specify number of reads and writes per second
 - Can use Auto Scaling
 - Good for predictable workloads
 - Consistent traffic or gradual changes



Amazon CloudFront Pricing

- **Traffic distribution** – data transfer and request pricing, varies across regions, and is based on the edge location from which the content is served
- **Requests** – the number and type of requests (HTTP or HTTPS) and the geographic region in which they are made
- **Data transfer out** – quantity of data transferred out of CloudFront edge locations
- There are additional chargeable items such as invalidation requests, field-level encryption requests, and custom SSL certificates



AWS Lambda Pricing

- **Number of requests**
- **Duration of request** – rounded up to the nearest millisecond
- Price is dependent on the amount of memory allocated to the function
- Free tier includes 1M free requests per month and 400,000 GB-seconds of compute time

AWS Pricing Calculator



AWS Support Plans





AWS Support Plans

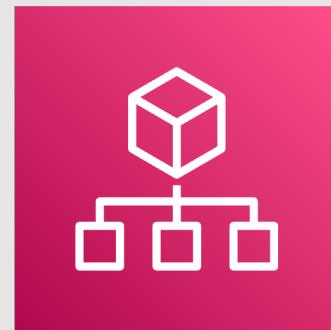
	Developer	Business	Enterprise
AWS Trusted Advisor Best Practice Checks	7 Core checks	Full set of checks	Full set of checks
Enhanced Technical Support	<ul style="list-style-type: none">Business hours email access to Cloud Support AssociatesUnlimited cases / 1 primary contact	<ul style="list-style-type: none">24x7 phone, email, and chat access to Cloud Support EngineersUnlimited cases / unlimited contacts (IAM supported)	<ul style="list-style-type: none">24x7 phone, email, and chat access to Cloud Support EngineersUnlimited cases / unlimited contacts (IAM supported)
Case Severity / Response Times*	<ul style="list-style-type: none">General guidance: < 24 hours**System impaired: < 12 hours**	<ul style="list-style-type: none">General guidance: < 24 hoursSystem impaired: < 12 hoursProduction system impaired: < 4 hoursProduction system down: < 1 hour	<ul style="list-style-type: none">General guidance: < 24 hoursSystem impaired: < 12 hoursProduction system impaired: < 4 hoursProduction system down: < 1 hourBusiness-critical system down: < 15 minutes



AWS Support Plans

	Developer	Business	Enterprise
Architectural Guidance	General	Contextual to your use-cases	Consultative review and guidance based on your applications
Programmatic Case Management		AWS Support API	AWS Support API
Proactive Programs and Services		Access to Infrastructure Event Management for additional fee	<ul style="list-style-type: none">• Infrastructure Event Management• Well-Architected Reviews• Access to proactive reviews, workshops, and deep dives
Technical Account Management			Designated Technical Account Manager (TAM) to proactively monitor your environment and assist with optimization and coordinate access to programs and AWS experts
Account Assistance			Concierge Support Team

Consolidated Billing





AWS Organizations

- Consolidated billing has the following benefits:
- One bill – You get one bill for multiple accounts
 - **Easy tracking** – You can track the charges across multiple accounts and download the combined cost and usage data
 - **Combined usage** – You can combine the usage across all accounts in the organization to share the volume pricing discounts and Reserved Instance discounts
 - **No extra fee** – Consolidated billing is offered at no additional cost



AWS Organizations

Tier Description	Price Per GB	Price Per TB
First 1 TB/month	\$0.10	\$100.00
Next 49 TB/month	\$0.08	\$80.00
Next 450 TB/month	\$0.06	\$60.00

Usage within Organization:

Account A (master) usage: 2 TB

Account B usage: 80 TB

Account C usage: 120 TB

Total: 202 TB

Calculation:

First 1 TB = \$100.00

Next 49 TB = \$3,920.00

Next 157 TB = \$9,120.00

Total cost = \$13,140.00

AWS Budgets





AWS Budgets

All budgets (1)	Cost budgets (1)	Usage budgets (0)	Reservation budgets (0)	Savings Plans budgets (0)		
Budget name	Type	Current	Budgeted	Forecasted	Current vs. budgeted	Forecasted vs. budgeted
MyBudget	Cost	\$13.20	\$100.00	\$129.47	<div style="width: 13.2%; background-color: #0070C0;"></div> 13.2%	<div style="width: 129.47%; background-color: #E74C3C;"></div> 129.47%

Set Custom Budgets - set custom usage and reservation budgets

Configure Alerts – receive alerts when you exceed or are forecast to exceed your alert thresholds

Integrated with other AWS services – Includes Cost Explorer Chatbot, and Service Catalog

AWS Cost Allocation Tags



AWS Cost Management Tools





AWS Cost Explorer

- The **AWS Cost Explorer** is a free tool that allows you to view charts of your costs
- You can view cost data for the past 13 months and forecast how much you are likely to spend over the next three months
- Cost Explorer can be used to discover patterns in how much you spend on AWS resources over time and to identify cost problem areas
- Cost Explorer can help you to identify service usage statistics such as:
 - Which services you use the most
 - View metrics for which AZ has the most traffic
 - Which linked account is used the most



AWS Cost & Usage Report

- Publish AWS billing reports to an Amazon S3 bucket
- Reports break down costs by:
 - Hour, day, month, product, product resource, tags
- Can update the report up to three times a day
- Create, retrieve, and delete your reports using the AWS CUR API Reference



AWS Price List API

- Query the prices of AWS services
- **Price List Service API** (AKA the Query API) – query with JSON
- **AWS Price List API** (AKA the Bulk API) – query with HTML
- Alerts via Amazon SNS when prices change

AWS Cost Management Tools



SECTION 16

Migration, Machine Learning and More

AWS Migration and Transfer Services





AWS Migration Tools



Region



AWS Migration Hub



Amazon S3



Amazon RDS



EC2 Instances



EFS File system



AWS Server Migration Service



AWS Database Migration Service



AWS DataSync



VPN or Direct
Connect connection

Corporate data center



Servers



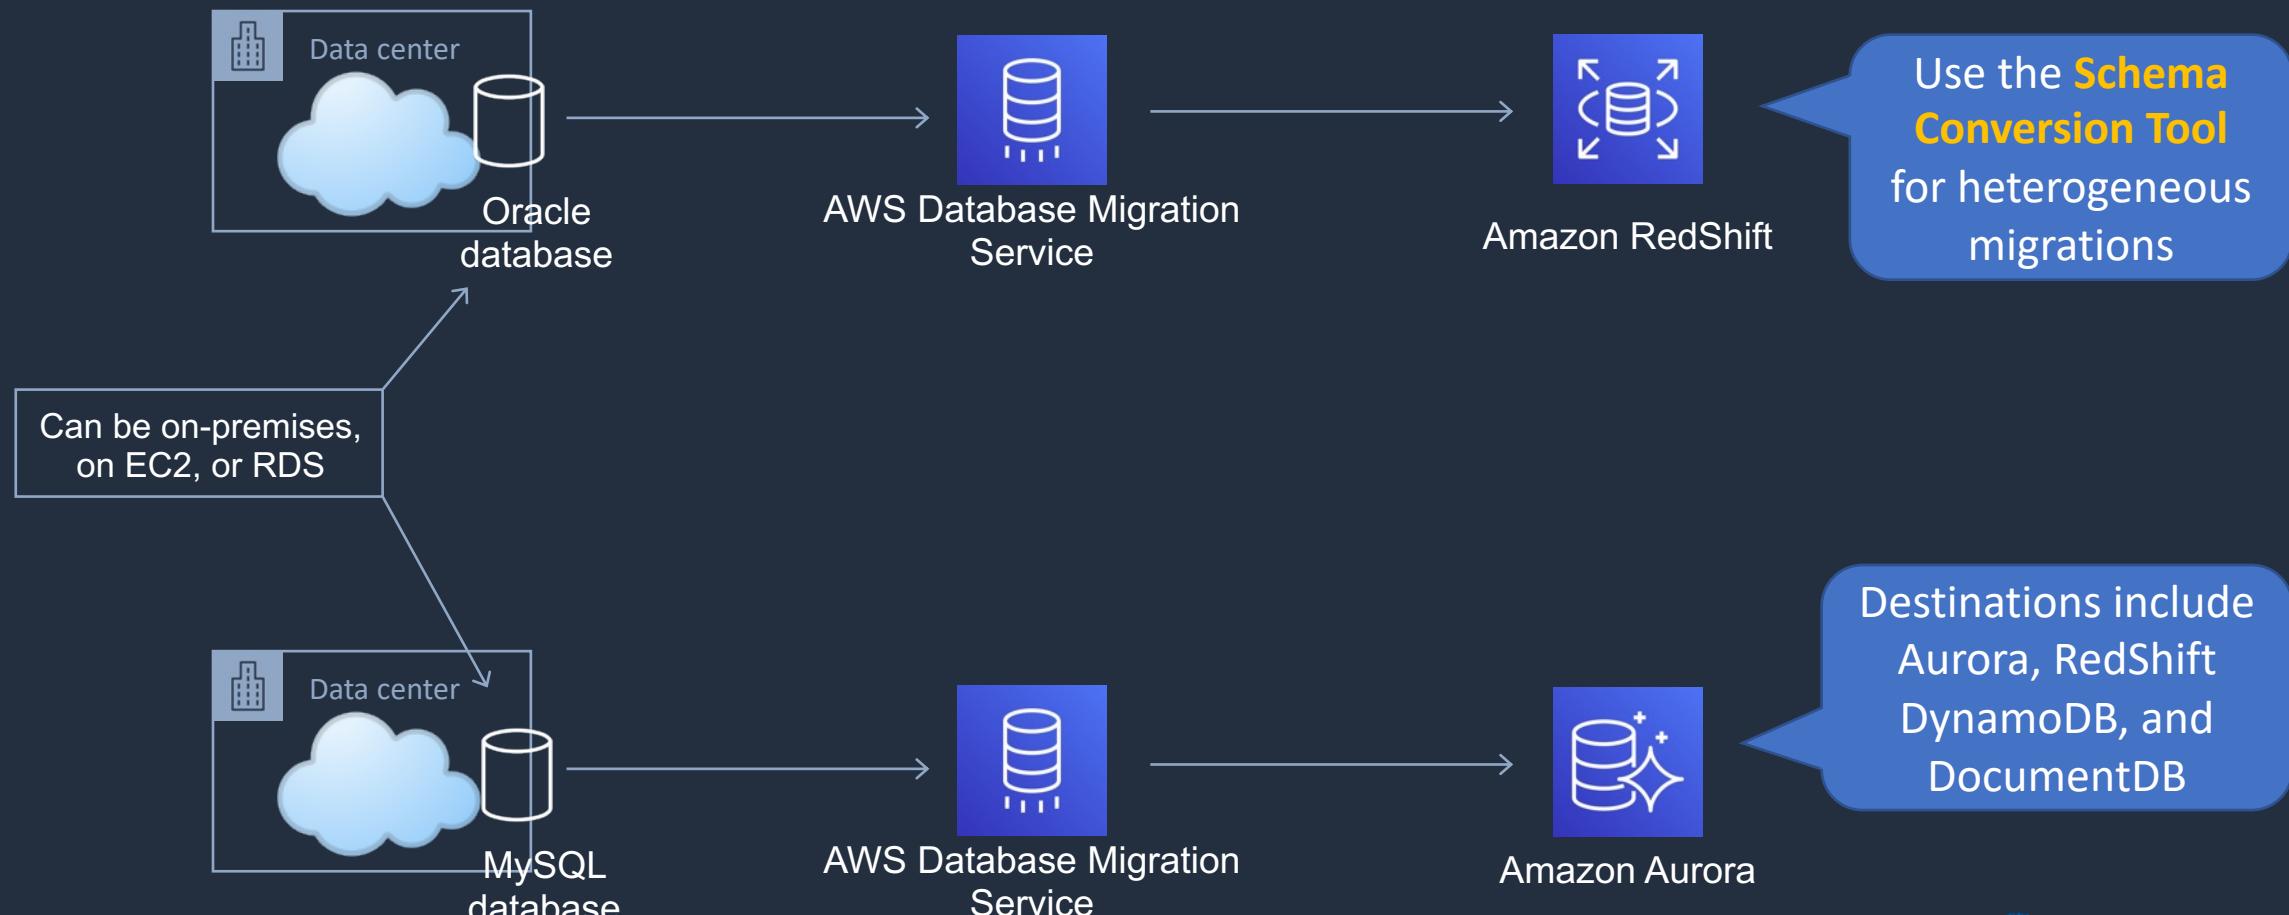
Database



NAS / File
Server



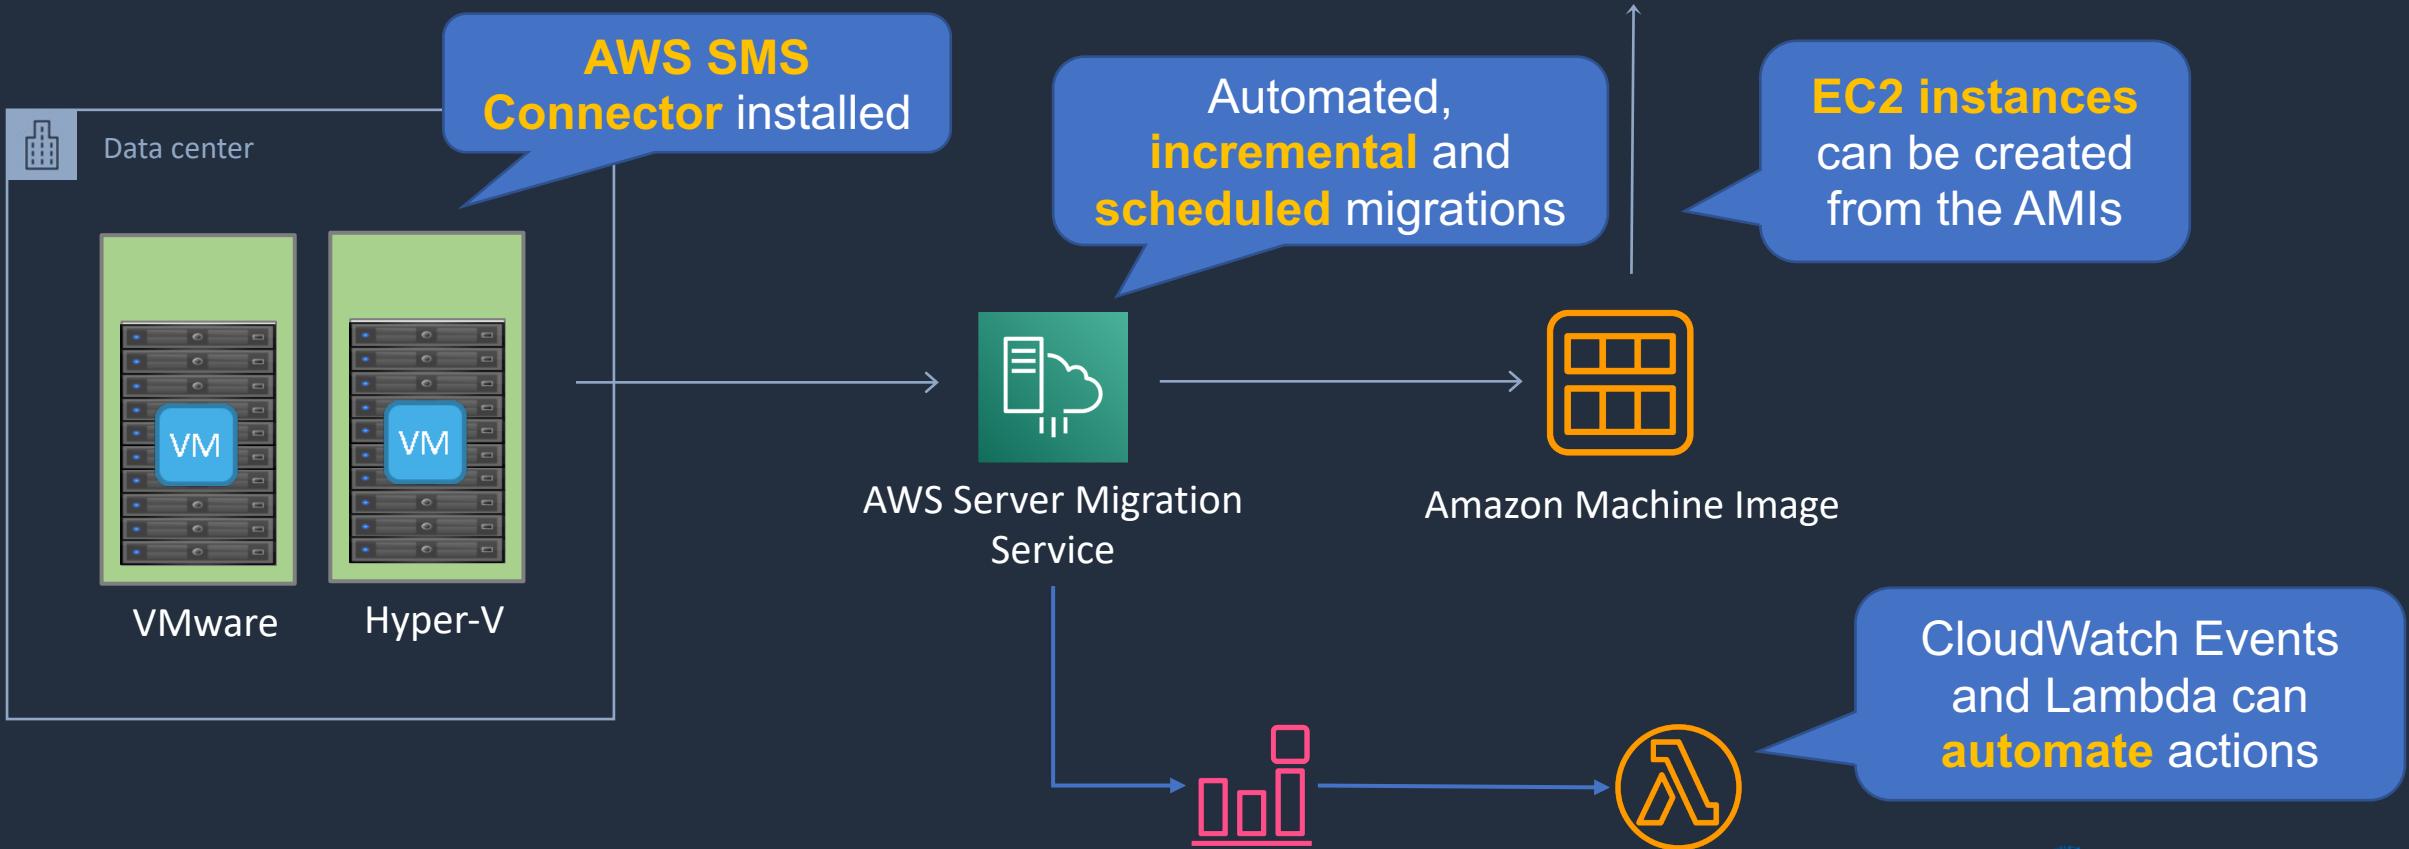
AWS Database Migration Service (DMS)





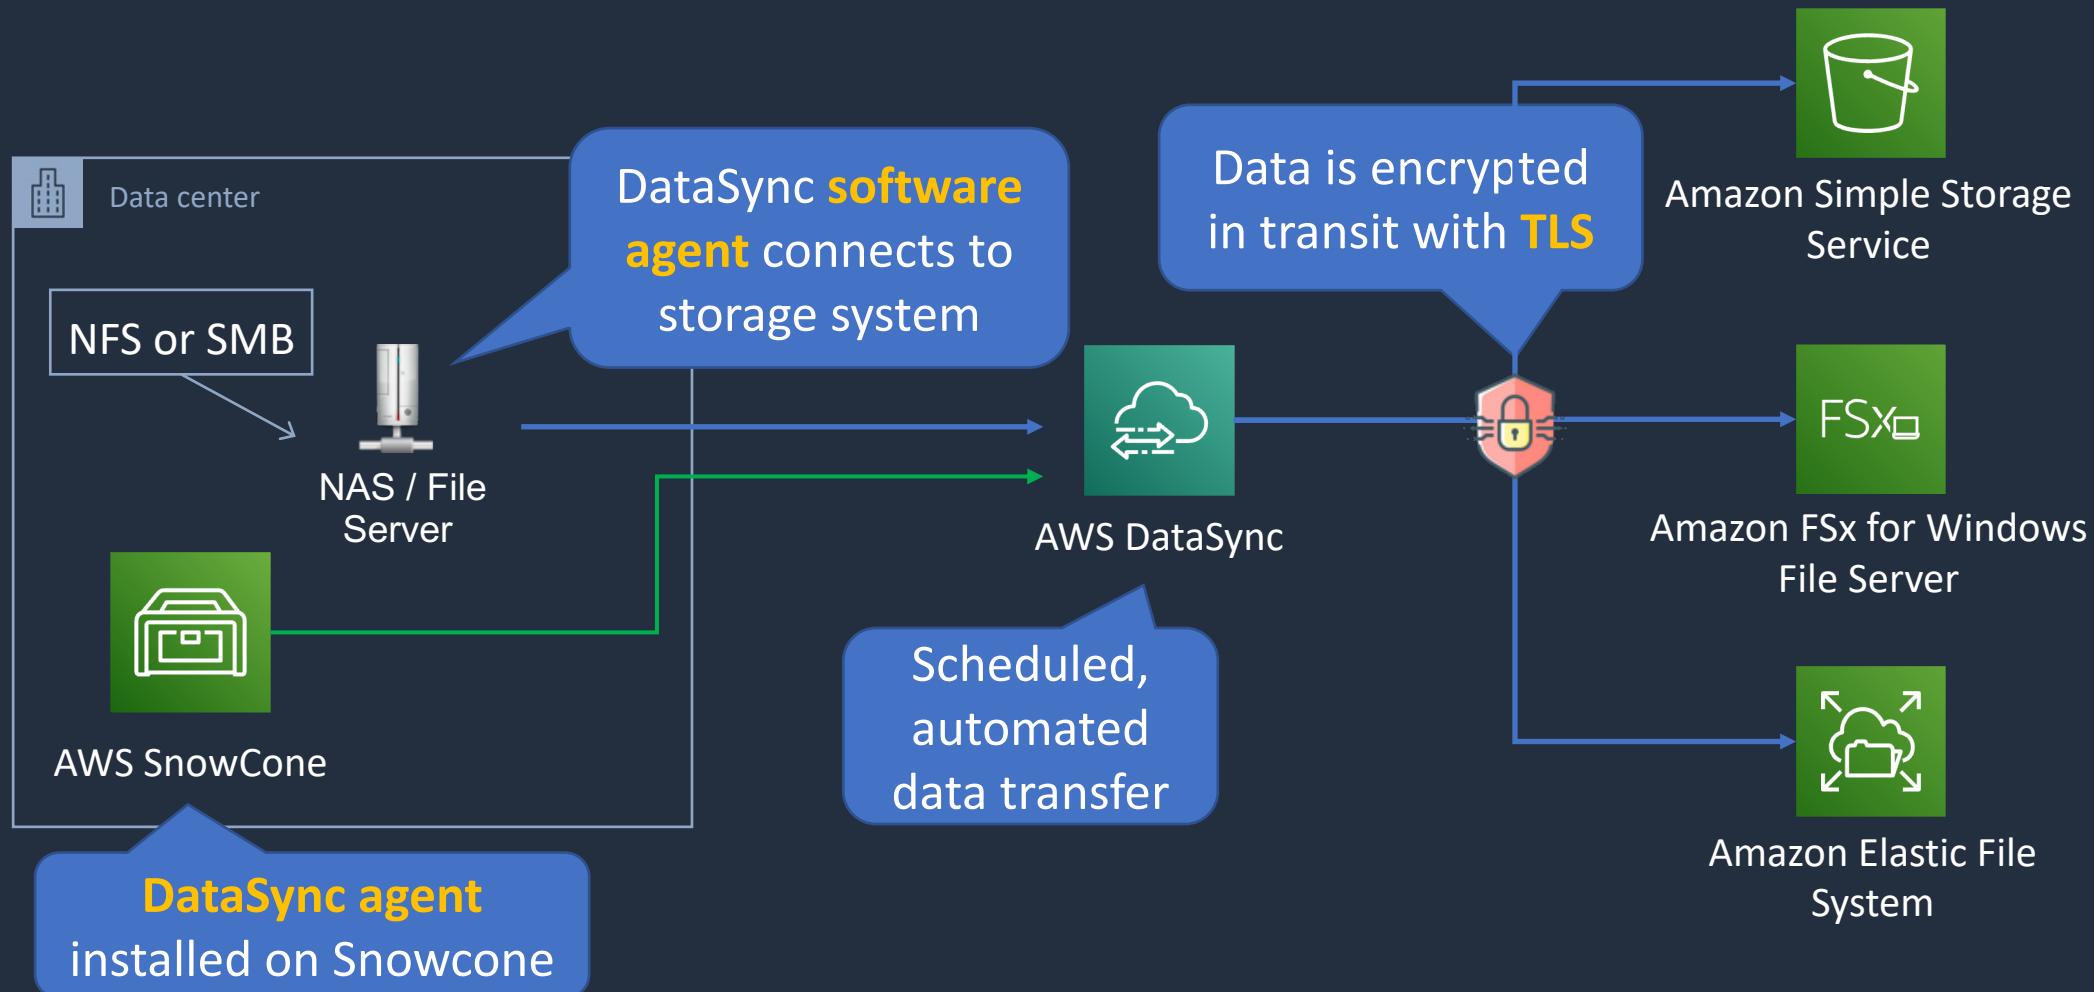
AWS Server Migration Service (SMS)

- AWS SMS migrates VMware vSphere, Microsoft Hyper-V/SCVMM, and Azure virtual machines to Amazon EC2





AWS DataSync





AWS Snowball Family

- **AWS Snowball and Snowmobile** are used for migrating large volumes of data to AWS
- **Snowball Edge Compute Optimized**
 - Provides block and object storage and optional GPU
 - Use for data collection, machine learning and processing, and storage in environments with intermittent connectivity (edge use cases)
- **Snowball Edge Storage Optimized**
 - Provides block storage and Amazon S3-compatible object storage
 - Use for local storage and large-scale data transfer
- **Snowcone**
 - Small device used for edge computing, storage and data transfer
 - Can transfer data offline or online with AWS DataSync agent



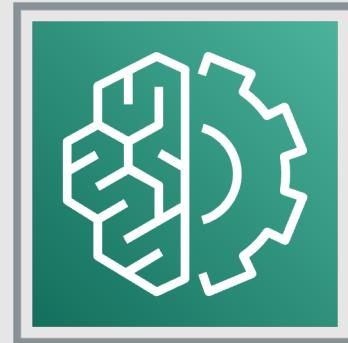


AWS Snowball Family

- Uses a secure storage device for physical transportation
- Snowball Client is software that is installed on a local computer and is used to identify, compress, encrypt, and transfer data
- Uses 256-bit encryption (managed with the AWS KMS) and tamper-resistant enclosures with TPM
- **Snowball** (80TB) (50TB) “petabyte scale”
- **Snowball Edge** (100TB) “petabyte scale”
- **Snowmobile** – “exabyte scale” with up to 100PB per Snowmobile



AWS Machine Learning Services





AWS Rekognition

Identify objects



Perform facial analysis

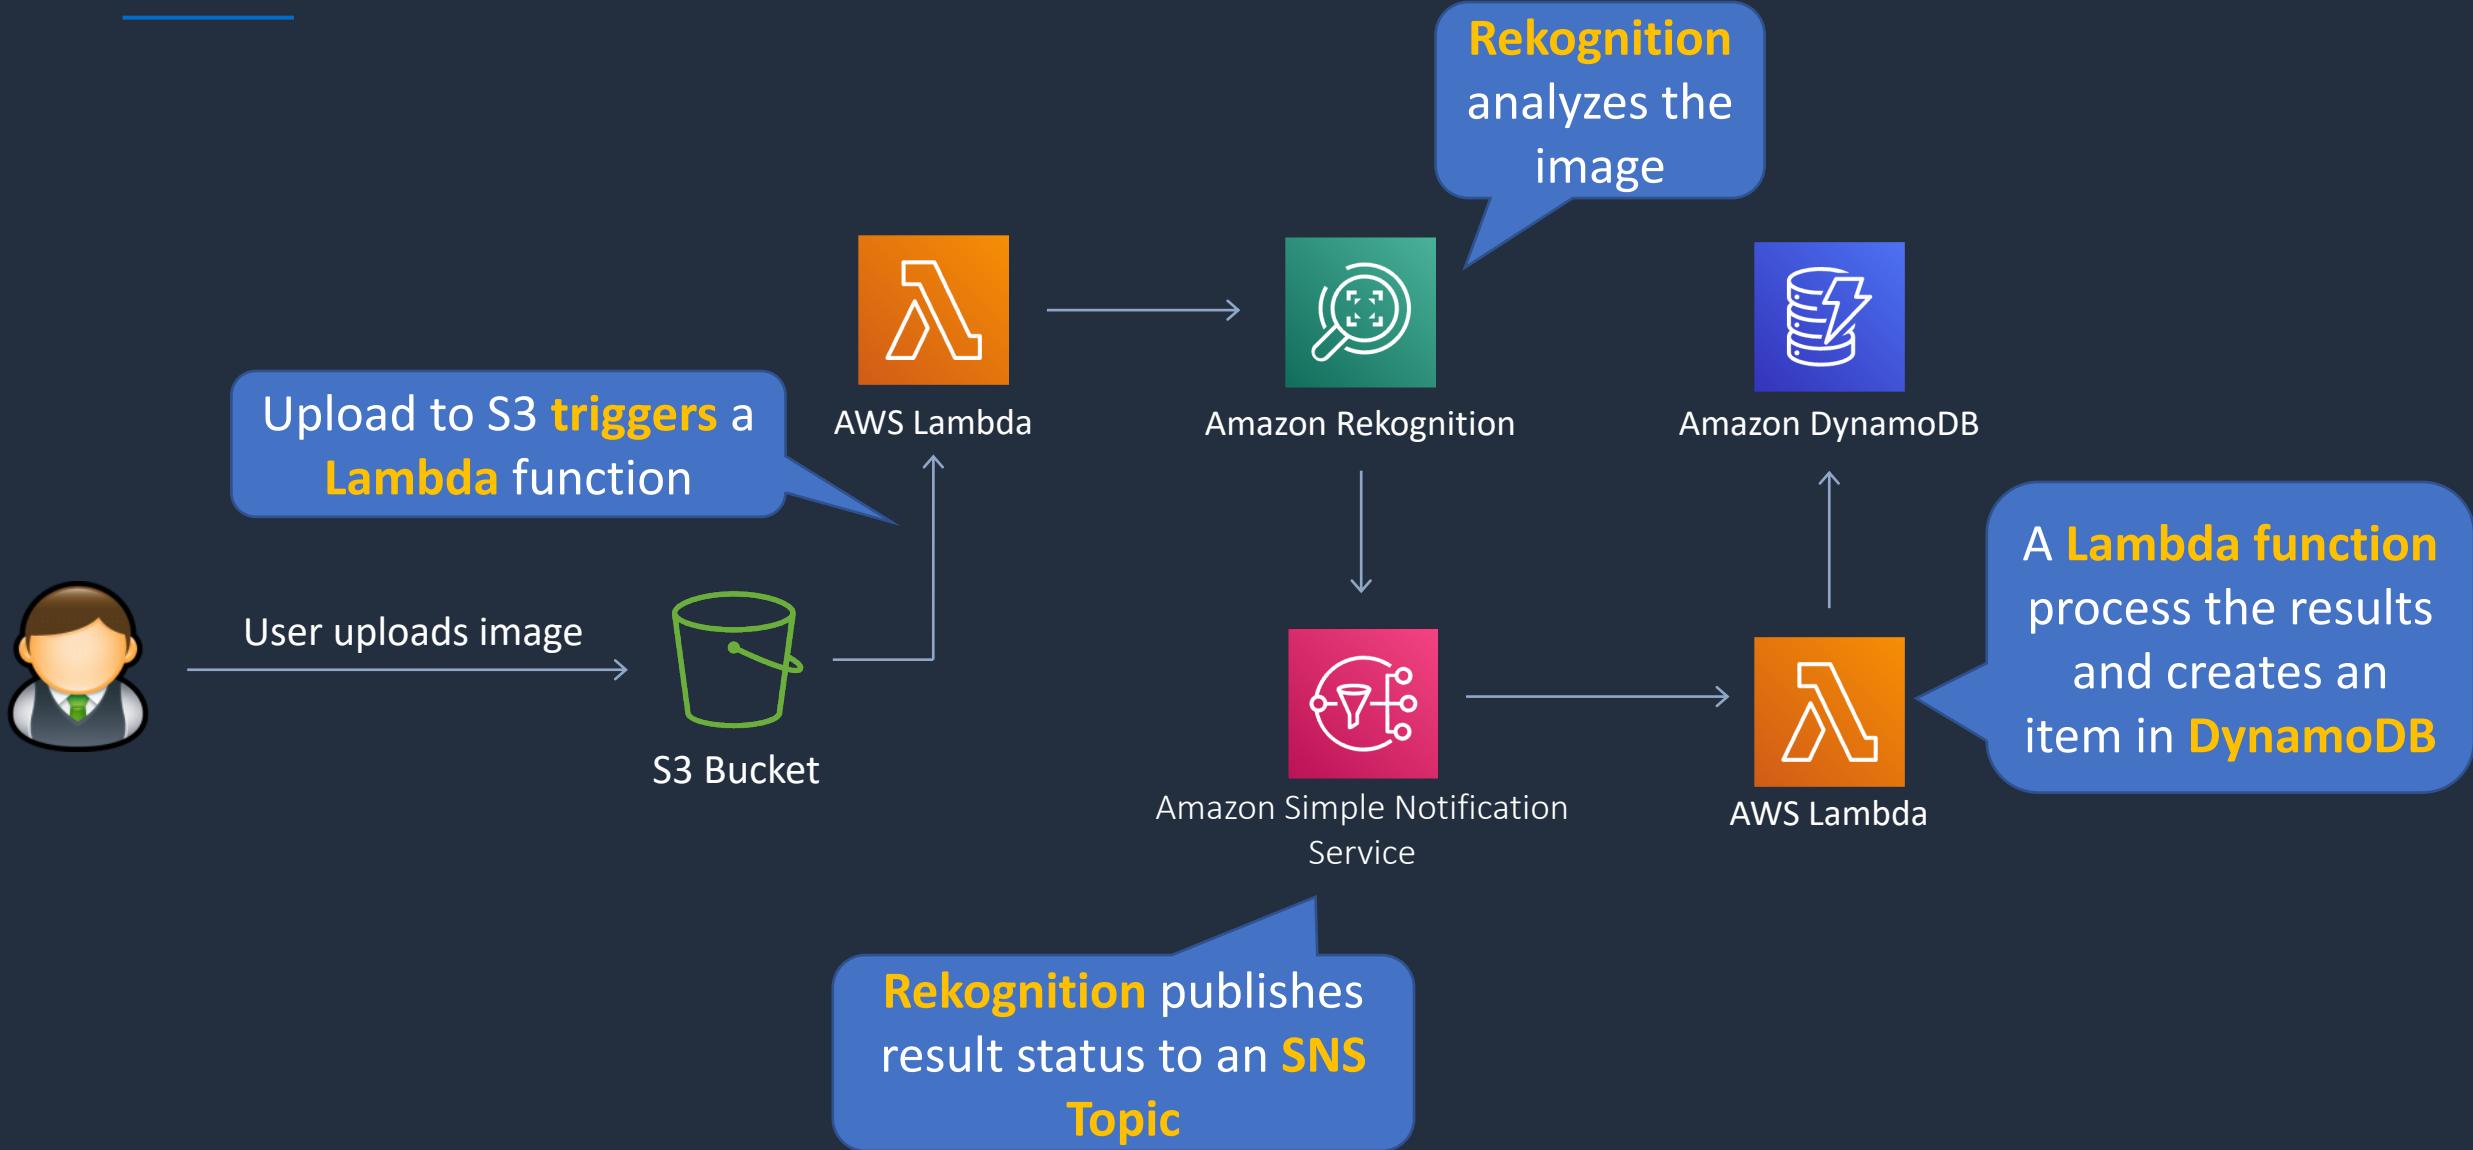


Celebrity recognition





AWS Rekognition in Event-Driven Architecture





AWS Rekognition

- Add image and video analysis to your applications
- Identify objects, people, text, scenes, and activities in images and videos
- Processes videos stored in an Amazon S3 bucket
- Publish completion status to Amazon SNS Topic



Amazon Transcribe

- Add speech to text capabilities to applications
- Recorded speech can be converted to text before it can be used in applications
- Uses a deep learning process called automatic speech recognition (ASR) to convert speech to text quickly and accurately



Amazon Translate

- Neural machine translation service that delivers fast, high-quality, and affordable language translation
- Uses deep learning models to deliver more accurate and more natural sounding translation
- Localize content such as websites and applications for your diverse users



Amazon Textract

- Automatically extract printed text, handwriting, and data from any document
- Features:
 - Optical character recognition (OCR)
 - Identifies relationships, structure, and text
 - Uses AI to extract text and structured data
 - Recognizes handwriting as well as printed text
 - Can extract text from documents such as PDFs, images, forms, and tables
 - Understands context. For example know what data to extract from a receipt or invoice



Amazon SageMaker

- Helps data scientists and developers to prepare, build, train, and deploy high-quality machine learning (ML) models
- ML development activities including:
 - Data preparation
 - Feature engineering
 - Statistical bias detection
 - Auto-ML
 - Training and tuning
 - Hosting
 - Monitoring
 - Workflows



Amazon Comprehend

- Natural-language processing (NLP) service
- Uses machine learning to uncover information in unstructured data
- Can identify critical elements in data, including references to language, people, and places, and the text files can be categorized by relevant topics
- In real time, you can automatically and accurately detect customer sentiment in your content



Amazon Lex

- Conversational AI for Chatbots
- Build conversational interfaces into any application using voice and text
- Build bots to increase contact center productivity, automate simple tasks, and drive operational efficiencies across the enterprise



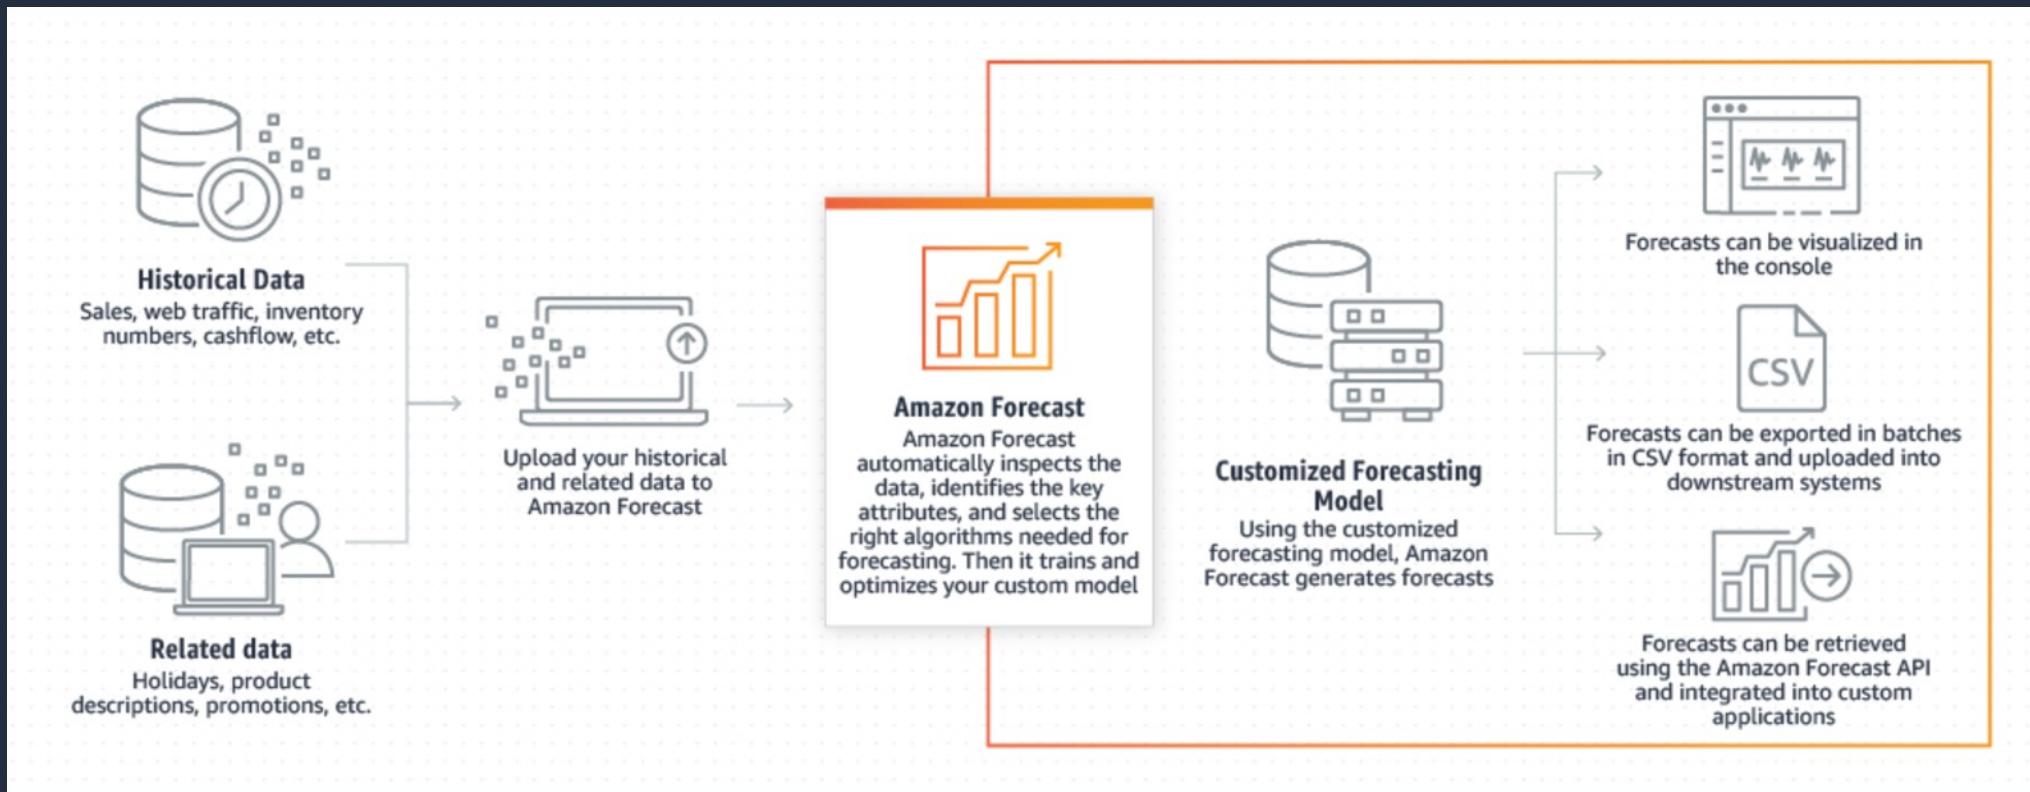
Amazon Polly

- Turns text into lifelike speech
- Create applications that talk, and build entirely new categories of speech-enabled products
- Text-to-Speech (TTS) service uses advanced deep learning technologies to synthesize natural sounding human speech



Amazon Forecast

- Time-series forecasting service
- Uses ML and built for business metrics analysis





Amazon DevOps Guru

- Cloud operations service for improving **application operational performance and availability**
- Detect behaviors that deviate from normal operating patterns
- Benefits:
 - Automatically detect operational issues
 - Resolve issues with ML-powered insights
 - Elastically scale operational analytics
 - Uses ML to reduce alarm noise

End User Computing



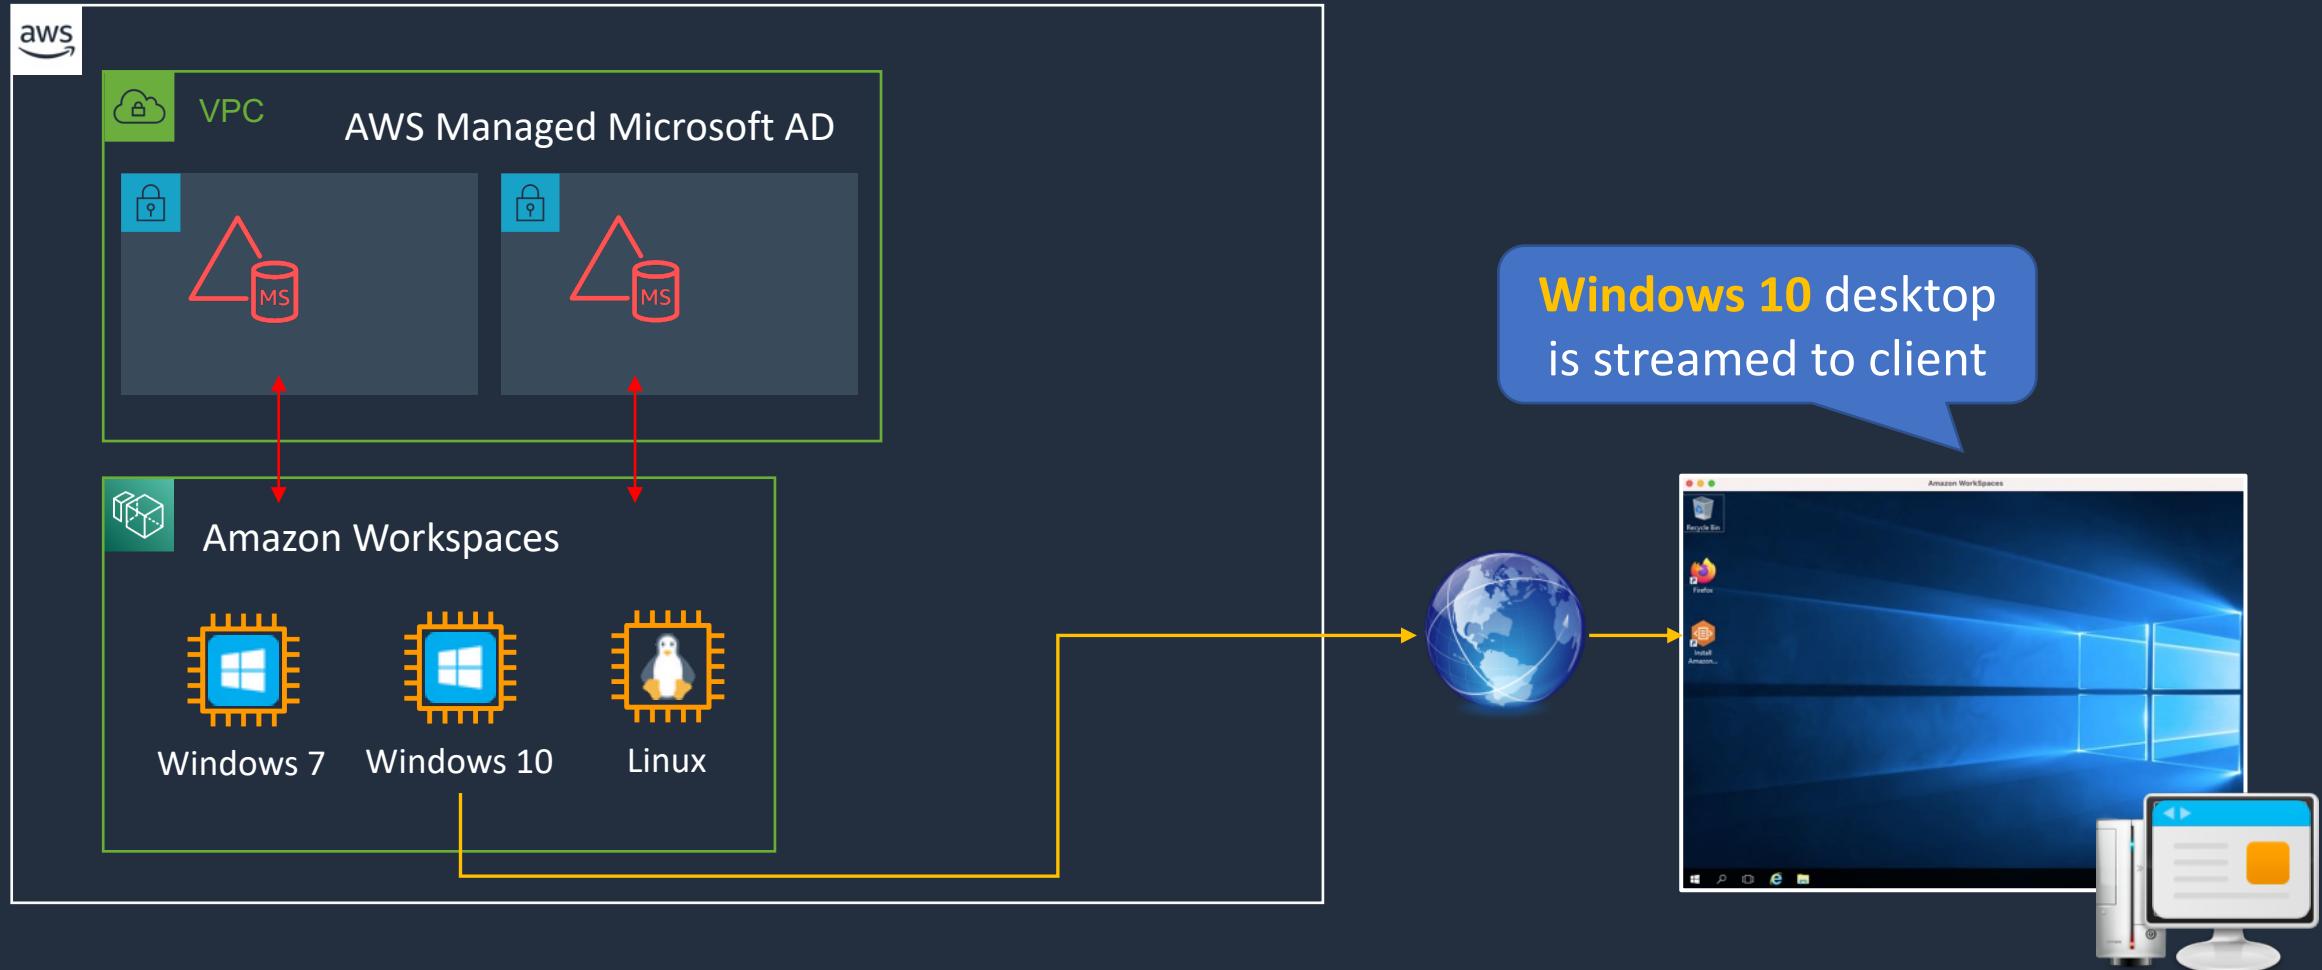


Amazon Workspaces

- Managed **Desktop-as-a-Service** (DaaS) solution
- Provision either Windows or Linux desktops
- Simplifies delivery of desktops compared to traditional virtual desktop infrastructure (VDI) deployments



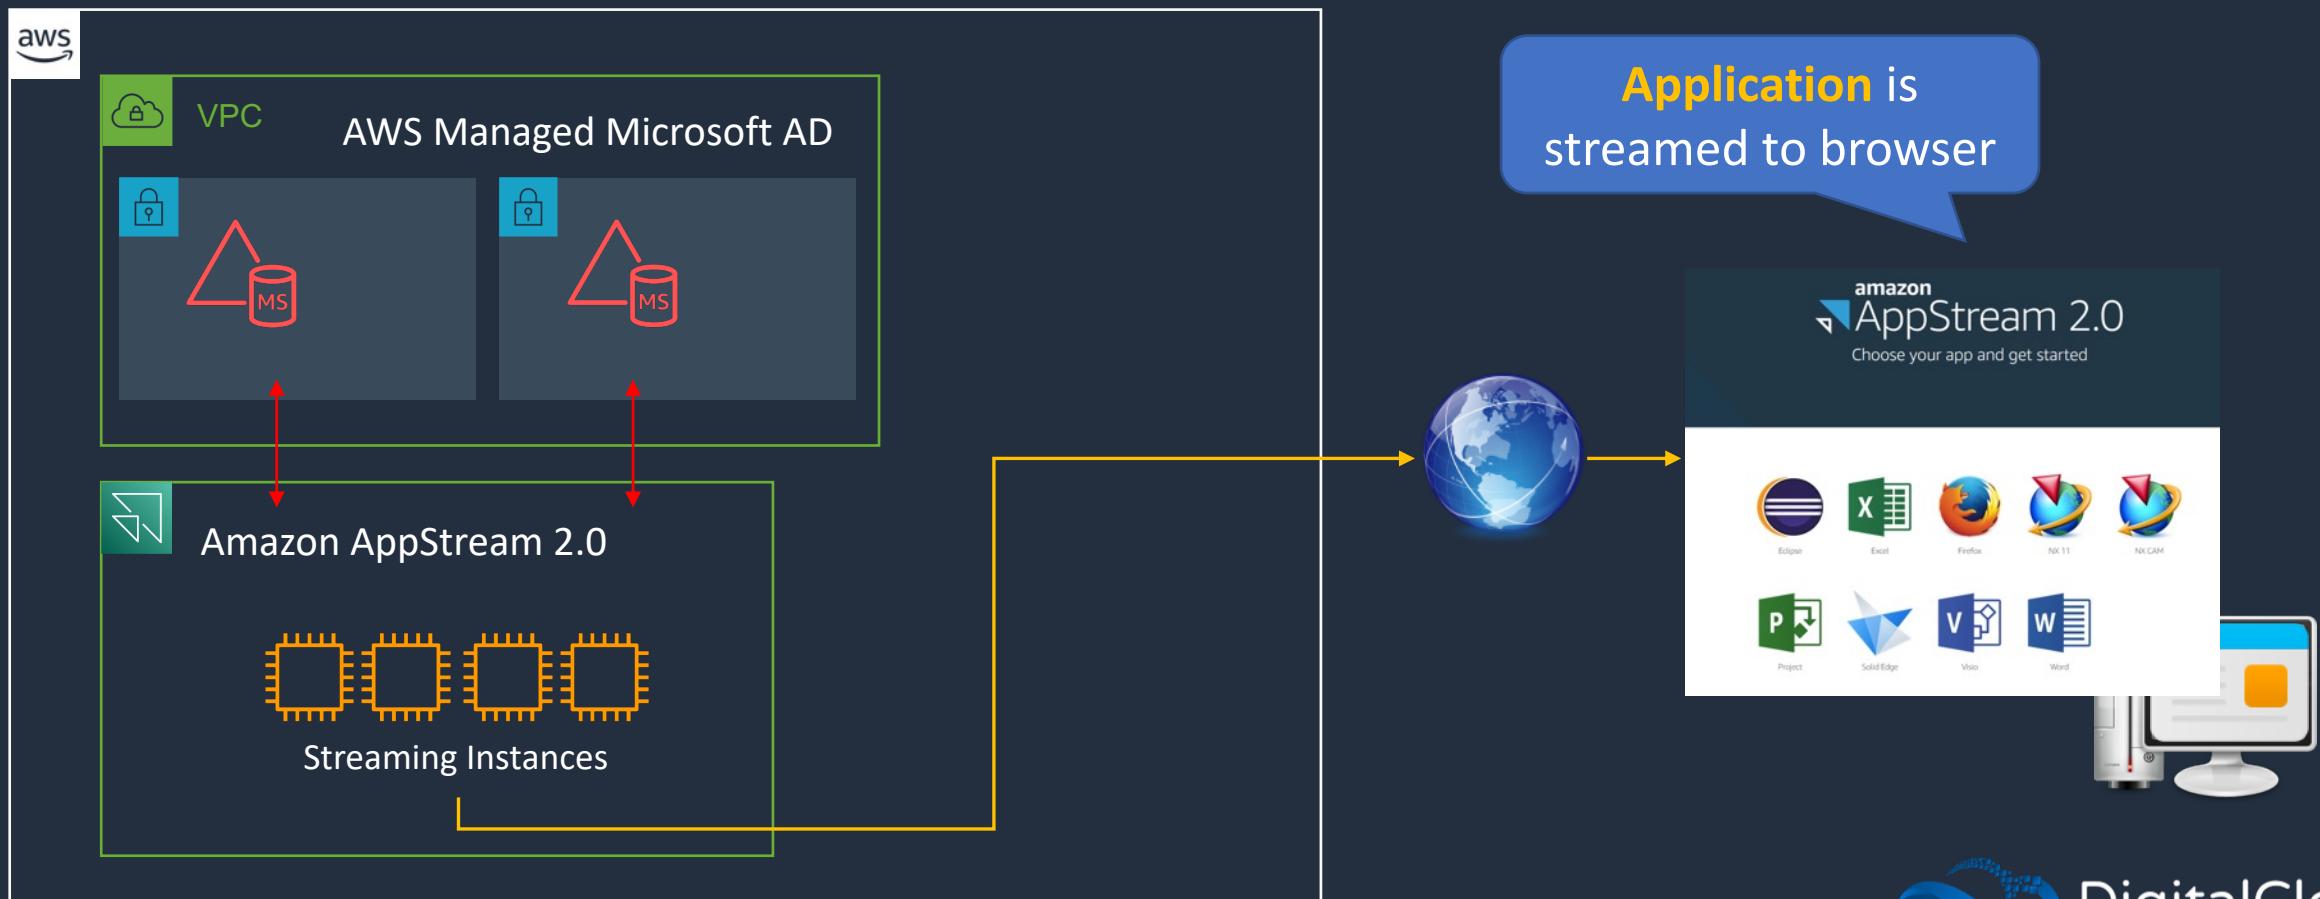
Amazon Workspaces





AWS AppStream 2.0

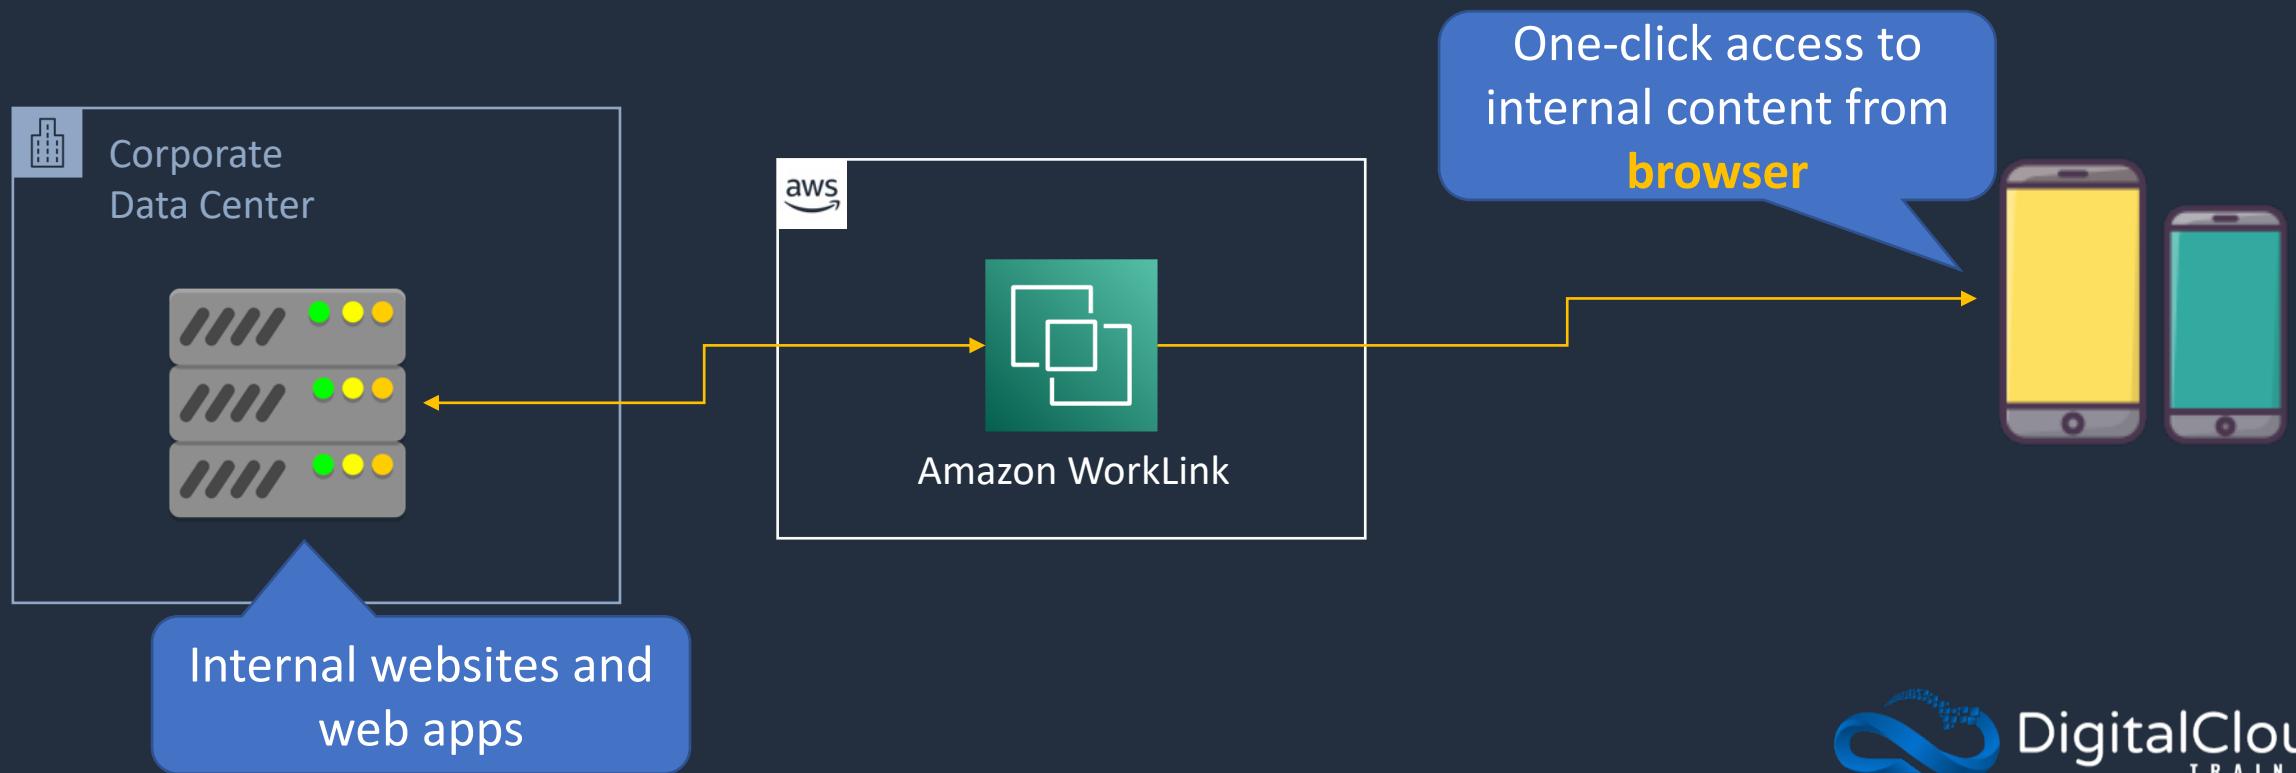
- Fully managed non-persistent application streaming service
- Alternative to popular products such as Citrix XenApp





AWS WorkLink

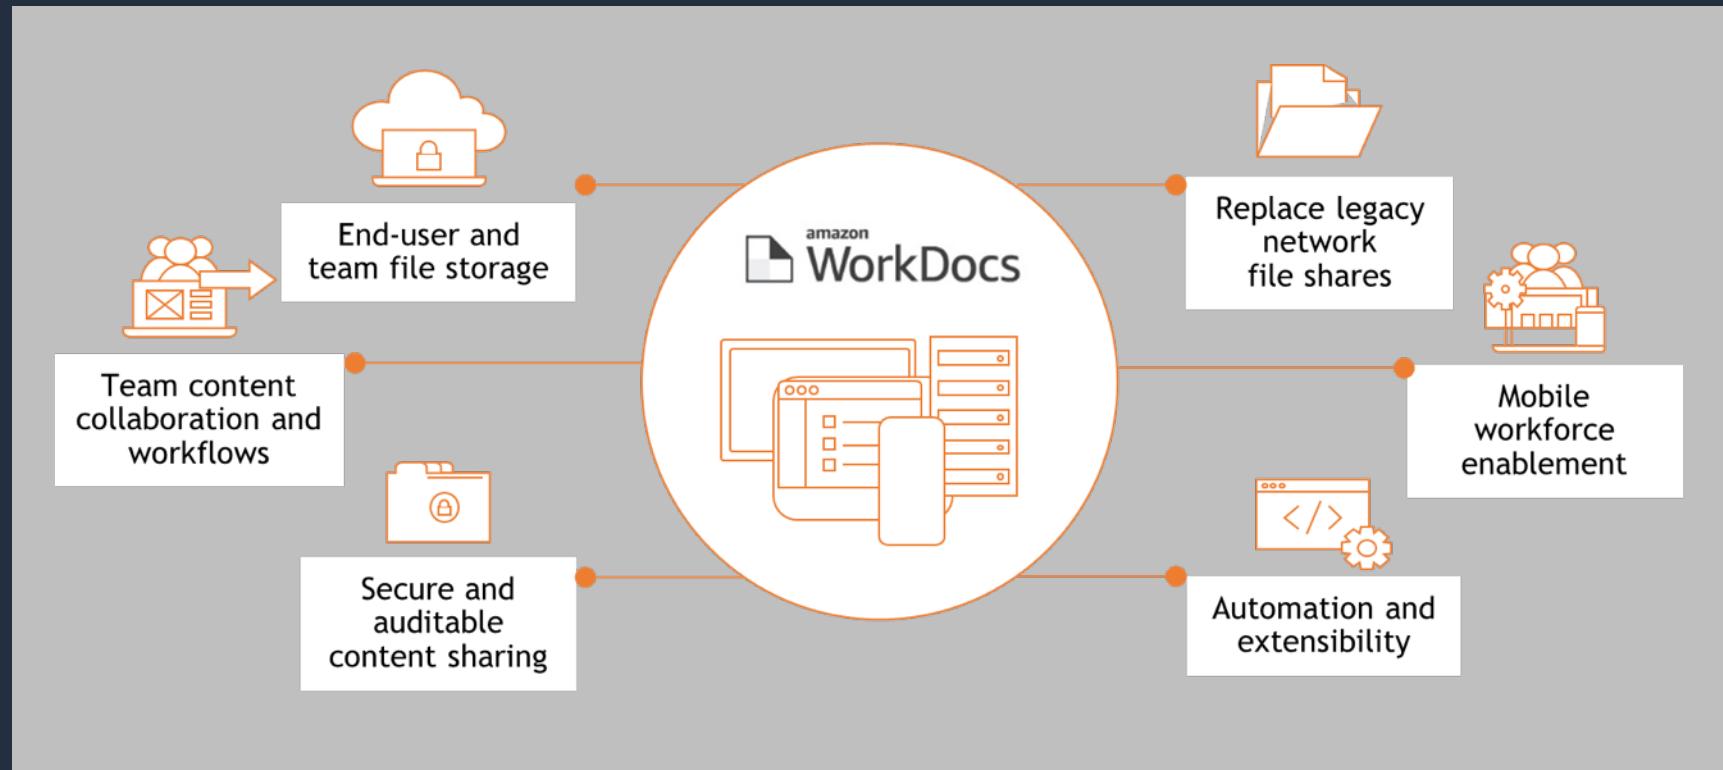
- Provides secure, one-click access to your internal websites and web apps using mobile phone browsers
- Does not require VPN client or App





AWS WorkDocs

- Fully managed, secure content creation, storage, and collaboration service
- Create, edit, and share content that's centrally stored on AWS



AWS IoT Core





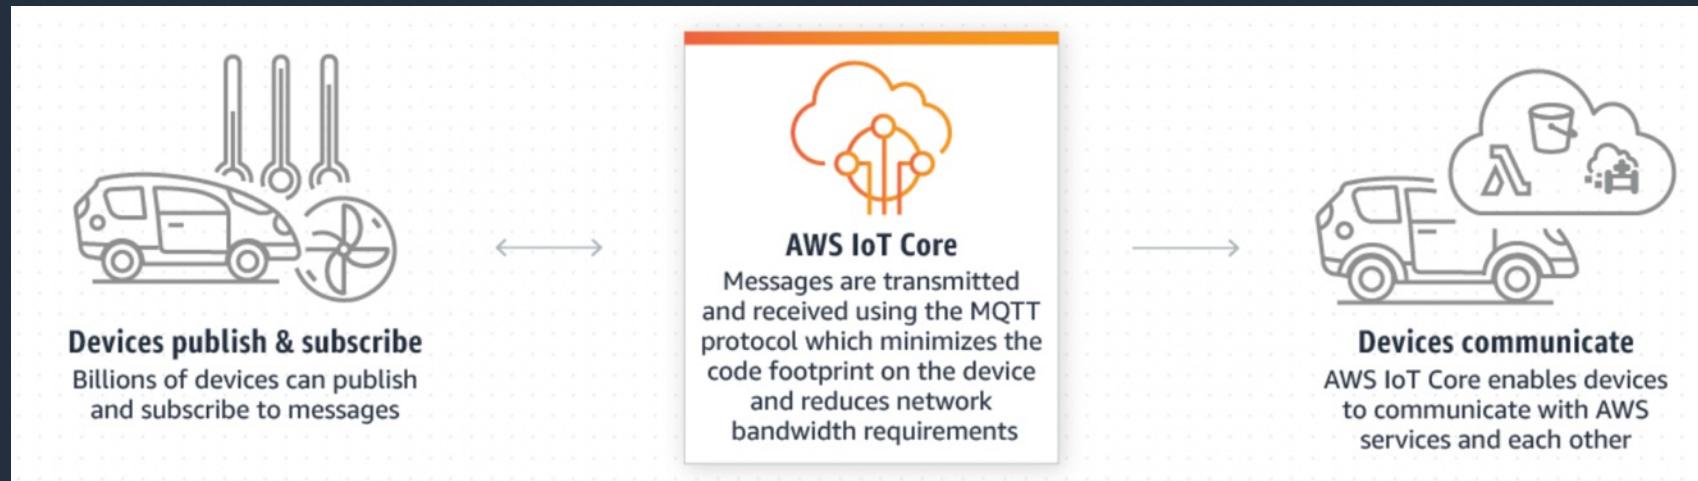
The Internet of Things (IoT)

- Describes the network of physical objects that are embedded with sensors or software
- Each IoT device can communicate and exchange data with other devices and systems
- Use cases include:
 - Smart home automation
 - Smart healthcare
 - Manufacturing
 - Agriculture

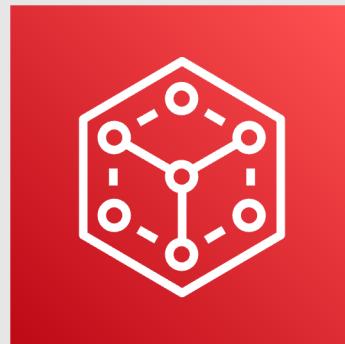


AWS IoT Core

- Lets you connect IoT devices to the AWS cloud without the need to provision or manage servers
- Can support billions of devices and trillions of messages



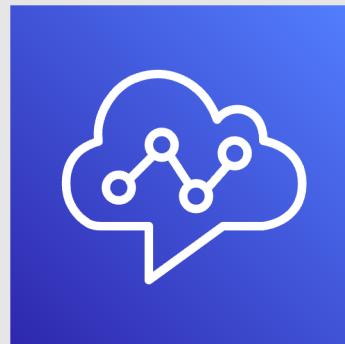
AWS Device Farm



AWS Knowledge Center



Amazon Connect





Amazon Connect

- Cloud Contact Center
- Facilitates human agents in helping customers
- Think human connection, not network connection!
- Features include telephony automation, chatbots, task management, and analytics

SECTION 17

Exam Preparation and Tips

Booking your Exam



Exam Preparation Tips





Exam Preparation Tips

- Dedicate regular time to learning
- Use the free study plan
- Use practice tests early and regularly
- Review knowledge areas where you score poorly
- Don't book the exam until you're ready
- Non-native English speakers can request an extension (extra 30 minutes)
- If you've taken another AWS exam before, use your 50% discount voucher

Exam Question Walkthrough

